
Non-Linear Least-Squares Minimization and Curve-Fitting for Python

Release 1.3.2

Matthew Newville, Till Stensitzki, Renee Otten, and others

Nov 05, 2024

CONTENTS

1	Getting started with Non-Linear Least-Squares Fitting	3
2	Downloading and Installation	7
2.1	Prerequisites	7
2.2	Downloads	7
2.3	Installation	8
2.4	Development Version	8
2.5	Testing	8
2.6	Acknowledgements	8
2.7	Copyright, Licensing, and Re-distribution	10
3	Getting Help	13
4	Frequently Asked Questions	15
4.1	What's the best way to ask for help or submit a bug report?	15
4.2	Why did my script break when upgrading from lmfit 0.8.3 to 0.9.0?	15
4.3	I get import errors from IPython	15
4.4	How can I fit multi-dimensional data?	15
4.5	How can I fit multiple data sets?	16
4.6	How can I fit complex data?	16
4.7	How should I cite LMFIT?	16
4.8	I get errors from NaN in my fit. What can I do?	17
4.9	Why are Parameter values sometimes stuck at initial values?	18
4.10	Why are uncertainties in Parameters sometimes not determined?	18
4.11	Can Parameters be used for Array Indices or Discrete Values?	18
5	Parameter and Parameters	23
5.1	The Parameter class	23
5.2	The Parameters class	25
5.3	The create_params() function	29
5.4	Simple Example	30
6	Performing Fits and Analyzing Outputs	33
6.1	The minimize() function	33
6.2	Writing a Fitting Function	35
6.3	Types of Data to Use for Fitting	37
6.4	Choosing Different Fitting Methods	37
6.5	MinimizerResult – the optimization result	39
6.6	Getting and Printing Fit Reports	44
6.7	Using a Iteration Callback Function	46
6.8	Using the Minimizer class	46

6.9	Minimizer.emcee() - calculating the posterior probability distribution of parameters	57
7	Modeling Data and Curve Fitting	65
7.1	Motivation and simple example: Fit data to Gaussian profile	65
7.2	The Model class	69
7.3	The ModelResult class	83
7.4	Composite Models : adding (or multiplying) Models	106
8	Built-in Fitting Models in the models module	113
8.1	Peak-like models	113
8.2	Linear and Polynomial Models	126
8.3	Periodic Models	130
8.4	Step-like models	130
8.5	Exponential and Power law models	132
8.6	Two dimensional Peak-like models	134
8.7	User-defined Models	134
8.8	Example 1: Fit Peak data to Gaussian, Lorentzian, and Voigt profiles	136
8.9	Example 2: Fit data to a Composite Model with pre-defined models	141
8.10	Example 3: Fitting Multiple Peaks – and using Prefixes	143
8.11	Example 4: Using a Spline Model	148
9	Calculation of confidence intervals	155
9.1	Method used for calculating confidence intervals	155
9.2	A basic example	155
9.3	Working without standard error estimates	156
9.4	Calculating and visualizing maps of χ^2	157
9.5	An advanced example for evaluating confidence intervals	161
9.6	Confidence Interval Functions	165
10	Bounds Implementation	169
11	Using Mathematical Constraints	171
11.1	Overview	171
11.2	Supported Operators, Functions, and Constants	172
11.3	Using Inequality Constraints	173
11.4	Advanced usage of Expressions in lmfit	173
12	Release Notes	175
12.1	Version 1.3.2 Release Notes (July 19, 2024)	175
12.2	Version 1.3.1 Release Notes (April 19, 2024)	175
12.3	Version 1.3.0 Release Notes (April 4, 2024)	176
12.4	Version 1.2.2 Release Notes (July 14, 2023)	177
12.5	Version 1.2.1 Release Notes (May 02, 2023)	177
12.6	Version 1.2.0 Release Notes (April 05, 2023)	177
12.7	Version 1.1.0 Release Notes (November 27, 2022)	178
12.8	Version 1.0.3 Release Notes (October 14, 2021)	179
12.9	Version 1.0.2 Release Notes (February 7, 2021)	180
12.10	Version 1.0.1 Release Notes	181
12.11	Version 1.0.0 Release Notes	182
12.12	Version 0.9.15 Release Notes	182
12.13	Version 0.9.14 Release Notes	183
12.14	Version 0.9.13 Release Notes	183
12.15	Version 0.9.12 Release Notes	184
12.16	Version 0.9.10 Release Notes	185
12.17	Version 0.9.9 Release Notes	185

12.18 Version 0.9.6 Release Notes	185
12.19 Version 0.9.5 Release Notes	186
12.20 Version 0.9.4 Release Notes	186
12.21 Version 0.9.3 Release Notes	186
12.22 Version 0.9.0 Release Notes	186
13 Examples gallery	189
13.1 Fit with Data in a pandas DataFrame	189
13.2 Using an ExpressionModel	191
13.3 Fit Using Inequality Constraint	192
13.4 Fit Using differential_evolution Algorithm	194
13.5 Fit Using Bounds	197
13.6 Fit with Algebraic Constraint	199
13.7 Fit Specifying Different Reduce Function	202
13.8 Building a lmfit model with SymPy	204
13.9 Fit comparing leastsq and basin hopping, or other methods	207
13.10 Fit Multiple Data Sets	210
13.11 Fit using the Model interface	213
13.12 Fit Specifying a Function to Compute the Jacobian	218
13.13 Outlier detection via leave-one-out	221
13.14 Emcee and the Model Interface	226
13.15 Complex Resonator Model	232
13.16 Model Selection using lmfit and emcee	237
13.17 Calculate Confidence Intervals	240
13.18 Fit Two Dimensional Peaks	245
13.19 Global minimization using the brute method (a.k.a. grid search)	253
14 Examples from the documentation	267
14.1 Examples from the documentation	267
Python Module Index	343
Index	345

Lmfit provides a high-level interface to non-linear optimization and curve fitting problems for Python. It builds on and extends many of the optimization methods of `scipy.optimize`. Initially inspired by (and named for) extending the Levenberg-Marquardt method from `scipy.optimize.leastsq`, Lmfit now provides a number of useful enhancements to optimization and data fitting problems, including:

- Using `Parameter` objects instead of plain floats as variables. A `Parameter` has a value that can be varied during the fit or kept at a fixed value. It can have upper and/or lower bounds. A Parameter can even have a value that is constrained by an algebraic expression of other Parameter values. As a Python object, a Parameter can also have attributes such as a standard error, after a fit that can estimate uncertainties.
- Ease of changing fitting algorithms. Once a fitting model is set up, one can change the fitting algorithm used to find the optimal solution without changing the objective function.
- Improved estimation of confidence intervals. While `scipy.optimize.leastsq` will automatically calculate uncertainties and correlations from the covariance matrix, the accuracy of these estimates is sometimes questionable. To help address this, Lmfit has functions to explicitly explore parameter space and determine confidence levels even for the most difficult cases. Additionally, Lmfit will use the `numdifftools` package (if installed) to estimate parameter uncertainties and correlations for algorithms that do not natively support this in SciPy.
- Improved curve-fitting with the `Model` class. This extends the capabilities of `scipy.optimize.curve_fit`, allowing you to turn a function that models your data into a Python class that helps you parametrize and fit data with that model.
- Many *built-in models* for common lineshapes are included and ready to use.

The Lmfit package is Free software, using an Open Source license. The software and this document are works in progress. If you are interested in participating in this effort please use the [Lmfit GitHub repository](#).

GETTING STARTED WITH NON-LINEAR LEAST-SQUARES FITTING

The lmfit package provides simple tools to help you build complex fitting models for non-linear least-squares problems and apply these models to real data. This section gives an overview of the concepts and describes how to set up and perform simple fits. Some basic knowledge of Python, NumPy, and modeling data are assumed – this is not a tutorial on why or how to perform a minimization or fit data, but is rather aimed at explaining how to use lmfit to do these things.

In order to do a non-linear least-squares fit of a model to data or for any other optimization problem, the main task is to write an *objective function* that takes the values of the fitting variables and calculates either a scalar value to be minimized or an array of values that are to be minimized, typically in the least-squares sense. For many data fitting processes, the latter approach is used, and the objective function should return an array of `(data-model)`, perhaps scaled by some weighting factor such as the inverse of the uncertainty in the data. For such a problem, the chi-square (χ^2) statistic is often defined as:

$$\chi^2 = \sum_i^N \frac{[y_i^{\text{meas}} - y_i^{\text{model}}(\mathbf{v})]^2}{\epsilon_i^2}$$

where y_i^{meas} is the set of measured data, $y_i^{\text{model}}(\mathbf{v})$ is the model calculation, \mathbf{v} is the set of variables in the model to be optimized in the fit, and ϵ_i is the estimated uncertainty in the data, respectively.

In a traditional non-linear fit, one writes an objective function that takes the variable values and calculates the residual array $y_i^{\text{meas}} - y_i^{\text{model}}(\mathbf{v})$, or the residual array scaled by the data uncertainties, $[y_i^{\text{meas}} - y_i^{\text{model}}(\mathbf{v})]/\epsilon_i$, or some other weighting factor.

As a simple concrete example, one might want to model data with a decaying sine wave, and so write an objective function like this:

```
from numpy import exp, sin

def residual(variables, x, data, uncertainty):
    """Model a decaying sine wave and subtract data."""
    amp = variables[0]
    phaseshift = variables[1]
    freq = variables[2]
    decay = variables[3]

    model = amp * sin(x*freq + phaseshift) * exp(-x*x*decay)

    return (data-model) / uncertainty
```

To perform the minimization with `scipy.optimize`, one would do this:

```

from numpy import linspace, random
from scipy.optimize import leastsq

# generate synthetic data with noise
x = linspace(0, 100)
noise = random.normal(size=x.size, scale=0.2)
data = 7.5 * sin(x*0.22 + 2.5) * exp(-x*x*0.01) + noise

# generate experimental uncertainties
uncertainty = abs(0.16 + random.normal(size=x.size, scale=0.05))

variables = [10.0, 0.2, 3.0, 0.007]
out = leastsq(residual, variables, args=(x, data, uncertainty))

```

Though it is wonderful to be able to use Python for such optimization problems, and the SciPy library is robust and easy to use, the approach here is not terribly different from how one would do the same fit in C or Fortran. There are several practical challenges to using this approach, including:

- The user has to keep track of the order of the variables, and their meaning – `variables[0]` is the `amplitude`, `variables[2]` is the `frequency`, and so on, although there is no intrinsic meaning to this order.
- If the user wants to fix a particular variable (*not* vary it in the fit), the residual function has to be altered to have fewer variables, and have the corresponding constant value passed in some other way. While reasonable for simple cases, this quickly becomes a significant work for more complex models, and greatly complicates modeling for people not intimately familiar with the details of the fitting code.
- There is no simple, robust way to put bounds on values for the variables, or enforce mathematical relationships between the variables. While some optimization methods in SciPy do provide bounds, they require bounds to be set for all variables with separate arrays that are in the same arbitrary order as variable values. Again, this is acceptable for small or one-off cases, but becomes painful if the fitting model needs to change.
- In some cases, constraints can be placed on Parameter values, but this is a pretty opaque and complex process.

While these shortcomings can be worked around with some work, they are all essentially due to the use of arrays or lists to hold the variables. This closely matches the implementation of the underlying Fortran code, but does not fit very well with Python's rich selection of objects and data structures. The key concept in lmfit is to define and use `Parameter` objects instead of plain floating point numbers as the variables for the fit. Using `Parameter` objects (or the closely related `Parameters` – a dictionary of `Parameter` objects), allows one to do the following:

- forget about the order of variables and refer to Parameters by meaningful names.
- place bounds on Parameters as attributes, without worrying about preserving the order of arrays for variables and boundaries, and without relying on the solver to support bounds itself.
- fix Parameters, without having to rewrite the objective function.
- place algebraic constraints on Parameters.

To illustrate the value of this approach, we can rewrite the above example for the decaying sine wave as:

```

from numpy import exp, sin

from lmfit import minimize, Parameters

def residual(params, x, data, uncertainty):
    amp = params['amp']
    phaseshift = params['phase']

```

(continues on next page)

(continued from previous page)

```

freq = params['frequency']
decay = params['decay']

model = amp * sin(x*freq + phaseshift) * exp(-x*x*decay)

return (data-model) / uncertainty

params = Parameters()
params.add('amp', value=10)
params.add('decay', value=0.007)
params.add('phase', value=0.2)
params.add('frequency', value=3.0)

out = minimize(residual, params, args=(x, data, uncertainty))

```

At first look, we simply replaced a list of values with a dictionary, so that we can access Parameters by name. Just by itself, this is better as it allows separation of the objective function from the code using it.

Note that creation of Parameters here could also be done as:

Added in version 1.2.0.

```

from lmfit import create_params

params = create_params(amp=10, decay=0.007, phase=0.2, frequency=3.0)

```

where keyword/value pairs set Parameter names and their initial values.

Either when using `create_param()` or `Parameters`, the resulting `params` object is an instance of `Parameters`, which acts like a dictionary, with keys being the Parameter name and values being individual `Parameter` objects. These `Parameter` objects hold the value and several other attributes that control how a Parameter acts. For example, Parameters can be fixed or bounded; setting attributes to control this behavior can be done during definition, as with:

```

params = Parameters()
params.add('amp', value=10, vary=False)
params.add('decay', value=0.007, min=0.0)
params.add('phase', value=0.2)
params.add('frequency', value=3.0, max=10)

```

Here `vary=False` will prevent the value from changing in the fit, and `min=0.0` will set a lower bound on that parameter's value. The same thing can be accomplished by providing a dictionary of attribute values to `create_params()`:

Added in version 1.2.0.

```

params = create_params(amp={'value': 10, 'vary': False},
                      decay={'value': 0.007, 'min': 0},
                      phase=0.2,
                      frequency={'value': 3.0, 'max': 10})

```

Parameter attributes can also be modified after they have been created:

```

params['amp'].vary = False
params['decay'].min = 0.10

```

Importantly, our objective function remains unchanged. This means the objective function can simply express the parametrized phenomenon to be calculated, accessing Parameter values by name and separating the choice of parameters to be varied in the fit.

The `params` object can be copied and modified to make many user-level changes to the model and fitting process. Of course, most of the information about how your data is modeled goes into the objective function, but the approach here allows some external control; that is, control by the `user` performing the fit, instead of by the author of the objective function.

Finally, in addition to the `Parameters` approach to fitting data, lmfit allows switching optimization methods without changing the objective function, provides tools for generating fitting reports, and provides a better determination of Parameters confidence levels.

DOWNLOADING AND INSTALLATION

2.1 Prerequisites

Lmfit works with [Python](#) versions 3.8 and higher. Version 0.9.15 is the final version to support Python 2.7.

Lmfit requires the following Python packages, with versions given:

- [NumPy](#) version 1.23 or higher.
- [SciPy](#) version 1.8 or higher.
- [asteval](#) version 1.0 or higher.
- [uncertainties](#) version 3.2.2 or higher.
- [dill](#) version 0.3.4 or higher.

All of these are readily available on PyPI, and are installed automatically if installing with `pip install lmfit`.

In order to run the test suite, the [pytest](#), [pytest-cov](#), and [flaky](#) packages are required. Some functionality requires the [emcee](#) (version 3+), [corner](#), [pandas](#), [Jupyter](#), [matplotlib](#), or [numdifftools](#) packages. These are not installed automatically, but we highly recommend each of them.

For building the documentation and generating the examples gallery, [matplotlib](#), [emcee](#) (version 3+), [corner](#), [Sphinx](#), [sphinx-gallery](#), [jupyter_sphinx](#), [ipykernel](#), [Pillow](#), and [SymPy](#) are required. For generating the PDF documentation, the Python packages [sphinxcontrib-svg2pdfconverter](#) and [cairosvg](#) are also required, as well as the LaTex package [Latexmk](#) (which is included by default in some LaTex distributions).

Please refer to `pyproject.toml` under `project.optional-dependencies` for a list of all dependencies that are needed if you want to participate in the development of lmfit. You can install all these dependencies automatically by doing `pip install lmfit[all]`, or select only a subset (e.g., `dev`, `doc`, or `test`).

Please note: the “original” `python setup.py install` is deprecated, but we will provide a shim `setup.py` file for as long as `Python` and/or `setuptools` allow the use of this legacy command.

2.2 Downloads

The latest stable version of lmfit is 1.3.2 and is available from [PyPI](#). Check the [Release Notes](#) for a list of changes compared to earlier releases.

2.3 Installation

The easiest way to install lmfit is with:

```
pip install lmfit
```

For Anaconda Python, lmfit is not an official package, but several Anaconda channels provide it, allowing installation with (for example):

```
conda install -c conda-forge lmfit
```

2.4 Development Version

To get the latest development version from the [lmfit GitHub repository](https://github.com/lmfit/lmfit-py), use:

```
git clone https://github.com/lmfit/lmfit-py.git
```

and install using:

```
pip install --upgrade build pip setuptools wheel
```

to install the required build dependencies and then do:

```
python -m build
pip install ".[all]"
```

to generate the wheel and install lmfit with all its dependencies.

We welcome all contributions to lmfit! If you cloned the repository for this purpose, please read [CONTRIBUTING.md](#) for more detailed instructions.

2.5 Testing

A battery of tests scripts that can be run with the [pytest](#) testing framework is distributed with lmfit in the `tests` folder. These are automatically run as part of the development process. For any release or any master branch from the git repository, running `pytest` should run all of these tests to completion without errors or failures.

Many of the examples in this documentation are distributed with lmfit in the `examples` folder, and should also run for you. Some of these examples assume that [matplotlib](#) has been installed and is working correctly.

2.6 Acknowledgements

Many people have contributed to lmfit. The attribution of credit in a project such as this is difficult to get perfect, and there are no doubt important contributions that are missing or under-represented here. Please consider this file as part of the code and documentation that may have bugs that need fixing.

Some of the largest and most important contributions (in approximate order

(continues on next page)

(continued from previous page)

of size of the contribution to the existing code) are from:

Matthew Newville wrote the original version and maintains the project.

Renee Otten wrote the brute force method, implemented the basin-hopping and AMPGO global solvers, implemented uncertainty calculations for scalar minimizers and has greatly improved the code, testing, and documentation and overall project.

Till Stensitzki wrote the improved estimates of confidence intervals, and contributed many tests, bug fixes, and documentation.

A. R. J. Nelson added differential_evolution, emcee, and greatly improved the code, docstrings, and overall project.

Antonino Ingargiola wrote much of the high level Model code and has provided many bug fixes and improvements.

Daniel B. Allan wrote much of the original version of the high level Model code, and many improvements to the testing and documentation.

Austen Fox fixed many of the built-in model functions and improved the testing and documentation of these.

Michal Rawlik added plotting capabilities for Models.

The method used for placing bounds on parameters was derived from the clear description in the MINUIT documentation, and adapted from J. J. Helmus's Python implementation in `leastsqbounds.py`.

E. O. Le Bigot wrote the uncertainties package, a version of which was used by lmfit for many years, and is now an external dependency.

The original AMPGO code came from Andrea Gavana and was adopted for lmfit.

The propagation of parameter uncertainties to uncertainties in a Model was adapted from the excellent description at <https://www.astro.rug.nl/software/kapteyn/kmpfittutorial.html#confidence-and-prediction-intervals>, which references the original work of: J. Wolberg, Data Analysis Using the Method of Least Squares, 2006, Springer.

Additional patches, bug fixes, and suggestions have come from Faustin Carter, Christoph Deil, Francois Boulogne, Thomas Caswell, Colin Brosseau, nmearl, Gustavo Pasquevich, Clemens Prescher, LiCode, Ben Gamari, Yoav Roam, Alexander Stark, Alexandre Beelen, Andrey Aristov, Nicholas Zobrist, Ethan Welty, Julius Zimmermann, Mark Dean, Arun Persaud, Ray Osborn, @lneuhaus, Marcel Stimberg, Yoshiera Huang, Leon Foks, Sebastian Weigand, Florian LB, Michael Hudson-Doyle, Ruben Verweij, @jedzill4, @spalato, Jens Hedegaard Nielsen, Martin Majli, Kristian Meyer, @azelcer, Ivan Usov, and many others.

(continues on next page)

(continued from previous page)

The lmfit code obviously depends on, and owes a very large debt to the code in `scipy.optimize`. Several discussions on the SciPy-user and lmfit mailing lists have also led to improvements in this code.

2.7 Copyright, Licensing, and Re-distribution

The LMFIT-py code is distributed under the following license:

BSD-3

Copyright 2022 Matthew Newville, The University of Chicago
Renee Otten, Brandeis University
Till Stensitzki, Freie Universität Berlin
A. R. J. Nelson, Australian Nuclear Science and Technology Organisation
Antonino Ingargiola, University of California, Los Angeles
Daniel B. Allen, Johns Hopkins University
Michal Rawlik, Eidgenössische Technische Hochschule, Zurich

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

1. Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its contributors may be used to endorse or promote products derived from this software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

Some code has been taken from the `scipy` library whose licence is below.

Copyright (c) 2001, 2002 Enthought, Inc.
All rights reserved.

(continues on next page)

(continued from previous page)

Copyright (c) 2003-2019 SciPy Developers.
All rights reserved.

Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions are met:

- a. Redistributions of source code must retain the above copyright notice,
this list of conditions and the following disclaimer.
- b. Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
- c. Neither the name of Enthought nor the names of the SciPy Developers
may be used to endorse or promote products derived from this software
without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS"
AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE
IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE
ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDERS OR CONTRIBUTORS
BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY,
OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF
SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS
INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN
CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE)
ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF
THE POSSIBILITY OF SUCH DAMAGE.

Some code has been taken from the AMPGO library of Andrea Gavana, which was
released under a MIT license.

CHAPTER
THREE

GETTING HELP

If you have questions, comments, or suggestions for LMFIT, please use the [mailing list](#). This provides an on-line conversation that is both archived well and can be searched easily with standard web searches. If you find a bug in the code or documentation, use [GitHub Issues](#) to submit a report. If you have an idea for how to solve the problem and are familiar with Python and GitHub, submitting a GitHub Pull Request would be greatly appreciated.

If you are unsure whether to use the mailing list or the Issue tracker, please start a conversation on the [mailing list](#). That is, the problem you're having may or may not be due to a bug. If it is due to a bug, creating an Issue from the conversation is easy. If it is not a bug, the problem will be discussed and then the Issue will be closed. While one *can* search through closed Issues on GitHub, these are not so easily searched, and the conversation is not easily useful to others later. Starting the conversation on the mailing list with "How do I do this?" or "Why didn't this work?" instead of "This should work and doesn't" is generally preferred, and will better help others with similar questions. Of course, there is not always an obvious way to decide if something is a Question or an Issue, and we will try our best to engage in all discussions.

FREQUENTLY ASKED QUESTIONS

A list of common questions.

4.1 What's the best way to ask for help or submit a bug report?

See [Getting Help](#).

4.2 Why did my script break when upgrading from lmfit 0.8.3 to 0.9.0?

See [Version 0.9.0 Release Notes](#).

4.3 I get import errors from IPython

If you see something like:

```
from IPython.html.widgets import Dropdown
ImportError: No module named 'widgets'
```

then you need to install the ipywidgets package, try: `pip install ipywidgets`.

4.4 How can I fit multi-dimensional data?

The fitting routines accept data arrays that are one-dimensional and double precision. So you need to convert the data and model (or the value returned by the objective function) to be one-dimensional. A simple way to do this is to use `numpy.ndarray.flatten`, for example:

```
def residual(params, x, data=None):
    ....
    resid = calculate_multidim_residual()
    return resid.flatten()
```

4.5 How can I fit multiple data sets?

As above, the fitting routines accept data arrays that are one-dimensional and double precision. So you need to convert the sets of data and models (or the value returned by the objective function) to be one-dimensional. A simple way to do this is to use `numpy.concatenate`. As an example, here is a residual function to simultaneously fit two lines to two different arrays. As a bonus, the two lines share the 'offset' parameter:

```
import numpy as np

def fit_function(params, x=None, dat1=None, dat2=None):
    model1 = params['offset'] + x * params['slope1']
    model2 = params['offset'] + x * params['slope2']

    resid1 = dat1 - model1
    resid2 = dat2 - model2
    return np.concatenate((resid1, resid2))
```

4.6 How can I fit complex data?

As with working with multi-dimensional data, you need to convert your data and model (or the value returned by the objective function) to be double precision, floating point numbers. The simplest approach is to use `numpy.ndarray.view`, perhaps like:

```
import numpy as np

def residual(params, x, data=None):
    ...
    resid = calculate_complex_residual()
    return resid.view(float)
```

Alternately, you can use the `lmfit.Model` class to wrap a fit function that returns a complex vector. It will automatically apply the above prescription when calculating the residual. The benefit to this method is that you also get access to the plot routines from the `ModelResult` class, which are also complex-aware.

4.7 How should I cite LMFIT?

For LMFIT 1.3.2, use <https://doi.org/10.5281/zenodo.12785036>.

For other versions, see <https://zenodo.org/records/12785036>

4.8 I get errors from NaN in my fit. What can I do?

The solvers used by lmfit use NaN (see <https://en.wikipedia.org/wiki/NaN>) values as signals that the calculation cannot continue. If any value in the residual array (typically $(\text{data}-\text{model})^*\text{weight}$) is NaN, then calculations of chi-square or comparisons with other residual arrays to try find a better fit will also give NaN and fail. There is no sensible way for lmfit or any of the optimization routines to know how to handle such NaN values. They indicate that numerical calculations are not sensible and must stop.

This means that if your objective function (if using `minimize`) or model function (if using `Model`) generates a NaN, the fit will stop immediately. If your objective or model function generates a NaN, you really must handle that.

4.8.1 nan_policy

If you are using `lmfit.Model` and the NaN values come from your data array and are meant to indicate missing values, or if you using `lmfit.minimize()` with the same basic intention, then it might be possible to get a successful fit in spite of the NaN values. To do this, you can add a `nan_policy='omit'` argument to `lmfit.minimize()`, or when creating a `lmfit.Model`, or when running `lmfit.Model.fit()`.

In order for this to be effective, the number of NaN values cannot ever change during the fit. If the NaN values come from the data and not the calculated model, that should be the case.

4.8.2 Common sources of NaN

If you are seeing errors due to NaN values, you will need to figure out where they are coming from and eliminate them. It is sometimes difficult to tell what causes NaN values. Keep in mind that all values should be assumed to be either scalar values or numpy arrays of double precision real numbers when fitting. Some of the most likely causes of NaNs are:

- taking `sqrt(x)` or `log(x)` where `x` is negative.
- doing `x**y` where `x` is negative. Since `y` is real, there will be a fractional component, and a negative number to a fractional exponent is not a real number.
- doing `x/y` where both `x` and `y` are 0.

If you use these very common constructs in your objective or model function, you should take some caution for what values you are passing these functions and operators. Many special functions have similar limitations and should also be viewed with some suspicion if NaNs are being generated.

A related problem is the generation of Inf (Infinity in floating point), which generally comes from `exp(x)` where `x` has values greater than 700 or so, so that the resulting value is greater than 1.e308. Inf is only slightly better than NaN. It will completely ruin the ability to do the fit. However, unlike NaN, it is also usually clear how to handle Inf, as you probably won't ever have values greater than 1.e308 and can therefore (usually) safely clip the argument passed to `exp()` to be smaller than about 700.

4.9 Why are Parameter values sometimes stuck at initial values?

In order for a Parameter to be optimized in a fit, changing its value must have an impact on the fit residual (data-model when curve fitting, for example). If a fit has not changed one or more of the Parameters, it means that changing those Parameters did not change the fit residual.

Normally (that is, unless you specifically provide a function for calculating the derivatives, in which case you probably would not be asking this question ;)), the fitting process begins by making a very small change to each Parameter value to determine which way and how large of a change to make for the parameter: This is the derivative or Jacobian (change in residual per change in parameter value). By default, the change made for each variable Parameter is to multiply its value by (1.0+1.0e-8) or so (unless the value is below about 1.e-15, in which case it adds 1.0e-8). If that small change does not change the residual, then the value of the Parameter will not be updated.

Parameter values that are “way off” are a common reason for Parameters being stuck at initial values. As an example, imagine fitting peak-like data with an x range of 0 to 10, peak centered at 6, and a width of 1 or 2 or so, as in the example at [Model - gaussian](#). A Gaussian function with an initial value of for the peak center at 5 and an initial width or 5 will almost certainly find a good fit. An initial value of the peak center of -50 will end up being stuck with a “bad fit” because a small change in Parameters will still lead the modeled Gaussian to have no intensity over the actual range of the data. You should make sure that initial values for Parameters are reasonable enough to actually effect the fit. As it turns out in the example linked to above, changing the center value to any value between about 0 and 10 (that is, the data range) will result to a good fit.

Another common cause for Parameters being stuck at initial values is when the initial value is at a boundary value. For this case, too, a small change in the initial value for the Parameter will still leave the value at the boundary value and not show any real change in the residual.

If you’re using bounds, make sure the initial values for the Parameters are not at the boundary values.

Finally, one reason for a Parameter to not change is that they are actually used as discrete values. This is discussed below in [Can Parameters be used for Array Indices or Discrete Values?](#).

4.10 Why are uncertainties in Parameters sometimes not determined?

In order for Parameter uncertainties to be estimated, each variable Parameter must actually change the fit, and cannot be stuck at an initial value or at a boundary value. See [Why are Parameter values sometimes stuck at initial values?](#) for why values may not change from their initial values.

4.11 Can Parameters be used for Array Indices or Discrete Values?

The short answer is “No”: variables in all of the fitting methods used in `lmfit` (and all of those available in `scipy.optimize`) are treated as continuous values, and represented as double precision floating point values. As an important example, you cannot have a variable that is somehow constrained to be an integer.

Still, it is a rather common question of how to fit data to a model that includes a breakpoint, perhaps

$$f(x; x_0, a, b, c) = \begin{cases} c & \text{for } x < x_0 \\ a + bx^2 & \text{for } x > x_0 \end{cases}$$

That you implement with a model function and use to fit data like this:

```

import numpy as np

import lmfit

def quad_off(x, x0, a, b, c):
    model = a + b * x**2
    model[np.where(x < x0)] = c
    return model

x0 = 19
b = 0.02
a = 2.0
xdat = np.linspace(0, 100, 101)
ydat = a + b * xdat**2
ydat[np.where(xdat < x0)] = a + b * x0**2
ydat += np.random.normal(scale=0.1, size=xdat.size)

mod = lmfit.Model(quad_off)
pars = mod.make_params(x0=22, a=1, b=1, c=1)

result = mod.fit(ydat, pars, x=xdat)
print(result.fit_report())

```

```

[[Model]]
    Model(quad_off)
[[Fit Statistics]]
    # fitting method    = leastsq
    # function evals   = 30
    # data points     = 101
    # variables       = 4
    chi-square        = 3.36908247
    reduced chi-square = 0.03473281
    Akaike info crit  = -335.448487
    Bayesian info crit = -324.988005
    R-squared         = 0.99999050
[[Variables]]
    x0: 21.9999940 +/- 1.7417e-05 (0.00%) (init = 22)
    a:  2.03865555 +/- 0.03854349 (1.89%) (init = 1)
    b:  0.01999489 +/- 7.5198e-06 (0.04%) (init = 1)
    c:  9.31956084 +/- 0.03973366 (0.43%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
    C(a, b) = -0.8368
    C(x0, a) = +0.1871
    C(x0, b) = -0.1508

```

This will not result in a very good fit, as the value for `x0` cannot be found by making a small change in its value. Specifically, `model[np.where(x < x0)]` will give the same result for `x0=22` and `x0=22.001`, and so that value is not changed during the fit.

There are a couple ways around this problem. First, you may be able to make the fit depend on `x0` in a way that is not just discrete. That depends on your model function. A second option is to treat the break not as a hard break but as a more gentle transition with a sigmoidal function, such as an error function. Like the break-point, these will go from 0

to 1, but more gently and with some finite value leaking into neighboring points. The amount of leakage or width of the step can also be adjusted.

A simple modification of the above to use an error function would look like this and give better fit results:

```
import numpy as np
from scipy.special import erf

import lmfit

def quad_off(x, x0, a, b, c):
    m1 = a + b * x**2
    m2 = c * np.ones(len(x))
    # step up from 0 to 1 at x0: (erf(x-x0)+1)/2
    # step down from 1 to 0 at x0: (1-erf(x-x0))/2
    model = m1 * (erf(x-x0)+1)/2 + m2 * (1-erf(x-x0))/2
    return model

x0 = 19
b = 0.02
a = 2.0
xdat = np.linspace(0, 100, 101)
ydat = a + b * xdat**2
ydat[np.where(xdat < x0)] = a + b * x0**2
ydat += np.random.normal(scale=0.1, size=xdat.size)

mod = lmfit.Model(quad_off)
pars = mod.make_params(x0=22, a=1, b=1, c=1)

result = mod.fit(ydat, pars, x=xdat)
print(result.fit_report())
```

```
[[Model]]
    Model(quad_off)
[[Fit Statistics]]
    # fitting method    = leastsq
    # function evals   = 62
    # data points      = 101
    # variables        = 4
    chi-square         = 0.83750892
    reduced chi-square = 0.00863411
    Akaike info crit  = -476.036831
    Bayesian info crit = -465.576349
    R-squared          = 0.99999764
[[Variables]]
    x0: 19.0868767 +/- 0.40511591 (2.12%) (init = 22)
    a:  1.98747903 +/- 0.01845982 (0.93%) (init = 1)
    b:  0.02000216 +/- 3.6534e-06 (0.02%) (init = 1)
    c:  9.19289010 +/- 0.02133535 (0.23%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
    C(a, b) = -0.8273
    C(x0, a) = +0.1560
```

(continues on next page)

(continued from previous page)

```
C(x0, b) = -0.1248
C(x0, c) = -0.1056
```

The natural width of the error function is about 2 x units, but you can adjust this, shortening it with `erf((x-x0)**2)` to give a sharper transition for example.

PARAMETER AND PARAMETERS

This chapter describes the `Parameter` object, which is a key concept of lmfit.

A `Parameter` is the quantity to be optimized in all minimization problems, replacing the plain floating point number used in the optimization routines from `scipy.optimize`. A `Parameter` has a value that can either be varied in the fit or held at a fixed value, and can have lower and/or upper bounds placed on the value. It can even have a value that is constrained by an algebraic expression of other Parameter values. Since `Parameter` objects live outside the core optimization routines, they can be used in **all** optimization routines from `scipy.optimize`. By using `Parameter` objects instead of plain variables, the objective function does not have to be modified to reflect every change of what is varied in the fit, or whether bounds can be applied. This simplifies the writing of models, allowing general models that describe the phenomenon and gives the user more flexibility in using and testing variations of that model.

Whereas a `Parameter` expands on an individual floating point variable, the optimization methods actually still need an ordered group of floating point variables. In the `scipy.optimize` routines this is required to be a one-dimensional `numpy.ndarray`. In lmfit, this one-dimensional array is replaced by a `Parameters` object, which works as an ordered dictionary of `Parameter` objects with a few additional features and methods. That is, while the concept of a `Parameter` is central to lmfit, one normally creates and interacts with a `Parameters` instance that contains many `Parameter` objects. For example, the objective functions you write for lmfit will take an instance of `Parameters` as its first argument. A table of parameter values, bounds, and other attributes can be printed using `Parameters.pretty_print()`.

5.1 The Parameter class

```
class Parameter(name, value=None, vary=True, min=-inf, max=inf, expr=None, brute_step=None,  
user_data=None)
```

A Parameter is an object that can be varied in a fit.

It is a central component of lmfit, and all minimization and modeling methods use Parameter objects.

A Parameter has a `name` attribute, and a scalar floating point `value`. It also has a `vary` attribute that describes whether the value should be varied during the minimization. Finite bounds can be placed on the Parameter's value by setting its `min` and/or `max` attributes. A Parameter can also have its value determined by a mathematical expression of other Parameter values held in the `expr` attribute. Additional attributes include `brute_step` used as the step size in a brute-force minimization, and `user_data` reserved exclusively for user's need.

After a minimization, a Parameter may also gain other attributes, including `stderr` holding the estimated standard error in the Parameter's value, and `correl`, a dictionary of correlation values with other Parameters used in the minimization.

Parameters

- `name` (`str`) – Name of the Parameter.
- `value` (`float`, *optional*) – Numerical Parameter value.

- **vary** (*bool*, *optional*) – Whether the Parameter is varied during a fit (default is True).
- **min** (*float*, *optional*) – Lower bound for value (default is `-numpy.inf`, no lower bound).
- **max** (*float*, *optional*) – Upper bound for value (default is `numpy.inf`, no upper bound).
- **expr** (*str*, *optional*) – Mathematical expression used to constrain the value during the fit (default is None).
- **brute_step** (*float*, *optional*) – Step size for grid points in the *brute* method (default is None).
- **user_data** (*optional*) – User-definable extra attribute used for a Parameter (default is None).

stderr

The estimated standard error for the best-fit value.

Type
`float`

correl

A dictionary of the correlation with the other fitted Parameters of the form:

```
{'decay': 0.404, 'phase': -0.020, 'frequency': 0.102}
```

Type
`dict`

See [Bounds Implementation](#) for details on the math used to implement the bounds with `min` and `max`.

The `expr` attribute can contain a mathematical expression that will be used to compute the value for the Parameter at each step in the fit. See [Using Mathematical Constraints](#) for more details and examples of this feature.

set(*value=None*, *vary=None*, *min=None*, *max=None*, *expr=None*, *brute_step=None*, *is_init_value=True*)

Set or update Parameter attributes.

Parameters

- **value** (*float*, *optional*) – Numerical Parameter value.
- **vary** (*bool*, *optional*) – Whether the Parameter is varied during a fit.
- **min** (*float*, *optional*) – Lower bound for value. To remove a lower bound you must use `-numpy.inf`.
- **max** (*float*, *optional*) – Upper bound for value. To remove an upper bound you must use `numpy.inf`.
- **expr** (*str*, *optional*) – Mathematical expression used to constrain the value during the fit. To remove a constraint you must supply an empty string.
- **brute_step** (*float*, *optional*) – Step size for grid points in the *brute* method. To remove the step size you must use 0.
- **is_init_value** (*bool*, *optional*) – Whether to set value as *init_value*, when setting value.

Notes

Each argument to `set()` has a default value of `None`, which will leave the current value for the attribute unchanged. Thus, to lift a lower or upper bound, passing in `None` will not work. Instead, you must set these to `-numpy.inf` or `numpy.inf`, as with:

```
par.set(min=None)      # leaves lower bound unchanged
par.set(min=-numpy.inf) # removes lower bound
```

Similarly, to clear an expression, pass a blank string, (not `None!`) as with:

```
par.set(expr=None)    # leaves expression unchanged
par.set(expr='')      # removes expression
```

Explicitly setting a value or setting `vary=True` will also clear the expression.

Finally, to clear the `brute_step` size, pass `0`, not `None`:

```
par.set(brute_step=None) # leaves brute_step unchanged
par.set(brute_step=0)    # removes brute_step
```

5.2 The Parameters class

`class Parameters(usersyms=None)`

A dictionary of Parameter objects.

It should contain all Parameter objects that are required to specify a fit model. All minimization and Model fitting routines in lmfit will use exactly one Parameters object, typically given as the first argument to the objective function.

All keys of a `Parameters()` instance must be strings and valid Python symbol names, so that the name must match `[a-zA-Z][a-zA-Z0-9_]*` and cannot be a Python reserved word.

All values of a `Parameters()` instance must be Parameter objects.

A `Parameters(xs)` instance includes an `asteval` Interpreter used for evaluation of constrained Parameters.

`Parameters()` support copying and pickling, and have methods to convert to and from serializations using json strings.

Parameters

`usersyms (dict, optional)` – Dictionary of symbols to add to the `asteval.Interpreter` (default is `None`).

`add(name, value=None, vary=True, min=-inf, max=inf, expr=None, brute_step=None)`

Add a Parameter.

Parameters

- `name (str or Parameter)` – If `name` refers to a Parameter object it will be added directly to the `Parameters` instance, otherwise a new Parameter object with name `string` is created before adding it. In both cases, `name` must match `[a-zA-Z][a-zA-Z0-9_]*` and cannot be a Python reserved word.
- `value (float, optional)` – Numerical Parameter value, typically the *initial value*.
- `vary (bool, optional)` – Whether the Parameter is varied during a fit (default is `True`).

- **min** (*float*, *optional*) – Lower bound for value (default is `-numpy.inf`, no lower bound).
- **max** (*float*, *optional*) – Upper bound for value (default is `numpy.inf`, no upper bound).
- **expr** (*str*, *optional*) – Mathematical expression used to constrain the value during the fit (default is `None`).
- **brute_step** (*float*, *optional*) – Step size for grid points in the *brute* method (default is `None`).

Examples

```
>>> params = Parameters()
>>> params.add('xvar', value=0.50, min=0, max=1)
>>> params.add('yvar', expr='1.0 - xvar')
```

which is equivalent to:

```
>>> params = Parameters()
>>> params['xvar'] = Parameter(name='xvar', value=0.50, min=0, max=1)
>>> params['yvar'] = Parameter(name='yvar', expr='1.0 - xvar')
```

`add_many(*parlist)`

Add many parameters, using a sequence of tuples.

Parameters

***parlist** (sequence of *tuple* or *Parameter*) – A sequence of tuples, or a sequence of *Parameter* instances. If it is a sequence of tuples, then each tuple must contain at least a *name*. The order in each tuple must be (*name*, *value*, *vary*, *min*, *max*, *expr*, *brute_step*).

Examples

```
>>> params = Parameters()
# add with tuples: (NAME VALUE VARY MIN MAX EXPR BRUTE_STEP)
>>> params.add_many(('amp', 10, True, None, None, None),
...                   ('cen', 4, True, 0.0, None, None, None),
...                   ('wid', 1, False, None, None, None),
...                   ('frac', 0.5))
# add a sequence of Parameters
>>> f = Parameter('par_f', 100)
>>> g = Parameter('par_g', 2.)
>>> params.add_many(f, g)
```

`pretty_print(oneline=False, colwidth=8, precision=4, fmt='g', columns=['value', 'min', 'max', 'stderr', 'vary', 'expr', 'brute_step'])`

Pretty-print of parameters data.

Parameters

- **oneline** (*bool*, *optional*) – If True prints a one-line parameters representation [False]
- **colwidth** (*int*, *optional*) – Column width for all columns specified in *columns* [8]

- **precision** (`int`, *optional*) – Number of digits to be printed after floating point [4]
- **fmt** (`{'g', 'e', 'f'}`, *optional*) – Single-character numeric formatter. Valid values are: ‘g’ floating point and exponential (default), ‘e’ exponential, or ‘f’ floating point.
- **columns** (`list` of `str`, *optional*) – List of `Parameter` attribute names to print (default is to show all attributes).

valuesdict()

Return an ordered dictionary of parameter values.

Returns

A dictionary of `name:value` pairs for each Parameter.

Return type

`dict`

create_uvars(`covar=None`)

Return a dict of uncertainties ufloats from the current Parameter values and stderr, and an optionally-supplied covariance matrix. Uncertainties in Parameters with constraint expressions will be calculated, propagating uncertainties (and including correlations)

Parameters

`covar` (*optional*) – Nvar x Nvar covariance matrix from fit

Return type

`dict` with keys of Parameter names and values of uncertainties.ufloats.

Notes

1. if covar is provide, it must correspond to the existing *variable* Parameters. If covar is given, the returned uncertainties ufloats will take the correlations into account when combining values.
2. See the uncertainties package documentation (<https://pythonhosted.org/uncertainties>) for more details.

dumps(`**kws`)

Represent Parameters as a JSON string.

Parameters

`**kws` (*optional*) – Keyword arguments that are passed to `json.dumps`.

Returns

JSON string representation of Parameters.

Return type

`str`

See also

`dump`, `loads`, `load`, `json.dumps`

dump(`fp`, `**kws`)

Write JSON representation of Parameters to a file-like object.

Parameters

- **fp** (`file-like object`) – An open and `.write()`-supporting file-like object.
- ****kws** (*optional*) – Keyword arguments that are passed to `dumps`.

Returns

Return value from `fp.write()`: the number of characters written.

Return type

`int`

↳ **See also**

`dumps`, `load`, `json.dump`

eval(expr)

Evaluate a statement using the `asteval` Interpreter.

Parameters

`expr (str)` – An expression containing parameter names and other symbols recognizable by the `asteval` Interpreter.

Returns

The result of evaluating the expression.

Return type

`float`

loads(s, **kws)

Load Parameters from a JSON string.

Parameters

`**kws (optional)` – Keyword arguments that are passed to `json.loads`.

Returns

Updated Parameters from the JSON string.

Return type

`Parameters`

Notes

Current Parameters will be cleared before loading the data from the JSON string.

↳ **See also**

`dump`, `dumps`, `load`, `json.loads`

load(fp, **kws)

Load JSON representation of Parameters from a file-like object.

Parameters

- `fp (file-like object)` – An open and `.read()`-supporting file-like object.
- `**kws (optional)` – Keyword arguments that are passed to `loads`.

Returns

Updated Parameters loaded from `fp`.

Return type

`Parameters`

See also

`dump`, `loads`, `json.load`

 **Warning**

Saving Parameters with user-added functions to the `_asteval` interpreter using :meth:`dump` and `dumps()` may not be easily recovered with the `load()` and `loads()`. See [Saving and Loading Models](#) for further discussion.

5.3 The `create_params()` function

Added in version 1.2.0.

The `create_params()` function is probably the easiest method for making `Parameters` objects, as it allows defining Parameter names by keyword with values either being the numerical initial value for the Parameter or being a dictionary with keyword/value pairs for `value` as well as other Parameter attribute such as `min`, `max`, `expr`, and so forth.

`create_params(**kws)`

Create `Imfit.Parameters` instance and set initial values and attributes.

Parameters

`**kws` – keywords are parameter names, value are dictionaries of Parameter values and attributes.

Return type

`Parameters` instance

Notes

1. keyword arguments will be used to create parameter names.
2. values can either be numbers (floats or integers) to set the parameter value, or can be dictionaries with any of the following keywords: `value`, `vary`, `min`, `max`, `expr`, `brute_step`, or `is_init_value` to set those parameter attributes.
3. for each parameter, `is_init_value` controls whether to set `init_value` when setting `value`, and defaults to `True`.

Examples

```
>>> params = create_params(amplitude=2, center=200,
                           sigma={'value': 3, 'min':0},
                           fwhm={'expr': '2.0*sigma'})
```

5.4 Simple Example

A basic example making use of `Parameters` and the `minimize()` function (discussed in the next chapter) might look like this:

```
# <examples/doc_parameters_basic.py>
import numpy as np

from lmfit import Minimizer, Parameters, create_params, report_fit

# create data to be fitted
x = np.linspace(0, 15, 301)
np.random.seed(2021)
data = (5.0 * np.sin(2.0*x - 0.1) * np.exp(-x*x*0.025) +
        np.random.normal(size=x.size, scale=0.2))

# define objective function: returns the array to be minimized
def fcn2min(params, x, data):
    """Model a decaying sine wave and subtract data."""
    amp = params['amp']
    shift = params['shift']
    omega = params['omega']
    decay = params['decay']
    model = amp * np.sin(x*omega + shift) * np.exp(-x*x*decay)
    return model - data

# create a set of Parameters
params = Parameters()
params.add('amp', value=10, min=0)
params.add('decay', value=0.1)
params.add('shift', value=0.0, min=-np.pi/2., max=np.pi/2.)
params.add('omega', value=3.0)

# ... or use
params = create_params(amp=dict(value=10, min=0),
                       decay=0.1,
                       omega=3,
                       shift=dict(value=0, min=-np.pi/2, max=np.pi/2))

# do fit, here with the default leastsq algorithm
minner = Minimizer(fcn2min, params, fcn_args=(x, data))
result = minner.minimize()

# calculate final result
final = data + result.residual

# write error report
report_fit(result)

# try to plot results
try:
```

(continues on next page)

(continued from previous page)

```
import matplotlib.pyplot as plt
plt.plot(x, data, '+')
plt.plot(x, final)
plt.show()
except ImportError:
    pass
# <end of examples/doc_parameters_basic.py>
```

Here, the objective function explicitly unpacks each Parameter value. This can be simplified using the [Parameters](#) `valuesdict()` method, which would make the objective function `fcn2min` above look like:

```
def fcn2min(params, x, data):
    """Model a decaying sine wave and subtract data."""
    v = params.valuesdict()

    model = v['amp'] * np.sin(x*v['omega'] + v['shift']) * np.exp(-x*x*v['decay'])
    return model - data
```

The results are identical, and the difference is a stylistic choice.

PERFORMING FITS AND ANALYZING OUTPUTS

As shown in the previous chapter, a simple fit can be performed with the `minimize()` function. For more sophisticated modeling, the `Minimizer` class can be used to gain a bit more control, especially when using complicated constraints or comparing results from related fits.

6.1 The `minimize()` function

The `minimize()` function is a wrapper around `Minimizer` for running an optimization problem. It takes an objective function (the function that calculates the array to be minimized), a `Parameters` object, and several optional arguments. See [Writing a Fitting Function](#) for details on writing the objective function.

```
minimize(fcn, params, method='leastsq', args=None, kws=None, iter_cb=None, scale_covar=True,  
nan_policy='raise', reduce_fcn=None, calc_covar=True, max_nfev=None, **fit_kws)
```

Perform the minimization of the objective function.

The `minimize` function takes an objective function to be minimized, a dictionary (`Parameters` ; `Parameters`) containing the model parameters, and several optional arguments including the fitting method.

Parameters

- `fcn (callable)` – Objective function to be minimized. When method is ‘`leastsq`’ or ‘`least_squares`’, the objective function should return an array of residuals (difference between model and data) to be minimized in a least-squares sense. With the scalar methods the objective function can either return the residuals array or a single scalar value. The function must have the signature:

```
fcn(params, *args, **kws)
```

- `params (Parameters)` – Contains the Parameters for the model.
- `method (str, optional)` – Name of the fitting method to use. Valid values are:
 - ‘`leastsq`’: Levenberg-Marquardt (default)
 - ‘`least_squares`’: Least-Squares minimization, using Trust Region Reflective method
 - ‘`differential_evolution`’: differential evolution
 - ‘`brute`’: brute force method
 - ‘`basinhopping`’: basinhopping
 - ‘`ampgo`’: Adaptive Memory Programming for Global Optimization
 - ‘`nelder`’: Nelder-Mead
 - ‘`lbfgsb`’: L-BFGS-B

- ‘`powell`’: Powell
- ‘`cg`’: Conjugate-Gradient
- ‘`newton`’: Newton-CG
- ‘`cobyla`’: Cobyla
- ‘`bfgs`’: BFGS
- ‘`trnc`’: Truncated Newton
- ‘`trust-ncg`’: Newton-CG trust-region
- ‘`trust-exact`’: nearly exact trust-region
- ‘`trust-krylov`’: Newton GLTR trust-region
- ‘`trust-constr`’: trust-region for constrained optimization
- ‘`dogleg`’: Dog-leg trust-region
- ‘`slsqp`’: Sequential Linear Squares Programming
- ‘`emcee`’: Maximum likelihood via Monte-Carlo Markov Chain
- ‘`shgo`’: Simplicial Homology Global Optimization
- ‘`dual_annealing`’: Dual Annealing optimization

In most cases, these methods wrap and use the method of the same name from `scipy.optimize`, or use `scipy.optimize.minimize` with the same `method` argument. Thus ‘`leastsq`’ will use `scipy.optimize.leastsq`, while ‘`powell`’ will use `scipy.optimize.minimizer(..., method='powell')`

For more details on the fitting methods please refer to the [SciPy docs](#).

- **`args` (`tuple`, *optional*)** – Positional arguments to pass to `fcn`.
- **`kws` (`dict`, *optional*)** – Keyword arguments to pass to `fcn`.
- **`iter_cb` (`callable`, *optional*)** – Function to be called at each fit iteration. This function should have the signature:

```
iter_cb(params, iter, resid, *args, **kws),
```

where `params` will have the current parameter values, `iter` the iteration number, `resid` the current residual array, and `*args` and `**kws` as passed to the objective function.

- **`scale_covar` (`bool`, *optional*)** – Whether to automatically scale the covariance matrix (default is `True`).
- **`nan_policy` (`{'raise', 'propagate', 'omit'}`, *optional*)** – Specifies action if `fcn` (or a Jacobian) returns NaN values. One of:
 - ‘`raise`’ : a `ValueError` is raised
 - ‘`propagate`’ : the values returned from `userfcn` are un-altered
 - ‘`omit`’ : non-finite values are filtered
- **`reduce_fcn` (`str` or `callable`, *optional*)** – Function to convert a residual array to a scalar value for the scalar minimizers. See Notes in *Minimizer*.
- **`calc_covar` (`bool`, *optional*)** – Whether to calculate the covariance matrix (default is `True`) for solvers other than ‘`leastsq`’ and ‘`least_squares`’. Requires the `numdifftools` package to be installed.

- **max_nfev** (*int or None, optional*) – Maximum number of function evaluations (default is None). The default value depends on the fitting method.
- ****fit_kws** (*dict, optional*) – Options to pass to the minimizer being used.

Returns

Object containing the optimized parameters and several goodness-of-fit statistics.

Return type

MinimizerResult

Changed in version 0.9.0: Return value changed to *MinimizerResult*.

Notes

The objective function should return the value to be minimized. For the Levenberg-Marquardt algorithm from `leastsq()`, this returned value must be an array, with a length greater than or equal to the number of fitting variables in the model. For the other methods, the return value can either be a scalar or an array. If an array is returned, the sum-of-squares of the array will be sent to the underlying fitting method, effectively doing a least-squares optimization of the return values.

A common use for `args` and `kws` would be to pass in other data needed to calculate the residual, including such things as the data array, dependent variable, uncertainties in the data, and other data structures for the model calculation.

On output, `params` will be unchanged. The best-fit values and, where appropriate, estimated uncertainties and correlations, will all be contained in the returned *MinimizerResult*. See *MinimizerResult – the optimization result* for further details.

This function is simply a wrapper around *Minimizer* and is equivalent to:

```
fitter = Minimizer(fcn, params, fcn_args=args, fcn_kws=kws,
                    iter_cb=iter_cb, scale_covar=scale_covar,
                    nan_policy=nan_policy, reduce_fcn=reduce_fcn,
                    calc_covar=calc_covar, **fit_kws)
fitter.minimize(method=method)
```

6.2 Writing a Fitting Function

An important component of a fit is writing a function to be minimized – the *objective function*. Since this function will be called by other routines, there are fairly stringent requirements for its call signature and return value. In principle, your function can be any Python callable, but it must look like this:

func(params, *args, **kws):

Calculate objective residual to be minimized from parameters.

Parameters

- **params** (*Parameters*) – Parameters.
- **args** – Positional arguments. Must match `args` argument to `minimize()`.
- **kws** – Keyword arguments. Must match `kws` argument to `minimize()`.

Returns

Residual array (generally data-model) to be minimized in the least-squares sense.

Return type

`numpy.ndarray`. The length of this array cannot change between calls.

A common use for the positional and keyword arguments would be to pass in other data needed to calculate the residual, including things as the data array, dependent variable, uncertainties in the data, and other data structures for the model calculation.

The objective function should return the value to be minimized. For the Levenberg-Marquardt algorithm from `leastsq()`, this returned value **must** be an array, with a length greater than or equal to the number of fitting variables in the model. For the other methods, the return value can either be a scalar or an array. If an array is returned, the sum of squares of the array will be sent to the underlying fitting method, effectively doing a least-squares optimization of the return values.

Since the function will be passed in a dictionary of `Parameters`, it is advisable to unpack these to get numerical values at the top of the function. A simple way to do this is with `Parameters.valuesdict()`, as shown below:

```
from numpy import exp, sign, sin, pi

def residual(pars, x, data=None, eps=None):
    # unpack parameters: extract .value attribute for each parameter
    parvals = pars.valuesdict()
    period = parvals['period']
    shift = parvals['shift']
    decay = parvals['decay']

    if abs(shift) > pi/2:
        shift = shift - sign(shift)*pi

    if abs(period) < 1.e-10:
        period = sign(period)*1.e-10

    model = parvals['amp'] * sin(shift + x/period) * exp(-x*x*decay*decay)

    if data is None:
        return model
    if eps is None:
        return model - data
    return (model-data) / eps
```

In this example, `x` is a positional (required) argument, while the `data` array is actually optional (so that the function returns the model calculation if the data is neglected). Also note that the model calculation will divide `x` by the value of the `period` Parameter. It might be wise to ensure this parameter cannot be 0. It would be possible to use bounds on the `Parameter` to do this:

```
params['period'] = Parameter(name='period', value=2, min=1.e-10)
```

but putting this directly in the function with:

```
if abs(period) < 1.e-10:
    period = sign(period)*1.e-10
```

is also a reasonable approach. Similarly, one could place bounds on the `decay` parameter to take values only between $-\pi/2$ and $\pi/2$.

6.3 Types of Data to Use for Fitting

Minimization methods assume that data is numerical. For all the fitting methods supported by lmfit, data and fitting parameters are also assumed to be continuous variables. As the routines make heavy use of numpy and scipy, the most natural data to use in fitting is then numpy nd-arrays. In fact, many of the underlying fitting algorithms - including the default `leastsq()` method - **require** the values in the residual array used for the minimization to be a 1-dimensional numpy array with data type (`dtype`) of “`float64`”: a 64-bit representation of a floating point number (sometimes called a “double precision float”).

Python is generally forgiving about data types, and in the scientific Python community there is a concept of an object being “array like” which essentially means that the can usually be coerced or interpreted as a numpy array, often with that object having an `__array__()` method specially designed for that conversion. Important examples of objects that can be considered “array like” include Lists and Tuples that contain only numbers, pandas Series, and HDF5 Datasets. Many objects from data-processing libraries like dask, xarray, zarr, and more are also “array like”.

Lmfit tries to be accommodating in the data that can be used in the fitting process. When using `Minimizer`, the data you pass in as extra arrays for the calculation of the residual array will not be altered, and can be used in your objective function in whatever form you send. Usually, “array like” data will work, but some care may be needed. In the example above, if `x` was not a numpy array but a list of numbers, this would give an error message like:

```
TypeError: unsupported operand type(s) for /: 'list' and 'float'
```

or:

```
TypeError: can't multiply sequence by non-int of type 'float'
```

because a list of numbers is only sometimes “array like”.

Sending in a “more array-like” object like a pandas Series will avoid many (though maybe not all!) such exceptions, but the resulting calculation returned from the function would then also be a pandas Series. Lmfit `minimize()` will always coerce the return value from the objective function into a 1-D numpy array with `dtype` of “`float64`”. This will usually “just work”, but there may be exceptions.

When in doubt, or if running it trouble, converting data to `float64` numpy arrays before being used in a fit is recommended. If using complex data or functions, a `dtype` of “`complex128`” will also always work, and will be converted to “`float64`” with `ndarray.view("float64")`. Numpy arrays of other `dtype` (say, “`int16`” or “`float32`”) should be used with caution. In particular, “`float32`” data should be avoided: Multiplying a “`float32`” array and a Python float will result in a “`float32`” array for example. As fitting variables may have small changes made to them, the results may be at or below “`float32`” precision, which will cause the fit to give up. For integer data, results are more sometimes promoted to “`float64`”, but many numpy ufuncs (say, `numpy.exp()`) will promote only to “`float32`”, so care is still needed.

See also [Data Types for data and independent data with Model](#) for discussion of data passed in for curve-fitting.

6.4 Choosing Different Fitting Methods

By default, the [Levenberg-Marquardt](#) algorithm is used for fitting. While often criticized, including the fact it finds a *local* minimum, this approach has some distinct advantages. These include being fast, and well-behaved for most curve-fitting needs, and making it easy to estimate uncertainties for and correlations between pairs of fit variables, as discussed in [MinimizerResult – the optimization result](#).

Alternative algorithms can also be used by providing the `method` keyword to the `minimize()` function or `Minimizer.minimize()` class as listed in the [Table of Supported Fitting Methods](#). If you have the `numdifftools` package installed, lmfit will try to estimate the covariance matrix and determine parameter uncertainties and correlations if `calc_covar` is `True` (default).

Table of Supported Fitting Methods:

Fitting Method	method arg to <code>minimize()</code> or <code>Minimizer.minimize()</code>
Levenberg-Marquardt	<code>leastsq</code> or <code>least_squares</code>
Nelder-Mead	<code>nelder</code>
L-BFGS-B	<code>lbfgsb</code>
Powell	<code>powell</code>
Conjugate Gradient	<code>cg</code>
Newton-CG	<code>newton</code>
COBYLA	<code>cobyla</code>
BFGS	<code>bfgsb</code>
Truncated Newton	<code>tnc</code>
Newton CG trust-region	<code>trust-ncg</code>
Exact trust-region	<code>trust-exact</code>
Newton GLTR trust-region	<code>trust-krylov</code>
Constrained trust-region	<code>trust-constr</code>
Dogleg	<code>dogleg</code>
Sequential Linear Squares Programming	<code>slsqp</code>
Differential Evolution	<code>differential_evolution</code>
Brute force method	<code>brute</code>
Basinhopping	<code>basinhopping</code>
Adaptive Memory Programming for Global Optimization	<code>ampgo</code>
Simplicial Homology Global Optimization	<code>shgo</code>
Dual Annealing	<code>dual_annealing</code>
Maximum likelihood via Monte-Carlo Markov Chain	<code>emcee</code>

Note

The objective function for the Levenberg-Marquardt method **must** return an array, with more elements than variables. All other methods can return either a scalar value or an array. The Monte-Carlo Markov Chain or `emcee` method has two different operating methods when the objective function returns a scalar value. See the documentation for `emcee`.

Warning

Much of this documentation assumes that the Levenberg-Marquardt (`leastsq`) method is used. Many of the fit statistics and estimates for uncertainties in parameters discussed in [MinimizerResult – the optimization result](#) are done only unconditionally for this (and the `least_squares`) method. Lmfit versions newer than 0.9.11 provide the capability to use `numdifftools` to estimate the covariance matrix and calculate parameter uncertainties and correlations for other methods as well.

6.5 MinimizerResult – the optimization result

An optimization with `minimize()` or `Minimizer.minimize()` will return a `MinimizerResult` object. This is an otherwise plain container object (that is, with no methods of its own) that simply holds the results of the minimization. These results will include several pieces of informational data such as status and error messages, fit statistics, and the updated parameters themselves.

Importantly, the parameters passed in to `Minimizer.minimize()` will not be changed. To find the best-fit values, uncertainties and so on for each parameter, one must use the `MinimizerResult.params` attribute. For example, to print the fitted values, bounds and other parameter attributes in a well-formatted text tables you can execute:

```
result.params.pretty_print()
```

with `results` being a `MinimizerResult` object. Note that the method `pretty_print()` accepts several arguments for customizing the output (e.g., column width, numeric format, etcetera).

```
class MinimizerResult(**kws)
```

The results of a minimization.

Minimization results include data such as status and error messages, fit statistics, and the updated (i.e., best-fit) parameters themselves in the `params` attribute.

The list of (possible) `MinimizerResult` attributes is given below:

residual

Residual array Resid_i . Return value of the objective function when using the best-fit values of the parameters.

Type

`numpy.ndarray`

params

The best-fit Parameters resulting from the fit.

Type

`Parameters`

uvars

Dictionary of uncertainties ufloats from Parameters

Type

`dict`

var_names

list of variable Parameter names used in optimization in the same order as the values in `init_vals` and `covar`.

Type

`list`

covar

Covariance matrix from minimization, with rows and columns corresponding to `var_names`. If uncertainties cannot be determined, this value will be `None`.

Type

`numpy.ndarray` or `None`

init_vals

List of initial values for variable parameters using `var_names`.

Type

`list`

init_values

Dictionary of initial values for variable parameters.

Type

`dict`

aborted

Whether the fit was aborted.

Type

`bool`

status

Termination status of the optimizer. Its value depends on the underlying solver. Refer to `message` for details.

Type

`int`

success

True if the fit succeeded, otherwise False. This is an optimistic view of success, meaning that the method finished without error.

Type

`bool`

errorbars

whether uncertainties were estimated for variable Parameters.

Type

`bool`

message

Message about fit success.

Type

`str`

ier

Integer error value from `scipy.optimize.leastsq` ('`leastsq`' method only).

Type

`int`

lmdif_message

Message from `scipy.optimize.leastsq` ('`leastsq`' method only).

Type

`str`

call_kws

Keyword arguments sent to underlying solver.

Type

`dict`

flatchain

A flatchain view of the sampling chain from the *emcee* method.

Type

`pandas.DataFrame`

nfev

Number of function evaluations.

Type

`int`

nvarys

Number of variables in fit: N_{varys} .

Type

`int`

ndata

Number of data points: N .

Type

`int`

nfree

Degrees of freedom in fit: $N - N_{\text{varys}}$.

Type

`int`

chisqr

Chi-square: $\chi^2 = \sum_i^N [\text{Resid}_i]^2$.

Type

`float`

redchi

Reduced chi-square: $\chi^2_\nu = \chi^2 / (N - N_{\text{varys}})$.

Type

`float`

aic

Akaike Information Criterion statistic: $N \ln(\chi^2/N) + 2N_{\text{varys}}$.

Type

`float`

bic

Bayesian Information Criterion statistic: $N \ln(\chi^2/N) + \ln(N)N_{\text{varys}}$.

Type

`float`

show_candidates()

`pretty_print()` representation of candidates from the *brute* fitting method.

6.5.1 Goodness-of-Fit Statistics

Table of Fit Results: These values, including the standard Goodness-of-Fit statistics, are all attributes of the `MinimizerResult` object returned by `minimize()` or `Minimizer.minimize()`.

Attribute Name	Description / Formula
<code>nfev</code>	number of function evaluations
<code>nvarys</code>	number of variables in fit N_{varys}
<code>ndata</code>	number of data points: N
<code>nfree</code>	degrees of freedom in fit: $N - N_{\text{varys}}$
<code>aborted</code>	boolean of whether the fit has been aborted.
<code>success</code>	boolean for a minimal test of whether the fit finished successfully
<code>errorbars</code>	boolean of whether error bars and uncertainty were estimated
<code>ier</code>	integer flag describing message from <code>leastsq</code> .
<code>message</code>	simple message from <code>leastsq</code>
<code>method</code>	name of fitting methods
<code>residual</code>	residual array, returned by the objective function: $\{\text{Resid}_i\}$
<code>chisqr</code>	chi-square: $\chi^2 = \sum_i^N [\text{Resid}_i]^2$
<code>redchi</code>	reduced chi-square: $\chi^2_\nu = \chi^2 / (N - N_{\text{varys}})$
<code>aic</code>	Akaike Information Criterion statistic (see below)
<code>bic</code>	Bayesian Information Criterion statistic (see below)
<code>params</code>	best-fit parameters after fit, with uncertainties is available
<code>var_names</code>	ordered list of variable parameter names used for <code>init_vals</code> and <code>covar</code>
<code>covar</code>	covariance matrix (with rows/columns using <code>var_names</code>)
<code>init_vals</code>	list of initial values for variable parameters
<code>init_values</code>	dictionary of initial values for variable Parameters.
<code>uvars</code>	dictionary of uncertainties <code>uvalues</code> for all Parameters.
<code>call_kwds</code>	dict of keyword arguments sent to underlying solver

Note that the calculation of chi-square and reduced chi-square assume that the returned residual function is scaled properly to the uncertainties in the data. For these statistics to be meaningful, the person writing the function to be minimized **must** scale them properly.

After a fit using the `leastsq()` or `least_squares()` method has completed successfully, standard errors for the fitted variables and correlations between pairs of fitted variables are automatically calculated from the covariance matrix. For other methods, the `calc_covar` parameter (default is `True`) in the `Minimizer` class determines whether or not to use the `numdifftools` package to estimate the covariance matrix. The standard error (estimated 1σ error-bar) goes into the `stderr` attribute of the Parameter. The correlations with all other variables will be put into the `correl` attribute of the Parameter – a dictionary with keys for all other Parameters and values of the corresponding correlation.

In some cases, it may not be possible to estimate the errors and correlations. For example, if a variable actually has no practical effect on the fit, it will likely cause the covariance matrix to be singular, making standard errors impossible to estimate. Placing bounds on varied Parameters makes it more likely that errors cannot be estimated, as being near the maximum or minimum value makes the covariance matrix singular. In these cases, the `errorbars` attribute of the fit result (`Minimizer` object) will be `False`.

6.5.2 Akaike and Bayesian Information Criteria

The `MinimizerResult` includes the traditional chi-square and reduced chi-square statistics:

$$\begin{aligned}\chi^2 &= \sum_i^N r_i^2 \\ \chi^2_\nu &= \chi^2 / (N - N_{\text{varys}})\end{aligned}$$

where r is the residual array returned by the objective function (likely to be `(data-model)/uncertainty` for data modeling usages), N is the number of data points (`ndata`), and N_{varys} is number of variable parameters.

Also included are the [Akaike Information Criterion](#), and [Bayesian Information Criterion](#) statistics, held in the `aic` and `bic` attributes, respectively. These give slightly different measures of the relative quality for a fit, trying to balance quality of fit with the number of variable parameters used in the fit. These are calculated as:

$$\begin{aligned}\text{aic} &= N \ln(\chi^2/N) + 2N_{\text{varys}} \\ \text{bic} &= N \ln(\chi^2/N) + \ln(N)N_{\text{varys}}\end{aligned}$$

When comparing fits with different numbers of varying parameters, one typically selects the model with lowest reduced chi-square, Akaike information criterion, and/or Bayesian information criterion. Generally, the Bayesian information criterion is considered the most conservative of these statistics.

6.5.3 Uncertainties in Variable Parameters, and their Correlations

As mentioned above, when a fit is complete the uncertainties for fitted Parameters as well as the correlations between pairs of Parameters are usually calculated. This happens automatically either when using the default `leastsq()` method, the `least_squares()` method, or for most other fitting methods if the highly-recommended `numdifftools` package is available. The estimated standard error (the 1σ uncertainty) for each variable Parameter will be contained in the `stderr`, while the `correl` attribute for each Parameter will contain a dictionary of the correlation with each other variable Parameter. These updated parameters with uncertainty and correlation information will be placed in `MinimizerResult.params`, so that you may access the best fit value, standard error and correlation. For a successful fit for which uncertainties and correlations can be calculated, the `MinimizerResult` will also have a `uvars` attribute that is a dictionary with keynames for each Parameter (includng constraints) and values of `Ufloats` from the `uncertainties` package using the best fit values, the standard error and the correlation between Parameters.

These estimates of the uncertainties are done by inverting the Hessian matrix which represents the second derivative of fit quality for each variable parameter. There are situations for which the uncertainties cannot be estimated, which generally indicates that this matrix cannot be inverted because one of the fit is not actually sensitive to one of the variables. This can happen if a Parameter is stuck at an upper or lower bound, if the variable is simply not used by the fit, or if the value for the variable is such that it has no real influence on the fit.

In principle, the scale of the uncertainties in the Parameters is closely tied to the goodness-of-fit statistics chi-square and reduced chi-square (`chi_sq` and `redchi`). The standard errors or 1σ uncertainties are those that increase chi-square by 1. Since a “good fit” should have `redchi` of around 1, this requires that the data uncertainties (and to some extent the sampling of the N data points) is correct. Unfortunately, it is often not the case that one has high-quality estimates of the data uncertainties (getting the data is hard enough!). Because of this common situation, the uncertainties reported and held in `stderr` are not those that increase chi-square by 1, but those that increase chi-square by reduced chi-square. This is equivalent to rescaling the uncertainty in the data such that reduced chi-square would be 1. To be clear, this rescaling is done by default because if reduced chi-square is far from 1, this rescaling often makes the reported uncertainties sensible, and if reduced chi-square is near 1 it does little harm. If you have good scaling of the data uncertainty and believe the scale of the residual array is correct, this automatic rescaling can be turned off using `scale_covar=False`.

Note that the simple (and fast!) approach to estimating uncertainties and correlations by inverting the second derivative matrix assumes that the components of the residual array (if, indeed, an array is used) are distributed around 0 with a

normal (Gaussian distribution), and that a map of probability distributions for pairs would be elliptical – the size of the ellipse gives the uncertainty itself and the eccentricity of the ellipse gives the correlation. This simple approach to assessing uncertainties ignores outliers, highly asymmetric uncertainties, or complex correlations between Parameters. In fact, it is not too hard to come up with problems where such effects are important. Our experience is that the automated results are usually the right scale and quite reasonable as initial estimates, but a more thorough exploration of the Parameter space using the tools described in [Minimizer.emcee\(\) - calculating the posterior probability distribution of parameters](#) and [An advanced example for evaluating confidence intervals](#) can give a more complete understanding of the distributions and relations between Parameters.

6.6 Getting and Printing Fit Reports

`fit_report(inpars, modelpars=None, show_correl=True, min_correl=0.1, sort_pars=False, correl_mode='list')`

Generate a report of the fitting results.

The report contains the best-fit values for the parameters and their uncertainties and correlations.

Parameters

- **inpars** ([Parameters](#)) – Input Parameters from fit or MinimizerResult returned from a fit.
- **modelpars** ([Parameters](#), *optional*) – Known Model Parameters.
- **show_correl** ([bool](#), *optional*) – Whether to show list of sorted correlations (default is True).
- **min_correl** ([float](#), *optional*) – Smallest correlation in absolute value to show (default is 0.1).
- **sort_pars** ([bool](#) or [callable](#), *optional*) – Whether to show parameter names sorted in alphanumerical order. If False (default), then the parameters will be listed in the order they were added to the Parameters dictionary. If callable, then this (one argument) function is used to extract a comparison key from each list element.
- **correl_mode** (`{'list', 'table'}` [str](#), *optional*) – Mode for how to show correlations. Can be either ‘list’ (default) to show a sorted (if `sort_pars` is True) list of correlation values, or ‘table’ to show a complete, formatted table of correlations.

Returns

Multi-line text of fit report.

Return type

[str](#)

An example using this to write out a fit report would be:

```
# <examples/doc_fitting_withreport.py>
from numpy import exp, linspace, pi, random, sign, sin

from lmfit import create_params, fit_report, minimize

p_true = create_params(amp=14.0, period=5.46, shift=0.123, decay=0.032)

def residual(pars, x, data=None):
    """Model a decaying sine wave and subtract data."""
    vals = pars.valuesdict()
    amp = vals['amp']
```

(continues on next page)

(continued from previous page)

```

per = vals['period']
shift = vals['shift']
decay = vals['decay']

if abs(shift) > pi/2:
    shift = shift - sign(shift)*pi
model = amp * sin(shift + x/per) * exp(-x*x*decay*decay)
if data is None:
    return model
return model - data

random.seed(0)
x = linspace(0.0, 250., 1001)
noise = random.normal(scale=0.7215, size=x.size)
data = residual(p_true, x) + noise

fit_params = create_params(amp=13, period=2, shift=0, decay=0.02)

out = minimize(residual, fit_params, args=(x,), kws={'data': data})

print(fit_report(out))
# <end examples/doc_fitting_withreport.py>

```

which would give as output:

```

[[Fit Statistics]]
# fitting method    = leastsq
# function evals   = 83
# data points      = 1001
# variables        = 4
chi-square         = 498.811759
reduced chi-square = 0.50031270
Akaike info crit   = -689.222517
Bayesian info crit = -669.587497
[[Variables]]
amp:    13.9121959 +/- 0.14120321 (1.01%) (init = 13)
period: 5.48507038 +/- 0.02666520 (0.49%) (init = 2)
shift:  0.16203673 +/- 0.01405662 (8.67%) (init = 0)
decay:  0.03264539 +/- 3.8015e-04 (1.16%) (init = 0.02)
[[Correlations]] (unreported correlations are < 0.100)
C(period, shift) = +0.7974
C(amp, decay)   = +0.5816
C(amp, shift)   = -0.2966
C(amp, period)  = -0.2432
C(shift, decay) = -0.1819
C(period, decay) = -0.1496

```

To be clear, you can get at all of these values from the fit result `out` and `out.params`. For example, a crude printout of the best fit variables and standard errors could be done as

```
print('-----')
```

(continues on next page)

(continued from previous page)

```
print('Parameter    Value        Stderr')
for name, param in out.params.items():
    print(f'{name:7s} {param.value:11.5f} {param.stderr:11.5f}')
```

```
-----
Parameter    Value        Stderr
amp          13.91220    0.14120
period       5.48507    0.02667
shift         0.16204    0.01406
decay         0.03265    0.00038
```

6.7 Using a Iteration Callback Function

An iteration callback function is a function to be called at each iteration, just after the objective function is called. The iteration callback allows user-supplied code to be run at each iteration, and can be used to abort a fit.

iter_cb(params, iter, resid, *args, **kws):

User-supplied function to be run at each iteration.

Parameters

- **params** (*Parameters*) – Parameters.
- **iter** (*int*) – Iteration number.
- **resid** (*numpy.ndarray*) – Residual array.
- **args** – Positional arguments. Must match args argument to *minimize()*
- **kws** – Keyword arguments. Must match kws argument to *minimize()*

Returns

Iteration abort flag.

Return type

None for normal behavior, any value like True to abort the fit.

Normally, the iteration callback would have no return value or return None. To abort a fit, have this function return a value that is True (including any non-zero integer). The fit will also abort if any exception is raised in the iteration callback. When a fit is aborted this way, the parameters will have the values from the last iteration. The fit statistics are not likely to be meaningful, and uncertainties will not be computed.

6.8 Using the Minimizer class

For full control of the fitting process, you will want to create a *Minimizer* object.

```
class Minimizer(userfcn, params, fcn_args=None, fcn_kws=None, iter_cb=None, scale_covar=True,
                nan_policy='raise', reduce_fcn=None, calc_covar=True, max_nfev=None, **kws)
```

A general minimizer for curve fitting and optimization.

Parameters

- **userfcn** (*callable*) – Objective function that returns the residual (difference between model and data) to be minimized in a least-squares sense. This function must have the signature:

```
userfcn(params, *fcn_args, **fcn_kws)
```

- **params** (*Parameters*) – Contains the Parameters for the model.
- **fcn_args** (*tuple*, *optional*) – Positional arguments to pass to *userfcn*.
- **fcn_kws** (*dict*, *optional*) – Keyword arguments to pass to *userfcn*.
- **iter_cb** (*callable*, *optional*) – Function to be called at each fit iteration. This function should have the signature:

```
iter_cb(params, iter, resid, *fcn_args, **fcn_kws)
```

where *params* will have the current parameter values, *iter* the iteration number, *resid* the current residual array, and **fcn_args* and ***fcn_kws* are passed to the objective function.

- **scale_covar** (*bool*, *optional*) – Whether to automatically scale the covariance matrix (default is True).
- **nan_policy** (*{'raise', 'propagate', 'omit'}*, *optional*) – Specifies action if *userfcn* (or a Jacobian) returns NaN values. One of:
 - *'raise'* : a *ValueError* is raised (default)
 - *'propagate'* : the values returned from *userfcn* are un-altered
 - *'omit'* : non-finite values are filtered
- **reduce_fcn** (*str or callable*, *optional*) – Function to convert a residual array to a scalar value for the scalar minimizers. Optional values are (where *r* is the residual array):
 - None : sum-of-squares of residual (default)

$$= (r * r).sum()$$
 - *'negentropy'* : neg entropy, using normal distribution

$$= \rho * \log(\rho).sum(), \text{ where } \rho = \exp(-r * r / 2) / (\sqrt{2 * \pi})$$
 - *'neglogcauchy'* : neg log likelihood, using Cauchy distribution

$$= -\log(1 / (\pi * (1 + r * r))).sum()$$
 - callable : must take one argument (*r*) and return a float.
- **calc_covar** (*bool*, *optional*) – Whether to calculate the covariance matrix (default is True) for solvers other than 'leastsq' and 'least_squares'. Requires the *numdifftools* package to be installed.
- **max_nfev** (*int or None*, *optional*) – Maximum number of function evaluations (default is None). The default value depends on the fitting method.
- ****kws** (*dict*, *optional*) – Options to pass to the minimizer being used.

Notes

The objective function should return the value to be minimized. For the Levenberg-Marquardt algorithm from `leastsq()` or `least_squares()`, this returned value must be an array, with a length greater than or equal to the number of fitting variables in the model. For the other methods, the return value can either be a scalar or an array. If an array is returned, the sum-of-squares of the array will be sent to the underlying fitting method, effectively doing a least-squares optimization of the return values. If the objective function returns non-finite values then a `ValueError` will be raised because the underlying solvers cannot deal with them.

A common use for the `fcn_args` and `fcn_kws` would be to pass in other data needed to calculate the residual, including such things as the data array, dependent variable, uncertainties in the data, and other data structures for the model calculation.

The Minimizer object has a few public methods:

`Minimizer.minimize(method='leastsq', params=None, **kws)`

Perform the minimization.

Parameters

- **method** (`str`, *optional*) – Name of the fitting method to use. Valid values are:
 - `'leastsq'`: Levenberg-Marquardt (default)
 - `'least_squares'`: Least-Squares minimization, using Trust Region Reflective method
 - `'differential_evolution'`: differential evolution
 - `'brute'`: brute force method
 - `'basinhopping'`: basinhopping
 - `'ampgo'`: Adaptive Memory Programming for Global Optimization
 - `'nelder'`: Nelder-Mead
 - `'lbfgsb'`: L-BFGS-B
 - `'powell'`: Powell
 - `'cg'`: Conjugate-Gradient
 - `'newton'`: Newton-CG
 - `'cobyla'`: Cobyla
 - `'bfgs'`: BFGS
 - `'tnc'`: Truncated Newton
 - `'trust-ncg'`: Newton-CG trust-region
 - `'trust-exact'`: nearly exact trust-region
 - `'trust-krylov'`: Newton GLTR trust-region
 - `'trust-constr'`: trust-region for constrained optimization
 - `'dogleg'`: Dog-leg trust-region
 - `'slsqp'`: Sequential Linear Squares Programming
 - `'emcee'`: Maximum likelihood via Monte-Carlo Markov Chain
 - `'shgo'`: Simplicial Homology Global Optimization
 - `'dual_annealing'`: Dual Annealing optimization

In most cases, these methods wrap and use the method with the same name from `scipy.optimize`, or use `scipy.optimize.minimize` with the same `method` argument. Thus ‘`leastsq`’ will use `scipy.optimize.leastsq`, while ‘`powell`’ will use `scipy.optimize.minimizer(..., method='powell')`.

For more details on the fitting methods please refer to the [SciPy documentation](#).

- **params** (`Parameters`, *optional*) – Parameters of the model to use as starting values.
- ****kws** (*optional*) – Additional arguments are passed to the underlying minimization method.

Returns

Object containing the optimized parameters and several goodness-of-fit statistics.

Return type

`MinimizerResult`

Changed in version 0.9.0: Return value changed to `MinimizerResult`.

`Minimizer.leastsq(params=None, max_nfev=None, **kws)`

Use Levenberg-Marquardt minimization to perform a fit.

It assumes that the input `Parameters` have been initialized, and a function to minimize has been properly set up. When possible, this calculates the estimated uncertainties and variable correlations from the covariance matrix.

This method calls `scipy.optimize.leastsq` and, by default, numerical derivatives are used.

Parameters

- **params** (`Parameters`, *optional*) – Parameters to use as starting point.
- **max_nfev** (`int` or `None`, *optional*) – Maximum number of function evaluations. Defaults to `2000*(nvars+1)`, where `nvars` is the number of variable parameters.
- ****kws** (`dict`, *optional*) – Minimizer options to pass to `scipy.optimize.leastsq`.

Returns

Object containing the optimized parameters and several goodness-of-fit statistics.

Return type

`MinimizerResult`

Changed in version 0.9.0: Return value changed to `MinimizerResult`.

`Minimizer.least_squares(params=None, max_nfev=None, **kws)`

Least-squares minimization using `scipy.optimize.least_squares`.

This method wraps `scipy.optimize.least_squares`, which has built-in support for bounds and robust loss functions. By default it uses the Trust Region Reflective algorithm with a linear loss function (i.e., the standard least-squares problem).

Parameters

- **params** (`Parameters`, *optional*) – Parameters to use as starting point.
- **max_nfev** (`int` or `None`, *optional*) – Maximum number of function evaluations. Defaults to `2000*(nvars+1)`, where `nvars` is the number of variable parameters.
- ****kws** (`dict`, *optional*) – Minimizer options to pass to `scipy.optimize.least_squares`.

Returns

Object containing the optimized parameters and several goodness-of-fit statistics.

Return type*MinimizerResult*

Changed in version 0.9.0: Return value changed to *MinimizerResult*.

`Minimizer.scalar_minimize(method='Nelder-Mead', params=None, max_nfev=None, **kws)`

Scalar minimization using `scipy.optimize.minimize`.

Perform fit with any of the scalar minimization algorithms supported by `scipy.optimize.minimize`. Default argument values are:

<code>scalar_minimize()</code> arg	Default Value	Description
<code>method</code>	'Nelder-Mead'	fitting method
<code>tol</code>	1.e-7	fitting and parameter tolerance
<code>hess</code>	None	Hessian of objective function

Parameters

- **method** (*str*, *optional*) – Name of the fitting method to use. One of:
 - 'Nelder-Mead' (default)
 - 'L-BFGS-B'
 - 'Powell'
 - 'CG'
 - 'Newton-CG'
 - 'COBYLA'
 - 'BFGS'
 - 'TNC'
 - 'trust-ncg'
 - 'trust-exact'
 - 'trust-krylov'
 - 'trust-constr'
 - 'dogleg'
 - 'SLSQP'
 - 'differential_evolution'
- **params** (*Parameters*, *optional*) – Parameters to use as starting point.
- **max_nfev** (*int* or *None*, *optional*) – Maximum number of function evaluations. Defaults to $2000^*(\text{nvarys}+1)$, where *nvarys* is the number of variable parameters.
- ****kws** (*dict*, *optional*) – Minimizer options pass to `scipy.optimize.minimize`.

Returns

Object containing the optimized parameters and several goodness-of-fit statistics.

Return type*MinimizerResult*

Changed in version 0.9.0: Return value changed to *MinimizerResult*.

Notes

If the objective function returns a NumPy array instead of the expected scalar, the sum-of-squares of the array will be used.

Note that bounds and constraints can be set on Parameters for any of these methods, so are not supported separately for those designed to use bounds. However, if you use the `differential_evolution` method you must specify finite (`min`, `max`) for each varying Parameter.

`Minimizer.prepare_fit(params=None)`

Prepare parameters for fitting.

Prepares and initializes model and Parameters for subsequent fitting. This routine prepares the conversion of Parameters into fit variables, organizes parameter bounds, and parses, “compiles” and checks constrain expressions. The method also creates and returns a new instance of a `MinimizerResult` object that contains the copy of the Parameters that will actually be varied in the fit.

Parameters

- `params (Parameters, optional)` – Contains the Parameters for the model; if None, then the Parameters used to initialize the Minimizer object are used.

Return type

`MinimizerResult`

Notes

This method is called directly by the fitting methods, and it is generally not necessary to call this function explicitly.

Changed in version 0.9.0: Return value changed to `MinimizerResult`.

`Minimizer.brute(params=None, Ns=20, keep=50, workers=1, max_nfev=None)`

Use the `brute` method to find the global minimum of a function.

The following parameters are passed to `scipy.optimize.brute` and cannot be changed:

<code>brute()</code> arg	Value	Description
<code>full_output</code>	1	Return the evaluation grid and the objective function’s values on it.
<code>finish</code>	None	No “polishing” function is to be used after the grid search.
<code>disp</code>	False	Do not print convergence messages (when <code>finish</code> is not None).

It assumes that the input Parameters have been initialized, and a function to minimize has been properly set up.

Parameters

- `params (Parameters, optional)` – Contains the Parameters for the model. If None, then the Parameters used to initialize the Minimizer object are used.
- `Ns (int, optional)` – Number of grid points along the axes, if not otherwise specified (see Notes).
- `keep (int, optional)` – Number of best candidates from the brute force method that are stored in the `candidates` attribute. If ‘all’, then all grid points from `scipy.optimize.brute` are stored as candidates.
- `workers (int or map-like callable, optional)` – For parallel evaluation of the grid (see `scipy.optimize.brute` for more details).

- **max_nfev** (*int or None, optional*) – Maximum number of function evaluations (default is None). Defaults to $200000^*(\text{nvarys}+1)$.

Returns

Object containing the parameters from the brute force method. The return values (`x0`, `fval`, `grid`, `Jout`) from `scipy.optimize.brute` are stored as `brute_<parname>` attributes. The *MinimizerResult* also contains the `:attr:candidates` attribute and `show_candidates()` method. The `candidates` attribute contains the parameters and `chisqr` from the brute force method as a namedtuple, `('Candidate', ['params', 'score'])` sorted on the (lowest) `chisqr` value. To access the values for a particular candidate one can use `result.candidate[#].params` or `result.candidate[#].score`, where a lower # represents a better candidate. The `show_candidates()` method uses the `pretty_print()` method to show a specific candidate-# or all candidates when no number is specified.

Return type

`MinimizerResult`

Added in version 0.9.6.

Notes

The `brute()` method evaluates the function at each point of a multidimensional grid of points. The grid points are generated from the parameter ranges using `Ns` and (optional) `brute_step`. The implementation in `scipy.optimize.brute` requires finite bounds and the `range` is specified as a two-tuple (`min`, `max`) or slice-object (`min`, `max`, `brute_step`). A slice-object is used directly, whereas a two-tuple is converted to a slice object that interpolates `Ns` points from `min` to `max`, inclusive.

In addition, the `brute()` method in lmfit, handles three other scenarios given below with their respective slice-object:

- **lower bound (`min`) and `brute_step` are specified:**

`range = (min, min + Ns * brute_step, brute_step).`

- **upper bound (`max`) and `brute_step` are specified:**

`range = (max - Ns * brute_step, max, brute_step).`

- **numerical value (`value`) and `brute_step` are specified:**

`range = (value - (Ns//2) * brute_step, value + (Ns//2) * brute_step, brute_step).`

For more information, check the examples in `examples/lmfit_brute_example.ipynb`.

`Minimizer.basin hopping(params=None, max_nfev=None, **kws)`

Use the `basinhopping` algorithm to find the global minimum.

This method calls `scipy.optimize.basin hopping` using the default arguments. The default minimizer is BFGS, but since lmfit supports parameter bounds for all minimizers, the user can choose any of the solvers present in `scipy.optimize.minimize`.

Parameters

- **params** (`Parameters, optional`) – Contains the Parameters for the model. If None, then the Parameters used to initialize the Minimizer object are used.
- **max_nfev** (*int or None, optional*) – Maximum number of function evaluations (default is None). Defaults to $200000^*(\text{nvarys}+1)$.
- ****kws** (`dict, optional`) – Minimizer options to pass to `scipy.optimize.basin hopping`.

Returns

Object containing the optimization results from the basinhopping algorithm.

Return type*MinimizerResult*

Added in version 0.9.10.

Minimizer.ampgo(*params=None*, *max_nfev=None*, ***kws*)

Find the global minimum of a multivariate function using AMPGO.

AMPGO stands for ‘Adaptive Memory Programming for Global Optimization’ and is an efficient algorithm to find the global minimum.

Parameters

- **params** (*Parameters*, *optional*) – Contains the Parameters for the model. If None, then the Parameters used to initialize the Minimizer object are used.
- **max_nfev** (*int*, *optional*) – Maximum number of total function evaluations. If None (default), the optimization will stop after *totaliter* number of iterations (see below)..
- ****kws** (*dict*, *optional*) – Minimizer options to pass to the ampgo algorithm, the options are listed below:

```

local: str, optional
    Name of the local minimization method. Valid options
    are:
    - `'L-BFGS-B'` (default)
    - `'Nelder-Mead'`
    - `'Powell'`
    - `'TNC'`
    - `'SLSQP'`
local_opts: dict, optional
    Options to pass to the local minimizer (default is
    None).
maxfunevals: int, optional
    Maximum number of function evaluations. If None
    (default), the optimization will stop after
    `totaliter` number of iterations (deprecated: use
    `max_nfev` instead).
totaliter: int, optional
    Maximum number of global iterations (default is 20).
maxiter: int, optional
    Maximum number of `Tabu Tunneling` iterations during
    each global iteration (default is 5).
glbtol: float, optional
    Tolerance whether or not to accept a solution after a
    tunneling phase (default is 1e-5).
eps1: float, optional
    Constant used to define an aspiration value for the
    objective function during the Tunneling phase (default
    is 0.02).
eps2: float, optional
    Perturbation factor used to move away from the latest
    local minimum at the start of a Tunneling phase
    (default is 0.1).
tabulistsize: int, optional
    Size of the (circular) tabu search list (default is 5).
tabustrategy: {'farthest', 'oldest'}, optional

```

(continues on next page)

(continued from previous page)

Strategy to use when the size of the tabu list exceeds `tabulistsize`. It can be `oldest` to drop the oldest point from the tabu list or `farthest` (default) to drop the element farthest from the last local minimum found.

`disp: bool, optional`
Set to True to print convergence messages (default is False).

Returns

Object containing the parameters from the `ampgo` method, with fit parameters, statistics and such. The return values (`x0`, `fval`, `eval`, `msg`, `tunnel`) are stored as `ampgo_<parname>` attributes.

Return type

MinimizerResult

Added in version 0.9.10.

Notes

The Python implementation was written by Andrea Gavana in 2014 (http://infinity77.net/global_optimization/index.html).

The details of the AMPGO algorithm are described in the paper “Adaptive Memory Programming for Constrained Global Optimization” located here:

[http://leeds-faculty.colorado.edu/glover/fred%20pubs/416%20-%20AMP%20\(TS\)%20for%20Constrained%20Global%20Opt%20w%20Lasdon%20et%20al%20.pdf](http://leeds-faculty.colorado.edu/glover/fred%20pubs/416%20-%20AMP%20(TS)%20for%20Constrained%20Global%20Opt%20w%20Lasdon%20et%20al%20.pdf)

`Minimizer.shgo(params=None, max_nfev=None, **kws)`

Use the *SHGO* algorithm to find the global minimum.

SHGO stands for “simplicial homology global optimization” and calls `scipy.optimize.shgo` using its default arguments.

Parameters

- `params` (*Parameters, optional*) – Contains the Parameters for the model. If None, then the Parameters used to initialize the `Minimizer` object are used.
- `max_nfev` (*int or None, optional*) – Maximum number of function evaluations. Defaults to `2000000*(nvars+1)`, where `nvars` is the number of variable parameters.
- `**kws` (*dict, optional*) – `Minimizer` options to pass to the SHGO algorithm.

Returns

Object containing the parameters from the SHGO method. The return values specific to `scipy.optimize.shgo` (`x`, `xl`, `fun`, `funl`, `nfev`, `nit`, `nlfev`, `nlhev`, and `nljev`) are stored as `shgo_<parname>` attributes.

Return type

MinimizerResult

Added in version 0.9.14.

`Minimizer.dual_annealing(params=None, max_nfev=None, **kws)`

Use the *dual_annealing* algorithm to find the global minimum.

This method calls `scipy.optimize.dual_annealing` using its default arguments.

Parameters

- **params** ([Parameters](#), *optional*) – Contains the Parameters for the model. If None, then the Parameters used to initialize the Minimizer object are used.
- **max_nfev** ([int](#) or [None](#), *optional*) – Maximum number of function evaluations. Defaults to $200000 * (\text{nvarys} + 1)$, where nvarys is the number of variables.
- ****kws** ([dict](#), *optional*) – Minimizer options to pass to the dual_annealing algorithm.

Returns

Object containing the parameters from the dual_annealing method. The return values specific to `scipy.optimize.dual_annealing(x, fun, nfev, nhev, njev, and nit)` are stored as `da_<parname>` attributes.

Return type

MinimizerResult

Added in version 0.9.14.

```
Minimizer.emcee(params=None, steps=1000, nwalkers=100, burn=0, thin=1, ntemps=1, pos=None,  
reuse_sampler=False, workers=1, float_behavior='posterior', is_weighted=True, seed=None,  
progress=True, run_mcmc_kwarg={})
```

Bayesian sampling of the posterior distribution.

The method uses the `emcee` Markov Chain Monte Carlo package and assumes that the prior is Uniform. You need to have `emcee` version 3 or newer installed to use this method.

Parameters

- **params** ([Parameters](#), *optional*) – Parameters to use as starting point. If this is not specified then the Parameters used to initialize the Minimizer object are used.
- **steps** ([int](#), *optional*) – How many samples you would like to draw from the posterior distribution for each of the walkers?
- **nwalkers** ([int](#), *optional*) – Should be set so $nwalkers >> \text{nvarys}$, where `nvarys` are the number of parameters being varied during the fit. ‘Walkers are the members of the ensemble. They are almost like separate Metropolis-Hastings chains but, of course, the proposal distribution for a given walker depends on the positions of all the other walkers in the ensemble.’ - from the `emcee` webpage.
- **burn** ([int](#), *optional*) – Discard this many samples from the start of the sampling regime.
- **thin** ([int](#), *optional*) – Only accept 1 in every `thin` samples.
- **ntemps** ([int](#), *deprecated*) – `ntemps` has no effect.
- **pos** ([numpy.ndarray](#), *optional*) – Specify the initial positions for the sampler, an ndarray of shape `(nwalkers, nvarys)`. You can also initialise using a previous chain of the same `nwalkers` and `nvarys`. Note that `nvarys` may be one larger than you expect it to be if your `userfcn` returns an array and `is_weighted=False`.
- **reuse_sampler** ([bool](#), *optional*) – Set to True if you have already run `emcee` with the `Minimizer` instance and want to continue to draw from its `sampler` (and so retain the chain history). If False, a new sampler is created. The keywords `nwalkers`, `pos`, and `params` will be ignored when this is set, as they will be set by the existing sampler. **Important:** the Parameters used to create the sampler must not change in-between calls to `emcee`. Alteration of Parameters would include changed `min`, `max`, `vary` and `expr` attributes. This may happen, for example, if you use an altered Parameters object and call the `minimize` method in-between calls to `emcee`.

- **workers** (*Pool-like or int, optional*) – For parallelization of sampling. It can be any Pool-like object with a map method that follows the same calling sequence as the built-in `map` function. If `int` is given as the argument, then a multiprocessing-based pool is spawned internally with the corresponding number of parallel processes. ‘`mpi4py`’-based parallelization and ‘`joblib`’-based parallelization pools can also be used here. **Note:** because of multi-processing overhead it may only be worth parallelising if the objective function is expensive to calculate, or if there are a large number of objective evaluations per step (`nwalkers` * `nvarys`).
- **float_behavior** (*str, optional*) – Meaning of float (scalar) output of objective function. Use ‘`posterior`’ if it returns a log-posterior probability or ‘`chi2`’ if it returns χ^2 . See Notes for further details.
- **is_weighted** (*bool, optional*) – Has your objective function been weighted by measurement uncertainties? If `is_weighted=True` then your objective function is assumed to return residuals that have been divided by the true measurement uncertainty (`data - model`) / `sigma`. If `is_weighted=False` then the objective function is assumed to return unweighted residuals, `data - model`. In this case `emcee` will employ a positive measurement uncertainty during the sampling. This measurement uncertainty will be present in the output `params` and `output` chain with the name `__lnsigma`. A side effect of this is that you cannot use this parameter name yourself. **Important:** this parameter only has any effect if your objective function returns an array. If your objective function returns a float, then this parameter is ignored. See Notes for more details.
- **seed** (*int or numpy.random.RandomState, optional*) – If `seed` is an `int`, a new `numpy.random.RandomState` instance is used, seeded with `seed`. If `seed` is already a `numpy.random.RandomState` instance, then that `numpy.random.RandomState` instance is used. Specify `seed` for repeatable minimizations.
- **progress** (*bool, optional*) – Print a progress bar to the console while running.
- **run_mcmc_kwargs** (*dict, optional*) – Additional (optional) keyword arguments that are passed to `emcee.EnsembleSampler.run_mcmc`.

Returns

`MinimizerResult` object containing updated `params`, `statistics`, etc. The updated `params` represent the median of the samples, while the uncertainties are half the difference of the 15.87 and 84.13 percentiles. The `MinimizerResult` contains a few additional attributes: `chain` contain the samples and has shape `((steps - burn) // thin, nwalkers, nvarys)`. `flatchain` is a `pandas.DataFrame` of the flattened chain, that can be accessed with `result.flatchain[parname]`. `lnprob` contains the log probability for each sample in `chain`. The sample with the highest probability corresponds to the maximum likelihood estimate. `acor` is an array containing the auto-correlation time for each parameter if the auto-correlation time can be computed from the chain. Finally, `acceptance_fraction` (an array of the fraction of steps accepted for each walker).

Return type

`MinimizerResult`

Notes

This method samples the posterior distribution of the parameters using Markov Chain Monte Carlo. It calculates the log-posterior probability of the model parameters, F , given the data, D , $\ln p(F_{true}|D)$. This ‘posterior probability’ is given by:

$$\ln p(F_{true}|D) \propto \ln p(D|F_{true}) + \ln p(F_{true})$$

where $\ln p(D|F_{true})$ is the ‘log-likelihood’ and $\ln p(F_{true})$ is the ‘log-prior’. The default log-prior encodes prior information known about the model that the log-prior probability is `-numpy.inf` (impossible) if any of the parameters is outside its limits, and is zero if all the parameters are inside their bounds (uniform prior). The log-likelihood function is¹:

$$\ln p(D|F_{true}) = -\frac{1}{2} \sum_n \left[\frac{(g_n(F_{true}) - D_n)^2}{s_n^2} + \ln(2\pi s_n^2) \right]$$

The first term represents the residual (g being the generative model, D_n the data and s_n the measurement uncertainty). This gives χ^2 when summed over all data points. The objective function may also return the log-posterior probability, $\ln p(F_{true}|D)$. Since the default log-prior term is zero, the objective function can also just return the log-likelihood, unless you wish to create a non-uniform prior.

If the objective function returns a float value, this is assumed by default to be the log-posterior probability, (`float_behavior` default is ‘posterior’). If your objective function returns χ^2 , then you should use `float_behavior='chi2'` instead.

By default objective functions may return an ndarray of (possibly weighted) residuals. In this case, use `is_weighted` to select whether these are correctly weighted by measurement uncertainty. Note that this ignores the second term above, so that to calculate a correct log-posterior probability value your objective function should return a float value. With `is_weighted=False` the data uncertainty, s_n , will be treated as a nuisance parameter to be marginalized out. This uses strictly positive uncertainty (homoscedasticity) for each data point, $s_n = \exp(_linsigma)$. `_linsigma` will be present in `MinimizerResult.params`, as well as `Minimizer.chain` and `nvarys` will be increased by one.

References

6.9 `Minimizer.emcee()` - calculating the posterior probability distribution of parameters

`Minimizer.emcee()` can be used to obtain the posterior probability distribution of parameters, given a set of experimental data. Note that this method does *not* actually perform a fit at all. Instead, it explores parameter space to determine the probability distributions for the parameters, but without an explicit goal of attempting to refine the solution. It should not be used for fitting, but it is a useful method to more thoroughly explore the parameter space around the solution after a fit has been done and thereby get an improved understanding of the probability distribution for the parameters. It may be able to refine your estimate of the most likely values for a set of parameters, but it will not iteratively find a good solution to the minimization problem. To use this method effectively, you should first use another minimization method and then use this method to explore the parameter space around those best-fit values.

To illustrate this, we’ll use an example problem of fitting data to function of a double exponential decay, including a modest amount of Gaussian noise to the data. Note that this example is the same problem used in [An advanced example for evaluating confidence intervals](#) for evaluating confidence intervals in the parameters, which is a similar goal to the one here.

¹ <https://emcee.readthedocs.io>

```
import matplotlib.pyplot as plt
import numpy as np

import lmfit

x = np.linspace(1, 10, 250)
np.random.seed(0)
y = 3.0 * np.exp(-x / 2) - 5.0 * np.exp(-(x - 0.1) / 10.) + 0.1 * np.random.randn(x.size)
```

Create a Parameter set for the initial guesses:

```
p = lmfit.Parameters()
p.add_many(('a1', 4.), ('a2', 4.), ('t1', 3.), ('t2', 3., True))

def residual(p):
    v = p.valuesdict()
    return v['a1'] * np.exp(-x / v['t1']) + v['a2'] * np.exp(-(x - 0.1) / v['t2']) - y
```


Solving with `minimize()` gives the Maximum Likelihood solution. Note that we use the robust Nelder-Mead method here. The default Levenberg-Marquardt method seems to have difficulty with exponential decays, though it can refine the solution if starting near the solution:

```
mi = lmfit.minimize(residual, p, method='nelder', nan_policy='omit')
lmfit.printfuncs.report_fit(mi.params, min_correl=0.5)
```

```
[[Variables]]
a1: 2.98623689 +/- 0.15010519 (5.03%) (init = 4)
a2: -4.33525597 +/- 0.11765824 (2.71%) (init = 4)
t1: 1.30993186 +/- 0.13449656 (10.27%) (init = 3)
t2: 11.8240752 +/- 0.47172610 (3.99%) (init = 3)
[[Correlations]] (unreported correlations are < 0.500)
C(a2, t2) = +0.9876
C(a2, t1) = -0.9278
C(t1, t2) = -0.8852
C(a1, t1) = -0.6093
```

and plotting the fit using the Maximum Likelihood solution gives the graph below:

```
plt.plot(x, y, 'o')
plt.plot(x, residual(mi.params) + y, label='best fit')
plt.legend()
plt.show()
```



Note that the fit here (for which the `numdifftools` package is installed) does estimate and report uncertainties in the parameters and correlations for the parameters, and reports the correlation of parameters `a2` and `t2` to be very high. As we'll see, these estimates are pretty good, but when faced with such high correlation, it can be helpful to get the full probability distribution for the parameters. MCMC methods are very good for this.

Furthermore, we wish to deal with the data uncertainty. This is called marginalisation of a nuisance parameter. `emcee` requires a function that returns the log-posterior probability. The log-posterior probability is a sum of the log-prior probability and log-likelihood functions. The log-prior probability is assumed to be zero if all the parameters are within their bounds and `-np.inf` if any of the parameters are outside their bounds.

If the objective function returns an array of unweighted residuals (i.e., `data-model`) as is the case here, you can use `is_weighted=False` as an argument for `emcee`. In that case, `emcee` will automatically add/use the `__lnsigma` parameter to estimate the true uncertainty in the data. To place boundaries on this parameter one can do:

```
mi.params.add('__lnsigma', value=np.log(0.1), min=np.log(0.001), max=np.log(2))
```

Now we have to set up the minimizer and do the sampling (again, just to be clear, this is *not* doing a fit):


```
res = lmfit.minimize(residual, method='emcee', nan_policy='omit', burn=300, steps=1000, thin=20,
                      params=mi.params, is_weighted=False, progress=False)
```

As mentioned in the Notes for `Minimizer.emcee()`, the `is_weighted` argument will be ignored if your objective function returns a float instead of an array. For the documentation we set `progress=False`; the default is to print a progress bar to the Terminal if the `tqdm` package is installed.

The success of the method (i.e., whether or not the sampling went well) can be assessed by checking the integrated autocorrelation time and/or the acceptance fraction of the walkers. For this specific example the autocorrelation time could not be estimated because the “chain is too short”. Instead, we plot the acceptance fraction per walker and its mean value suggests that the sampling worked as intended (as a rule of thumb the value should be between 0.2 and

0.5).

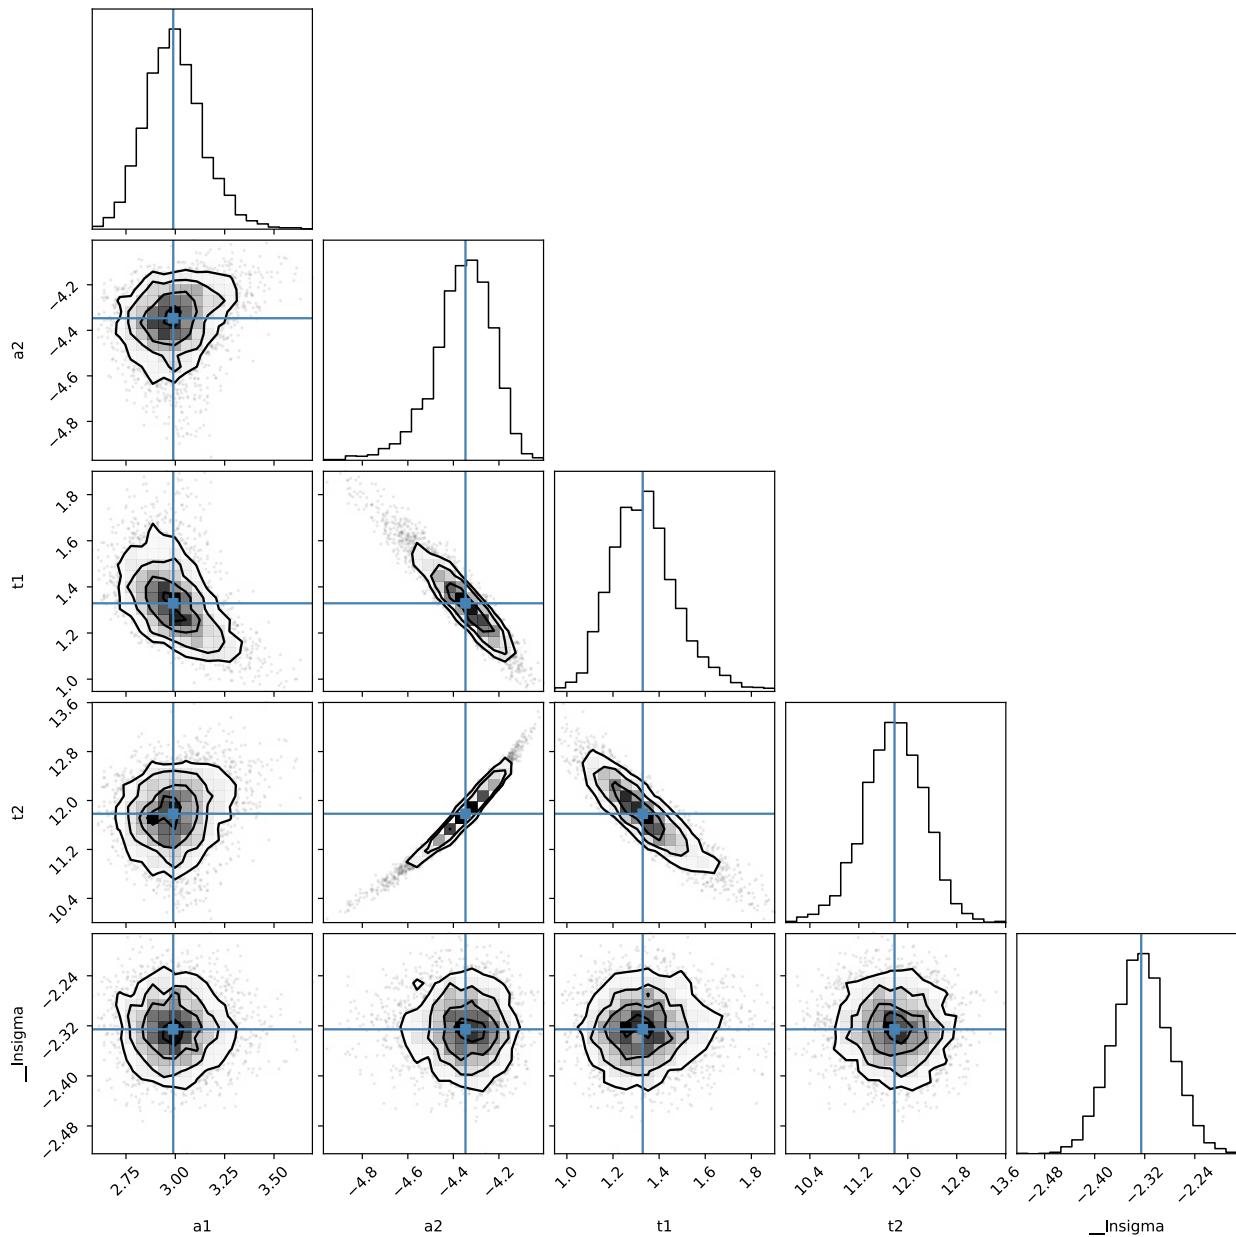
```
plt.plot(res.acceptance_fraction, 'o')
plt.xlabel('walker')
plt.ylabel('acceptance fraction')
plt.show()
```



With the results from emcee, we can visualize the posterior distributions for the parameters using the corner package:

```
import corner

emcee_plot = corner.corner(res.flatchain, labels=res.var_names,
                           truths=list(res.params.valuesdict().values()))
```



The values reported in the `MinimizerResult` are the medians of the probability distributions and a 1σ quantile, estimated as half the difference between the 15.8 and 84.2 percentiles. Printing these values:

```
print('median of posterior probability distribution')
print('-----')
lmfit.report_fit(res.params)
```

```
median of posterior probability distribution
-----
[[Variables]]
a1: 2.98945718 +/- 0.14033921 (4.69%) (init = 2.986237)
a2: -4.34687243 +/- 0.12131092 (2.79%) (init = -4.335256)
t1: 1.32883916 +/- 0.13766047 (10.36%) (init = 1.309932)
```

(continues on next page)

(continued from previous page)

```

t2:           11.7836194 +/- 0.47719763 (4.05%) (init = 11.82408)
__lnsigma: -2.32559226 +/- 0.04542650 (1.95%) (init = -2.302585)
[[Correlations]] (unreported correlations are < 0.100)
C(a2, t2) = +0.9811
C(a2, t1) = -0.9377
C(t1, t2) = -0.8943
C(a1, t1) = -0.5076
C(a1, a2) = +0.2140
C(a1, t2) = +0.1777

```

You can see that this recovered the right uncertainty level on the data. Note that these values agree pretty well with the results, uncertainties and correlations found by the fit and using `numdifftools` to estimate the covariance matrix. That is, even though the parameters `a2`, `t1`, and `t2` are all highly correlated and do not display perfectly Gaussian probability distributions, the probability distributions found by explicitly sampling the parameter space are not so far from elliptical as to make the simple (and much faster) estimates from inverting the covariance matrix completely invalid.

As mentioned above, the result from `emcee` reports the median values, which are not necessarily the same as the Maximum Likelihood Estimate. To obtain the values for the Maximum Likelihood Estimation (MLE) we find the location in the chain with the highest probability:

```

highest_prob = np.argmax(res.lnprob)
hp_loc = np.unravel_index(highest_prob, res.lnprob.shape)
mle_soln = res.chain[hp_loc]
for i, par in enumerate(p):
    p[par].value = mle_soln[i]

print('\nMaximum Likelihood Estimation from emcee      ')
print('-----')
print('Parameter  MLE Value  Median Value  Uncertainty')
fmt = '  {:5s}  {:11.5f}  {:11.5f}  {:11.5f}'.format
for name, param in p.items():
    print(fmt(name, param.value, res.params[name].value,
              res.params[name].stderr))

```

Maximum Likelihood Estimation from emcee

Parameter	MLE Value	Median Value	Uncertainty
a1	2.93839	2.98946	0.14034
a2	-4.35274	-4.34687	0.12131
t1	1.34310	1.32884	0.13766
t2	11.78782	11.78362	0.47720

Here the difference between MLE and median value are seen to be below 0.5%, and well within the estimated $1-\sigma$ uncertainty.

Finally, we can use the samples from `emcee` to work out the 1 - and 2 - σ error estimates.

```

print('\nError estimates from emcee:')
print('-----')
print('Parameter  -2sigma  -1sigma  median  +1sigma  +2sigma')

```

(continues on next page)

(continued from previous page)

```

for name in p.keys():
    quantiles = np.percentile(res.flatchain[name],
                               [2.275, 15.865, 50, 84.135, 97.275])
    median = quantiles[2]
    err_m2 = quantiles[0] - median
    err_m1 = quantiles[1] - median
    err_p1 = quantiles[3] - median
    err_p2 = quantiles[4] - median
    fmt = ' {:5s}  {:8.4f} {:8.4f} {:8.4f} {:8.4f} {:8.4f}'.format
    print(fmt(name, err_m2, err_m1, median, err_p1, err_p2))

```

Error estimates from emcee:

Parameter	-2sigma	-1sigma	median	+1sigma	+2sigma
a1	-0.2656	-0.1362	2.9895	0.1445	0.3141
a2	-0.3209	-0.1309	-4.3469	0.1118	0.1985
t1	-0.2377	-0.1305	1.3288	0.1448	0.3278
t2	-1.0677	-0.4807	11.7836	0.4739	0.8990

And we see that the initial estimates for the $1-\sigma$ standard error using `numdifftools` was not too bad. We'll return to this example problem in [An advanced example for evaluating confidence intervals](#) and use a different method to calculate the 1- and 2- σ error bars.

MODELING DATA AND CURVE FITTING

A common use of least-squares minimization is *curve fitting*, where one has a parametrized model function meant to explain some phenomena and wants to adjust the numerical values for the model so that it most closely matches some data. With `scipy`, such problems are typically solved with `scipy.optimize.curve_fit`, which is a wrapper around `scipy.optimize.least_squares`. While `minimize()` can be used for curve-fitting problems, it is more general and not aimed specifically at this common use-case.

The `Model` class in lmfit provides a simple and flexible approach to curve-fitting problems. Like `scipy.optimize.curve_fit`, a `Model` uses a *model function* – a function that is meant to calculate a model for some phenomenon – and then uses that to best match an array of supplied data. Beyond that similarity, the `Model` interface is rather different from `scipy.optimize.curve_fit`. These differences include a) being class-based and so having multiple methods for working with `Model`'s, b) using `:class:`~lmfit.parameter.Parameters`, and all of the advantages they provide, and c) being easy to combine into composite models.

In addition to allowing you to turn any model function into a curve-fitting model, lmfit also provides canonical definitions for many known line shapes such as Gaussian or Lorentzian peaks and Exponential decays that are widely used in many scientific domains. These are available in the `models` module that will be discussed in more detail in the next chapter (*Built-in Fitting Models in the models module*). We mention it here as you may want to consult that list before writing your own model. For now, we focus on turning Python functions into high-level fitting models with the `Model` class, and using these to fit data.

7.1 Motivation and simple example: Fit data to Gaussian profile

We start with a simple and common example of fitting data to a Gaussian peak. As we will see, there is a built-in `GaussianModel` class that can help do this, but here we'll build our own. We start with a simple definition of the model function:

```
from numpy import exp, linspace, random

def gaussian(x, amp, cen, wid):
    return amp * exp(-(x-cen)**2 / wid)
```

With that function, we can easily calculate a model for our data, and the goal here is to use this function to fit to data $y(x)$ represented by the arrays `y` and `x`. With `scipy.optimize.curve_fit`, this would be:

```
from scipy.optimize import curve_fit

x = linspace(-10, 10, 101)
y = gaussian(x, 2.33, 0.21, 1.51) + random.normal(0, 0.2, x.size)
```

(continues on next page)

(continued from previous page)

```
init_vals = [1, 0, 1] # for [amp, cen, wid]
best_vals, covar = curve_fit(gaussian, x, y, p0=init_vals)
```

That is, we create data (maybe adding a little noise), make an initial guess of the model values, and run `scipy.optimize.curve_fit` with the model function, data arrays, and initial guesses. The results returned are the optimal values for the parameters and the covariance matrix.

With lmfit, we create a `Model` that wraps the `gaussian` model function, which automatically generates the appropriate residual function, and determines the corresponding parameter names from the function signature itself:

```
from lmfit import Model

gmodel = Model(gaussian)
print('parameter names: {gmodel.param_names}')
print('independent variables: {gmodel.independent_vars}')
```

```
parameter names: ['amp', 'cen', 'wid']
independent variables: ['x']
```

As you can see, the Model `gmodel` determined the names of the parameters and the independent variables. By default, the first argument of the function is taken as the independent variable, held in `independent_vars`, and the rest of the functions positional arguments (and, in certain cases, keyword arguments – see below) are used for Parameter names. Thus, for the `gaussian` function above, the independent variable is `x`, and the parameters are named `amp`, `cen`, and `wid`. These are all taken directly from the signature of the model function. As discussed below, you can modify the default assignment of independent variable / arguments and specify yourself what the independent variable.

Although the model determines what the parameters should be named, `Parameters` are *not* created when the model is created. That is, the model knows the parameters names, but nothing about the scale and range of your data. You will have to make Parameters for a Model yourself. To help you create Parameters for a Model, each model has a `make_params()` method that will generate parameters with the expected names. You can use this method or create Parameters some other way (say, with `create_params()`), but you will have to create Parameters and assign initial values for all Parameters. You can also assign other attributes when doing this:

```
params = gmodel.make_params()
```

This creates the `Parameters` but does not automatically give them initial values since it has no idea what the scale should be. If left unspecified, the initial values will be `-Inf`, which will generally fail to give useful results. You can set initial values for parameters with keyword arguments to `make_params()`:

```
params = gmodel.make_params(cen=0.3, amp=3, wid=1.25)
```

or assign them (and other parameter properties) after the `Parameters` class has been created.

A `Model` has several methods associated with it. For example, one can use the `eval()` method to evaluate the model or the `fit()` method to fit data to this model with a `Parameter` object. Both of these methods can take explicit keyword arguments for the parameter values. For example, one could use `eval()` to calculate the predicted function:

```
x_eval = linspace(0, 10, 201)
y_eval = gmodel.eval(params, x=x_eval)
```

or with:

```
y_eval = gmodel.eval(x=x_eval, cen=6.5, amp=100, wid=2.0)
```

Admittedly, this is a slightly long-winded way to calculate a Gaussian function, given that you could have called your `gaussian` function directly. But now that the model is set up, we can use its `fit()` method to fit this model to data, as with:

```
result = gmodel.fit(y, params, x=x)
```

or with:

```
result = gmodel.fit(y, x=x, cen=0.5, amp=10, wid=2.0)
```

Putting everything together, included in the `examples` folder with the source code, is:

```
# <examples/doc_model_gaussian.py>
import matplotlib.pyplot as plt
from numpy import exp, loadtxt, pi, sqrt

from lmfit import Model

data = loadtxt('model1d_gauss.dat')
x = data[:, 0]
y = data[:, 1]

def gaussian(x, amp, cen, wid):
    """1-d gaussian: gaussian(x, amp, cen, wid)"""
    return (amp / (sqrt(pi) * wid)) * exp(-(x-cen)**2 / (2*wid**2))

gmodel = Model(gaussian)
result = gmodel.fit(y, x=x, amp=5, cen=5, wid=1)

print(result.fit_report())

plt.plot(x, y, 'o')
plt.plot(x, result.init_fit, '--', label='initial fit')
plt.plot(x, result.best_fit, '-', label='best fit')
plt.legend()
plt.show()
# <end examples/doc_model_gaussian.py>
```

which is pretty compact and to the point. The returned `result` will be a `ModelResult` object. As we will see below, this has many components, including a `fit_report()` method, which will show:


```
[[Model]]
  Model(gaussian)
[[Fit Statistics]]
  # fitting method   = leastsq
  # function evals  = 33
  # data points    = 101
  # variables      = 3
  chi-square        = 3.40883599
  reduced chi-square = 0.03478404
  Akaike info crit  = -336.263713
  Bayesian info crit = -328.418352
```

(continues on next page)

(continued from previous page)

```
R-squared          = 0.98533348
[[Variables]]
amp: 8.88021893 +/- 0.11359522 (1.28%) (init = 5)
cen: 5.65866102 +/- 0.01030495 (0.18%) (init = 5)
wid: 0.69765478 +/- 0.01030505 (1.48%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(amp, wid) = +0.5774
```

As the script shows, the result will also have `init_fit` for the fit with the initial parameter values and a `best_fit` for the fit with the best fit parameter values. These can be used to generate the following plot:



which shows the data in blue dots, the best fit as a solid green line, and the initial fit as a dashed orange line.

Note that the model fitting was really performed with:

```
gmodel = Model(gaussian)
result = gmodel.fit(y, params, x=x, amp=5, cen=5, wid=1)
```

These lines clearly express that we want to turn the `gaussian` function into a fitting model, and then fit the $y(x)$ data to this model, starting with values of 5 for `amp`, 5 for `cen` and 1 for `wid`. In addition, all the other features of lmfit are included: `Parameters` can have bounds and constraints and the result is a rich object that can be reused to explore the model fit in detail.

7.2 The Model class

The `Model` class provides a general way to wrap a pre-defined function as a fitting model.

```
class Model(func, independent_vars=None, param_names=None, nan_policy='raise', prefix='', name=None, **kws)
```

Create a model from a user-supplied model function.

The model function will normally take an independent variable (generally, the first argument) and a series of arguments that are meant to be parameters for the model. It will return an array of data to model some data as for a curve-fitting problem.

Parameters

- `func (callable)` – Function to be wrapped.
- `independent_vars (list of str, optional)` – Arguments to `func` that are independent variables (default is None).
- `param_names (list of str, optional)` – Names of arguments to `func` that are to be made into parameters (default is None).
- `nan_policy ({'raise', 'propagate', 'omit'}, optional)` – How to handle NaN and missing values in data. See Notes below.
- `prefix (str, optional)` – Prefix used for the model.
- `name (str, optional)` – Name for the model. When None (default) the name is the same as the model function (`func`).
- `**kws (dict, optional)` – Additional keyword arguments to pass to model function.

Notes

1. The model function must return an array that will be the same size as the data being modeled.
2. Parameter names are inferred from the function arguments by default, and a residual function is automatically constructed.
3. Specifying `independent_vars` here will explicitly name the independent variables for the Model. in contrast, `param_names` is meant to help infer Parameter names for keyword arguments defined with `**kws` in the Model function.
4. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - ‘raise’ : raise a `ValueError` (default)
 - ‘propagate’ : do nothing
 - ‘omit’ : drop missing data

Examples

The model function will normally take an independent variable (generally, the first argument) and a series of arguments that are meant to be parameters for the model. Thus, a simple peak using a Gaussian defined as:

```
>>> import numpy as np
>>> def gaussian(x, amp, cen, wid):
...     return amp * np.exp(-(x-cen)**2 / wid)
```

can be turned into a Model with:

```
>>> gmodel = Model(gaussian)
```

this will automatically discover the names of the independent variables and parameters:

```
>>> print(gmodel.param_names, gmodel.independent_vars)
['amp', 'cen', 'wid'], ['x']
```

7.2.1 Model class Methods

`Model.eval(params=None, **kwargs)`

Evaluate the model with supplied parameters and keyword arguments.

Parameters

- `params (Parameters, optional)` – Parameters to use in Model.
- `**kwargs (optional)` – Additional keyword arguments to pass to model function.

Returns

Value of model given the parameters and other arguments.

Return type

`numpy.ndarray, float, int or complex`

Notes

1. if `params` is `None`, the values for all parameters are expected to be provided as keyword arguments.
2. If `params` is given, and a keyword argument for a parameter value is also given, the keyword argument will be used in place of the value in the value in `params`.
3. all non-parameter arguments for the model function, **including all the independent variables** will need to be passed in using keyword arguments.
4. The return types are generally `numpy.ndarray`, but may depends on the model function and input independent variables. That is, return values may be Python `float`, `int`, or `complex` values.

`Model.fit(data, params=None, weights=None, method='leastsq', iter_cb=None, scale_covar=True, verbose=False, fit_kwds=None, nan_policy=None, calc_covar=True, max_nfev=None, coerce_farray=True, **kwargs)`

Fit the model to the data using the supplied Parameters.

Parameters

- `data (array_like)` – Array of data to be fit.
- `params (Parameters, optional)` – Parameters to use in fit (default is `None`).

- **weights** (*array_like, optional*) – Weights to use for the calculation of the fit residual [i.e., $weights^*(data-fit)$]. Default is None; must have the same size as *data*.
- **method** (*str, optional*) – Name of fitting method to use (default is ‘*leastsq*’).
- **iter_cb** (*callable, optional*) – Callback function to call at each iteration (default is None).
- **scale_covar** (*bool, optional*) – Whether to automatically scale the covariance matrix when calculating uncertainties (default is True).
- **verbose** (*bool, optional*) – Whether to print a message when a new parameter is added because of a hint (default is True).
- **fit_kws** (*dict, optional*) – Options to pass to the minimizer being used.
- **nan_policy** (*{'raise', 'propagate', 'omit'}*, *optional*) – What to do when encountering NaNs when fitting Model.
- **calc_covar** (*bool, optional*) – Whether to calculate the covariance matrix (default is True) for solvers other than ‘*leastsq*’ and ‘*least_squares*’. Requires the *numdifftools* package to be installed.
- **max_nfev** (*int or None, optional*) – Maximum number of function evaluations (default is None). The default value depends on the fitting method.
- **coerce_farray** (*bool, optional*) – Whether to coerce data and independent data to be ndarrays with dtype of float64 (or complex128). If set to False, data and independent data are not coerced at all, but the output of the model function will be. (default is True)
- ****kwargs** (*optional*) – Arguments to pass to the model function, possibly overriding parameters.

Return type

ModelResult

Notes

1. if *params* is None, the values for all parameters are expected to be provided as keyword arguments. Mixing *params* and keyword arguments is deprecated (see *Model.eval*).
2. all non-parameter arguments for the model function, **including all the independent variables** will need to be passed in using keyword arguments.
3. Parameters are copied on input, so that the original Parameter objects are unchanged, and the updated values are in the returned *ModelResult*.

Examples

Take *t* to be the independent variable and *data* to be the curve we will fit. Use keyword arguments to set initial guesses:

```
>>> result = my_model.fit(data, tau=5, N=3, t=t)
```

Or, for more control, pass a Parameters object.

```
>>> result = my_model.fit(data, params, t=t)
```

Model.guess(*data*, *x*, ***kws*)

Guess starting values for the parameters of a Model.

This is not implemented for all models, but is available for many of the built-in models.

Parameters

- **data** (*array_like*) – Array of data (i.e., y-values) to use to guess parameter values.
- **x** (*array_like*) – Array of values for the independent variable (i.e., x-values).
- ****kws** (*optional*) – Additional keyword arguments, passed to model function.

Returns

Initial, guessed values for the parameters of a Model.

Return type

Parameters

Raises

NotImplementedError – If the *guess* method is not implemented for a Model.

Notes

Should be implemented for each model subclass to run *self.make_params()*, update starting values and return a Parameters object.

Changed in version 1.0.3: Argument *x* is now explicitly required to estimate starting values.

Model.make_params(*verbose=False*, ***kwargs*)

Create a Parameters object for a Model.

Parameters

- **verbose** (*bool*, *optional*) – Whether to print out messages (default is False).
- ****kwargs** (*optional*) –

Parameter names and initial values or dictionaries of
values and attributes.

Returns

params – Parameters object for the Model.

Return type

Parameters

Notes

1. Parameter values can be numbers (floats or ints) to set the parameter value, or dictionaries with any of the following keywords: `value`, `vary`, `min`, `max`, `expr`, `brute_step`, `is_init_value` to set those parameter attributes.
2. This method will also apply any default values or parameter hints that may have been defined for the model.

Example

```
>>> gmodel = GaussianModel(prefix='peak_') + LinearModel(prefix='bkg_')
>>> gmodel.make_params(peak_center=3200, bkg_offset=0, bkg_slope=0,
...                      peak_amplitdue=dict(value=100, min=2),
...                      peak_sigma=dict(value=25, min=0, max=1000))
```

`Model.set_param_hint(name, **kwargs)`

Set *hints* to use when creating parameters with `make_params()`.

The given hint can include optional bounds and constraints (`value`, `vary`, `min`, `max`, `expr`), which will be used by `Model.make_params()` when building default parameters.

While this can be used to set initial values, `Model.make_params` or the function `create_params` should be preferred for creating parameters with initial values.

The intended use here is to control how a Model should create parameters, such as setting bounds that are required by the mathematics of the model (for example, that a peak width cannot be negative), or to define common constrained parameters.

Parameters

- **name** (`str`) – Parameter name, can include the models `prefix` or not.
- ****kwargs (optional)** – Arbitrary keyword arguments, needs to be a Parameter attribute.
Can be any of the following:
 - **value**
[float, optional] Numerical Parameter value.
 - **vary**
[bool, optional] Whether the Parameter is varied during a fit (default is True).
 - **min**
[float, optional] Lower bound for value (default is `-numpy.inf`, no lower bound).
 - **max**
[float, optional] Upper bound for value (default is `numpy.inf`, no upper bound).
 - **expr**
[str, optional] Mathematical expression used to constrain the value during the fit.

Example

```
>>> model = GaussianModel()
>>> model.set_param_hint('sigma', min=0)
```

See [Using parameter hints](#).

`Model.print_param_hints(colwidth=8)`

Print a nicely aligned text-table of parameter hints.

Parameters

`colwidth (int, optional)` – Width of each column, except for first and last columns.

See [Using parameter hints](#).

`Model.post_fit(fitresult)`

function that is called just after fit, can be overloaded by subclasses to add non-fitting ‘calculated parameters’

See [Using uncertainties in the fitted parameters for post-fit calculations](#).

7.2.2 Model class Attributes

`func`

The model function used to calculate the model.

`independent_vars`

List of strings for names of the independent variables.

`nan_policy`

Describes what to do for NaNs that indicate missing values in the data. The choices are:

- 'raise': Raise a `ValueError` (default)
- 'propagate': Do not check for NaNs or missing values. The fit will try to ignore them.
- 'omit': Remove NaNs or missing observations in data. If pandas is installed, `pandas.isnull()` is used, otherwise `numpy.isnan()` is used.

`name`

Name of the model, used only in the string representation of the model. By default this will be taken from the model function.

`opts`

Extra keyword arguments to pass to model function. Normally this will be determined internally and should not be changed.

`param_hints`

Dictionary of parameter hints. See [Using parameter hints](#).

`param_names`

List of strings of parameter names.

`prefix`

Prefix used for name-mangling of parameter names. The default is ''. If a particular `Model` has arguments `amplitude`, `center`, and `sigma`, these would become the parameter names. Using a prefix of 'g1_' would convert these parameter names to `g1_amplitude`, `g1_center`, and `g1_sigma`. This can be essential to avoid name collision in composite models.

7.2.3 Determining parameter names and independent variables for a function

The `Model` created from the supplied function `func` will guess which function arguments of the model function should be made into `Parameters` object, and which should be considered non-varying `independent variables` that need to be specified when evaluating or fitting with the Model.

By convention and by default, the first argument to the model function is taken to be the independent variable – this is often the `x` value for the model. You can explicitly set this to be a different functional argument, and will need to do so if the independent variable is not first in the list, or if there is more than one independent variable. Since curve fitting most commonly fits some function $y(x)$, it is common for the independent variable to be a numpy float ndarray of the same length as the data to be fit. While this is common, it is not actually required. The independent variable does not need to be an ndarray, and can be any Python data type, and it is fine to have multiple independent variables.

Changed in version 1.2.3.

By convention and default, the positional arguments (that is, those without default values specified in the function signature) other than the first argument are marked as being Parameters – these will have an invalid default value (or `-Inf`) that must be supplied before using the Parameter. Keyword arguments to the model function that have numerical values will also be marked as being Parameters, and the supplied numerical value will be used as the default value for that Parameter.

Keyword arguments to the model function that have non-numerical values (including `None`, `False`, and `True`, but also strings) will be marked as being independent variables and their default value will be stored and used unless overwritten. There is some ambiguity for function arguments that have a value supplied in the function signature of `None`, `False`, or `True`. These function arguments are tagged as independent variables, but can be made into a Parameter with `Model.make_params()`, in which case, they will be treated as variables. An example is given below.

7.2.4 Explicitly specifying `independent_vars`

As we saw for the Gaussian example above, creating a `Model` from a function is fairly easy. Let's try another one:

```
import numpy as np
from lmfit import Model

def decay(t, tau, N):
    return N*np.exp(-t/tau)

decay_model = Model(decay)
print(f'independent variables: {decay_model.independent_vars}')

params = decay_model.make_params()
print('\nParameters:')
for pname, par in params.items():
    print(pname, par)
```

```
independent variables: ['t']

Parameters:
tau <Parameter 'tau', value=-inf, bounds=[-inf:inf]>
N <Parameter 'N', value=-inf, bounds=[-inf:inf]>
```

Here, `t` is assumed to be the independent variable because it is the first argument to the function. The other function arguments are used to create parameters for the model.

If you want `tau` to be the independent variable in the above example, you can say so:

```
decay_model = Model(decay, independent_vars=['tau'])
print(f'independent variables: {decay_model.independent_vars}')

params = decay_model.make_params()
print('\nParameters:')
for pname, par in params.items():
    print(pname, par)
```

```
independent variables: ['tau']
```

(continues on next page)

(continued from previous page)

Parameters:

```
t <Parameter 't', value=-inf, bounds=[-inf:inf]>
N <Parameter 'N', value=-inf, bounds=[-inf:inf]>
```

You can also supply multiple values for multi-dimensional functions with multiple independent variables. In fact, the meaning of *independent variable* here is simple, and based on how it treats arguments of the function you are modeling:

independent variable

A function argument that is not a parameter or otherwise part of the model, and that will be required to be explicitly provided as a keyword argument for each fit with `Model.fit()` or evaluation with `Model.eval()`.

Note that independent variables are not required to be arrays, or even floating point numbers. Dictionaries or Objects from user-defined classes can be independent variables.

7.2.5 Functions with keyword arguments

If the model function had keyword parameters, these would be turned into Parameters if the supplied default value was a valid number:

```
def decay2(t, tau, N=10, check_positive=False):
    if check_positive:
        arg = abs(t)/max(1.e-9, abs(tau))
    else:
        arg = t/tau
    return N*np.exp(arg)

mod = Model(decay2)
print(f'independent variables: {mod.independent_vars}')

params = mod.make_params()
print('Parameters:')
for pname, par in params.items():
    print(pname, par)
```

```
independent variables: ['t', 'check_positive']
Parameters:
tau <Parameter 'tau', value=-inf, bounds=[-inf:inf]>
N <Parameter 'N', value=10, bounds=[-inf:inf]>
```

Here, `check_positive` is listed as an independent variable, because it has a default value of `False`.

The function argument `N` is expected to be a Parameter, because it has a numerical default value. This value will be kept as the default value when building the Parameters. Also, note that the default value for `tau` is `-np.inf`, which will need to be fixed before really be used.

Function arguments with default values of `None`, `False`, and `True` are slightly ambiguous whether they should be considered independent variables or Parameters. That is, in many contexts (including basic arithmetic) `True` acts like the integer 1, and `False` acts like 0. A default value of `None` might be used to signal some default behavior. For example, the builtin `lmfit.lineshapes.voigt()` function is defined as

```
def voigt(x, amplitude=1.0, center=0.0, sigma=1.0, gamma=None):
    """Return a 1-dimensional Voigt function.
```

(continues on next page)

(continued from previous page)

```

    ...
    """
if gamma is None:
    gamma = sigma
z = (x-center + 1j*gamma) / max(tiny, (sigma*s2))
return amplitude*real(wofz(z)) / max(tiny, (sigma*s2pi))

mod = Model(voigt)
print(f'independent variables: {mod.independent_vars}')

params = mod.make_params(amplitude=10, center=5, sigma=2)

# or
params2 = mod.make_params(amplitude=10, center=5, sigma=2, gamma=1.5)

```

```
independent variables: ['x', 'gamma']
```

In this case, `gamma` will be listed as an independent variable. But because its default value is `None` (or `True` or `False`) it can be made a Parameter that can be varied independently as shown above. Of course, whether this makes sense depends on the details of the implementation of the model function: it can be sensible (and even intended) for `gamma` in the `voigt` function above to be a numerical value that could be a variable Parameter, but it is not sensible for the `check_positive` argument in `decay2` to be variable that can take any numeric value.

7.2.6 Defining a prefix for the Parameters

As we will see below when combining models, it is sometimes necessary to decorate the parameter names in the model, but still have them be correctly used in the underlying model function. This would be necessary, for example, if two parameters in a composite model (see [Composite Models : adding \(or multiplying\) Models](#) or examples in the next chapter) would have the same name. To avoid this, we can add a `prefix` to the `Model` which will automatically do this mapping for us.

```

def myfunc(x, amplitude=1, center=0, sigma=1):
    # function definition, for now just ``pass``
    pass

mod = Model(myfunc, prefix='f1_')
params = mod.make_params()
print('Parameters:')
for pname, par in params.items():
    print(pname, par)

```

```
Parameters:
f1_amplitude <Parameter 'f1_amplitude', value=1, bounds=[-inf:inf]>
f1_center <Parameter 'f1_center', value=0, bounds=[-inf:inf]>
f1_sigma <Parameter 'f1_sigma', value=1, bounds=[-inf:inf]>
```

You would refer to these parameters as `f1_amplitude` and so forth, and the model will know to map these to the `amplitude` argument of `myfunc`.

7.2.7 Initializing model parameter values

As mentioned above, creating a model does not automatically create the corresponding `Parameters`. These can be created with either the `create_params()` function, or the `Model.make_params()` method of the corresponding instance of `Model`.

When creating Parameters, each parameter is created with invalid initial value of `-Inf` if it is not set explicitly. That is to say, parameter values **must** be initialized in order for the model to evaluate a finite result or used in a fit. There are a few different ways to do this:

1. You can supply initial values in the definition of the model function.
2. You can initialize the parameters when creating parameters with `Model.make_params()`.
3. You can create a `Parameters` object with `Parameters` or `create_params()`.
4. You can supply initial values for the parameters when calling `Model.eval()` or `Model.fit()` methods.

Generally, using the `Model.make_params()` method is recommended. The methods described above can be mixed, allowing you to overwrite initial values at any point in the process of defining and using the model.

Initializing values in the function definition

To supply initial values for parameters in the definition of the model function, you can simply supply a default value:

```
def myfunc(x, a=1, b=0):  
    return a*x + 10*a - b
```

instead of using:

```
def myfunc(x, a, b):  
    return a*x + 10*a - b
```

This has the advantage of working at the function level – all parameters with keywords can be treated as options. It also means that some default initial value will always be available for the parameter.

Initializing values with `Model.make_params()`

When creating parameters with `Model.make_params()` you can specify initial values. To do this, use keyword arguments for the parameter names. You can either set initial values as numbers (floats or ints) or as dictionaries with keywords of (`value`, `vary`, `min`, `max`, `expr`, `brute_step`, and `is_init_value`) to specify these parameter attributes.

```
mod = Model(myfunc)  
  
# simply supply initial values  
pars = mod.make_params(a=3, b=0.5)  
  
# supply initial values, attributes for bounds, etcetera:  
pars_bound = mod.make_params(a=dict(value=3, min=0),  
                           b=dict(value=0.5, vary=False))
```

Creating a Parameters object directly

You can also create your own Parameters directly using `create_params()`. This is independent of using the `Model` class, but is essentially equivalent to `Model.make_params()` except with less checking of errors for model prefixes and so on.

```
from lmfit import create_params

mod = Model(myfunc)

# simply supply initial values
pars = create_params(a=3, b=0.5)

# supply initial values and attributes for bounds, etc:
pars_bounded = create_params(a=dict(value=3, min=0),
                             b=dict(value=0.5, vary=False))
```

Because less error checking is done, `Model.make_params()` should probably be preferred when using Models.

Initializing parameter values for a model with keyword arguments

Finally, you can explicitly supply initial values when using a model. That is, as with `Model.make_params()`, you can include values as keyword arguments to either the `Model.eval()` or `Model.fit()` methods:

```
x = linspace(0, 10, 100)
y_eval = mod.eval(x=x, a=7.0, b=-2.0)
y_sim = y_eval + random.normal(0, 0.2, x.size)
out = mod.fit(y_sim, pars, x=x, a=3.0, b=0.0)
```

These approaches to initialization provide many opportunities for setting initial values for parameters. The methods can be combined, so that you can set parameter hints but then change the initial value explicitly with `Model.fit()`.

7.2.8 Using parameter hints

After a model has been created, but prior to creating parameters with `Model.make_params()`, you can define parameter hints for that model. This allows you to set other parameter attributes for bounds, whether it is varied in the fit, or set a default constraint expression for a parameter. You can also set the initial value, but that is not really the intention of the method, which is to really let you say that about the idealized Model, for example that some values may not make sense for some parameters, or that some parameters might be a small change from another parameter and so be fixed or constrained by default.

To set a parameter hint, you can use `Model.set_param_hint()`, as with:

```
mod = Model(myfunc)
mod.set_param_hint('bounded_parameter', min=0, max=1.0)
pars = mod.make_params()
```

Parameter hints are discussed in more detail in section [Using parameter hints](#).

Parameter hints are stored in a model's `param_hints` attribute, which is simply a nested dictionary:

```
print('Parameter hints:')
for pname, par in mod.param_hints.items():
    print(pname, par)
```

Parameter hints:

```
bounded_parameter {'min': 0, 'max': 1.0}
```

You can change this dictionary directly or use the `Model.set_param_hint()` method. Either way, these parameter hints are used by `Model.make_params()` when making parameters.

Parameter hints also allow you to create new parameters. This can be useful to make derived parameters with constraint expressions. For example to get the full-width at half maximum of a Gaussian model, one could use a parameter hint of:

```
mod = Model(gaussian)
mod.set_param_hint('wid', min=0)
mod.set_param_hint('fwhm', expr='2.3548*wid')
params = mod.make_params(amp={'value': 10, 'min':0.1, 'max':2000},
                        cen=5.5, wid=1.25)
params.pretty_print()
```

Name	Value	Min	Max	Stderr	Vary	Expr	Brute_Step
amp	10	0.1	2000	None	True	None	None
cen	5.5	-inf	inf	None	True	None	None
fwhm	2.944	-inf	inf	None	False	2.3548*wid	None
wid	1.25	0	inf	None	True	None	None

With that definition, the value (and uncertainty) of the `fwhm` parameter will be reported in the output of any fit done with that model.

7.2.9 Data Types for data and independent data with Model

The model as defined by your model function will use the independent variable(s) you specify to best match the data you provide. The model is meant to be an abstract representation for data, but when you do a fit with `Model.fit()`, you really need to pass in values for the data to be modeled and the independent data used to calculate that data.

As discussed in [Types of Data to Use for Fitting](#), the mathematical solvers used by `lmfit` all work exclusively with 1-dimensional numpy arrays of datatype (dtype) “float64”. The value of the calculation `(model-data)*weights` using the calculation of your model function, and the data and weights you pass in **will always be coerced** to an 1-dimensional ndarray with dtype “float64” when it is passed to the solver. If it cannot be coerced, an error will occur and the fit will be aborted.

That coercion will usually work for “array like” data that is not already a float64 ndarray. But, depending on the model function, the calculations within the model function may not always work well for some “array like” data types - especially independent data that are in list of numbers and ndarrays of type “float32” or “int16” or less precision.

To be clear, independent data for models using `Model` are meant to be truly independent, and not **not** required to be strictly numerical or objects that are easily converted to arrays of numbers. This could, for example, be a dictionary, an instance of a user-defined class, or other type of structured data. You can use independent data any way you want in your model function. But, as with almost all the examples given here, independent data is often also a 1-dimensional array of values, say `x`, and a simple view of the fit would be to plot the data as `y` as a function of `x`. Again, this is not required, but it is very common, especially for novice users.

By default, all data and independent data passed to `Model.fit()` that is “array like” - a list or tuple of numbers, a `pandas.Series`, and `h5py.Dataset`, or any object that has an `__array__()` method – will be converted to a “float64” ndarray before the fit begins. If the array-like data is complex, it will be converted to a “complex128” ndarray, which will always work too. This conversion before the fit begins ensures that the model function sees only “float64 ndarrays”, and nearly guarantees that data type conversion will not cause problems for the fit. But it also means that

if you have passed a `pandas.Series` as data or independent data, not all of the methods or attributes of that `Series` will be available by default within the model function.

Added in version 1.2.2.

This coercion can be turned off with the `coerce_farray` option to `Model.fit()`. When set to `False`, neither the data nor the independent data will be coerced from their original data type, and the user will be responsible to arrange for the calculation and return value from the model function to be allow a proper and accurate conversion to a “`float64`” `ndarray`.

See also [Types of Data to Use for Fitting](#) for general advise and recommendations on types of data to use when fitting data.

7.2.10 Saving and Loading Models

It is sometimes desirable to save a `Model` for later use outside of the code used to define the model. Lmfit provides a `save_model()` function that will save a `Model` to a file. There is also a companion `load_model()` function that can read this file and reconstruct a `Model` from it.

Saving a model turns out to be somewhat challenging. The main issue is that Python is not normally able to *serialize* a function (such as the model function making up the heart of the Model) in a way that can be reconstructed into a callable Python object. The `dill` package can sometimes serialize functions, but with the limitation that it can be used only in the same version of Python. In addition, class methods used as model functions will not retain the rest of the class attributes and methods, and so may not be usable. With all those warnings, it should be emphasized that if you are willing to save or reuse the definition of the model function as Python code, then saving the Parameters and rest of the components that make up a model presents no problem.

If the `dill` package is installed, the model function will also be saved using it. But because saving the model function is not always reliable, saving a model will always save the *name* of the model function. The `load_model()` takes an optional `funcdefs` argument that can contain a dictionary of function definitions with the function names as keys and function objects as values. If one of the dictionary keys matches the saved name, the corresponding function object will be used as the model function. If it is not found by name, and if `dill` was used to save the model, and if `dill` is available at run-time, the `dill`-encoded function will try to be used. Note that this approach will generally allow you to save a model that can be used by another installation of the same version of Python, but may not work across Python versions. For preserving fits for extended periods of time (say, archiving for documentation of scientific results), we strongly encourage you to save the full Python code used for the model function and fit process.

`save_model(model, fname)`

Save a Model to a file.

Parameters

- `model` (`Model`) – Model to be saved.
- `fname` (`str`) – Name of file for saved Model.

`load_model(fname, funcdefs=None)`

Load a saved Model from a file.

Parameters

- `fname` (`str`) – Name of file containing saved Model.
- `funcdefs` (`dict`, *optional*) – Dictionary of custom function names and definitions.

Returns

Model object loaded from file.

Return type

`Model`

As a simple example, one can save a model as:

```
# <examples/doc_model_savemodel.py>
import numpy as np

from lmfit.model import Model, save_model

def mysine(x, amp, freq, shift):
    return amp * np.sin(x*freq + shift)

sinemodel = Model(mysine)
pars = sinemodel.make_params(amp=1, freq=0.25, shift=0)

save_model(sinemodel, 'sinemodel.sav')
# <end examples/doc_model_savemodel.py>
```

To load that later, one might do:

```
# <examples/doc_model_loadmodel.py>
import os
import sys

import matplotlib.pyplot as plt
import numpy as np

from lmfit.model import load_model

if not os.path.exists('sinemodel.sav'):
    os.system(f"{sys.executable} doc_model_savemodel.py")

def mysine(x, amp, freq, shift):
    return amp * np.sin(x*freq + shift)

data = np.loadtxt('sinedata.dat')
x = data[:, 0]
y = data[:, 1]

model = load_model('sinemodel.sav', funcdefs={'mysine': mysine})
params = model.make_params(amp=dict(value=3, min=0),
                           freq=0.52,
                           shift=dict(value=0, min=-1, max=1))

result = model.fit(y, params, x=x)
print(result.fit_report())

plt.plot(x, y, 'o')
plt.plot(x, result.best_fit, '-')
plt.show()
# <end examples/doc_model_loadmodel.py>
```

See also [Saving and Loading ModelResults](#).

7.3 The ModelResult class

A `ModelResult` (which had been called `ModelFit` prior to version 0.9) is the object returned by `Model.fit()`. It is a subclass of `Minimizer`, and so contains many of the fit results. Of course, it knows the `Model` and the set of `Parameters` used in the fit, and it has methods to evaluate the model, to fit the data (or re-fit the data with changes to the parameters, or fit with different or modified data) and to print out a report for that fit.

While a `Model` encapsulates your model function, it is fairly abstract and does not contain the parameters or data used in a particular fit. A `ModelResult` *does* contain parameters and data as well as methods to alter and re-do fits. Thus the `Model` is the idealized model while the `ModelResult` is the messier, more complex (but perhaps more useful) object that represents a fit with a set of parameters to data with a model.

A `ModelResult` has several attributes holding values for fit results, and several methods for working with fits. These include statistics inherited from `Minimizer` useful for comparing different models, including `chisqr`, `redchi`, `aic`, and `bic`.

```
class ModelResult(model, params, data=None, weights=None, method='leastsq', fcn_args=None,
                  fcn_kws=None, iter_cb=None, scale_covar=True, nan_policy='raise', calc_covar=True,
                  max_nfev=None, **fit_kws)
```

Result from the Model fit.

This has many attributes and methods for viewing and working with the results of a fit using `Model`. It inherits from `Minimizer`, so that it can be used to modify and re-run the fit for the `Model`.

Parameters

- `model` (`Model`) – Model to use.
- `params` (`Parameters`) – Parameters with initial values for model.
- `data` (`array_like`, *optional*) – Data to be modeled.
- `weights` (`array_like`, *optional*) – Weights to multiply (`data`-`model`) for fit residual.
- `method` (`str`, *optional*) – Name of minimization method to use (default is '`leastsq`').
- `fcn_args` (`sequence`, *optional*) – Positional arguments to send to model function.
- `fcn_dict` (`dict`, *optional*) – Keyword arguments to send to model function.
- `iter_cb` (`callable`, *optional*) – Function to call on each iteration of fit.
- `scale_covar` (`bool`, *optional*) – Whether to scale covariance matrix for uncertainty evaluation.
- `nan_policy` (`{'raise', 'propagate', 'omit'}`, *optional*) – What to do when encountering NaNs when fitting `Model`.
- `calc_covar` (`bool`, *optional*) – Whether to calculate the covariance matrix (default is `True`) for solvers other than '`leastsq`' and '`least_squares`'. Requires the `numdifftools` package to be installed.
- `max_nfev` (`int` or `None`, *optional*) – Maximum number of function evaluations (default is `None`). The default value depends on the fitting method.
- `**fit_kws` (*optional*) – Keyword arguments to send to minimization routine.

7.3.1 ModelResult methods

`ModelResult.eval(params=None, **kwargs)`

Evaluate model function.

Parameters

- `params (Parameters, optional)` – Parameters to use.
- `**kwargs (optional)` – Options to send to Model.eval().

Returns

Array or value for the evaluated model.

Return type

`numpy.ndarray, float, int, or complex`

`ModelResult.eval_components(params=None, **kwargs)`

Evaluate each component of a composite model function.

Parameters

- `params (Parameters, optional)` – Parameters, defaults to ModelResult.params.
- `**kwargs (optional)` – Keyword arguments to pass to model function.

Returns

Keys are prefixes of component models, and values are the estimated model value for each component of the model.

Return type

`dict`

`ModelResult.fit(data=None, params=None, weights=None, method=None, nan_policy=None, **kwargs)`

Re-perform fit for a Model, given data and params.

Parameters

- `data (array_like, optional)` – Data to be modeled.
- `params (Parameters, optional)` – Parameters with initial values for model.
- `weights (array_like, optional)` – Weights to multiply (data-model) for fit residual.
- `method (str, optional)` – Name of minimization method to use (default is ‘leastsq’).
- `nan_policy ({'raise', 'propagate', 'omit'}, optional)` – What to do when encountering NaNs when fitting Model.
- `**kwargs (optional)` – Keyword arguments to send to minimization routine.

`ModelResult.fit_report(modelpars=None, show_correl=True, min_correl=0.1, sort_pars=False, correl_mode='list')`

Return a printable fit report.

The report contains fit statistics and best-fit values with uncertainties and correlations.

Parameters

- `modelpars (Parameters, optional)` – Known Model Parameters.
- `show_correl (bool, optional)` – Whether to show list of sorted correlations (default is True).

- **min_correl** (*float, optional*) – Smallest correlation in absolute value to show (default is 0.1).
- **sort_pars** (*callable, optional*) – Whether to show parameter names sorted in alphanumerical order (default is False). If False, then the parameters will be listed in the order as they were added to the Parameters dictionary. If callable, then this (one argument) function is used to extract a comparison key from each list element.
- **correl_mode** ({'list', 'table'} *str, optional*) – Mode for how to show correlations. Can be either ‘list’ (default) to show a sorted (if **sort_pars** is True) list of correlation values, or ‘table’ to show a complete, formatted table of correlations.

Returns

Multi-line text of fit report.

Return type

str

ModelResult.summary()

Return a dictionary with statistics and attributes of a ModelResult.

Returns

Dictionary of statistics and many attributes from a ModelResult.

Return type

dict

Notes

1. values for data arrays are not included.
2. The result summary dictionary will include the following entries:

```
model, method, ndata, nvarys, nfree, chisqr, redchi, aic, bic, rsquared, nfev, max_nfev, aborted, errorbars, success, message, lmdif_message, ier, nan_policy, scale_covar, calc_covar, ci_out, col_deriv, flatchain, call_kws, var_names, user_options, kws, init_values, best_values, and params.
```

where ‘params’ is a list of parameter “states”: tuples with entries of (*name, value, vary, expr, min, max, brute_step, stderr, correl, init_value, user_data*).

3. The result will include only plain Python objects, and so should be easily serializable with JSON or similar tools.

ModelResult.conf_interval(**kwargs)

Calculate the confidence intervals for the variable parameters.

Confidence intervals are calculated using the `confidence.conf_interval()` function and keyword arguments (***kwargs*) are passed to that function. The result is stored in the `ci_out` attribute so that it can be accessed without recalculating them.

ModelResult.ci_report(*with_offset=True, ndigits=5, **kwargs*)

Return a formatted text report of the confidence intervals.

Parameters

- **with_offset** (*bool, optional*) – Whether to subtract best value from all other values (default is True).
- **ndigits** (*int, optional*) – Number of significant digits to show (default is 5).

- ****kwargs** (*optional*) – Keyword arguments that are passed to the `conf_interval` function.

Returns

Text of formatted report on confidence intervals.

Return type

`str`

`ModelResult.eval_uncertainty(params=None, sigma=1, dscale=0.01, **kwargs)`

Evaluate the uncertainty of the *model function*.

This can be used to give confidence bands for the model from the uncertainties in the best-fit parameters.

Parameters

- **params** (`Parameters`, *optional*) – Parameters, defaults to `ModelResult.params`.
- **sigma** (`float`, *optional*) – Confidence level, i.e. how many sigma (default is 1).
- **dscale** (`float`, *optional*) – Scale for derivative steps (default is 0.01).
- ****kwargs** (*optional*) – Values of options, independent variables, etcetera.

Returns

Uncertainty at each value of the model.

Return type

`numpy.ndarray`

Notes

1. This is based on the excellent and clear example from <https://www.astro.rug.nl/software/kapteyn/kmpfittutorial.html#confidence-and-prediction-intervals>, which references the original work of: J. Wolberg, Data Analysis Using the Method of Least Squares, 2006, Springer
2. The value of `sigma` is number of `sigma` values, and is converted to a probability. Values of 1, 2, or 3 give probabilities of 0.6827, 0.9545, and 0.9973, respectively. If the `sigma` value is < 1, it is interpreted as the probability itself. That is, `sigma=1` and `sigma=0.6827` will give the same results, within precision errors.
3. The derivatives are calculated by stepping each Parameter from its best value to to `+- stderr*dscale`, where `dscale` can be passed in and defaults to 0.01.
4. Sets attributes of `dely` for the uncertainty of the model (which will be the same as the array returned by this method) and `dely_comps`, a dictionary of `dely` for each component.
5. Sets the attribute of `dely_predicted` for the ‘predicted interval’, the sigma-scaled quadrature sum of the uncertainty interval `dely` and reduced chi-square. This should give an idea of the expected range in the data.

Examples

```
>>> out = model.fit(data, params, x=x)
>>> dely = out.eval_uncertainty(x=x)
>>> plt.plot(x, data)
>>> plt.plot(x, out.best_fit)
>>> plt.fill_between(x, out.best_fit-dely,
...                   out.best_fit+dely, color='#888888')
```

```
ModelResult.plot(datafmt='o', fitfmt='-', initfmt='--', xlabel=None, ylabel=None, yerr=None, numpoints=None,
fig=None, data_kws=None, fit_kws=None, init_kws=None, ax_res_kws=None,
ax_fit_kws=None, fig_kws=None, show_init=False, parse_complex='abs', title=None)
```

Plot the fit results and residuals using matplotlib.

The method will produce a matplotlib figure (if package available) with both results of the fit and the residuals plotted. If the fit model included weights, errorbars will also be plotted. To show the initial conditions for the fit, pass the argument `show_init=True`.

Parameters

- `datafmt (str, optional)` – Matplotlib format string for data points.
- `fitfmt (str, optional)` – Matplotlib format string for fitted curve.
- `initfmt (str, optional)` – Matplotlib format string for initial conditions for the fit.
- `xlabel (str, optional)` – Matplotlib format string for labeling the x-axis.
- `ylabel (str, optional)` – Matplotlib format string for labeling the y-axis.
- `yerr (numpy.ndarray, optional)` – Array of uncertainties for data array.
- `numpoints (int, optional)` – If provided, the final and initial fit curves are evaluated not only at data points, but refined to contain `numpoints` points in total.
- `fig (matplotlib.figure.Figure, optional)` – The figure to plot on. The default is None, which means use the current pyplot figure or create one if there is none.
- `data_kws (dict, optional)` – Keyword arguments passed to the plot function for data points.
- `fit_kws (dict, optional)` – Keyword arguments passed to the plot function for fitted curve.
- `init_kws (dict, optional)` – Keyword arguments passed to the plot function for the initial conditions of the fit.
- `ax_res_kws (dict, optional)` – Keyword arguments for the axes for the residuals plot.
- `ax_fit_kws (dict, optional)` – Keyword arguments for the axes for the fit plot.
- `fig_kws (dict, optional)` – Keyword arguments for a new figure, if a new one is created.
- `show_init (bool, optional)` – Whether to show the initial conditions for the fit (default is False).
- `parse_complex ({'abs', 'real', 'imag', 'angle'}, optional)` – How to reduce complex data for plotting. Options are one of: 'abs' (default), 'real', 'imag', or 'angle', which correspond to the NumPy functions with the same name.
- `title (str, optional)` – Matplotlib format string for figure title.

Return type

`matplotlib.figure.Figure`


See also

`ModelResult.plot_fit`

Plot the fit results using matplotlib.

`ModelResult.plot_residuals`

Plot the fit residuals using matplotlib.



Notes

The method combines `ModelResult.plot_fit` and `ModelResult.plot_residuals`.

If `yerr` is specified or if the fit model included weights, then `matplotlib.axes.Axes.errorbar` is used to plot the data. If `yerr` is not specified and the fit includes weights, `yerr` set to `1/self.weights`.

If model returns complex data, `yerr` is treated the same way that weights are in this case.

If `fig` is `None` then `matplotlib.pyplot.figure(**fig_kws)` is called, otherwise `fig_kws` is ignored.

```
ModelResult.plot_fit(ax=None, datafmt='o', fitfmt='-', initfmt='--', xlabel=None, ylabel=None, yerr=None,
                      numpoints=None, data_kws=None, fit_kws=None, init_kws=None, ax_kws=None,
                      show_init=False, parse_complex='abs', title=None)
```

Plot the fit results using matplotlib, if available.

The plot will include the data points, the initial fit curve (optional, with `show_init=True`), and the best-fit curve. If the fit model included weights or if `yerr` is specified, errorbars will also be plotted.

Parameters

- `ax (matplotlib.axes.Axes, optional)` – The axes to plot on. The default is `None`, which means use the current pyplot axis or create one if there is none.
- `datafmt (str, optional)` – Matplotlib format string for data points.
- `fitfmt (str, optional)` – Matplotlib format string for fitted curve.
- `initfmt (str, optional)` – Matplotlib format string for initial conditions for the fit.
- `xlabel (str, optional)` – Matplotlib format string for labeling the x-axis.
- `ylabel (str, optional)` – Matplotlib format string for labeling the y-axis.
- `yerr (numpy.ndarray, optional)` – Array of uncertainties for data array.
- `numpoints (int, optional)` – If provided, the final and initial fit curves are evaluated not only at data points, but refined to contain `numpoints` points in total.
- `data_kws (dict, optional)` – Keyword arguments passed to the plot function for data points.
- `fit_kws (dict, optional)` – Keyword arguments passed to the plot function for fitted curve.
- `init_kws (dict, optional)` – Keyword arguments passed to the plot function for the initial conditions of the fit.
- `ax_kws (dict, optional)` – Keyword arguments for a new axis, if a new one is created.
- `show_init (bool, optional)` – Whether to show the initial conditions for the fit (default is `False`).
- `parse_complex ({'abs', 'real', 'imag', 'angle'}, optional)` – How to reduce complex data for plotting. Options are one of: `'abs'` (default), `'real'`, `'imag'`, or `'angle'`, which correspond to the NumPy functions with the same name.
- `title (str, optional)` – Matplotlib format string for figure title.

Return type

`matplotlib.axes.Axes`

See also

`ModelErrorResult.plot_residuals`

Plot the fit residuals using matplotlib.

`ModelErrorResult.plot`

Plot the fit results and residuals using matplotlib.

Notes

For details about plot format strings and keyword arguments see documentation of `matplotlib.axes.Axes.plot`.

If `yerr` is specified or if the fit model included weights, then `matplotlib.axes.Axes.errorbar` is used to plot the data. If `yerr` is not specified and the fit includes weights, `yerr` set to `1/weights`.

If model returns complex data, `yerr` is treated the same way that weights are in this case.

If `ax` is `None` then `matplotlib.pyplot.gca(**ax_kws)` is called.

`ModelErrorResult.plot_residuals(ax=None, datafmt='o', yerr=None, data_kws=None, fit_kws=None, ax_kws=None, parse_complex='abs', title=None)`

Plot the fit residuals using matplotlib, if available.

If `yerr` is supplied or if the model included weights, errorbars will also be plotted.

Parameters

- `ax` (`matplotlib.axes.Axes`, `optional`) – The axes to plot on. The default is `None`, which means use the current pyplot axis or create one if there is none.
- `datafmt` (`str`, `optional`) – Matplotlib format string for data points.
- `yerr` (`numpy.ndarray`, `optional`) – Array of uncertainties for data array.
- `data_kws` (`dict`, `optional`) – Keyword arguments passed to the plot function for data points.
- `fit_kws` (`dict`, `optional`) – Keyword arguments passed to the plot function for fitted curve.
- `ax_kws` (`dict`, `optional`) – Keyword arguments for a new axis, if a new one is created.
- `parse_complex` (`{'abs', 'real', 'imag', 'angle'}`, `optional`) – How to reduce complex data for plotting. Options are one of: `'abs'` (default), `'real'`, `'imag'`, or `'angle'`, which correspond to the NumPy functions with the same name.
- `title` (`str`, `optional`) – Matplotlib format string for figure title.

Return type

`matplotlib.axes.Axes`

See also

`ModelErrorResult.plot_fit`

Plot the fit results using matplotlib.

`ModelErrorResult.plot`

Plot the fit results and residuals using matplotlib.

Notes

For details about plot format strings and keyword arguments see documentation of `matplotlib.axes.Axes.plot`.

If `yerr` is specified or if the fit model included weights, then `matplotlib.axes.Axes.errorbar` is used to plot the data. If `yerr` is not specified and the fit includes weights, `yerr` set to `1/self.weights`.

If model returns complex data, `yerr` is treated the same way that weights are in this case.

If `ax` is `None` then `matplotlib.pyplot.gca(**ax_kws)` is called.

`ModelResult.ITER_CB()`

Optional callable function, to be called at each fit iteration. This must take arguments of `(params, iter, resid, *args, **kws)`, where `params` will have the current parameter values, `iter` the iteration, `resid` the current residual array, and `*args` and `**kws` as passed to the objective function. See [Using a Iteration Callback Function](#).

`ModelResult.JACFCN()`

Optional callable function, to be called to calculate Jacobian array.

7.3.2 ModelResult attributes

A `ModelResult` will take all of the attributes of `MinimizerResult`, and several more. Here, we arrange them into categories.

Parameters and Variables

`best_values`

Dictionary with parameter names as keys, and best-fit values as values.

`init_params`

Initial parameters, as passed to `Model.fit()`.

`init_values`

Dictionary with parameter names as keys, and initial values as values.

`init_vals`

list of values for the variable parameters.

`params`

Parameters used in fit; will contain the best-fit values and uncertainties if available.

`uvars`

Dictionary of `uncertainties` ufloats from Parameters.

`var_names`

List of variable Parameter names used in optimization in the same order as the values in `init_vals` and `covar`.

Fit Arrays and Model

best_fit

numpy.ndarray result of model function, evaluated at provided independent variables and with best-fit parameters.

covar

numpy.ndarray (square) covariance matrix returned from fit.

data

numpy.ndarray of data to compare to model.

dely

numpy.ndarray of estimated uncertainties or *confidence interval* in the y values of the model from [ModelResult.eval_uncertainty\(\)](#) (see [Calculating uncertainties in the model function](#)).

dely_comps

a dictionary of estimated uncertainties in the y values of the model components, from [ModelResult.eval_uncertainty\(\)](#) (see [Calculating uncertainties in the model function](#)).

dely_predicted

numpy.ndarray of estimated expected uncertainties in the data or *predicted interval* in the y values of the model from [ModelResult.eval_uncertainty\(\)](#) (see [Calculating uncertainties in the model function](#)).

init_fit

numpy.ndarray result of model function, evaluated at provided independent variables and with initial parameters.

residual

numpy.ndarray for residual.

weights

numpy.ndarray (or `None`) of weighting values to be used in fit. If not `None`, it will be used as a multiplicative factor of the residual array, so that `weights*(data - fit)` is minimized in the least-squares sense.

components

List of components of the [Model](#).

Fit Status

aborted

Whether the fit was aborted.

errorbars

Boolean for whether error bars were estimated by fit.

flatchain

A `pandas.DataFrame` view of the sampling chain if the `emcee` method is used.

ier

Integer returned code from `scipy.optimize.leastsq`.

lmdif_message

String message returned from `scipy.optimize.leastsq`.

message

String message returned from `minimize()`.

method

String naming fitting method for `minimize()`.

call_kws

Dict of keyword arguments actually send to underlying solver with `minimize()`.

model

Instance of `Model` used for model.

scale_covar

Boolean flag for whether to automatically scale covariance matrix.

userargs

positional arguments passed to `Model.fit()`, a tuple of (y, weights)

userkws

keyword arguments passed to `Model.fit()`, a dict, which will have independent data arrays such as x.

Fit Statistics

aic

Floating point best-fit Akaike Information Criterion statistic (see [MinimizerResult – the optimization result](#)).

bic

Floating point best-fit Bayesian Information Criterion statistic (see [MinimizerResult – the optimization result](#)).

chisqr

Floating point best-fit chi-square statistic (see [MinimizerResult – the optimization result](#)).

ci_out

Confidence interval data (see [Calculation of confidence intervals](#)) or None if the confidence intervals have not been calculated.

ndata

Integer number of data points.

nfev

Integer number of function evaluations used for fit.

nfree

Integer number of free parameters in fit.

nvarys

Integer number of independent, freely varying variables in fit.

redchi

Floating point reduced chi-square statistic (see [MinimizerResult – the optimization result](#)).

rsquared

Floating point R^2 statistic, defined for data y and best-fit model f as

$$R^2 = 1 - \frac{\sum_i (y_i - f_i)^2}{\sum_i (y_i - \bar{y})^2}$$

success

Boolean value of whether fit succeeded. This is an optimistic view of success, meaning that the method finished without error.


7.3.3 Calculating uncertainties in the model function

We return to the first example above and ask not only for the uncertainties in the fitted parameters but for the range of values that those uncertainties mean for the model function itself. We can use the `ModelResult.eval_uncertainty()` method of the model result object to evaluate the uncertainty in the model with a specified level for σ .

That is, adding:

```
dely = result.eval_uncertainty(sigma=3)
plt.fill_between(x, result.best_fit-dely, result.best_fit+dely, color="#ABABAB",
                  label=r'3-$\sigma$ uncertainty band')
```

to the example fit to the Gaussian at the beginning of this chapter will give 3σ bands for the best-fit Gaussian, and produce the figure below.



Added in version 1.0.4.

If the model is a composite built from multiple components, the `ModelResult.eval_uncertainty()` method will evaluate the uncertainty of both the full model (often the sum of multiple components) as well as the uncertainty in each component. The uncertainty of the full model will be held in `result.dely`, and the uncertainties for each component will be held in the dictionary `result.dely_comps`, with keys that are the component prefixes.

An example script shows how the uncertainties in components of a composite model can be calculated and used:

```

# <examples/doc_model_uncertainty2.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit.models import ExponentialModel, GaussianModel

dat = np.loadtxt('NIST_Gauss2.dat')
x = dat[:, 1]
y = dat[:, 0]

model = (GaussianModel(prefix='g1_') +
          GaussianModel(prefix='g2_') +
          ExponentialModel(prefix='bkg_'))

params = model.make_params(bkg_amplitude=100, bkg_decay=80,
                           g1_amplitude=3000,
                           g1_center=100,
                           g1_sigma=10,
                           g2_amplitude=3000,
                           g2_center=150,
                           g2_sigma=10)

result = model.fit(y, params, x=x)
print(result.fit_report(min_correl=0.5))

comps = result.eval_components(x=x)
dely = result.eval_uncertainty(sigma=3)

fig, axes = plt.subplots(2, 2, figsize=(12.8, 9.6))

axes[0][0].plot(x, y, 'o', color="#99002299", markersize=3, label='data')
axes[0][0].plot(x, result.best_fit, '--', label='best fit')
axes[0][0].plot(x, result.init_fit, '---', label='initial fit')
axes[0][0].set_title('data, initial fit, and best-fit')
axes[0][0].legend()

axes[0][1].plot(x, y, 'o', color="#99002299", markersize=3, label='data')
axes[0][1].plot(x, result.best_fit, '--', label='best fit')
axes[0][1].fill_between(x, result.best_fit-dely, result.best_fit+dely,
                       color="#8A8A8A", label=r'3-$\sigma$ band')
axes[0][1].set_title('data, best-fit, and uncertainty band')
axes[0][1].legend()

axes[1][0].plot(x, result.best_fit, '--', label=r'best fit, 3-$\sigma$ band')
axes[1][0].fill_between(x,
                       result.best_fit-dely,
                       result.best_fit+dely,
                       color="#8A8A8A")

axes[1][0].plot(x, comps['bkg_'], label=r'background, 3-$\sigma$ band')
axes[1][0].fill_between(x,
                       comps['bkg_']-result.dely_comps['bkg_'],
                       comps['bkg_']+result.dely_comps['bkg_'],

```

(continues on next page)

(continued from previous page)

```

color="#8A8A8A")

axes[1][0].plot(x, comps['g1_'], label=r'Gaussian #1, 3-$\sigma$ band')
axes[1][0].fill_between(x,
                       comps['g1_']-result.dely_comps['g1_'],
                       comps['g1_']+result.dely_comps['g1_'],
                       color="#8A8A8A")

axes[1][0].plot(x, comps['g2_'], label=r'Gaussian #2, 3-$\sigma$ band')
axes[1][0].fill_between(x,
                       comps['g2_']-result.dely_comps['g2_'],
                       comps['g2_']+result.dely_comps['g2_'],
                       color="#8A8A8A")
axes[1][0].set_title('model components with uncertainty bands')
axes[1][0].legend()

axes[1][1].plot(x, result.best_fit, '--', label='best fit')
axes[1][1].plot(x, 10*result.dely, label=r'3-$\sigma$ total (x10)')
axes[1][1].plot(x, 10*result.dely_comps['bkg_'], label=r'3-$\sigma$ background (x10)')
axes[1][1].plot(x, 10*result.dely_comps['g1_'], label=r'3-$\sigma$ Gaussian #1 (x10)')
axes[1][1].plot(x, 10*result.dely_comps['g2_'], label=r'3-$\sigma$ Gaussian #2 (x10)')
axes[1][1].set_title('uncertainties for model components')
axes[1][1].legend()

plt.show()
# <end examples/doc_model_uncertainty2.py>

```

```

[[Model]]
((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
→ prefix='bkg_'))
[[Fit Statistics]]
# fitting method   = leastsq
# function evals   = 55
# data points      = 250
# variables        = 8
chi-square         = 1247.52821
reduced chi-square = 5.15507524
Akaike info crit   = 417.864631
Bayesian info crit = 446.036318
R-squared          = 0.99648654
[[Variables]]
g1_amplitude:    4257.77399 +/- 42.3838008 (1.00%) (init = 3000)
g1_center:        107.030957 +/- 0.15006868 (0.14%) (init = 100)
g1_sigma:         16.6725789 +/- 0.16048222 (0.96%) (init = 10)
g2_amplitude:    2493.41715 +/- 36.1696228 (1.45%) (init = 3000)
g2_center:        153.270104 +/- 0.19466723 (0.13%) (init = 150)
g2_sigma:         13.8069453 +/- 0.18680099 (1.35%) (init = 10)
bkg_amplitude:   99.0183280 +/- 0.53748639 (0.54%) (init = 100)
bkg_decay:        90.9508824 +/- 1.10310769 (1.21%) (init = 80)
g1_fwhm:          39.2609222 +/- 0.37790675 (0.96%) == '2.3548200*g1_sigma'
g1_height:        101.880228 +/- 0.59217122 (0.58%) == '0.3989423*g1_amplitude/max(1e-
→15, g1_sigma)'

```

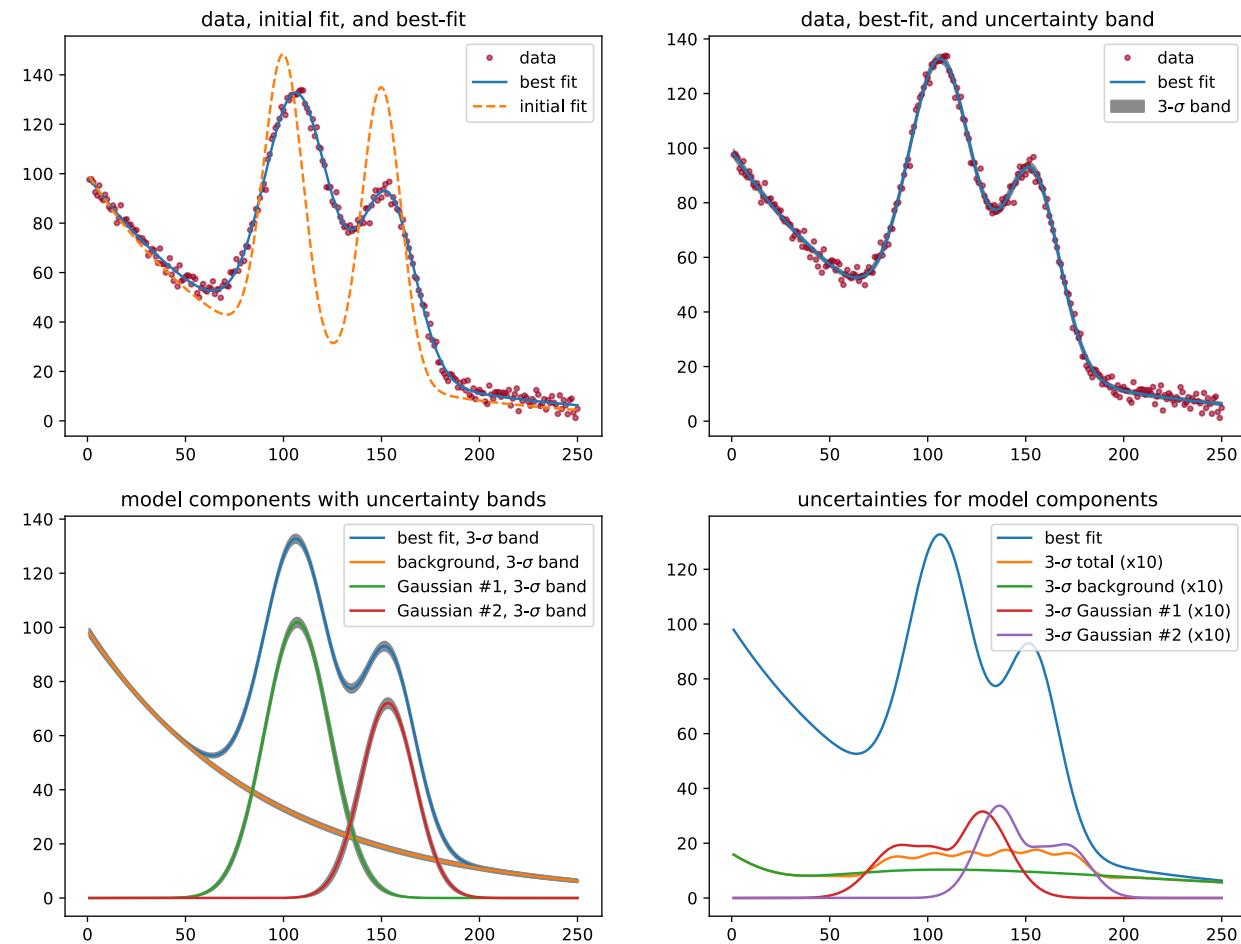
(continues on next page)

(continued from previous page)

```

g2_fwhm:      32.5128710 +/- 0.43988270 (1.35%) == '2.3548200*g2_sigma'
g2_height:    72.0455936 +/- 0.61721901 (0.86%) == '0.3989423*g2_amplitude/max(1e-
              -15, g2_sigma)'
[[Correlations]] (unreported correlations are < 0.500)
C(g1_amplitude, g1_sigma) = +0.8243
C(g2_amplitude, g2_sigma) = +0.8154
C(bkg_amplitude, bkg_decay) = -0.6946
C(g1_sigma, g2_center) = +0.6842
C(g1_center, g2_amplitude) = -0.6689
C(g1_center, g2_sigma) = -0.6520
C(g1_amplitude, g2_center) = +0.6477
C(g1_center, g2_center) = +0.6205
C(g1_center, g1_sigma) = +0.5075
C(g1_amplitude, bkg_decay) = -0.5074

```



Added in version 1.2.3.

In addition to the “confidence interval” `result.dely`, the `ModelResult.eval_uncertainty()` method will also estimate the “predicted interval” – the expected range for data matching the model, and hold this in `result.dely_predicted`. An example script showing both confidence and predicted intervals is shown below:

```
# <examples/doc_model_uncertainty_pred.py>
```

(continues on next page)

(continued from previous page)

```

# Here we reproduce the results presented in Section 5.1 of the
# nonlinear regression lecture by Julien Chiquet
# https://jchiquet.github.io/MAP566/docs/regression/map566-lecture-nonlinear-regression.
#html

import matplotlib.pyplot as plt
import numpy as np

from lmfit import Model

# data from the R base package datasets for the waiting time between eruptions
# and the duration of the eruption for the Old Faithful geyser in Yellowstone
# National Park, Wyoming, USA.

tlen = np.array([3.6, 1.8, 3.333, 2.283, 4.533, 2.883, 4.7, 3.6, 1.95, 4.35,
                 1.833, 3.917, 4.2, 1.75, 4.7, 2.167, 1.75, 4.8, 1.6, 4.25,
                 1.8, 1.75, 3.45, 3.067, 4.533, 3.6, 1.967, 4.083, 3.85, 4.433,
                 4.3, 4.467, 3.367, 4.033, 3.833, 2.017, 1.867, 4.833, 1.833,
                 4.783, 4.35, 1.883, 4.567, 1.75, 4.533, 3.317, 3.833, 2.1,
                 4.633, 2, 4.8, 4.716, 1.833, 4.833, 1.733, 4.883, 3.717,
                 1.667, 4.567, 4.317, 2.233, 4.5, 1.75, 4.8, 1.817, 4.4, 4.167,
                 4.7, 2.067, 4.7, 4.033, 1.967, 4.5, 4, 1.983, 5.067, 2.017,
                 4.567, 3.883, 3.6, 4.133, 4.333, 4.1, 2.633, 4.067, 4.933,
                 3.95, 4.517, 2.167, 4, 2.2, 4.333, 1.867, 4.817, 1.833, 4.3,
                 4.667, 3.75, 1.867, 4.9, 2.483, 4.367, 2.1, 4.5, 4.05, 1.867,
                 4.7, 1.783, 4.85, 3.683, 4.733, 2.3, 4.9, 4.417, 1.7, 4.633,
                 2.317, 4.6, 1.817, 4.417, 2.617, 4.067, 4.25, 1.967, 4.6,
                 3.767, 1.917, 4.5, 2.267, 4.65, 1.867, 4.167, 2.8, 4.333,
                 1.833, 4.383, 1.883, 4.933, 2.033, 3.733, 4.233, 2.233, 4.533,
                 4.817, 4.333, 1.983, 4.633, 2.017, 5.1, 1.8, 5.033, 4, 2.4,
                 4.6, 3.567, 4, 4.5, 4.083, 1.8, 3.967, 2.2, 4.15, 2, 3.833,
                 3.5, 4.583, 2.367, 5, 1.933, 4.617, 1.917, 2.083, 4.583,
                 3.333, 4.167, 4.333, 4.5, 2.417, 4, 4.167, 1.883, 4.583, 4.25,
                 3.767, 2.033, 4.433, 4.083, 1.833, 4.417, 2.183, 4.8, 1.833,
                 4.8, 4.1, 3.966, 4.233, 3.5, 4.366, 2.25, 4.667, 2.1, 4.35,
                 4.133, 1.867, 4.6, 1.783, 4.367, 3.85, 1.933, 4.5, 2.383, 4.7,
                 1.867, 3.833, 3.417, 4.233, 2.4, 4.8, 2, 4.15, 1.867, 4.267,
                 1.75, 4.483, 4, 4.117, 4.083, 4.267, 3.917, 4.55, 4.083,
                 2.417, 4.183, 2.217, 4.45, 1.883, 1.85, 4.283, 3.95, 2.333,
                 4.15, 2.35, 4.933, 2.9, 4.583, 3.833, 2.083, 4.367, 2.133,
                 4.35, 2.2, 4.45, 3.567, 4.5, 4.15, 3.817, 3.917, 4.45, 2,
                 4.283, 4.767, 4.533, 1.85, 4.25, 1.983, 2.25, 4.75, 4.117,
                 2.15, 4.417, 1.817, 4.467])

delay = np.array([79, 54, 74, 62, 85, 55, 88, 85, 51, 85, 54, 84, 78, 47, 83,
                  52, 62, 84, 52, 79, 51, 47, 78, 69, 74, 83, 55, 76, 78, 79,
                  73, 77, 66, 80, 74, 52, 48, 80, 59, 90, 80, 58, 84, 58, 73,
                  83, 64, 53, 82, 59, 75, 90, 54, 80, 54, 83, 71, 64, 77, 81,
                  59, 84, 48, 82, 60, 92, 78, 78, 65, 73, 82, 56, 79, 71, 62,
                  76, 60, 78, 76, 83, 75, 82, 70, 65, 73, 88, 76, 80, 48, 86,
                  60, 90, 50, 78, 63, 72, 84, 75, 51, 82, 62, 88, 49, 83, 81,

```

(continues on next page)

(continued from previous page)

```

47, 84, 52, 86, 81, 75, 59, 89, 79, 59, 81, 50, 85, 59, 87,
53, 69, 77, 56, 88, 81, 45, 82, 55, 90, 45, 83, 56, 89, 46,
82, 51, 86, 53, 79, 81, 60, 82, 77, 76, 59, 80, 49, 96, 53,
77, 77, 65, 81, 71, 70, 81, 93, 53, 89, 45, 86, 58, 78, 66,
76, 63, 88, 52, 93, 49, 57, 77, 68, 81, 81, 73, 50, 85, 74,
55, 77, 83, 83, 51, 78, 84, 46, 83, 55, 81, 57, 76, 84, 77,
81, 87, 77, 51, 78, 60, 82, 91, 53, 78, 46, 77, 84, 49, 83,
71, 80, 49, 75, 64, 76, 53, 94, 55, 76, 50, 82, 54, 75, 78,
79, 78, 78, 70, 79, 70, 54, 86, 50, 90, 54, 54, 77, 79, 64,
75, 47, 86, 63, 85, 82, 57, 82, 67, 74, 54, 83, 73, 73, 88,
80, 71, 83, 56, 79, 78, 84, 58, 83, 43, 60, 75, 81, 46, 90,
46, 74])

# model with a modified logistic function + offset
def logistic_func(t, amp, off, tau, gamma):
    return (amp-off)/(1+np.exp(-gamma*(t-tau))) + off

# set up model and initial parameter values
mod = Model(logistic_func)
pars = mod.make_params(amp=90, gamma=2, tau=2, off=50)

# run fit, and print report
result = mod.fit(delay, pars, t=tlen)
print(result.fit_report())

# make finely spaced grid of duration values, extending past data range
tfine = np.linspace(1, 6, 101)
yfine = result.eval(t=tfine)

# now calculate uncertainty interval and predicted interval for sigma=2, 95% level
efine = result.eval_uncertainty(t=tfine, sigma=2)
pfine = result.dely_predicted

plt.plot(tlen, delay, '.', label='observations')
plt.plot(tlen, result.best_fit, 'o', label='best fit')
plt.plot(tfine, yfine, '-', label='fine grid model')
plt.fill_between(tfine, yfine-efine, yfine+efine,
                 color="#c0c0c0", label=r'$2\sigma$ confidence interval')


plt.fill_between(tfine, yfine-pfine, yfine+pfine,
                 color="#d0d0a060", label=r'$2\sigma$ predicted interval')

plt.xlabel('duration (minutes)')
plt.ylabel('delay to next eruption (minutes)')
plt.legend()
plt.show()

# <end examples/doc_model_uncertainty_pred.py>

```

```
[[Model]]
  Model(logistic_func)
[[Fit Statistics]]
  # fitting method = leastsq
  # function evals = 60
  # data points   = 272
  # variables     = 4
  chi-square      = 8469.42359
  reduced chi-square = 31.6023268
  Akaike info crit = 943.249061
  Bayesian info crit = 957.672270
  R-squared        = 0.83090615
[[Variables]]
  amp: 82.4660404 +/- 0.99770817 (1.21%) (init = 90)
  off: 51.3218514 +/- 1.83125060 (3.57%) (init = 50)
  tau: 3.05525924 +/- 0.11068420 (3.62%) (init = 2)
  gamma: 2.25386376 +/- 0.43544396 (19.32%) (init = 2)
[[Correlations]] (unreported correlations are < 0.100)
  C(off, gamma) = +0.8634
  C(amp, gamma) = -0.7706
  C(off, tau)   = +0.7277
  C(amp, off)   = -0.5381
  C(tau, gamma) = +0.4699
```



7.3.4 Using uncertainties in the fitted parameters for post-fit calculations

Added in version 1.2.2.

As with the previous section, after a fit is complete, you may want to do some further calculations with the resulting Parameter values. Since these Parameters will have not only best-fit values but also usually have uncertainties, it is desirable for subsequent calculations to be able to propagate those uncertainties to any resulting calculated value. In addition, it is common for Parameters to have finite - and sometimes large - correlations which should be taken into account in such calculations.

The `ModelResult.uvars` will be a dictionary with keys for all variable Parameters and values that are `uvalues` from the `uncertainties` package. When used in mathematical calculations with basic Python operators or `numpy` functions, these `uvalues` will automatically propagate their uncertainties to the resulting calculation, and taking into account the full covariance matrix describing the correlation between values.

This readily allows “derived Parameters” to be evaluated just after the fit. In fact, it might be useful to have a Model always do such a calculation just after the fit. The `Model.post_fit()` method allows exactly that: you can overwrite this otherwise empty method for any Model. It takes one argument: the `ModelResult` instance just after the actual fit has run (and before `Model.fit()` returns) and can be used to add Parameters or do other post-fit processing.

The following example script shows two different methods for calculating a centroid value for two peaks, either by doing the calculation directly after the fit with the `result.uvars` or by capturing this in a `Model.post_fit()` method that would be run for all instances of that model. It also demonstrates that taking correlations between Parameters into account when performing calculations can have a noticeable influence on the resulting uncertainties.

```
# <examples/doc_parameters_uvars.py>
import numpy as np

from lmfit.models import ExponentialModel, GaussianModel

dat = np.loadtxt('NIST_Gauss2.dat')
x = dat[:, 1]
y = dat[:, 0]

bkg = ExponentialModel(prefix='exp_')
gauss1 = GaussianModel(prefix='g1_')
gauss2 = GaussianModel(prefix='g2_')

mod = gauss1 + gauss2 + bkg

pars = mod.make_params(g1_center=105, g1_amplitude=4000, g1_sigma={'value': 15, 'min': 0}
                       ,
                       g2_center=155, g2_amplitude=2000, g2_sigma={'value': 15, 'min': 0}
                       ,
                       exp_amplitude=100, exp_decay=100)

out = mod.fit(y, pars, x=x)

print(out.fit_report(min_correl=0.5))

# Area and Centroid of the combined peaks
# option 1: from output uncertainties values

uvar_g1amp = out.uvars['g1_amplitude']
uvar_g2amp = out.uvars['g2_amplitude']
```

(continues on next page)

(continued from previous page)

```

uvar_g1cen = out.uvars['g1_center']
uvar_g2cen = out.uvars['g2_center']

print("### Peak Area: ")
print(f"Peak1: {out.params['g1_amplitude'].value:.5f}+/-{out.params['g1_amplitude'].\
→stderr:.5f}")
print(f"Peak2: {out.params['g2_amplitude'].value:.5f}+/-{out.params['g2_amplitude'].\
→stderr:.5f}")
print(f"Sum : {(uvar_g1amp + uvar_g2amp):.5f}")

area_quadrature = np.sqrt(out.params['g1_amplitude'].stderr**2 + out.params['g2_amplitude']\
→.stderr**2)

print(f"Correlation of peak1 and peak2 areas: {out.params['g1_amplitude'].\
→correl['g2_amplitude']:.5f}")
print(f"Stderr Quadrature sum for peak1 and peak2 area2: {area_quadrature:.5f}")
print(f"Stderr of peak1 + peak2 area, with correlation: {(uvar_g1amp+uvar_g2amp).std_\
→dev:.5f}")

print("### Peak Centroid: ")

centroid = (uvar_g1amp * uvar_g1cen + uvar_g2amp * uvar_g2cen) / (uvar_g1amp + uvar_\
→g2amp)

print(f"Peak Centroid: {centroid:.5f}")

# option 2: define corresponding variables after fit

def post_fit(result):
    "example post fit function"
    result.params.add('post_area', expr='g1_amplitude + g2_amplitude')
    result.params.add('post_centroid', expr='(g1_amplitude*g1_center+ g2_amplitude*g2_\
→center)/(g1_amplitude + g2_amplitude)')

mod.post_fit = post_fit

out = mod.fit(y, pars, x=x)
print(out.fit_report(min_correl=0.5))

# option 3: define corresponding variables before fit
pars.add('area', expr='g1_amplitude + g2_amplitude')
pars.add('centroid', expr='(g1_amplitude*g1_center+ g2_amplitude*g2_center)/(g1_\
→amplitude + g2_amplitude)')

out = mod.fit(y, pars, x=x)
print(out.fit_report(min_correl=0.5))
# <end examples/doc_parameters_uvars.py>

```

```

[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
  ↪ prefix='exp_'))
[[Fit Statistics]]
# fitting method = leastsq
# function evals = 46
# data points = 250
# variables = 8
chi-square = 1247.52821
reduced chi-square = 5.15507524
Akaike info crit = 417.864631
Bayesian info crit = 446.036318
R-squared = 0.99648654
[[Variables]]
g1_amplitude: 4257.77390 +/- 42.3837959 (1.00%) (init = 4000)
g1_center: 107.030957 +/- 0.15006877 (0.14%) (init = 105)
g1_sigma: 16.6725785 +/- 0.16048211 (0.96%) (init = 15)
g2_amplitude: 2493.41716 +/- 36.1697062 (1.45%) (init = 2000)
g2_center: 153.270104 +/- 0.19466774 (0.13%) (init = 155)
g2_sigma: 13.8069453 +/- 0.18680153 (1.35%) (init = 15)
exp_amplitude: 99.0183277 +/- 0.53748650 (0.54%) (init = 100)
exp_decay: 90.9508838 +/- 1.10310767 (1.21%) (init = 100)
g1_fwhm: 39.2609214 +/- 0.37790648 (0.96%) == '2.3548200*g1_sigma'
g1_height: 101.880229 +/- 0.59217051 (0.58%) == '0.3989423*g1_amplitude/max(1e-
  ↪ 15, g1_sigma)'
g2_fwhm: 32.5128709 +/- 0.43988398 (1.35%) == '2.3548200*g2_sigma'
g2_height: 72.0455939 +/- 0.61721985 (0.86%) == '0.3989423*g2_amplitude/max(1e-
  ↪ 15, g2_sigma)'
[[Correlations]] (unreported correlations are < 0.500)
C(g1_amplitude, g1_sigma) = +0.8243
C(g2_amplitude, g2_sigma) = +0.8154
C(exp_amplitude, exp_decay) = -0.6946
C(g1_sigma, g2_center) = +0.6842
C(g1_center, g2_amplitude) = -0.6689
C(g1_center, g2_sigma) = -0.6520
C(g1_amplitude, g2_center) = +0.6477
C(g1_center, g2_center) = +0.6205
C(g1_center, g1_sigma) = +0.5075
C(g1_amplitude, exp_decay) = -0.5074
### Peak Area:
Peak1: 4257.77390+/-42.38380
Peak2: 2493.41716+/-36.16971
Sum : 6751.19106+/-46.50802
Correlation of peak1 and peak2 areas: -0.30712
Stderr Quadrature sum for peak1 and peak2 area2: 55.71924
Stderr of peak1 + peak2 area, with correlation: 46.50802
### Peak Centroid:
Peak Centroid: 124.10846+/-0.14074
[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
  ↪ prefix='exp_'))
[[Fit Statistics]]
# fitting method = leastsq

```

(continues on next page)

(continued from previous page)

```

# function evals      = 46
# data points        = 250
# variables          = 8
chi-square           = 1247.52821
reduced chi-square   = 5.15507524
Akaike info crit    = 417.864631
Bayesian info crit  = 446.036318
R-squared             = 0.99648654
[[Variables]]
g1_amplitude:      4257.77390 +/- 42.3837959 (1.00%) (init = 4000)
g1_center:          107.030957 +/- 0.15006877 (0.14%) (init = 105)
g1_sigma:           16.6725785 +/- 0.16048211 (0.96%) (init = 15)
g2_amplitude:       2493.41716 +/- 36.1697062 (1.45%) (init = 2000)
g2_center:          153.270104 +/- 0.19466774 (0.13%) (init = 155)
g2_sigma:           13.8069453 +/- 0.18680153 (1.35%) (init = 15)
exp_amplitude:     99.0183277 +/- 0.53748650 (0.54%) (init = 100)
exp_decay:          90.9508838 +/- 1.10310767 (1.21%) (init = 100)
g1_fwhm:            39.2609214 +/- 0.37790648 (0.96%) == '2.3548200*g1_sigma'
g1_height:          101.880229 +/- 0.59217051 (0.58%) == '0.3989423*g1_amplitude/max(1e-
-15, g1_sigma)'
g2_fwhm:            32.5128709 +/- 0.43988398 (1.35%) == '2.3548200*g2_sigma'
g2_height:          72.0455939 +/- 0.61721985 (0.86%) == '0.3989423*g2_amplitude/max(1e-
-15, g2_sigma)'
post_area:          6751.19106 +/- 46.5080249 (0.69%) == 'g1_amplitude + g2_amplitude'
post_centroid:      124.108459 +/- 0.14074482 (0.11%) == '(g1_amplitude*g1_center+ g2_
-amplitude*g2_center)/(g1_amplitude + g2_amplitude)'
[[Correlations]] (unreported correlations are < 0.500)
C(g1_amplitude, g1_sigma) = +0.8243
C(g2_amplitude, g2_sigma) = +0.8154
C(exp_amplitude, exp_decay) = -0.6946
C(g1_sigma, g2_center) = +0.6842
C(g1_center, g2_amplitude) = -0.6689
C(g1_center, g2_sigma) = -0.6520
C(g1_amplitude, g2_center) = +0.6477
C(g1_center, g2_center) = +0.6205
C(g1_center, g1_sigma) = +0.5075
C(g1_amplitude, exp_decay) = -0.5074
[[Model]]
((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
- prefix='exp_'))
[[Fit Statistics]]
# fitting method = leastsq
# function evals = 46
# data points   = 250
# variables     = 8
chi-square       = 1247.52821
reduced chi-square = 5.15507524
Akaike info crit = 417.864631
Bayesian info crit = 446.036318
R-squared         = 0.99648654
[[Variables]]
g1_amplitude: 4257.77390 +/- 42.3837959 (1.00%) (init = 4000)

```

(continues on next page)

(continued from previous page)

```

g1_center:      107.030957 +/- 0.15006877 (0.14%) (init = 105)
g1_sigma:       16.6725785 +/- 0.16048211 (0.96%) (init = 15)
g2_amplitude:   2493.41716 +/- 36.1697062 (1.45%) (init = 2000)
g2_center:      153.270104 +/- 0.19466774 (0.13%) (init = 155)
g2_sigma:        13.8069453 +/- 0.18680153 (1.35%) (init = 15)
exp_amplitude:  99.0183277 +/- 0.53748650 (0.54%) (init = 100)
exp_decay:      90.9508838 +/- 1.10310767 (1.21%) (init = 100)
g1_fwhm:        39.2609214 +/- 0.37790648 (0.96%) == '2.3548200*g1_sigma'
g1_height:      101.880229 +/- 0.59217051 (0.58%) == '0.3989423*g1_amplitude/max(1e-
-15, g1_sigma)'
g2_fwhm:        32.5128709 +/- 0.43988398 (1.35%) == '2.3548200*g2_sigma'
g2_height:      72.0455939 +/- 0.61721985 (0.86%) == '0.3989423*g2_amplitude/max(1e-
-15, g2_sigma)'
area:           6751.19106 +/- 46.5080249 (0.69%) == 'g1_amplitude + g2_amplitude'
centroid:       124.108459 +/- 0.14074482 (0.11%) == '(g1_amplitude*g1_center+ g2_
-amplitude*g2_center)/(g1_amplitude + g2_amplitude)'
post_area:      6751.19106 +/- 46.5080249 (0.69%) == 'g1_amplitude + g2_amplitude'
post_centroid:  124.108459 +/- 0.14074482 (0.11%) == '(g1_amplitude*g1_center+ g2_
-amplitude*g2_center)/(g1_amplitude + g2_amplitude)'

[[Correlations]] (unreported correlations are < 0.500)
C(g1_amplitude, g1_sigma) = +0.8243
C(g2_amplitude, g2_sigma) = +0.8154
C(exp_amplitude, exp_decay) = -0.6946
C(g1_sigma, g2_center) = +0.6842
C(g1_center, g2_amplitude) = -0.6689
C(g1_center, g2_sigma) = -0.6520
C(g1_amplitude, g2_center) = +0.6477
C(g1_center, g2_center) = +0.6205
C(g1_center, g1_sigma) = +0.5075
C(g1_amplitude, exp_decay) = -0.5074

```

Note that the `Model.post_fit()` does not need to be limited to this use case of adding derived Parameters.

7.3.5 Saving and Loading ModelResults

As with saving models (see section [Saving and Loading Models](#)), it is sometimes desirable to save a `ModelResult`, either for later use or to organize and compare different fit results. Lmfit provides a `save_modelresult()` function that will save a `ModelResult` to a file. There is also a companion `load_modelresult()` function that can read this file and reconstruct a `ModelResult` from it.

As discussed in section [Saving and Loading Models](#), there are challenges to saving model functions that may make it difficult to restore a saved a `ModelResult` in a way that can be used to perform a fit. Use of the optional `funcdefs` argument is generally the most reliable way to ensure that a loaded `ModelResult` can be used to evaluate the model function or redo the fit.

`save_modelresult(modelresult, fname)`

Save a ModelResult to a file.

Parameters

- `modelresult (ModelResult)` – ModelResult to be saved.
- `fname (str)` – Name of file for saved ModelResult.

load_modelresult(*fname*, *funcdefs*=None)

Load a saved ModelResult from a file.

Parameters

- **fname** (*str*) – Name of file containing saved ModelResult.
- **funcdefs** (*dict*, *optional*) – Dictionary of custom function names and definitions.

Returns

ModelResult object loaded from file.

Return type

ModelResult

An example of saving a *ModelResult* is:

```
# <examples/doc_model_savemodelresult.py>
import numpy as np

from lmfit.model import save_modelresult
from lmfit.models import GaussianModel

data = np.loadtxt('model1d_gauss.dat')
x = data[:, 0]
y = data[:, 1]

gmodel = GaussianModel()
result = gmodel.fit(y, x=x, amplitude=5, center=5, sigma=1)

save_modelresult(result, 'gauss_modelresult.sav')

print(result.fit_report())
# <end examples/doc_model_savemodelresult.py>
```

To load that later, one might do:

```
# <examples/doc_model_loadmodelresult.py>
import os
import sys

import matplotlib.pyplot as plt
import numpy as np

from lmfit.model import load_modelresult

if not os.path.exists('gauss_modelresult.sav'):
    os.system(f'{sys.executable} doc_model_savemodelresult.py')

data = np.loadtxt('model1d_gauss.dat')
x = data[:, 0]
y = data[:, 1]

result = load_modelresult('gauss_modelresult.sav')
print(result.fit_report())
```

(continues on next page)

(continued from previous page)

```
plt.plot(x, y, 'o')
plt.plot(x, result.best_fit, '-')
plt.show()
# <end examples/doc_model_loadmodelresult.py>
```

7.4 Composite Models : adding (or multiplying) Models

One of the more interesting features of the `Model` class is that Models can be added together or combined with basic algebraic operations (add, subtract, multiply, and divide) to give a composite model. The composite model will have parameters from each of the component models, with all parameters being available to influence the whole model. This ability to combine models will become even more useful in the next chapter, when pre-built subclasses of `Model` are discussed. For now, we'll consider a simple example, and build a model of a Gaussian plus a line, as to model a peak with a background. For such a simple problem, we could just build a model that included both components:

```
def gaussian_plus_line(x, amp, cen, wid, slope, intercept):
    """line + 1-d gaussian"""

    gauss = (amp / (sqrt(2*pi) * wid)) * exp(-(x-cen)**2 / (2*wid**2))
    line = slope*x + intercept
    return gauss + line
```

and use that with:

```
mod = Model(gaussian_plus_line)
```

But we already had a function for a gaussian function, and maybe we'll discover that a linear background isn't sufficient which would mean the model function would have to be changed.

Instead, lmfit allows models to be combined into a `CompositeModel`. As an alternative to including a linear background in our model function, we could define a linear function:

```
def line(x, slope, intercept):
    """a line"""

    return slope*x + intercept
```

and build a composite model with just:

```
mod = Model(gaussian) + Model(line)
```

This model has parameters for both component models, and can be used as:

```
# <examples/doc_model_two_components.py>
import matplotlib.pyplot as plt
from numpy import exp, loadtxt, pi, sqrt

from lmfit import Model

data = loadtxt('model1d_gauss.dat')
x = data[:, 0]
y = data[:, 1] + 0.25*x - 1.0
```

(continues on next page)

(continued from previous page)

```

def gaussian(x, amp, cen, wid):
    """1-d gaussian: gaussian(x, amp, cen, wid)"""
    return (amp / (sqrt(2*pi) * wid)) * exp(-(x-cen)**2 / (2*wid**2))

def line(x, slope, intercept):
    """a line"""
    return slope*x + intercept

mod = Model(gaussian) + Model(line)
pars = mod.make_params(amp=5, cen=5, wid={'value': 1, 'min': 0},
                      slope=0, intercept=1)

result = mod.fit(y, pars, x=x)
print(result.fit_report())

plt.plot(x, y, 'o')
plt.plot(x, result.init_fit, '--', label='initial fit')
plt.plot(x, result.best_fit, ' ', label='best fit')
plt.legend()
plt.show()
# <end examples/doc_model_two_components.py>

```

which prints out the results:

```

[[Model]]
(Model(gaussian) + Model(line))
[[Fit Statistics]]
# fitting method    = leastsq
# function evals   = 55
# data points      = 101
# variables        = 5
chi-square         = 2.57855517
reduced chi-square = 0.02685995
Akaike info crit   = -360.457020
Bayesian info crit = -347.381417
R-squared          = 0.99194643
[[Variables]]
amp:      8.45930976 +/- 0.12414531 (1.47%) (init = 5)
cen:      5.65547889 +/- 0.00917673 (0.16%) (init = 5)
wid:      0.67545513 +/- 0.00991697 (1.47%) (init = 1)
slope:    0.26484403 +/- 0.00574892 (2.17%) (init = 0)
intercept: -0.96860189 +/- 0.03352202 (3.46%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(slope, intercept) = -0.7954
C(amp, wid)        = +0.6664
C(amp, intercept)  = -0.2216
C(amp, slope)       = -0.1692
C(cen, slope)       = -0.1618
C(wid, intercept)  = -0.1477


```

(continues on next page)

(continued from previous page)

```
C(cen, intercept) = +0.1287
C(wid, slope) = -0.1127
```

and shows the plot on the left.



On the left, data is shown in blue dots, the total fit is shown in solid green line, and the initial fit is shown as a orange dashed line. The figure on the right shows again the data in blue dots, the Gaussian component as a orange dashed line and the linear component as a green dashed line. It is created using the following code:

```
comps = result.eval_components()
plt.plot(x, y, 'o')
plt.plot(x, comps['gaussian'], '--', label='Gaussian component')
plt.plot(x, comps['line'], '--', label='Line component')
```

The components were generated after the fit using the `ModelResult.eval_components()` method of the `result`, which returns a dictionary of the components, using keys of the model name (or `prefix` if that is set). This will use the parameter values in `result.params` and the independent variables (`x`) used during the fit. Note that while the `ModelResult` held in `result` does store the best parameters and the best estimate of the model in `result.best_fit`, the original model and parameters in `pars` are left unaltered.

You can apply this composite model to other data sets, or evaluate the model at other values of `x`. You may want to do this to give a finer or coarser spacing of data point, or to extrapolate the model outside the fitting range. This can be done with:

```
xwide = linspace(-5, 25, 3001)
predicted = mod.eval(result.params, x=xwide)
```

In this example, the argument names for the model functions do not overlap. If they had, the `prefix` argument to `Model` would have allowed us to identify which parameter went with which component model. As we will see in the next chapter, using composite models with the built-in models provides a simple way to build up complex models.

```
class CompositeModel(left, right, op[, **kws])
```

Combine two models (`left` and `right`) with binary operator (`op`).

Normally, one does not have to explicitly create a `CompositeModel`, but can use normal Python operators `+`, `-`, `*`, and `/` to combine components as in:

```
>>> mod = Model(fcn1) + Model(fcn2) * Model(fcn3)
```

Parameters

- **left** (`Model`) – Left-hand model.
- **right** (`Model`) – Right-hand model.
- **op** (`callable binary operator`) – Operator to combine *left* and *right* models.
- ****kws** (*optional*) – Additional keywords are passed to `Model` when creating this new model.

Notes

The two models can use different independent variables.

Note that when using built-in Python binary operators, a `CompositeModel` will automatically be constructed for you. That is, doing:

```
mod = Model(fcn1) + Model(fcn2) * Model(fcn3)
```

will create a `CompositeModel`. Here, `left` will be `Model(fcn1)`, `op` will be `operator.add()`, and `right` will be another `CompositeModel` that has a `left` attribute of `Model(fcn2)`, an `op` of `operator.mul()`, and a `right` of `Model(fcn3)`.

To use a binary operator other than `+`, `-`, `*`, or `/` you can explicitly create a `CompositeModel` with the appropriate binary operator. For example, to convolve two models, you could define a simple convolution function, perhaps as:

```
import numpy as np

def convolve(dat, kernel):
    """simple convolution of two arrays"""
    npts = min(len(dat), len(kernel))
    pad = np.ones(npts)
    tmp = np.concatenate((pad*dat[0], dat, pad*dat[-1]))
    out = np.convolve(tmp, kernel, mode='valid')
    noff = int(len(out) - npts) / 2
    return (out[noff:])[npts:]
```

which extends the data in both directions so that the convolving kernel function gives a valid result over the data range. Because this function takes two array arguments and returns an array, it can be used as the binary operator. A full script using this technique is here:

```
# <examples/doc_model_composite.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit import CompositeModel, Model
from lmfit.lineshapes import gaussian, step

# create data from broadened step
x = np.linspace(0, 10, 201)
y = step(x, amplitude=12.5, center=4.5, sigma=0.88, form='erf')
np.random.seed(0)
y = y + np.random.normal(scale=0.35, size=x.size)
```

(continues on next page)

(continued from previous page)

```

def jump(x, mid):
    """Heaviside step function."""
    o = np.zeros(x.size)
    imid = max(np.where(x <= mid)[0])
    o[imid:] = 1.0
    return o

def convolve(arr, kernel):
    """Simple convolution of two arrays."""
    npts = min(arr.size, kernel.size)
    pad = np.ones(npts)
    tmp = np.concatenate((pad*arr[0], arr, pad*arr[-1]))
    out = np.convolve(tmp, kernel, mode='valid')
    noff = int((len(out) - npts) / 2)
    return out[noff:noff+npts]

# create Composite Model using the custom convolution operator
mod = CompositeModel(Model(jump), Model(gaussian), convolve)

# create parameters for model. Note that 'mid' and 'center' will be highly
# correlated. Since 'mid' is used as an integer index, it will be very
# hard to fit, so we fix its value
pars = mod.make_params(amplitude=dict(value=1, min=0),
                       center=3.5,
                       sigma=dict(value=1.5, min=0),
                       mid=dict(value=4, vary=False))

# fit this model to data array y
result = mod.fit(y, params=pars, x=x)

print(result.fit_report())

# generate components
comps = result.eval_components(x=x)

# plot results
fig, axes = plt.subplots(1, 2, figsize=(12.8, 4.8))

axes[0].plot(x, y, 'bo')
axes[0].plot(x, result.init_fit, 'k--', label='initial fit')
axes[0].plot(x, result.best_fit, 'r-', label='best fit')
axes[0].legend()

axes[1].plot(x, y, 'bo')
axes[1].plot(x, 10*comps['jump'], 'k--', label='Jump component')
axes[1].plot(x, 10*comps['gaussian'], 'r-', label='Gaussian component')
axes[1].legend()


plt.show()
# <end examples/doc_model_composite.py>

```

which prints out the results:

```
[[Model]]
  (Model(jump) <function convolve at 0x17db29f80> Model(gaussian))
[[Fit Statistics]]
  # fitting method    = leastsq
  # function evals   = 33
  # data points      = 201
  # variables        = 3
  chi-square         = 24.7562335
  reduced chi-square = 0.12503148
  Akaike info crit  = -414.939746
  Bayesian info crit = -405.029832
  R-squared          = 0.99632577
[[Variables]]
  mid:      4 (fixed)
  amplitude: 0.62508458 +/- 0.00189732 (0.30%) (init = 1)
  center:   5.50853669 +/- 0.00973231 (0.18%) (init = 3.5)
  sigma:    0.59576097 +/- 0.01348579 (2.26%) (init = 1.5)
[[Correlations]] (unreported correlations are < 0.100)
  C(amplitude, center) = +0.3292
  C(amplitude, sigma)  = +0.2680
```

and shows the plots:



Using composite models with built-in or custom operators allows you to build complex models from testable sub-components.

BUILT-IN FITTING MODELS IN THE MODELS MODULE

Lmfit provides several built-in fitting models in the `models` module. These pre-defined models each subclass from the `Model` class of the previous chapter and wrap relatively well-known functional forms, such as Gaussian, Lorentzian, and Exponential that are used in a wide range of scientific domains. In fact, all the models are based on simple, plain Python functions defined in the `lineshapes` module. In addition to wrapping a function into a `Model`, these models also provide a `guess()` method that is intended to give a reasonable set of starting values from a data array that closely approximates the data to be fit.

As shown in the previous chapter, a key feature of the `Model` class is that models can easily be combined to give a composite `CompositeModel`. Thus, while some of the models listed here may seem pretty trivial (notably, `ConstantModel` and `LinearModel`), the main point of having these is to be able to use them in composite models. For example, a Lorentzian plus a linear background might be represented as:

```
from lmfit.models import LinearModel, LorentzianModel

peak = LorentzianModel()
background = LinearModel()
model = peak + background
```

Almost all the models listed below are one-dimensional, with an independent variable named `x`. Many of these models represent a function with a distinct peak, and so share common features. To maintain uniformity, common parameter names are used whenever possible. Thus, most models have a parameter called `amplitude` that represents the overall intensity (or area of) a peak or function and a `sigma` parameter that gives a characteristic width.

After a list of built-in models, a few examples of their use are given.

8.1 Peak-like models

There are many peak-like models available. These include `GaussianModel`, `LorentzianModel`, `VoigtModel`, `PseudoVoigtModel`, and some less commonly used variations. Most of these models are *unit-normalized* and share the same parameter names so that you can easily switch between models and interpret the results. The `amplitude` parameter is the multiplicative factor for the unit-normalized peak lineshape, and so will represent the strength of that peak or the area under that curve. The `center` parameter will be the centroid `x` value. The `sigma` parameter is the characteristic width of the peak, with many functions using $(x - \mu)/\sigma$ where μ is the centroid value. Most of these peak functions will have two additional parameters derived from and constrained by the other parameters. The first of these is `fwhm` which will hold the estimated “Full Width at Half Max” for the peak, which is often easier to compare between different models than `sigma`. The second of these is `height` which will contain the maximum value of the peak, typically the value at $x = \mu$. Finally, each of these models has a `guess()` method that uses data to make a fairly crude but usually sufficient guess for the value of `amplitude`, `center`, and `sigma`, and sets a lower bound of 0 on the value of `sigma`.

8.1.1 GaussianModel

```
class GaussianModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on a Gaussian or normal distribution lineshape.

The model has three Parameters: *amplitude*, *center*, and *sigma*. In addition, parameters *fwhm* and *height* are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma) = \frac{A}{\sigma\sqrt{2\pi}} e^{[-(x-\mu)^2/2\sigma^2]}$$

where the parameter *amplitude* corresponds to *A*, *center* to μ , and *sigma* to σ . The full width at half maximum is $2\sigma\sqrt{2\ln 2}$, approximately 2.3548σ .

For more information, see: https://en.wikipedia.org/wiki/Normal_distribution

Parameters

- **independent_vars** (*list of str*, optional) – Arguments to the model function that are independent variables default is ['x']).
- **prefix** (*str*, optional) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** ({'raise', 'propagate', 'omit'}, optional) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (optional) – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:

- 'raise' : raise a *ValueError* (default)
- 'propagate' : do nothing
- 'omit' : drop missing data

8.1.2 LorentzianModel

```
class LorentzianModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on a Lorentzian or Cauchy-Lorentz distribution function.

The model has three Parameters: *amplitude*, *center*, and *sigma*. In addition, parameters *fwhm* and *height* are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma) = \frac{A}{\pi} \left[\frac{\sigma}{(x - \mu)^2 + \sigma^2} \right]$$

where the parameter *amplitude* corresponds to *A*, *center* to μ , and *sigma* to σ . The full width at half maximum is 2σ .

For more information, see: https://en.wikipedia.org/wiki/Cauchy_distribution

Parameters

- **independent_vars** (*list of str*, optional) – Arguments to the model function that are independent variables default is ['x']).

- **prefix**(*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy**({'raise', 'propagate', 'omit'}, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs**(*optional*) – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:

- 'raise' : raise a *ValueError* (default)
- 'propagate' : do nothing
- 'omit' : drop missing data

8.1.3 SplitLorentzianModel

```
class SplitLorentzianModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on a Lorentzian or Cauchy-Lorentz distribution function.

The model has four parameters: *amplitude*, *center*, *sigma*, and *sigma_r*. In addition, parameters *FWHM* and *height* are included as constraints to report full width at half maximum and maximum peak height, respectively.

'Split' means that the width of the distribution is different between left and right slopes.

$$f(x; A, \mu, \sigma, \sigma_r) = \frac{2A}{\pi(\sigma + \sigma_r)} \left[\frac{\sigma^2}{(x - \mu)^2 + \sigma^2} * H(\mu - x) + \frac{\sigma_r^2}{(x - \mu)^2 + \sigma_r^2} * H(x - \mu) \right]$$

where the parameter *amplitude* corresponds to *A*, *center* to *μ*, *sigma* to *σ*, *sigma_l* to *σ_l*, and *H(x)* is a Heaviside step function:

$$H(x) = 0|x < 0, 1|x \geq 0$$

The full width at half maximum is $\sigma_l + \sigma_r$. Just as with the Lorentzian model, integral of this function from *-inf* to *+inf* equals to *amplitude*.

For more information, see: https://en.wikipedia.org/wiki/Cauchy_distribution

Parameters

- **independent_vars**(*list* of *str*, *optional*) – Arguments to the model function that are independent variables default is ['x']).
- **prefix**(*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy**({'raise', 'propagate', 'omit'}, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs**(*optional*) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘raise’ : raise a `ValueError` (default)
- ‘propagate’ : do nothing
- ‘omit’ : drop missing data

8.1.4 VoigtModel

```
class VoigtModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on a Voigt distribution function.

The model has four Parameters: *amplitude*, *center*, *sigma*, and *gamma*. By default, *gamma* is constrained to have a value equal to *sigma*, though it can be varied independently. In addition, parameters *fwlm* and *height* are included as constraints to report full width at half maximum and maximum peak height, respectively. The definition for the Voigt function used here is:

$$f(x; A, \mu, \sigma, \gamma) = \frac{A \operatorname{Re}[w(z)]}{\sigma \sqrt{2\pi}}$$

where

$$\begin{aligned} z &= \frac{x - \mu + i\gamma}{\sigma \sqrt{2}} \\ w(z) &= e^{-z^2} \operatorname{erfc}(-iz) \end{aligned}$$

and `erfc()` is the complementary error function. As above, *amplitude* corresponds to *A*, *center* to μ , and *sigma* to σ . The parameter *gamma* corresponds to γ . If *gamma* is kept at the default value (constrained to *sigma*), the full width at half maximum is approximately 3.6013σ .

For more information, see: https://en.wikipedia.org/wiki/Voigt_profile

Parameters

- `independent_vars` (list of str, optional) – Arguments to the model function that are independent variables default is ['x']).
- `prefix(str, optional)` – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- `nan_policy` ({‘raise’, ‘propagate’, ‘omit’}, optional) – How to handle NaN and missing values in data. See Notes below.
- `**kwargs` (optional) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘raise’ : raise a `ValueError` (default)
- ‘propagate’ : do nothing
- ‘omit’ : drop missing data

8.1.5 PseudoVoigtModel

```
class PseudoVoigtModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on a pseudo-Voigt distribution function.

This is a weighted sum of a Gaussian and Lorentzian distribution function that share values for *amplitude* (A), *center* (μ), and full width at half maximum *fwlm* (and so has constrained values of *sigma* (σ) and *height* (maximum peak height). The parameter *fraction* (α) controls the relative weight of the Gaussian and Lorentzian components, giving the full definition of:

$$f(x; A, \mu, \sigma, \alpha) = \frac{(1 - \alpha)A}{\sigma_g \sqrt{2\pi}} e^{[-(x - \mu)^2 / 2\sigma_g^2]} + \frac{\alpha A}{\pi} \left[\frac{\sigma}{(x - \mu)^2 + \sigma^2} \right]$$

where $\sigma_g = \sigma / \sqrt{2 \ln 2}$ so that the full width at half maximum of each component and of the sum is 2σ . The `guess()` function always sets the starting value for *fraction* at 0.5.

For more information, see: https://en.wikipedia.org/wiki/Voigt_profile#Pseudo-Voigt_Approximation

Parameters

- **independent_vars** (`list` of `str`, optional) – Arguments to the model function that are independent variables default is `['x']`.
- **prefix** (`str`, optional) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`{'raise', 'propagate', 'omit'}`, optional) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (optional) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- `'raise'` : raise a `ValueError` (default)
- `'propagate'` : do nothing
- `'omit'` : drop missing data

8.1.6 MoffatModel

```
class MoffatModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on the Moffat distribution function.

The model has four Parameters: *amplitude* (A), *center* (μ), a width parameter *sigma* (σ), and an exponent *beta* (β). In addition, parameters *fwlm* and *height* are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma, \beta) = A \left[\left(\frac{x - \mu}{\sigma} \right)^2 + 1 \right]^{-\beta}$$

the full width at half maximum is $2\sigma\sqrt{2^{1/\beta} - 1}$. The `guess()` function always sets the starting value for *beta* to 1.

Note that for ($\beta = 1$) the Moffat has a Lorentzian shape. For more information, see: https://en.wikipedia.org/wiki/Moffat_distribution

Parameters

- **independent_vars** (list of str, optional) – Arguments to the model function that are independent variables default is ['x']).
- **prefix(str, optional)** – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** ({'raise', 'propagate', 'omit'}, optional) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs (optional)** – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:

- 'raise' : raise a *ValueError* (default)
- 'propagate' : do nothing
- 'omit' : drop missing data

8.1.7 Pearson4Model

```
class Pearson4Model(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on a Pearson IV distribution.

The model has five parameters: *amplitude* (*A*), *center* (μ), *sigma* (σ), *expon* (*m*) and *skew* (ν). In addition, parameters *frwhm*, *height* and *position* are included as constraints to report estimates for the approximate full width at half maximum (20% error), the peak height, and the peak position (the position of the maximal function value), respectively. The *frwhm* value has an error of about 20% in the parameter range *expon*: (0.5, 1000], *skew*: [-1000, 1000].

$$f(x; A, \mu, \sigma, m, \nu) = A \frac{\left| \frac{\Gamma(m+i\frac{\nu}{2})}{\Gamma(m)} \right|^2}{\sigma \beta(m - \frac{1}{2}, \frac{1}{2})} \left[1 + \frac{(x - \mu)^2}{\sigma^2} \right]^{-m} \exp \left(-\nu \arctan \left(\frac{x - \mu}{\sigma} \right) \right)$$

where β is the beta function (see `scipy.special.beta`). The `guess()` function always gives a starting value of 1.5 for *expon*, and 0 for *skew*.

For more information, see: https://en.wikipedia.org/wiki/Pearson_distribution#The_Pearson_type_IV_distribution

Parameters

- **independent_vars** (list of str, optional) – Arguments to the model function that are independent variables default is ['x']).
- **prefix(str, optional)** – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** ({'raise', 'propagate', 'omit'}, optional) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs (optional)** – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘`raise`’ : raise a `ValueError` (default)
- ‘`propagate`’ : do nothing
- ‘`omit`’ : drop missing data

8.1.8 Pearson7Model

```
class Pearson7Model(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on a Pearson VII distribution.

The model has four parameters: *amplitude* (A), *center* (μ), *sigma* (σ), and *exponent* (m). In addition, parameters *fwhm* and *height* are included as constraints to report estimates for the full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma, m) = \frac{A}{\sigma \beta(m - \frac{1}{2}, \frac{1}{2})} \left[1 + \frac{(x - \mu)^2}{\sigma^2} \right]^{-m}$$

where β is the beta function (see `scipy.special.beta`). The `guess()` function always gives a starting value for *exponent* of 1.5. In addition, parameters *fwhm* and *height* are included as constraints to report full width at half maximum and maximum peak height, respectively.

For more information, see: https://en.wikipedia.org/wiki/Pearson_distribution#The_Pearson_type_VII_distribution

Parameters

- `independent_vars` (`list` of `str`, *optional*) – Arguments to the model function that are independent variables default is `['x']`.
- `prefix(str, optional)` – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- `nan_policy({'raise', 'propagate', 'omit'}, optional)` – How to handle NaN and missing values in data. See Notes below.
- `**kwargs (optional)` – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘`raise`’ : raise a `ValueError` (default)
- ‘`propagate`’ : do nothing
- ‘`omit`’ : drop missing data

8.1.9 StudentsTModel

```
class StudentsTModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on a Student's t-distribution function.

The model has three Parameters: *amplitude* (A), *center* (μ), and *sigma* (σ). In addition, parameters *fw hm* and *height* are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma) = \frac{A \Gamma(\frac{\sigma+1}{2})}{\sqrt{\sigma \pi} \Gamma(\frac{\sigma}{2})} \left[1 + \frac{(x - \mu)^2}{\sigma^2} \right]^{-\frac{\sigma+1}{2}}$$

where $\Gamma(x)$ is the gamma function.

For more information, see: https://en.wikipedia.org/wiki/Student%27s_t-distribution

Parameters

- **independent_vars** (`list of str`, optional) – Arguments to the model function that are independent variables default is ['x']).
- **prefix** (`str, optional`) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`{'raise', 'propagate', 'omit'}`, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:

- 'raise' : raise a *ValueError* (default)
- 'propagate' : do nothing
- 'omit' : drop missing data

8.1.10 BreitWignerModel

```
class BreitWignerModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on a Breit-Wigner-Fano function.

The model has four Parameters: *amplitude* (A), *center* (μ), *sigma* (σ), and *q* (q).

$$f(x; A, \mu, \sigma, q) = \frac{A(q\sigma/2 + x - \mu)^2}{(\sigma/2)^2 + (x - \mu)^2}$$

For more information, see: https://en.wikipedia.org/wiki/Fano_resonance

Parameters

- **independent_vars** (`list of str`, optional) – Arguments to the model function that are independent variables default is ['x']).
- **prefix** (`str, optional`) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`{'raise', 'propagate', 'omit'}`, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘`raise`’ : raise a `ValueError` (default)
- ‘`propagate`’ : do nothing
- ‘`omit`’ : drop missing data

8.1.11 LognormalModel

class LognormalModel(`independent_vars=['x']`, `prefix=''`, `nan_policy='raise'`, `kwargs`)**

A model based on the Log-normal distribution function.

The model has three Parameters *amplitude* (A), *center* (μ), and *sigma* (σ). In addition, parameters *FWHM* and *height* are included as constraints to report estimates of full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma) = \frac{A}{\sigma\sqrt{2\pi}} \frac{e^{-(\ln(x)-\mu)^2/2\sigma^2}}{x}$$

For more information, see: <https://en.wikipedia.org/wiki/Lognormal>

Parameters

- `independent_vars` (`list` of `str`, optional) – Arguments to the model function that are independent variables default is `['x']`.
- `prefix(str, optional)` – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- `nan_policy` (`{'raise', 'propagate', 'omit'}`, `optional`) – How to handle NaN and missing values in data. See Notes below.
- `**kwargs` (`optional`) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘`raise`’ : raise a `ValueError` (default)
- ‘`propagate`’ : do nothing
- ‘`omit`’ : drop missing data

8.1.12 DampedOscillatorModel

class DampedOscillatorModel(`independent_vars=['x']`, `prefix=''`, `nan_policy='raise'`, `kwargs`)**

A model based on the Damped Harmonic Oscillator Amplitude.

The model has three Parameters: *amplitude* (A), *center* (μ), and *sigma* (σ). In addition, the parameter *height* is included as a constraint to report the maximum peak height.

$$f(x; A, \mu, \sigma) = \frac{A}{\sqrt{[1 - (x/\mu)^2]^2 + (2\sigma x/\mu)^2}}$$

For more information, see: https://en.wikipedia.org/wiki/Harmonic_oscillator#Amplitude_part

Parameters

- **independent_vars** (`list` of `str`, optional) – Arguments to the model function that are independent variables default is `['x']`).
- **prefix** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`{'raise', 'propagate', 'omit'}`, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- `'raise'` : raise a `ValueError` (default)
- `'propagate'` : do nothing
- `'omit'` : drop missing data

8.1.13 DampedHarmonicOscillatorModel

```
class DampedHarmonicOscillatorModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on a variation of the Damped Harmonic Oscillator.

The model follows the definition given in DAVE/PAN (see: <https://www.ncnr.nist.gov/dave>) and has four Parameters: *amplitude* (A), *center* (μ), *sigma* (σ), and *gamma* (γ). In addition, parameters *FWHM* and *height* are included as constraints to report estimates for full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma, \gamma) = \frac{A\sigma}{\pi[1 - \exp(-x/\gamma)]} \left[\frac{1}{(x - \mu)^2 + \sigma^2} - \frac{1}{(x + \mu)^2 + \sigma^2} \right]$$

where $\gamma = kT$, k is the Boltzmann constant in $eV K^{-1}$, and T is the temperature in K .

For more information, see: https://en.wikipedia.org/wiki/Harmonic_oscillator

Parameters

- **independent_vars** (`list` of `str`, optional) – Arguments to the model function that are independent variables default is `['x']`).
- **prefix** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`{'raise', 'propagate', 'omit'}`, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘`raise`’ : raise a `ValueError` (default)
- ‘`propagate`’ : do nothing
- ‘`omit`’ : drop missing data

8.1.14 ExponentialGaussianModel

class ExponentialGaussianModel(`independent_vars=['x']`, `prefix=''`, `nan_policy='raise'`, `kwargs`)**

A model of an Exponentially modified Gaussian distribution.

The model has four Parameters: *amplitude* (A), *center* (μ), *sigma* (σ), and *gamma* (γ).

$$f(x; A, \mu, \sigma, \gamma) = \frac{A\gamma}{2} \exp[\gamma(\mu - x + \gamma\sigma^2/2)] \operatorname{erfc}\left(\frac{\mu + \gamma\sigma^2 - x}{\sqrt{2}\sigma}\right)$$

where `erfc()` is the complementary error function.

For more information, see: https://en.wikipedia.org/wiki/Exponentially_modified_Gaussian_distribution

Parameters

- **independent_vars** (`list of str`, optional) – Arguments to the model function that are independent variables default is `['x']`.
- **prefix(`str`, optional)** – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`{'raise', 'propagate', 'omit'}`, `optional`) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (`optional`) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘`raise`’ : raise a `ValueError` (default)
- ‘`propagate`’ : do nothing
- ‘`omit`’ : drop missing data

8.1.15 SkewedGaussianModel

class SkewedGaussianModel(`independent_vars=['x']`, `prefix=''`, `nan_policy='raise'`, `kwargs`)**

A skewed Gaussian model, using a skewed normal distribution.

The model has four Parameters: *amplitude* (A), *center* (μ), *sigma* (σ), and *gamma* (γ).

$$f(x; A, \mu, \sigma, \gamma) = \frac{A}{\sigma\sqrt{2\pi}} e^{[-(x-\mu)^2/2\sigma^2]} \left\{ 1 + \operatorname{erf}\left[\frac{\gamma(x-\mu)}{\sigma\sqrt{2}}\right] \right\}$$

where `erf()` is the error function.

For more information, see: https://en.wikipedia.org/wiki/Skew_normal_distribution

Parameters

- **independent_vars** (`list` of `str`, optional) – Arguments to the model function that are independent variables default is `['x']`.
- **prefix** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`{'raise', 'propagate', 'omit'}`, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- `'raise'` : raise a `ValueError` (default)
- `'propagate'` : do nothing
- `'omit'` : drop missing data

8.1.16 SkewedVoigtModel

class SkewedVoigtModel(`independent_vars=['x']`, `prefix=''`, `nan_policy='raise'`, `**kwargs`)

A skewed Voigt model, modified using a skewed normal distribution.

The model has five Parameters *amplitude* (A), *center* (μ), *sigma* (σ), and *gamma* (γ), as usual for a Voigt distribution, and adds a new Parameter *skew*.

$$f(x; A, \mu, \sigma, \gamma, \text{skew}) = \text{Voigt}(x; A, \mu, \sigma, \gamma) \left\{ 1 + \text{erf} \left[\frac{\text{skew}(x - \mu)}{\sigma \sqrt{2}} \right] \right\}$$

where `erf()` is the error function.

For more information, see: https://en.wikipedia.org/wiki/Skew_normal_distribution

Parameters

- **independent_vars** (`list` of `str`, optional) – Arguments to the model function that are independent variables default is `['x']`.
- **prefix** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`{'raise', 'propagate', 'omit'}`, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘`raise`’ : raise a `ValueError` (default)
- ‘`propagate`’ : do nothing
- ‘`omit`’ : drop missing data

8.1.17 ThermalDistributionModel

```
class ThermalDistributionModel(independent_vars=['x', 'form'], prefix='', nan_policy='raise', form='bose',
**kwargs)
```

Return a thermal distribution function.

Variable `form` defines the kind of distribution as below with three Parameters: *amplitude* (`A`), *center* (`x0`), and `kt` (`kt`). The following distributions are available:

- ‘`bose`’ : Bose-Einstein distribution (default)
- ‘`maxwell`’ : Maxwell-Boltzmann distribution
- ‘`fermi`’ : Fermi-Dirac distribution

The functional forms are defined as:

$$\begin{aligned} f(x; A, x_0, kt, \text{form} = \text{'bose'}) &= \frac{1}{A \exp(\frac{x-x_0}{kt}) - 1} \\ f(x; A, x_0, kt, \text{form} = \text{'maxwell'}) &= \frac{1}{A \exp(\frac{x-x_0}{kt})} \\ f(x; A, x_0, kt, \text{form} = \text{'fermi'}) &= \frac{1}{A \exp(\frac{x-x_0}{kt}) + 1} \end{aligned}$$

Notes

- `kt` should be defined in the same units as `x` ($k_B = 8.617 \times 10^{-5}$ eV/K).
- set `kt < 0` to implement the energy loss convention common in scattering research.

For more information, see: <http://hyperphysics.phy-astr.gsu.edu/hbase/quantum/disfcn.html>

Parameters

- `independent_vars` (`list` of `str`, *optional*) – Arguments to the model function that are independent variables default is `['x']`.
- `prefix` (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- `nan_policy` (`{'raise', 'propagate', 'omit'}`, *optional*) – How to handle NaN and missing values in data. See Notes below.
- `**kwargs` (*optional*) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘`raise`’ : raise a `ValueError` (default)
- ‘`propagate`’ : do nothing
- ‘`omit`’ : drop missing data

8.1.18 DoniachModel

```
class DoniachModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model of an Doniach Sunjic asymmetric lineshape.

This model is used in photo-emission and has four Parameters: *amplitude* (A), *center* (μ), *sigma* (σ), and *gamma* (γ). In addition, parameter *height* is included as a constraint to report maximum peak height.

$$f(x; A, \mu, \sigma, \gamma) = \frac{A}{\sigma^{1-\gamma}} \frac{\cos[\pi\gamma/2 + (1-\gamma)\arctan((x-\mu)/\sigma)]}{[1+(x-\mu)/\sigma]^{(1-\gamma)/2}}$$

For more information, see: https://www.casaxps.com/help_manual/line_shapes.htm

Parameters

- `independent_vars` (`list` of `str`, optional) – Arguments to the model function that are independent variables default is ['x']).
- `prefix(str, optional)` – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- `nan_policy` ({‘`raise`’, ‘`propagate`’, ‘`omit`’}, optional) – How to handle NaN and missing values in data. See Notes below.
- `**kwargs` (optional) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘`raise`’ : raise a `ValueError` (default)
- ‘`propagate`’ : do nothing
- ‘`omit`’ : drop missing data

8.2 Linear and Polynomial Models

These models correspond to polynomials of some degree. Of course, lmfit is a very inefficient way to do linear regression (see `numpy.polyfit` or `scipy.stats.linregress`), but these models may be useful as one of many components of a composite model. The SplineModel below corresponds to a cubic spline.

8.2.1 ConstantModel

```
class ConstantModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

Constant model, with a single Parameter: c .

Note that this is ‘constant’ in the sense of having no dependence on the independent variable x , not in the sense of being non-varying. To be clear, c will be a Parameter that will be varied in the fit (by default, of course).

Parameters

- **independent_vars** (*list of str, optional*) – Arguments to the model function that are independent variables default is ['x']).
- **prefix** (*str, optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (*{'raise', 'propagate', 'omit'}*, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘raise’ : raise a *ValueError* (default)
- ‘propagate’ : do nothing
- ‘omit’ : drop missing data

8.2.2 LinearModel

```
class LinearModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

Linear model, with two Parameters: *intercept* and *slope*.

Defined as:

$$f(x; m, b) = mx + b$$

with *slope* for m and *intercept* for b .

Parameters

- **independent_vars** (*list of str, optional*) – Arguments to the model function that are independent variables default is ['x']).
- **prefix** (*str, optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (*{'raise', 'propagate', 'omit'}*, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘raise’ : raise a `ValueError` (default)
- ‘propagate’ : do nothing
- ‘omit’ : drop missing data

8.2.3 QuadraticModel

```
class QuadraticModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A quadratic model, with three Parameters: a , b , and c .

Defined as:

$$f(x; a, b, c) = ax^2 + bx + c$$

Parameters

- `independent_vars` (`list` of `str`, optional) – Arguments to the model function that are independent variables default is ['x']).
- `prefix(str, optional)` – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- `nan_policy` ({‘raise’, ‘propagate’, ‘omit’}, optional) – How to handle NaN and missing values in data. See Notes below.
- `**kwargs (optional)` – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘raise’ : raise a `ValueError` (default)
- ‘propagate’ : do nothing
- ‘omit’ : drop missing data

8.2.4 PolynomialModel

```
class PolynomialModel(degree=7, independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A polynomial model with up to 7 Parameters, specified by `degree`.

$$f(x; c_0, c_1, \dots, c_7) = \sum_{i=0,7} c_i x^i$$

with parameters $c0, c1, \dots, c7$. The supplied `degree` will specify how many of these are actual variable parameters. This uses `numpy.polyval` for its calculation of the polynomial.

Parameters

- `independent_vars` (`list` of `str`, optional) – Arguments to the model function that are independent variables default is ['x']).

- **prefix** (*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (*{'raise', 'propagate', 'omit'}*, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘raise’ : raise a *ValueError* (default)
- ‘propagate’ : do nothing
- ‘omit’ : drop missing data

8.2.5 SplineModel

```
class SplineModel(xknots, independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A 1-D cubic spline model with a variable number of *knots* and parameters $s0, s1, \dots, sN$, for N knots.

The user must supply a list or ndarray *xknots*: the *x* values for the ‘knots’ which control the flexibility of the spline function.

The parameters $s0, \dots, sN$ (where N is the size of *xknots*) will correspond to the *y* values for the spline knots at the $x=xknots$ positions where the highest order derivative will be discontinuous. The resulting curve will not necessarily pass through these knot points, but for finely-spaced knots, the spline parameter values will be very close to the *y* values of the resulting curve.

The maximum number of knots supported is 300.

Using the *guess()* method to initialize parameter values is highly recommended.

Parameters

- **xknots** (*list* of floats or *ndarray*, required) – x-values of knots for spline.
- **independent_vars** (*list* of *str*, *optional*) – Arguments to the model function that are independent variables default is ['x']).
- **prefix** (*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (*{'raise', 'propagate', 'omit'}*, *optional*) – How to handle NaN and missing values in data. See Notes below.

Notes

1. There must be at least 4 knot points, and not more than 300.

2. **nan_policy** sets what to do when a NaN or missing value is seen in

the data. Should be one of:

- ‘raise’ : raise a *ValueError* (default)
- ‘propagate’ : do nothing
- ‘omit’ : drop missing data

8.3 Periodic Models

These models correspond to periodic functions.

8.3.1 SineModel

```
class SineModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on a sinusoidal lineshape.

The model has three Parameters: *amplitude*, *frequency*, and *shift*.

$$f(x; A, \phi, f) = A \sin(fx + \phi)$$

where the parameter *amplitude* corresponds to *A*, *frequency* to *f*, and *shift* to ϕ . All are constrained to be non-negative, and *shift* additionally to be smaller than 2π .

Parameters

- **independent_vars** (`list of str`, optional) – Arguments to the model function that are independent variables default is `['x']`.
- **prefix** (`str, optional`) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`{'raise', 'propagate', 'omit'}`, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:

- `'raise'` : raise a *ValueError* (default)
- `'propagate'` : do nothing
- `'omit'` : drop missing data

8.4 Step-like models

Two models represent step-like functions, and share many characteristics.

8.4.1 StepModel

```
class StepModel(independent_vars=['x', 'form'], prefix='', nan_policy='raise', form='linear', **kwargs)
```

A model based on a Step function.

The model has three Parameters: *amplitude* (*A*), *center* (μ), and *sigma* (σ).

There are four choices for *form*:

- `'linear'` (default)
- `'atan'` or `'arctan'` for an arc-tangent function

- ‘erf’ for an error function
- ‘logistic’ for a logistic function (for more information, see: https://en.wikipedia.org/wiki/Logistic_function)

The step function starts with a value 0 and ends with a value of $\text{sign}(\sigma)A$ rising or falling to $A/2$ at μ , with σ setting the characteristic width of the step. The functional forms are defined as:

$$\begin{aligned}
 f(x; A, \mu, \sigma, \text{form} = \text{'linear'}) &= A \min [1, \max (0, \alpha + 1/2)] \\
 f(x; A, \mu, \sigma, \text{form} = \text{'arctan'}) &= A[1/2 + \arctan(\alpha)/\pi] \\
 f(x; A, \mu, \sigma, \text{form} = \text{'erf'}) &= A[1 + \text{erf}(\alpha)]/2 \\
 f(x; A, \mu, \sigma, \text{form} = \text{'logistic'}) &= A \left[1 - \frac{1}{1 + e^{\alpha}} \right]
 \end{aligned}$$

where $\alpha = (x - \mu)/\sigma$.

Note that $\sigma > 0$ gives a rising step, while $\sigma < 0$ gives a falling step.

Parameters

- **independent_vars** (list of str, optional) – Arguments to the model function that are independent variables default is ['x']).
- **prefix(str, optional)** – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** ({'raise', 'propagate', 'omit'}, optional) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs (optional)** – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘raise’ : raise a *ValueError* (default)
- ‘propagate’ : do nothing
- ‘omit’ : drop missing data

8.4.2 RectangleModel

```
class RectangleModel(independent_vars=['x', 'form'], prefix='', nan_policy='raise', form='linear', **kwargs)
```

A model based on a Step-up and Step-down function.

The model has five Parameters: *amplitude* (A), *center1* (μ_1), *center2* (μ_2), *sigma1* (σ_1), and *sigma2* (σ_2).

There are four choices for *form*, which is used for both the Step up and the Step down:

- ‘linear’ (default)
- ‘atan’ or ‘arctan’ for an arc-tangent function
- ‘erf’ for an error function
- ‘logistic’ for a logistic function (for more information, see: https://en.wikipedia.org/wiki/Logistic_function)

The function starts with a value 0 and transitions to a value of A , taking the value $A/2$ at μ_1 , with σ_1 setting the characteristic width. The function then transitions again to the value $A/2$ at μ_2 , with σ_2 setting the characteristic width. The functional forms are defined as:

$$\begin{aligned}
 f(x; A, \mu, \sigma, \text{form} = \text{'linear'}) &= A \{ \min [1, \max (-1, \alpha_1)] + \min [1, \max (-1, \alpha_2)] \} / 2 \\
 f(x; A, \mu, \sigma, \text{form} = \text{'arctan'}) &= A [\arctan(\alpha_1) + \arctan(\alpha_2)] / \pi \\
 f(x; A, \mu, \sigma, \text{form} = \text{'erf'}) &= A [\text{erf}(\alpha_1) + \text{erf}(\alpha_2)] / 2 \\
 f(x; A, \mu, \sigma, \text{form} = \text{'logistic'}) &= A \left[1 - \frac{1}{1 + e^{\alpha_1}} - \frac{1}{1 + e^{\alpha_2}} \right]
 \end{aligned}$$

where $\alpha_1 = (x - \mu_1) / \sigma_1$ and $\alpha_2 = -(x - \mu_2) / \sigma_2$.

Note that, unlike a StepModel, $\sigma_1 > 0$ is enforced, giving a rising initial step, and $\sigma_2 > 0$ gives a falling final step.

Parameters

- **independent_vars** (`list` of `str`, optional) – Arguments to the model function that are independent variables default is `['x']`.
- **prefix** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`{'raise', 'propagate', 'omit'}`, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:

- `'raise'` : raise a `ValueError` (default)
- `'propagate'` : do nothing
- `'omit'` : drop missing data

8.5 Exponential and Power law models

8.5.1 ExponentialModel

```
class ExponentialModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on an exponential decay function.

The model has two Parameters: *amplitude* (A) and *decay* (τ) and is defined as:

$$f(x; A, \tau) = Ae^{-x/\tau}$$

For more information, see: https://en.wikipedia.org/wiki/Exponential_decay

Parameters

- **independent_vars** (`list` of `str`, optional) – Arguments to the model function that are independent variables default is `['x']`.
- **prefix** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.

- **nan_policy** (`{'raise', 'propagate', 'omit'}`, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘raise’ : raise a *ValueError* (default)
- ‘propagate’ : do nothing
- ‘omit’ : drop missing data

8.5.2 PowerLawModel

```
class PowerLawModel(independent_vars=['x'], prefix='', nan_policy='raise', **kwargs)
```

A model based on a Power Law.

The model has two Parameters: *amplitude* (*A*) and *exponent* (*k*) and is defined as:

$$f(x; A, k) = Ax^k$$

For more information, see: https://en.wikipedia.org/wiki/Power_law

Parameters

- **independent_vars** (`list` of `str`, *optional*) – Arguments to the model function that are independent variables default is `['x']`.
- **prefix** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`{'raise', 'propagate', 'omit'}`, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘raise’ : raise a *ValueError* (default)
- ‘propagate’ : do nothing
- ‘omit’ : drop missing data

8.6 Two dimensional Peak-like models

The one example of a two-dimensional peak is a two-dimensional Gaussian.

8.6.1 Gaussian2dModel

```
class Gaussian2dModel(independent_vars=['x', 'y'], prefix='', nan_policy='raise', **kwargs)
```

A model based on a two-dimensional Gaussian function.

The model has two independent variables x and y and five Parameters: *amplitude*, *centerx*, *sigmax*, *centery*, and *sigmay*. In addition, parameters *fhwmx*, *fhwmy*, and *height* are included as constraints to report the maximum peak height and the two full width at half maxima, respectively.

$$f(x, y; A, \mu_x, \sigma_x, \mu_y, \sigma_y) = Ag(x; A = 1, \mu_x, \sigma_x)g(y; A = 1, \mu_y, \sigma_y)$$

where subfunction $g(x; A, \mu, \sigma)$ is a Gaussian lineshape:

$$g(x; A, \mu, \sigma) = \frac{A}{\sigma\sqrt{2\pi}}e^{[-(x-\mu)^2/2\sigma^2]}.$$

Parameters

- **independent_vars** (*list of str, optional*) – Arguments to the model function that are independent variables default is ['x', 'y']).
- **prefix** (*str, optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (*{'raise', 'propagate', 'omit'}*, *optional*) – How to handle NaN and missing values in data. See Notes below.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:

- 'raise' : raise a *ValueError* (default)
- 'propagate' : do nothing
- 'omit' : drop missing data

8.7 User-defined Models

As shown in the previous chapter (*Modeling Data and Curve Fitting*), it is fairly straightforward to build fitting models from parametrized Python functions. The number of model classes listed so far in the present chapter should make it clear that this process is not too difficult. Still, it is sometimes desirable to build models from a user-supplied function. This may be especially true if model-building is built-in to some larger library or application for fitting in which the user may not be able to easily build and use a new model from Python code.

The *ExpressionModel* allows a model to be built from a user-supplied expression. This uses the *asteval* module also used for mathematical constraints as discussed in *Using Mathematical Constraints*.

8.7.1 ExpressionModel

```
class ExpressionModel(expr, independent_vars=None, init_script=None, nan_policy='raise', **kws)
```

ExpressionModel class.

Generate a model from user-supplied expression.

Parameters

- **expr** (*str*) – Mathematical expression for model.
- **independent_vars** (*list* of *str* or *None*, optional) – Variable names to use as independent variables.
- **init_script** (*str* or *None*, optional) – Initial script to run in asteval interpreter.
- **nan_policy** (*{'raise', 'propagate', 'omit'}*, optional) – How to handle NaN and missing values in data. See Notes below.
- ****kws** (optional) – Keyword arguments to pass to Model.

Notes

1. each instance of ExpressionModel will create and use its own version of an asteval interpreter.
2. *prefix* is **not supported** for ExpressionModel.
3. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - ‘raise’ : raise a *ValueError* (default)
 - ‘propagate’ : do nothing
 - ‘omit’ : drop missing data

Since the point of this model is that an arbitrary expression will be supplied, the determination of what are the parameter names for the model happens when the model is created. To do this, the expression is parsed, and all symbol names are found. Names that are already known (there are over 500 function and value names in the asteval namespace, including most Python built-ins, more than 200 functions inherited from NumPy, and more than 20 common lineshapes defined in the `lineshapes` module) are not converted to parameters. Unrecognized names are expected to be names of either parameters or independent variables. If `independent_vars` is the default value of `None`, and if the expression contains a variable named `x`, that will be used as the independent variable. Otherwise, `independent_vars` must be given.

For example, if one creates an `ExpressionModel` as:

```
from lmfit.models import ExpressionModel

mod = ExpressionModel('off + amp * exp(-x/x0) * sin(x*phase)')
```

The name `exp` will be recognized as the exponent function, so the model will be interpreted to have parameters named `off`, `amp`, `x0` and `phase`. In addition, `x` will be assumed to be the sole independent variable. In general, there is no obvious way to set default parameter values or parameter hints for bounds, so this will have to be handled explicitly.

To evaluate this model, you might do the following:

```
from numpy import exp, linspace, sin

x = linspace(0, 10, 501)
params = mod.make_params(off=0.25, amp=1.0, x0=2.0, phase=0.04)
y = mod.eval(params, x=x)
```

While many custom models can be built with a single line expression (especially since the names of the lineshapes like `gaussian`, `lorentzian` and so on, as well as many NumPy functions, are available), more complex models will inevitably require multiple line functions. You can include such Python code with the `init_script` argument. The text of this script is evaluated when the model is initialized (and before the actual expression is parsed), so that you can define functions to be used in your expression.

As a probably unphysical example, to make a model that is the derivative of a Gaussian function times the logarithm of a Lorentzian function you may could to define this in a script:

```
script = """
def mycurve(x, amp, cen, sig):
    loren = lorentzian(x, amplitude=amp, center=cen, sigma=sig)
    gauss = gaussian(x, amplitude=amp, center=cen, sigma=sig)
    return log(loren) * gradient(gauss) / gradient(x)
"""


```

and then use this with `ExpressionModel` as:

```
mod = ExpressionModel('mycurve(x, height, mid, wid)', init_script=script,
                      independent_vars=['x'])
```

As above, this will interpret the parameter names to be `height`, `mid`, and `wid`, and build a model that can be used to fit data.

8.8 Example 1: Fit Peak data to Gaussian, Lorentzian, and Voigt profiles

Here, we will fit data to three similar lineshapes, in order to decide which might be the better model. We will start with a Gaussian profile, as in the previous chapter, but use the built-in `GaussianModel` instead of writing one ourselves. This is a slightly different version from the one in previous example in that the parameter names are different, and have built-in default values. We will simply use:

```
from numpy import loadtxt

from lmfit.models import GaussianModel

data = loadtxt('test_peak.dat')
x = data[:, 0]
y = data[:, 1]

mod = GaussianModel()

pars = mod.guess(y, x=x)
out = mod.fit(y, pars, x=x)

print(out.fit_report(min_correl=0.25))
```

which prints out the results:

```
[[Model]]
    Model(gaussian)
[[Fit Statistics]]
```

(continues on next page)


(continued from previous page)

```

# fitting method      = leastsq
# function evals    = 25
# data points       = 401
# variables         = 3
chi-square          = 29.9943157
reduced chi-square = 0.07536260
Akaike info crit   = -1033.77437
Bayesian info crit = -1021.79248
R-squared           = 0.99045513
[[Variables]]
amplitude: 30.3135789 +/- 0.15712752 (0.52%) (init = 43.62238)
center:    9.24277046 +/- 0.00737497 (0.08%) (init = 9.25)
sigma:     1.23218496 +/- 0.00737506 (0.60%) (init = 1.35)
fwhm:      2.90157379 +/- 0.01736695 (0.60%) == '2.3548200*sigma'
height:    9.81457271 +/- 0.05087308 (0.52%) == '0.3989423*amplitude/max(1e-15,_
sigma)'
[[Correlations]] (unreported correlations are < 0.250)
C(amplitude, sigma) = +0.5774

```

We see a few interesting differences from the results of the previous chapter. First, the parameter names are longer. Second, there are `fwhm` and `height` parameters, to give the full-width-at-half-maximum and maximum peak height, respectively. And third, the automated initial guesses are pretty good. A plot of the fit:



shows a decent match to the data – the fit worked with no explicit setting of initial parameter values. Looking more closely, the fit is not perfect, especially in the tails of the peak, suggesting that a different peak shape, with longer tails, should be used. Perhaps a Lorentzian would be better? To do this, we simply replace `GaussianModel` with `LorentzianModel` to get a `LorentzianModel`:

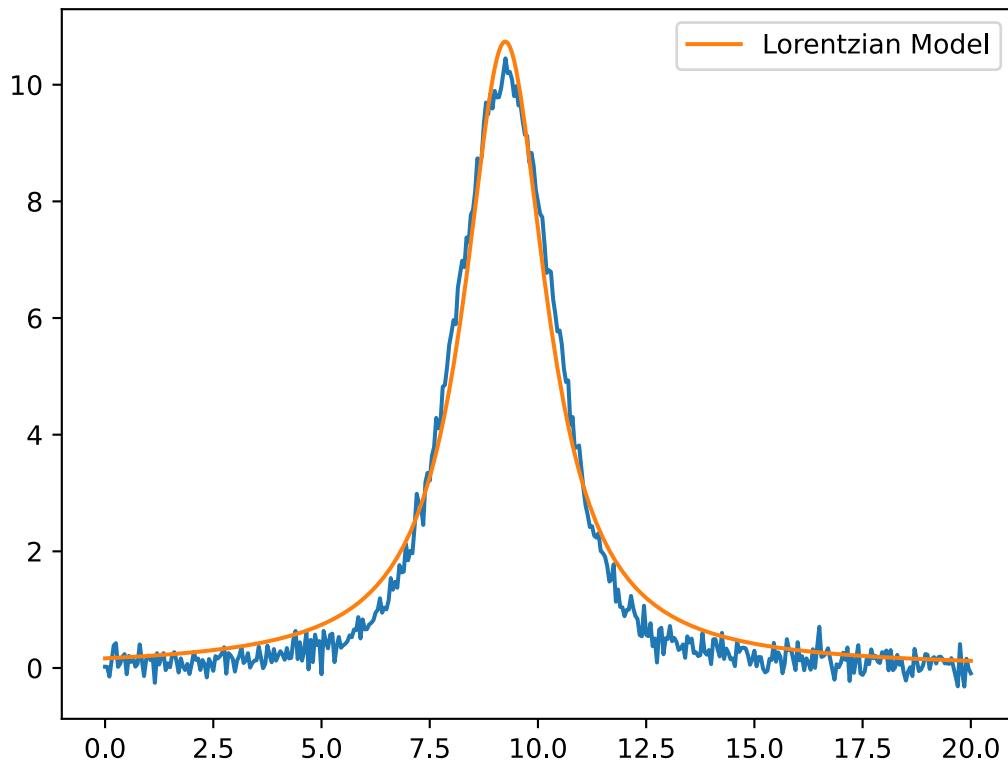
```
from lmfit.models import LorentzianModel

mod = LorentzianModel()
```

with the rest of the script as above. Perhaps predictably, the first thing we try gives results that are worse by comparing the fit statistics:

```
[[Model]]
    Model(lorentzian)
[[Fit Statistics]]
    # fitting method    = leastsq
    # function evals   = 25
    # data points      = 401
    # variables        = 3
    chi-square         = 53.7535387
    reduced chi-square = 0.13505914
    Akaike info crit   = -799.830322
    Bayesian info crit = -787.848438
    R-squared          = 0.98289441
[[Variables]]
    amplitude: 38.9726380 +/- 0.31386754 (0.81%) (init = 54.52798)
    center:    9.24439393 +/- 0.00927645 (0.10%) (init = 9.25)
    sigma:     1.15483177 +/- 0.01315708 (1.14%) (init = 1.35)
    fwhm:      2.30966354 +/- 0.02631416 (1.14%) == '2.0000000*sigma'
    height:    10.7421504 +/- 0.08634317 (0.80%) == '0.3183099*amplitude/max(1e-15,_
    sigma)'
[[Correlations]] (unreported correlations are < 0.250)
    C(amplitude, sigma) = +0.7087
```

and also by visual inspection of the fit to the data (figure below).



The tails are now too big, and the value for χ^2 almost doubled. A Voigt model does a better job. Using `VoigtModel`, this is as simple as using:

```
from lmfit.models import VoigtModel
mod = VoigtModel()
```

with all the rest of the script as above. This gives:

```
[[Model]]
    Model(voigt)
[[Fit Statistics]]
    # fitting method    = leastsq
    # function evals   = 25
    # data points     = 401
    # variables       = 3
    chi-square        = 14.5448627
    reduced chi-square = 0.03654488
    Akaike info crit  = -1324.00615
    Bayesian info crit = -1312.02427
    R-squared         = 0.99537150
[[Variables]]
    amplitude:  35.7553799 +/- 0.13861559 (0.39%) (init = 65.43358)
    center:     9.24411179 +/- 0.00505496 (0.05%) (init = 9.25)
    sigma:      0.73015485 +/- 0.00368473 (0.50%) (init = 0.8775)
    gamma:      0.73015485 +/- 0.00368473 (0.50%) == 'sigma'
    fwhm:       2.62949983 +/- 0.01326979 (0.50%) == '1.0692*gamma+sqrt(0.
    ↵8664*gamma**2+5.545083*sigma**2)'
```


(continues on next page)

(continued from previous page)

```

height: 10.2204068 +/- 0.03959933 (0.39%) == '(amplitude/(max(1e-15,
sigma*sqrt(2*pi)))*real(wofz((1j*gamma)/(max(1e-15, sigma*sqrt(2))))))'
[[Correlations]] (unreported correlations are < 0.250)
C(amplitude, sigma) = +0.6513
    
```

which has a much better value for χ^2 and the other goodness-of-fit measures, and an obviously better match to the data as seen in the figure below (left).



Fit to peak with Voigt model (left) and Voigt model with gamma varying independently of sigma (right).

Can we do better? The Voigt function has a γ parameter (`gamma`) that can be distinct from `sigma`. The default behavior used above constrains `gamma` to have exactly the same value as `sigma`. If we allow these to vary separately, does the fit improve? To do this, we have to change the `gamma` parameter from a constrained expression and give it a starting value using something like:

```

mod = VoigtModel()
pars = mod.guess(y, x=x)
pars['gamma'].set(value=0.7, vary=True, expr='')
    
```

which gives:

```

[[Model]]
    Model(voigt)
[[Fit Statistics]]
    # fitting method    = leastsq
    # function evals   = 26
    # data points     = 401
    # variables       = 4
    chi-square        = 10.9301767
    reduced chi-square = 0.02753193
    Akaike info crit  = -1436.57602
    Bayesian info crit = -1420.60017
    R-squared          = 0.99652177
[[Variables]]
    amplitude: 34.1914716 +/- 0.17946974 (0.52%) (init = 65.43358)
    center:    9.24374846 +/- 0.00441904 (0.05%) (init = 9.25)
    sigma:     0.89518951 +/- 0.01415479 (1.58%) (init = 0.8775)
    
```

(continues on next page)

(continued from previous page)

```

gamma:      0.52540156 +/- 0.01857994 (3.54%) (init = 0.7)
fwhm:      2.72573678 +/- 0.01363994 (0.50%) == '1.0692*gamma+sqrt(0.
-8664*gamma**2+5.545083*sigma**2)'
height:    10.0872197 +/- 0.03482126 (0.35%) == '(amplitude/(max(1e-15,
-sigma*sqrt(2*pi)))*real(wofz((1j*gamma)/(max(1e-15, sigma*sqrt(2))))))'
[[Correlations]] (unreported correlations are < 0.250)
C(sigma, gamma) = -0.9285
C(amplitude, gamma) = +0.8210
C(amplitude, sigma) = -0.6512

```

and the fit shown on the right above.

Comparing the two fits with the Voigt function, we see that χ^2 is definitely improved with a separately varying `gamma` parameter. In addition, the two values for `gamma` and `sigma` differ significantly – well outside the estimated uncertainties. More compelling, reduced χ^2 is improved even though a fourth variable has been added to the fit. In the simplest statistical sense, this suggests that `gamma` is a significant variable in the model. In addition, we can use both the Akaike or Bayesian Information Criteria (see [Akaike and Bayesian Information Criteria](#)) to assess how likely the model with variable `gamma` is to explain the data than the model with `gamma` fixed to the value of `sigma`. According to theory, $\exp(-(AIC1 - AIC0)/2)$ gives the probability that a model with `AIC1` is more likely than a model with `AIC0`. For the two models here, with `AIC` values of -1436 and -1324 (Note: if we had more carefully set the value for `weights` based on the noise in the data, these values might be positive, but the difference would be roughly the same), this says that the model with `gamma` fixed to `sigma` has a probability less than 5.e-25 of being the better model.

8.9 Example 2: Fit data to a Composite Model with pre-defined models

Here, we repeat the point made at the end of the last chapter that instances of `Model` class can be added together to make a *composite model*. By using the large number of built-in models available, it is therefore very simple to build models that contain multiple peaks and various backgrounds. An example of a simple fit to a noisy step function plus a constant:

```

# <examples/doc_builtinmodels_stepmodel.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit.models import LinearModel, StepModel

x = np.linspace(0, 10, 201)
y = np.ones_like(x)
y[:48] = 0.0
y[48:77] = np.arange(77-48)/(77.0-48)
np.random.seed(0)
y = 110.2 * (y + 9e-3*np.random.randn(x.size)) + 12.0 + 2.22*x

step_mod = StepModel(form='erf', prefix='step_')
line_mod = LinearModel(prefix='line_')

pars = line_mod.make_params(intercept=y.min(), slope=0)
pars += step_mod.guess(y, x=x, center=2.5)

```

(continues on next page)

(continued from previous page)

```

mod = step_mod + line_mod
out = mod.fit(y, pars, x=x)

print(out.fit_report())

plt.plot(x, y)
plt.plot(x, out.init_fit, '--', label='initial fit')
plt.plot(x, out.best_fit, ' ', label='best fit')
plt.legend()
plt.show()
# <end examples/doc_builtinmodels_stepmodel.py>

```


After constructing step-like data, we first create a `StepModel` telling it to use the `erf` form (see details above), and a `ConstantModel`. We set initial values, in one case using the data and `guess()` method for the initial step function parameters, and `make_params()` arguments for the linear component. After making a composite model, we run `fit()` and report the results, which gives:

```

[[Model]]
(Model(step, prefix='step_', form='erf') + Model(linear, prefix='line_'))
[[Fit Statistics]]
# fitting method    = leastsq
# function evals   = 55
# data points      = 201
# variables        = 5
chi-square         = 593.709621
reduced chi-square = 3.02913072
Akaike info crit   = 227.700173
Bayesian info crit = 244.216698
R-squared          = 0.99897798
[[Variables]]
line_slope:      1.87162051 +/- 0.09318576 (4.98%) (init = 0)
line_intercept: 12.0964556 +/- 0.27606000 (2.28%) (init = 11.58574)
step_amplitude: 112.858605 +/- 0.65391558 (0.58%) (init = 134.7378)
step_center:    3.13494787 +/- 0.00516601 (0.16%) (init = 2.5)
step_sigma:     0.67393526 +/- 0.01091153 (1.62%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(line_slope, step_amplitude) = -0.8791
C(step_amplitude, step_sigma) = +0.5643
C(line_slope, step_sigma) = -0.4569
C(line_intercept, step_center) = +0.4269
C(line_slope, line_intercept) = -0.3093
C(line_slope, step_center) = -0.2338
C(line_intercept, step_sigma) = -0.1372
C(line_intercept, step_amplitude) = -0.1173
C(step_amplitude, step_center) = +0.1095

```

with a plot of



8.10 Example 3: Fitting Multiple Peaks – and using Prefixes

As shown above, many of the models have similar parameter names. For composite models, this could lead to a problem of having parameters for different parts of the model having the same name. To overcome this, each `Model` can have a `prefix` attribute (normally set to a blank string) that will be put at the beginning of each parameter name. To illustrate, we fit one of the classic datasets from the NIST StRD suite involving a decaying exponential and two Gaussians.

```
# <examples/doc_builtinmodels_nistgauss.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit.models import ExponentialModel, GaussianModel

dat = np.loadtxt('NIST_Gauss2.dat')
x = dat[:, 1]
y = dat[:, 0]

exp_mod = ExponentialModel(prefix='exp_')
pars = exp_mod.guess(y, x=x)

gauss1 = GaussianModel(prefix='g1_')
pars.update(gauss1.make_params(center=dict(value=105, min=75, max=125),
                               sigma=dict(value=15, min=0),
                               amplitude=dict(value=2000, min=0)))

gauss2 = GaussianModel(prefix='g2_')
```

(continues on next page)

(continued from previous page)

```

pars.update(gauss2.make_params(center=dict(value=155, min=125, max=175),
                                sigma=dict(value=15, min=0),
                                amplitude=dict(value=2000, min=0)))

mod = gauss1 + gauss2 + exp_mod

init = mod.eval(pars, x=x)
out = mod.fit(y, pars, x=x)

print(out.fit_report(correl_mode='table'))

fig, axes = plt.subplots(1, 2, figsize=(12.8, 4.8))
axes[0].plot(x, y)
axes[0].plot(x, init, '--', label='initial fit')
axes[0].plot(x, out.best_fit, '-', label='best fit')
axes[0].legend()

comps = out.eval_components(x=x)
axes[1].plot(x, y)
axes[1].plot(x, comps['g1_'], '--', label='Gaussian component 1')
axes[1].plot(x, comps['g2_'], '--', label='Gaussian component 2')
axes[1].plot(x, comps['exp_'], '--', label='Exponential component')
axes[1].legend()

plt.show()
# <end examples/doc_builtinmodels_nistgauss.py>

```

where we give a separate prefix to each model (they all have an `amplitude` parameter). The `prefix` values are attached transparently to the models.

Note that the calls to `make_param()` used the bare name, without the prefix. We could have used the prefixes, but because we used the individual model `gauss1` and `gauss2`, there was no need.

Note also in the example here that we explicitly set bounds on many of the parameter values.

The fit results printed out are:

```

[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
  ↪ prefix='exp_'))
[[Fit Statistics]]
  # fitting method    = leastsq
  # function evals   = 46
  # data points      = 250
  # variables        = 8
  chi-square         = 1247.52821
  reduced chi-square = 5.15507524
  Akaike info crit  = 417.864631
  Bayesian info crit = 446.036318
  R-squared          = 0.99648654
[[Variables]]
  exp_amplitude: 99.0183278 +/- 0.53748593 (0.54%) (init = 162.2102)
  exp_decay:     90.9508853 +/- 1.10310778 (1.21%) (init = 93.24905)
  g1_amplitude: 4257.77360 +/- 42.3836478 (1.00%) (init = 2000)

```

(continues on next page)

(continued from previous page)

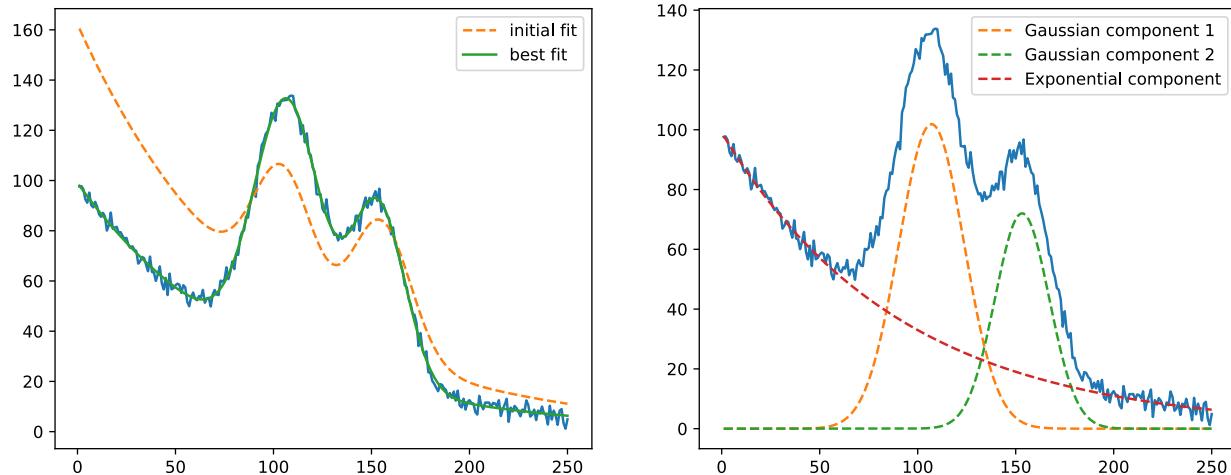
```

g1_center:      107.030956 +/- 0.15006851 (0.14%) (init = 105)
g1_sigma:       16.6725772 +/- 0.16048381 (0.96%) (init = 15)
g1_fwhm:        39.2609181 +/- 0.37791049 (0.96%) == '2.3548200*g1_sigma'
g1_height:      101.880230 +/- 0.59217173 (0.58%) == '0.3989423*g1_amplitude/max(1e-
˓→15, g1_sigma)'
g2_amplitude:   2493.41735 +/- 36.1697789 (1.45%) (init = 2000)
g2_center:      153.270102 +/- 0.19466802 (0.13%) (init = 155)
g2_sigma:        13.8069464 +/- 0.18679695 (1.35%) (init = 15)
g2_fwhm:        32.5128735 +/- 0.43987320 (1.35%) == '2.3548200*g2_sigma'
g2_height:      72.0455941 +/- 0.61722243 (0.86%) == '0.3989423*g2_amplitude/max(1e-
˓→15, g2_sigma)'

[[Correlations]] (unreported correlations are < 0.100)
C(g1_amplitude, g1_sigma)      = +0.8243
C(g2_amplitude, g2_sigma)      = +0.8154
C(exp_amplitude, exp_decay)   = -0.6946
C(g1_sigma, g2_center)         = +0.6842
C(g1_center, g2_amplitude)     = -0.6689
C(g1_center, g2_sigma)         = -0.6520
C(g1_amplitude, g2_center)     = +0.6477
C(g1_center, g2_center)        = +0.6205
C(g1_center, g1_sigma)         = +0.5075
C(exp_decay, g1_amplitude)     = -0.5074
C(g1_sigma, g2_amplitude)      = -0.4915
C(g2_center, g2_sigma)         = -0.4889
C(g1_sigma, g2_sigma)          = -0.4826
C(g2_amplitude, g2_center)     = -0.4763
C(exp_decay, g2_amplitude)     = -0.4270
C(g1_amplitude, g1_center)     = +0.4183
C(g1_amplitude, g2_sigma)      = -0.4010
C(g1_amplitude, g2_amplitude)   = -0.3071
C(exp_amplitude, g2_amplitude) = +0.2821
C(exp_decay, g1_sigma)          = -0.2520
C(exp_decay, g2_sigma)          = -0.2329
C(exp_amplitude, g2_sigma)      = +0.1714
C(exp_decay, g2_center)         = -0.1514
C(exp_amplitude, g1_amplitude)  = +0.1478
C(exp_decay, g1_center)         = +0.1055

```

We get a very good fit to this problem (described at the NIST site as of average difficulty, but the tests there are generally deliberately challenging) by applying reasonable initial guesses and putting modest but explicit bounds on the parameter values. The overall fit is shown on the left, with its individual components displayed on the right:



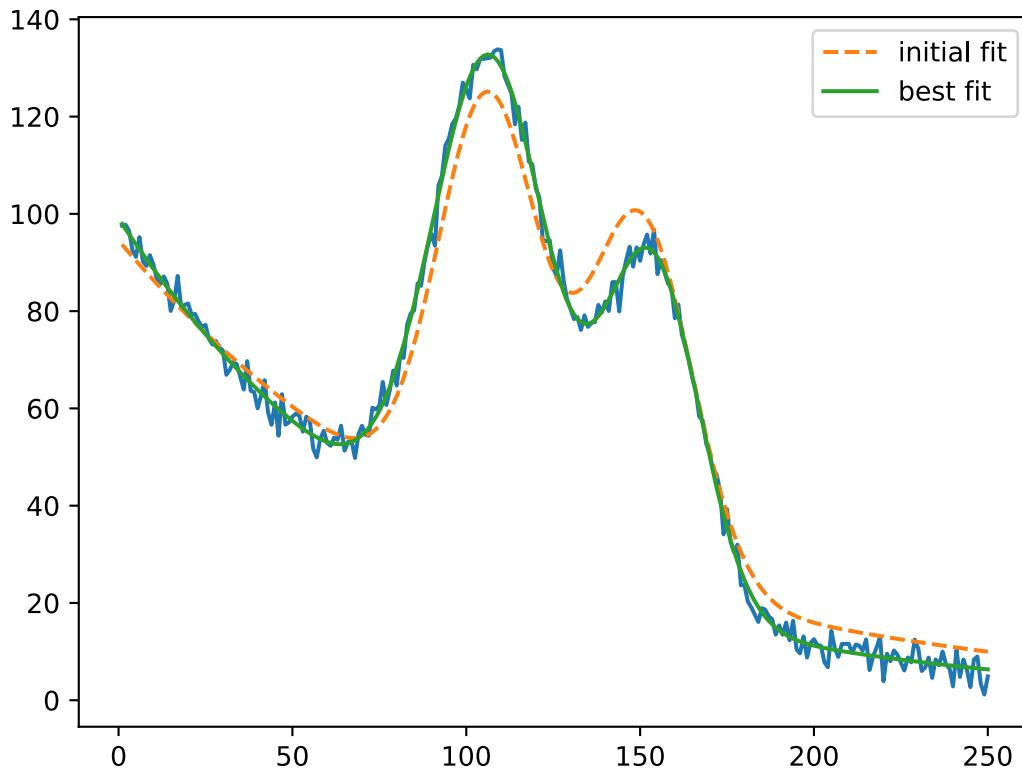
One final point on setting initial values. From looking at the data itself, we can see the two Gaussian peaks are reasonably well separated but do overlap. Furthermore, we can tell that the initial guess for the decaying exponential component was poorly estimated because we used the full data range. We can simplify the initial parameter values by using this, and by defining an `index_of()` function to limit the data range. That is, with:

```
def index_of(arrval, value):
    """Return index of array *at or below* value."""
    if value < min(arrval):
        return 0
    return max(np.where(arrval <= value)[0])

ix1 = index_of(x, 75)
ix2 = index_of(x, 135)
ix3 = index_of(x, 175)

exp_mod.guess(y[:ix1], x=x[:ix1])
gauss1.guess(y[ix1:ix2], x=x[ix1:ix2])
gauss2.guess(y[ix2:ix3], x=x[ix2:ix3])
```

we can get a better initial estimate (see below).



The fit converges to the same answer, giving to identical values (to the precision printed out in the report), but in fewer steps, and without any bounds on parameters at all:

```
[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
  ↪ prefix='exp_'))
[[Fit Statistics]]
  # fitting method    = leastsq
  # function evals   = 37
  # data points      = 250
  # variables        = 8
  chi-square         = 1247.52821
  reduced chi-square = 5.15507524
  Akaike info crit   = 417.864631
  Bayesian info crit = 446.036318
  R-squared          = 0.99648654
[[Variables]]
  exp_amplitude: 99.0183265 +/- 0.53748764 (0.54%) (init = 94.53724)
  exp_decay:     90.9508884 +/- 1.10310753 (1.21%) (init = 111.1985)
  g1_amplitude:  4257.77384 +/- 42.3839276 (1.00%) (init = 3189.648)
  g1_center:     107.030957 +/- 0.15006934 (0.14%) (init = 106.5)
  g1_sigma:       16.6725783 +/- 0.16048220 (0.96%) (init = 14.5)
  g1_fwhm:        39.2609209 +/- 0.37790669 (0.96%) == '2.3548200*g1_sigma'
  g1_height:     101.880228 +/- 0.59216965 (0.58%) == '0.3989423*g1_amplitude/max(1e-
  ↪ 15, g1_sigma)'
  g2_amplitude:  2493.41698 +/- 36.1699974 (1.45%) (init = 2818.337)
  g2_center:     153.270103 +/- 0.19466966 (0.13%) (init = 150)
  g2_sigma:       13.8069440 +/- 0.18680331 (1.35%) (init = 15)
```

(continues on next page)

(continued from previous page)

```

g2_fwhm:           32.5128679 +/- 0.43988818 (1.35%) == '2.3548200*g2_sigma'
g2_height:         72.0455954 +/- 0.61722288 (0.86%) == '0.3989423*g2_amplitude/max(1e-
                   -15, g2_sigma)'
[[Correlations]] (unreported correlations are < 0.100)
C(g1_amplitude, g1_sigma)      = +0.8243
C(g2_amplitude, g2_sigma)      = +0.8154
C(exp_amplitude, exp_decay)   = -0.6946
C(g1_sigma, g2_center)         = +0.6842
C(g1_center, g2_amplitude)     = -0.6689
C(g1_center, g2_sigma)         = -0.6521
C(g1_amplitude, g2_center)     = +0.6477
C(g1_center, g2_center)        = +0.6205
C(g1_center, g1_sigma)         = +0.5075
C(exp_decay, g1_amplitude)    = -0.5074
C(g1_sigma, g2_amplitude)      = -0.4914
C(g2_center, g2_sigma)         = -0.4890
C(g1_sigma, g2_sigma)          = -0.4826
C(g2_amplitude, g2_center)     = -0.4763
C(exp_decay, g2_amplitude)     = -0.4270
C(g1_amplitude, g1_center)     = +0.4183
C(g1_amplitude, g2_sigma)      = -0.4011
C(g1_amplitude, g2_amplitude)   = -0.3071
C(exp_amplitude, g2_amplitude) = +0.2821
C(exp_decay, g1_sigma)         = -0.2520
C(exp_decay, g2_sigma)         = -0.2329
C(exp_amplitude, g2_sigma)     = +0.1714
C(exp_decay, g2_center)        = -0.1514
C(exp_amplitude, g1_amplitude) = +0.1478
C(exp_decay, g1_center)        = +0.1055

```

This script is in the file `doc_builtinmodels_nistgauss2.py` in the examples folder, and the figure above shows an improved initial estimate of the data.

8.11 Example 4: Using a Spline Model

In the example above, the two peaks might represent the interesting part of the data, and the exponential decay could be viewed a “background” which might be due to other physical effects or part of some response of the instrumentation used to make the measurement. That is, the background might be well-understood to be modeled as an exponential decay, as in the example above and so easily included in the full analysis. As the results above show, there is some – but not huge – correlation of the parameters between the peak amplitudes and the decay of the exponential function. That means that it is helpful to include all of those components in a single fit, as the uncertainties in the peak amplitudes (which would be interpreted as “line strength” or “area”) will reflect some of the uncertainty in how well we modeled the background.

Sometimes a background is more complex or at least has a less obvious functional form. In these cases, it can be useful to use a *spline* to model part of the curve. Just for completeness, a spline is a piecewise continuous polynomial function (typically made of cubic polynomials) that has a series of x values known as “knots” at which the highest order derivative is allowed to be discontinuous. By adding more knots, the spline function has more flexibility to follow a particular function.

As an example (see the example file “`doc_builtinmodels_splinemodel.py`”), we start with data with a single peak and a background that is hard to characterize clearly as a simple decay, oscillatory structure.

```

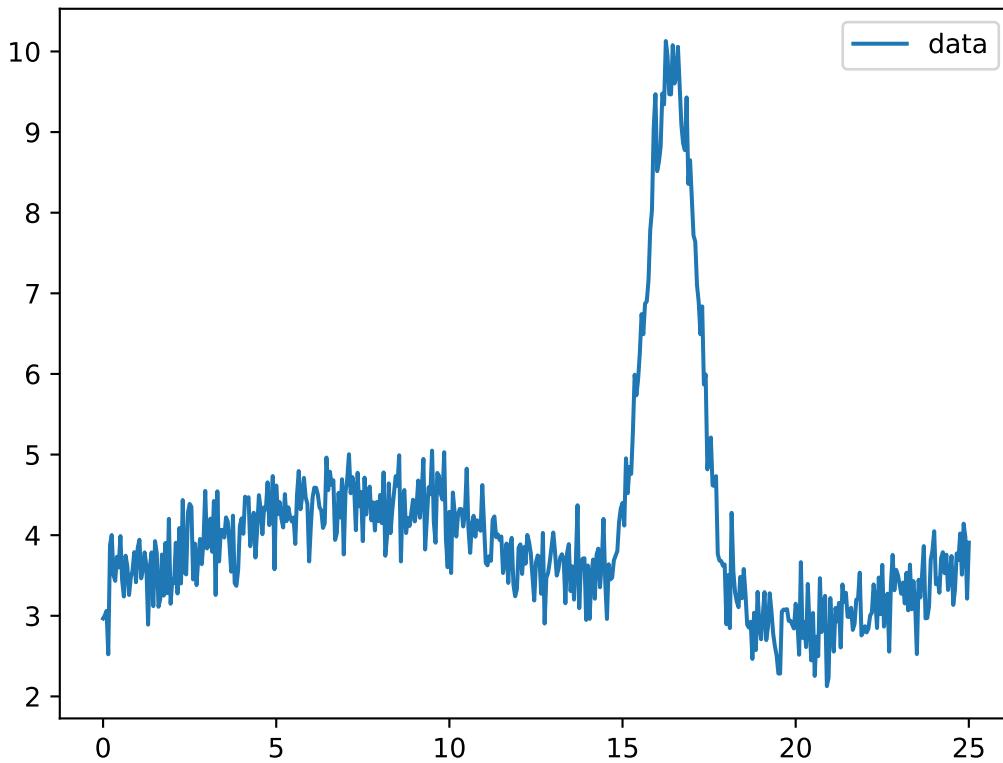
import numpy as np
import matplotlib.pyplot as plt
from lmfit.models import SplineModel, GaussianModel

data = np.loadtxt('test_splinepeak.dat')
x = data[:, 0]
y = data[:, 1]

plt.plot(x, y, label='data')
plt.legend()
plt.show()

```

which shows (figure below):



There is definitely a peak there, so we could start with building a model for a Gaussian peak, say with:

```

model = GaussianModel(prefix='peak_')
params = model.make_params(amplitude=8, center=16, sigma=1)

```

To account for that changing background, we'll use a spline, but need to know where to put the “knots”. Picking points away from the peak makes sense – we don't want to fit the peak – but we want it to have some flexibility near the peak. Let's try spacing knot points at $x=1, 3, \dots, 13$, then skip over the peak at around $x=16$ and then pick up knots points at $x=19, 21, 23, 25$.

```

knot_xvals = np.array([1, 3, 5, 7, 9, 11, 13, 19, 21, 23, 25])

bkg = SplineModel(prefix='bkg_', xknots=knot_xvals)
params.update(bkg.guess(y, x))

```

Note that we used `bkg.guess()` to guess the initial values of the spline parameters and then update the `params` Parameters object with these 11 parameters to account for the spline. These will be very close to the `y` values at the knot `x` values. The precise definition of the spline knot parameters is not “the `y`-values through which the resulting spline curve goes”, but these values are pretty good estimates for the resulting spline values. You’ll see below that these initial values are close.

With a spline background defined, we can create a composite model, and run a fit.

```
model = model + bkg

params['peak_amplitude'].min = 0
params['peak_center'].min = 10
params['peak_center'].max = 20

out = model.fit(y, params, x=x)
print(out.fit_report(min_correl=0.3))
```

You’ll see that we first set some “sanity bounds” on the peak parameters to prevent the peak from going completely wrong. This really is not necessary in this case, but it is often a reasonable thing to do - the general advice for this is to be generous in the bounds, not overly restrictive.

This fit will print out a report of

```
[[Model]]
    (Model(gaussian, prefix='peak_') + Model(spline_model, prefix='bkg_'))
[[Fit Statistics]]
    # fitting method    = leastsq
    # function evals   = 92
    # data points      = 501
    # variables        = 14
    chi-square         = 52.6611549
    reduced chi-square = 0.10813379
    Akaike info crit   = -1100.61674
    Bayesian info crit = -1041.58425
    R-squared          = 0.94690612
[[Variables]]
    peak_amplitude: 12.2231135 +/- 0.29554108 (2.42%) (init = 8)
    peak_center:    16.4280869 +/- 0.01091051 (0.07%) (init = 16)
    peak_sigma:     0.72096400 +/- 0.01336667 (1.85%) (init = 1)
    peak_fwhm:      1.69774046 +/- 0.03147610 (1.85%) == '2.3548200*peak_sigma'
    peak_height:    6.76360674 +/- 0.09854044 (1.46%) == '0.3989423*peak_amplitude/
    -max(1e-15, peak_sigma)'
    bkg_s0:         3.51175736 +/- 0.04941392 (1.41%) (init = 3.787995)
    bkg_s1:         3.72930068 +/- 0.09558236 (2.56%) (init = 3.959487)
    bkg_s2:         4.26846495 +/- 0.12650286 (2.96%) (init = 4.384009)
    bkg_s3:         4.42375490 +/- 0.10170203 (2.30%) (init = 4.431971)
    bkg_s4:         4.49590448 +/- 0.10615552 (2.36%) (init = 4.243976)
    bkg_s5:         3.96515315 +/- 0.09336555 (2.35%) (init = 4.115153)
    bkg_s6:         3.35531899 +/- 0.12669985 (3.78%) (init = 3.965325)
    bkg_s7:         2.89909752 +/- 0.16190211 (5.58%) (init = 2.788437)
    bkg_s8:         2.82656963 +/- 0.13445495 (4.76%) (init = 2.984317)
    bkg_s9:         3.43338680 +/- 0.15987281 (4.66%) (init = 3.383491)
    bkg_s10:        3.73024843 +/- 0.12096865 (3.24%) (init = 3.791937)
[[Correlations]] (unreported correlations are < 0.300)
    C(bkg_s7, bkg_s8) = -0.8192
```

(continues on next page)

(continued from previous page)


C(peak_amplitude, peak_sigma) =	+0.7987
C(bkg_s8, bkg_s9) =	-0.7063
C(bkg_s5, bkg_s6) =	-0.6950
C(peak_amplitude, bkg_s7) =	-0.6878
C(bkg_s2, bkg_s3) =	-0.6672
C(bkg_s9, bkg_s10) =	-0.6060
C(bkg_s3, bkg_s4) =	-0.5743
C(bkg_s1, bkg_s2) =	-0.5646
C(bkg_s4, bkg_s5) =	-0.5542
C(bkg_s7, bkg_s9) =	+0.5216
C(peak_sigma, bkg_s7) =	-0.5193
C(peak_amplitude, bkg_s8) =	+0.5185
C(bkg_s0, bkg_s1) =	+0.4448
C(peak_sigma, bkg_s8) =	+0.3733
C(peak_center, bkg_s6) =	+0.3599
C(bkg_s4, bkg_s6) =	+0.3597
C(bkg_s0, bkg_s2) =	-0.3595
C(bkg_s2, bkg_s4) =	+0.3504
C(bkg_s8, bkg_s10) =	+0.3455
C(bkg_s6, bkg_s7) =	-0.3332
C(peak_center, bkg_s7) =	-0.3301
C(peak_amplitude, bkg_s9) =	-0.3206

from this we can make a few observations. First, the correlation between the “spline” parameters and the “peak parameters” is noticeable, but not extremely high – that’s good, and the estimated uncertainties do account for this correlation. The spline components are correlated with each other (especially with the N-1 and N+1 spline parameter). Second, we can see that the initial values for the background spline parameters are pretty good.

We can plot the results and fit components with

```
comps = out.eval_components()
plt.plot(x, out.best_fit, label='best fit')
plt.plot(x, comps['bkg_'], label='background')
plt.plot(x, comps['peak_'], label='peak')
plt.legend()
```


which will generate the plot shown below:



If we're interested in seeing the locations of the knots, you might do

```
knot_yvals = np.array([o.value for o in out.params.values() if o.name.startswith('bkg')] )
plt.plot(knot_xvals, knot_yvals, 'o', color='black', label='spline knots values')
```

which will generate be shown as



You might be interested in trying to assess what impact the select of the knots has on the resulting peak intensity. For example, you might try some of the following set of knot values:

```
knot_xvals1 = np.array([1, 3, 5, 7, 9, 11, 13, 19, 21, 23, 25])
knot_xvals2 = np.array([1, 3, 5, 7, 9, 11, 13, 16, 19, 21, 23, 25])
knot_xvals3 = np.array([1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25])
```

and re-run the fit with these different sets of knot points. The results are shown in the table below.

Table of Peak amplitudes with varying spline points

spline x points	N	Peak amplitude value and uncertainty
knot_xvals1	11	12.223 (0.295)
knot_xvals2	12	11.746 (0.594)
knot_xvals3	13	12.052 (0.872)

Adding more spline points, especially near the peak center around $x=16.4$, can impact the measurement of the amplitude but the uncertainty increases dramatically enough to mostly cover the same range of values. This is a interesting case of adding more parameters to a fit and having the uncertainties in the fitted parameters getting worse. The interested reader is encouraged to explore the fit reports and plot these different case.

Finally, the basic case above used 11 spline points to fit the baseline. In fact, it would be reasonable to ask whether that is enough parameters to fit the full spectra. By imposing that there is also a Gaussian peak nearby makes the spline fit only the background, but without the Gaussian, the spline could fit the full curve. By way of example, we'll just try increasing the number of spline points to fit this data

```
plt.plot(x, y, 'o', label='data')
for nknots in (10, 15, 20, 25):
```

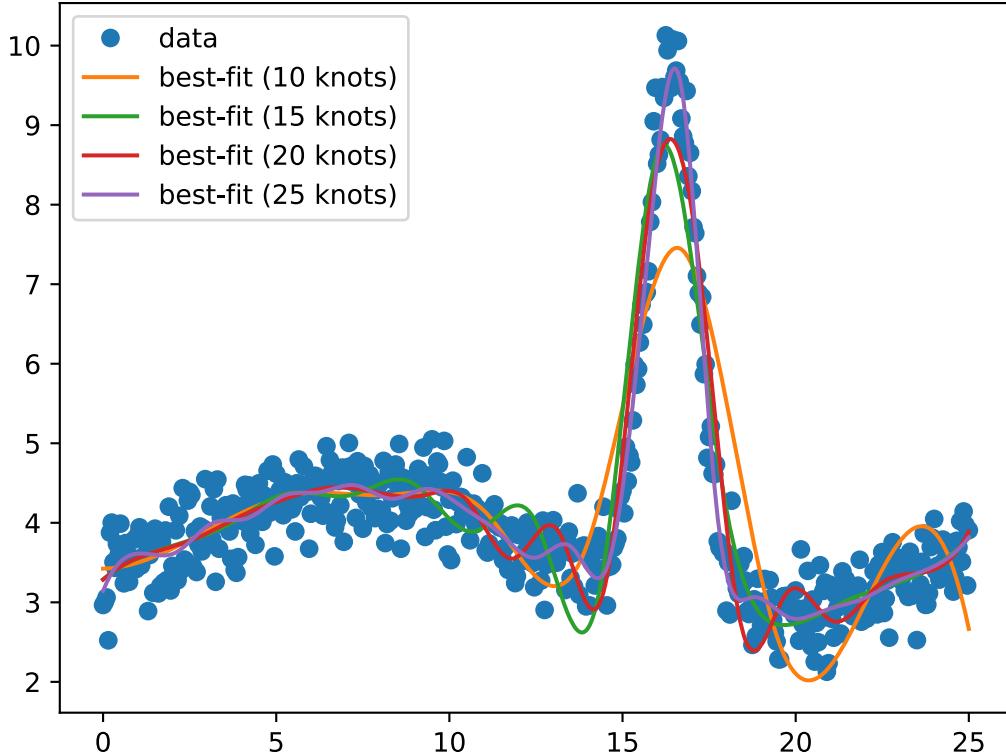
(continues on next page)

(continued from previous page)

```
model = SplineModel(prefix='bkg_', xknots=np.linspace(0, 25, nknots))
params = model.guess(y, x)
out = model.fit(y, params, x=x)
plt.plot(x, out.best_fit, label=f'best-fit ({nknots} knots)')

plt.legend()
plt.show()
```

which will show the fit below:



By itself, 10 knots does not give a very good fit, but 25 knots or more does give a very good fit to the peak. This should give some confidence that the fit with 11 parameters for the background spline is acceptable, but also give some reason to be careful in selecting the number of spline points to use.

CALCULATION OF CONFIDENCE INTERVALS

The lmfit `confidence` module allows you to explicitly calculate confidence intervals for variable parameters. For most models, it is not necessary since the estimation of the standard error from the estimated covariance matrix is normally quite good.

But for some models, the sum of two exponentials for example, the approximation begins to fail. For this case, lmfit has the function `conf_interval()` to calculate confidence intervals directly. This is substantially slower than using the errors estimated from the covariance matrix, but the results are more robust.

Please note that `conf_interval()` is not suited to work with fit results obtained by `emcee`.

9.1 Method used for calculating confidence intervals

The F-test is used to compare our null model, which is the best fit we have found, with an alternate model, where one of the parameters is fixed to a specific value. The value is changed until the difference between χ_0^2 and χ_f^2 can't be explained by the loss of a degree of freedom within a certain confidence.

$$F(P_{fix}, N - P) = \left(\frac{\chi_f^2}{\chi_0^2} - 1 \right) \frac{N - P}{P_{fix}}$$

N is the number of data points and P the number of parameters of the null model. P_{fix} is the number of fixed parameters (or to be more clear, the difference of number of parameters between our null model and the alternate model).

Adding a log-likelihood method is under consideration.

9.2 A basic example

First we create an example problem:

```
import numpy as np

import lmfit

x = np.linspace(0.3, 10, 100)
np.random.seed(0)
y = 1/(0.1*x) + 2 + 0.1*np.random.randn(x.size)
pars = lmfit.Parameters()
pars.add_many(('a', 0.1), ('b', 1))
```

(continues on next page)

(continued from previous page)

```
def residual(p):
    return 1/(p['a']**x) + p['b'] - y
```

before we can generate the confidence intervals, we have to run a fit, so that the automated estimate of the standard errors can be used as a starting point:

```
mini = lmfit.Minimizer(residual, pars)
result = mini.minimize()

print(lmfit.fit_report(result.params))
```

```
[[Variables]]
a: 0.09943896 +/- 1.9322e-04 (0.19%) (init = 0.1)
b: 1.98476942 +/- 0.01222678 (0.62%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(a, b) = +0.6008
```

Now it is just a simple function call to calculate the confidence intervals:

```
ci = lmfit.conf_interval(mini, result)
lmfit.printfuncs.report_ci(ci)
```

	99.73%	95.45%	68.27%	_BEST_	68.27%	95.45%	99.73%
a:	-0.00059	-0.00039	-0.00019	0.09944	+0.00019	+0.00039	+0.00060
b:	-0.03764	-0.02477	-0.01229	1.98477	+0.01229	+0.02477	+0.03764

This shows the best-fit values for the parameters in the `_BEST_` column, and parameter values that are at the varying confidence levels given by steps in σ . As we can see, the estimated error is almost the same, and the uncertainties are well behaved: Going from $1-\sigma$ (68% confidence) to $3-\sigma$ (99.7% confidence) uncertainties is fairly linear. It can also be seen that the errors are fairly symmetric around the best fit value. For this problem, it is not necessary to calculate confidence intervals, and the estimates of the uncertainties from the covariance matrix are sufficient.

9.3 Working without standard error estimates

Sometimes the estimation of the standard errors from the covariance matrix fails, especially if values are near given bounds. Hence, to find the confidence intervals in these cases, it is necessary to set the errors by hand. Note that the standard error is only used to find an upper limit for each value, hence the exact value is not important.

To set the step-size to 10% of the initial value we loop through all parameters and set it manually:

```
for p in result.params:
    result.params[p].stderr = abs(result.params[p].value * 0.1)
```

9.4 Calculating and visualizing maps of χ^2

The estimated values for the $1 - \sigma$ standard error calculated by default for each fit include the effects of correlation between pairs of variables, but assumes the uncertainties are symmetric. While it doesn't exactly say what the values of the $n - \sigma$ uncertainties would be, the implication is that the $n - \sigma$ error is simply $n^2\sigma$.

The `conf_interval()` function described above improves on these automatically (and quickly) calculated uncertainties by explicitly finding $n - \sigma$ confidence levels in both directions – it does not assume that the uncertainties are symmetric. This function also takes into account the correlations between pairs of variables, but it does not convey this information very well.

For even further exploration of the confidence levels of parameter values, it can be useful to calculate maps of χ^2 values for pairs of variables around their best fit values and visualize these as contour plots. Typically, pairs of variables will have elliptical contours of constant $n - \sigma$ level, with highly-correlated pairs of variables having high ratios of major and minor axes.

The `conf_interval2d()` can calculate 2-d arrays or maps of either probability or $\delta\chi^2 = \chi^2 - \chi^2_{\text{best}}$ for any pair of variables. Visualizing these can help better understand the nature of the uncertainties and correlations between parameters. To illustrate this, we'll start with an example fit to data that we deliberately add components not accounted for in the model, and with slightly non-Gaussian noise – a constructed but “real-world” example:

```
# <examples/doc_confidence_chi2_maps.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit import conf_interval, conf_interval2d, report_ci
from lmfit.lineshapes import gaussian
from lmfit.models import GaussianModel, LinearModel

sigma_levels = [1, 2, 3]

rng = np.random.default_rng(seed=102)

#####
# set up data -- deliberately adding imperfections and
# a small amount of non-Gaussian noise
npts = 501
x = np.linspace(1, 100, num=npts)
noise = rng.normal(scale=0.3, size=npts) + 0.2*rng.f(3, 9, size=npts)
y = (gaussian(x, amplitude=83, center=47., sigma=5.)
     + 0.02*x + 4 + 0.25*np.cos((x-20)/8.0) + noise)

mod = GaussianModel() + LinearModel()
params = mod.make_params(amplitude=100, center=50, sigma=5,
                        slope=0, intercept=2)
out = mod.fit(y, params, x=x)
print(out.fit_report())

#####
# run conf_interval, print report
sigma_levels = [1, 2, 3]
ci = conf_interval(out, out, sigmas=sigma_levels)

print("## Confidence Report:")
```

(continues on next page)

(continued from previous page)

report_ci(ci)

```
[[Model]]
  (Model(gaussian) + Model(linear))
[[Fit Statistics]]
  # fitting method = leastsq
  # function evals = 31
  # data points   = 501
  # variables     = 5
  chi-square      = 103.861381
  reduced chi-square = 0.20939794
  Akaike info crit = -778.348033
  Bayesian info crit = -757.265003
  R-squared        = 0.93782756
[[Variables]]
  amplitude: 78.8171374 +/- 1.21910939 (1.55%) (init = 100)
  center:    47.0751649 +/- 0.07576660 (0.16%) (init = 50)
  sigma:     4.93298753 +/- 0.07984021 (1.62%) (init = 5)
  slope:     0.01839006 +/- 7.1957e-04 (3.91%) (init = 0)
  intercept: 4.39234411 +/- 0.04420227 (1.01%) (init = 0)
  fwhm:      11.6162977 +/- 0.18800933 (1.62%) == '2.3548200*sigma'
  height:    6.37412722 +/- 0.08603873 (1.35%) == '0.3989423*amplitude/max(1e-15,_
  ↵sigma)'

[[Correlations]] (unreported correlations are < 0.100)
  C(slope, intercept) = -0.8421
  C(amplitude, sigma) = +0.6371
  C(amplitude, intercept) = -0.3373
  C(sigma, intercept) = -0.2149
  C(center, slope) = -0.1026
```

```
## Confidence Report:
      99.73%   95.45%   68.27%   _BEST_   68.27%   95.45%   99.73%
amplitude: -3.62610  -2.41983  -1.21237  78.81714  +1.22111  +2.45479  +3.70515
center   : -0.22849  -0.15214  -0.07584  47.07516  +0.07587  +0.15225  +0.22873
sigma    : -0.23335  -0.15640  -0.07870  4.93299  +0.08000  +0.16158  +0.24509
slope    : -0.00217  -0.00144  -0.00072  0.01839  +0.00072  +0.00144  +0.00217
intercept: -0.13326  -0.08860  -0.04423  4.39234  +0.04421  +0.08854  +0.13312
```

The reports show that we obtained a pretty good fit, and that the automated estimates of the uncertainties are actually pretty good – agreeing to the second decimal place. But we also see that some of the uncertainties do become noticeably asymmetric at high $n - \sigma$ levels.

We'll plot this data and fit, and then further explore these uncertainties using `conf_interval2d()`. Please note that `conf_interval2d()` shows the confidence intervals obtained considering the standard errors, not those obtained by '`conf_interval(out, out, sigmas=sigma_levels)`'.

```
#####
# plot initial fit
colors = ('#2030b0', '#b02030', '#207070')
fig, axes = plt.subplots(2, 3, figsize=(15, 9.5))

axes[0, 0].plot(x, y, 'o', markersize=3, label='data', color=colors[0])
```

(continues on next page)

(continued from previous page)

```

axes[0, 0].plot(x, out.best_fit, label='fit', color=colors[1])
axes[0, 0].set_xlabel('x')
axes[0, 0].set_ylabel('y')
axes[0, 0].legend()

aix, aiy = 0, 0
nsamples = 50
explicitly_calculate_sigma = True

for pairs in (('sigma', 'amplitude'), ('intercept', 'amplitude'),
              ('slope', 'intercept'), ('slope', 'center'), ('sigma', 'center')):

    xpar, ypar = pairs
    if explicitly_calculate_sigma:
        print(f"Generating chi-square map for {pairs}")
        c_x, c_y, chi2_mat = conf_interval2d(out, out, xpar, ypar,
                                              nsamples, nsamples, nsigma=3.5,
                                              chi2_out=True)
        # explicitly calculate sigma matrix: sigma increases chi_square
        # from chi_square_best
        # to chi_square + sigma**2 * reduced_chi_square
        # so: sigma = sqrt((chi2-chi2_best)/ reduced_chi_square)
        chi2_min = chi2_mat.min()
        sigma_mat = np.sqrt((chi2_mat-chi2_min)/out.redchi)
    else:
        print(f"Generating sigma map for {pairs}")
        # or, you could just calculate the matrix of probabilities as:
        c_x, c_y, sigma_mat = conf_interval2d(out, out, xpar, ypar,
                                              nsamples, nsamples, nsigma=3.5)

    aix += 1
    if aix == 2:
        aix = 0
        aiy += 1
    ax = axes[aix, aiy]

    cnt = ax.contour(c_x, c_y, sigma_mat, levels=sigma_levels, colors=colors,
                      linestyles='--')
    ax.clabel(cnt, inline=True, fmt=r"\sigma=%.\0f\%", fontsize=13)

    # draw boxes for estimated uncertainties:
    # dotted : scaled stderr from initial fit
    # dashed : values found from conf_interval()
    xv = out.params[xpar].value
    xs = out.params[xpar].stderr
    yv = out.params[ypar].value
    ys = out.params[ypar].stderr

    cix = ci[xpar]
    ciy = ci[ypar]
    nc = len(sigma_levels)
    for i in sigma_levels:

```

(continues on next page)

(continued from previous page)

```
# dotted line: scaled stderr
ax.plot((xv-i*xs, xv+i*xs, xv-i*xs, xv-i*xs),
         (yv-i*ys, yv-i*ys, yv+i*ys, yv-i*ys),
         linestyle='dotted', color=colors[i-1])

# dashed line: refined uncertainties from conf_interval
xsp, xsm = cix[nc+i][1], cix[nc-i][1]
ysp, ysm = ciy[nc+i][1], ciy[nc-i][1]
ax.plot((xsm, xsp, xsp, xsm), (ysm, ysm, ysp, ysp),
         linestyle='dashed', color=colors[i-1])

ax.set_xlabel(xpar)
ax.set_ylabel(ypar)
ax.grid(True, color='#d0d0d0')
plt.show()
# <end examples/doc_confidence_chi2_maps.py>
```


Generating chi-square map for ('sigma', 'amplitude')

Generating chi-square map for ('intercept', 'amplitude')

Generating chi-square map for ('slope', 'intercept')

Generating chi-square map for ('slope', 'center')

Generating chi-square map for ('sigma', 'center')



Here we made contours for the $n - \sigma$ levels from the 2-D array of χ^2 by noting that the $n - \sigma$ level will have χ^2 increased by $n^2 \chi_{\nu}^2$ where χ_{ν}^2 is reduced chi-square.

The dotted boxes show both the scaled values of the standard errors from the initial fit, and the dashed boxes show the confidence levels from `conf_interval()`. You can see that the notion of increasing χ^2 by χ_{ν}^2 works very well, and that there is a small asymmetry in the uncertainties for the `amplitude` and `sigma` parameters.

9.5 An advanced example for evaluating confidence intervals

Now we look at a problem where calculating the error from approximated covariance can lead to misleading result – the same double exponential problem shown in [Minimizer.emcee\(\) - calculating the posterior probability distribution of parameters](#). In fact such a problem is particularly hard for the Levenberg-Marquardt method, so we first estimate the results using the slower but robust Nelder-Mead method. We can then compare the uncertainties computed (if the `numdifftools` package is installed) with those estimated using Levenberg-Marquardt around the previously found solution. We can also compare to the results of using `emcee`.

```
# <examples/doc_confidence_advanced.py>
import matplotlib.pyplot as plt
import numpy as np

import lmfit

x = np.linspace(1, 10, 250)
np.random.seed(0)
y = 3.0*np.exp(-x/2) - 5.0*np.exp(-(x-0.1)/10.) + 0.1*np.random.randn(x.size)
```

(continues on next page)

(continued from previous page)

```

p = lmfit.create_params(a1=4, a2=4, t1=3, t2=3)

def residual(p):
    return p['a1']*np.exp(-x/p['t1']) + p['a2']*np.exp(-(x-0.1)/p['t2']) - y

# create Minimizer
mini = lmfit.Minimizer(residual, p, nan_policy='propagate')

# first solve with Nelder-Mead algorithm
out1 = mini.minimize(method='Nelder')

# then solve with Levenberg-Marquardt using the
# Nelder-Mead solution as a starting point
out2 = mini.minimize(method='leastsq', params=out1.params)

lmfit.report_fit(out2.params, min_correl=0.5)

ci, trace = lmfit.conf_interval(mini, out2, sigmas=[1, 2], trace=True)
lmfit.printfuncs.report_ci(ci)

# plot data and best fit
plt.figure()
plt.plot(x, y)
plt.plot(x, residual(out2.params) + y, '-')
plt.show()

# plot confidence intervals (a1 vs t2 and a2 vs t2)
fig, axes = plt.subplots(1, 2, figsize=(12.8, 4.8))
cx, cy, grid = lmfit.conf_interval2d(mini, out2, 'a1', 't2', 30, 30)
ctp = axes[0].contourf(cx, cy, grid, np.linspace(0, 1, 11))
fig.colorbar(ctp, ax=axes[0])
axes[0].set_xlabel('a1')
axes[0].set_ylabel('t2')

cx, cy, grid = lmfit.conf_interval2d(mini, out2, 'a2', 't2', 30, 30)
ctp = axes[1].contourf(cx, cy, grid, np.linspace(0, 1, 11))
fig.colorbar(ctp, ax=axes[1])
axes[1].set_xlabel('a2')
axes[1].set_ylabel('t2')
plt.show()

# plot dependence between two parameters
fig, axes = plt.subplots(1, 2, figsize=(12.8, 4.8))
cx1, cy1, prob = trace['a1']['a1'], trace['a1']['t2'], trace['a1']['prob']
cx2, cy2, prob2 = trace['t2']['t2'], trace['t2']['a1'], trace['t2']['prob']

axes[0].scatter(cx1, cy1, c=prob, s=30)
axes[0].set_xlabel('a1')
axes[0].set_ylabel('t2')

```

(continues on next page)

(continued from previous page)

```
axes[1].scatter(cx2, cy2, c=prob2, s=30)
axes[1].set_xlabel('t2')
axes[1].set_ylabel('a1')
plt.show()
# <end examples/doc_confidence_advanced.py>
```


which will report:

```
[[Variables]]
a1: 2.98622095 +/- 0.14867027 (4.98%) (init = 2.986237)
a2: -4.33526363 +/- 0.11527574 (2.66%) (init = -4.335256)
t1: 1.30994276 +/- 0.13121215 (10.02%) (init = 1.309932)
t2: 11.8240337 +/- 0.46316956 (3.92%) (init = 11.82408)
[[Correlations]] (unreported correlations are < 0.500)
C(a2, t2) = +0.9871
C(a2, t1) = -0.9246
C(t1, t2) = -0.8805
C(a1, t1) = -0.5988

      95.45%   68.27%   _BEST_   68.27%   95.45%
a1: -0.27285  -0.14165  2.98622  +0.16354  +0.36343
a2: -0.30440  -0.13219  -4.33526  +0.10689  +0.19684
t1: -0.23392  -0.12494  1.30994  +0.14660  +0.32369
t2: -1.01937  -0.48813  11.82403  +0.46045  +0.90439
```

Again we called `conf_interval()`, this time with tracing and only for 1- and 2- σ . Comparing these two different estimates, we see that the estimate for `a1` is reasonably well approximated from the covariance matrix, but the estimates for `a2` and especially for `t1`, and `t2` are very asymmetric and that going from 1 σ (68% confidence) to 2 σ (95% confidence) is not very predictable.

Plots of the confidence region are shown in the figures below for `a1` and `t2` (left), and `a2` and `t2` (right):



Neither of these plots is very much like an ellipse, which is implicitly assumed by the approach using the covariance matrix. The plots actually look quite a bit like those found with MCMC and shown in the “corner plot” in [Minimizer.emcee\(\) - calculating the posterior probability distribution of parameters](#). In fact, comparing the confidence interval results here with the results for the 1- and 2- σ error estimated with `emcee`, we can see that the agreement is

pretty good and that the asymmetry in the parameter distributions are reflected well in the asymmetry of the uncertainties.

The trace returned as the optional second argument from `conf_interval()` contains a dictionary for each variable parameter. The values are dictionaries with arrays of values for each variable, and an array of corresponding probabilities for the corresponding cumulative variables. This can be used to show the dependence between two parameters:

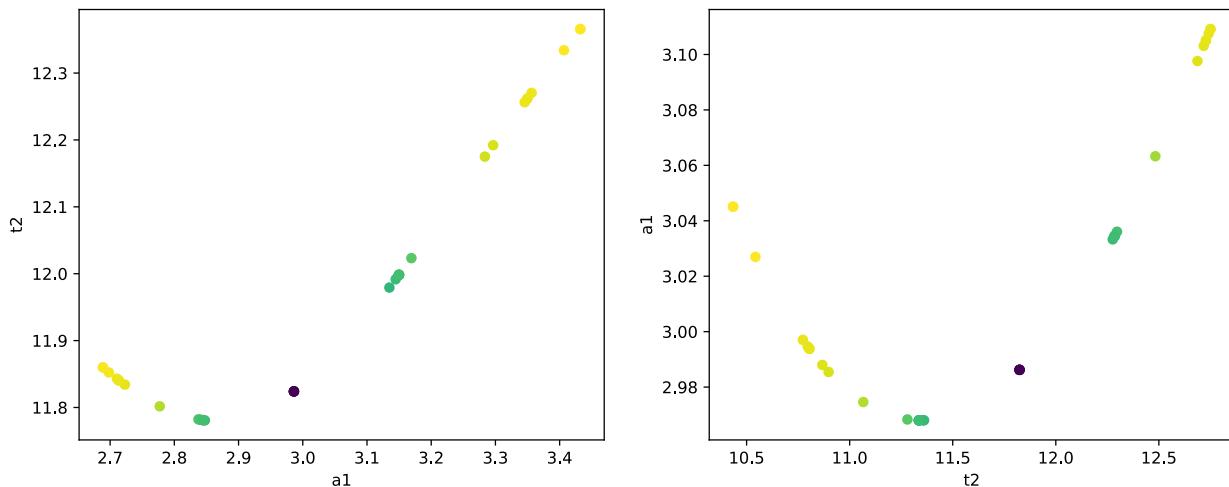
```
fig, axes = plt.subplots(1, 2, figsize=(12.8, 4.8))
cx1, cy1, prob = trace['a1']['a1'], trace['a1']['t2'], trace['a1']['prob']
cx2, cy2, prob2 = trace['t2']['t2'], trace['t2']['a1'], trace['t2']['prob']

axes[0].scatter(cx1, cy1, c=prob, s=30)
axes[0].set_xlabel('a1')
axes[0].set_ylabel('t2')

axes[1].scatter(cx2, cy2, c=prob2, s=30)
axes[1].set_xlabel('t2')
axes[1].set_ylabel('a1')

plt.show()
```

which shows the trace of values:



As an alternative/complement to the confidence intervals, the `Minimizer.emcee()` method uses Markov Chain Monte Carlo to sample the posterior probability distribution. These distributions demonstrate the range of solutions that the data supports and we refer to [Minimizer.emcee\(\) - calculating the posterior probability distribution of parameters](#) where this methodology was used on the same problem.

Credible intervals (the Bayesian equivalent of the frequentist confidence interval) can be obtained with this method. MCMC can be used for model selection, to determine outliers, to marginalize over nuisance parameters, etcetera. For example, you may have fractionally underestimated the uncertainties on a dataset. MCMC can be used to estimate the true level of uncertainty on each data point. A tutorial on the possibilities offered by MCMC can be found at¹.

¹ <https://jakevdp.github.io/blog/2014/03/11/frequentism-and-bayesianism-a-practical-intro/>

9.6 Confidence Interval Functions

```
conf_interval(minimizer, result, p_names=None, sigmas=None, trace=False, maxiter=200, verbose=False, prob_func=None, min_rel_change=1e-05)
```

Calculate the confidence interval (CI) for parameters.

The parameter for which the CI is calculated will be varied, while the remaining parameters are re-optimized to minimize the chi-square. The resulting chi-square is used to calculate the probability with a given statistic (e.g., F-test). This function uses a 1d-rootfinder from SciPy to find the values resulting in the searched confidence region.

Parameters

- **minimizer** (`Minimizer`) – The minimizer to use, holding objective function.
- **result** (`MinimizerResult`) – The result of running `minimize()`.
- **p_names** (`list`, *optional*) – Names of the parameters for which the CI is calculated. If `None` (default), the CI is calculated for every parameter.
- **sigmas** (`list`, *optional*) – The sigma-levels to find (default is `[1, 2, 3]`). See Notes below.
- **trace** (`bool`, *optional*) – Defaults to `False`; if `True`, each result of a probability calculation is saved along with the parameter. This can be used to plot so-called “profile traces”.
- **maxiter** (`int`, *optional*) – Maximum of iteration to find an upper limit (default is 200).
- **verbose** (`bool`, *optional*) – Print extra debugging information (default is `False`).
- **prob_func** (`None` or *callable*, *optional*) – Function to calculate the probability from the optimized chi-square. Default is `None` and uses the built-in function `f_compare` (i.e., F-test).
- **min_rel_change** (`float`, *optional*) – Minimum relative change in probability (default is `1e-5`).

Returns

- **output** (`dict`) – A dictionary containing a list of `(sigma, vals)`-tuples for each parameter.
- **trace_dict** (`dict`, *optional*) – Only if `trace` is `True`. Is a dictionary, the key is the parameter which was fixed. The values are again a dict with the names as keys, but with an additional key ‘prob’. Each contains an array of the corresponding values.

See also

[conf_interval2d](#)

Notes

The values for `sigma` are taken as the number of standard deviations for a normal distribution and converted to probabilities. That is, the default `sigma=[1, 2, 3]` will use probabilities of 0.6827, 0.9545, and 0.9973. If any of the sigma values is less than 1, that will be interpreted as a probability. That is, a value of 1 and 0.6827 will give the same results, within precision.

Examples

```
>>> from lmfit.printfuncs import *
>>> mini = minimize(some_func, params)
>>> mini.leastsq()
True
>>> report_errors(params)
... #report
>>> ci = conf_interval(mini)
>>> report_ci(ci)
... #report
```

Now with quantiles for the sigmas and using the trace.

```
>>> ci, trace = conf_interval(mini, sigmas=[0.5, 1, 2, 3], trace=True)
>>> fixed = trace['para1']['para1']
>>> free = trace['para1']['not_para1']
>>> prob = trace['para1']['prob']
```

This makes it possible to plot the dependence between free and fixed parameters.

```
conf_interval2d(minimizer, result, x_name, y_name, nx=10, ny=10, limits=None, prob_func=None, nsigma=5,
                 chi2_out=False)
```

Calculate confidence regions for two fixed parameters.

The method itself is explained in `conf_interval`: here we are fixing two parameters.

Parameters

- `minimizer` (`Minimizer`) – The minimizer to use, holding objective function.
- `result` (`MinimizerResult`) – The result of running `minimize()`.
- `x_name` (`str`) – The name of the parameter which will be the x direction.
- `y_name` (`str`) – The name of the parameter which will be the y direction.
- `nx` (`int, optional`) – Number of points in the x direction (default is 10).
- `ny` (`int, optional`) – Number of points in the y direction (default is 10).
- `limits` (`tuple, optional`) – Should have the form `((x_upper, x_lower), (y_upper, y_lower))`. If not given, the default is `nsigma*stderr` in each direction.
- `prob_func` (`None or callable, deprecated`) – Starting with version 1.2, this argument is unused and has no effect.
- `nsigma` (`float or int, optional`) – Multiplier of `stderr` for limits (default is 5).
- `chi2_out` (`bool`) – Whether to return chi-square at each coordinate instead of probability.

Returns

- **x** (`numpy.ndarray`) – X-coordinates (same shape as *nx*).
- **y** (`numpy.ndarray`) – Y-coordinates (same shape as *ny*).
- **grid** (`numpy.ndarray`) – 2-D array (with shape (*nx*, *ny*)) containing the calculated probabilities or chi-square.

See also

`conf_interval`

Examples

```
>>> mini = Minimizer(some_func, params)
>>> result = mini.leastsq()
>>> x, y, gr = conf_interval2d(mini, result, 'para1','para2')
>>> plt.contour(x,y,gr)
```

ci_report(*ci*, *with_offset=True*, *ndigits=5*)

Return text of a report for confidence intervals.

Parameters

- **ci** (`dict`) – The result of `conf_interval()`: a dictionary containing a list of (*sigma*, *vals*)-tuples for each parameter.
- **with_offset** (`bool`, *optional*) – Whether to subtract best value from all other values (default is True).
- **ndigits** (`int`, *optional*) – Number of significant digits to show (default is 5).

Returns

Text of formatted report on confidence intervals.

Return type

`str`

BOUNDS IMPLEMENTATION

This section describes the implementation of Parameter bounds. The [MINPACK-1](#) implementation used in `scipy.optimize.leastsq` for the Levenberg-Marquardt algorithm does not explicitly support bounds on parameters, and expects to be able to fully explore the available range of values for any Parameter. Simply placing hard constraints (that is, resetting the value when it exceeds the desired bounds) prevents the algorithm from determining the partial derivatives, and leads to unstable results.

Instead of placing such hard constraints, bounded parameters are mathematically transformed using the formulation devised (and documented) for [MINUIT](#). This is implemented following (and borrowing heavily from) the `leastsqbound` from J. J. Helmus. Parameter values are mapped from internally used, freely variable values P_{internal} to bounded parameters P_{bounded} . When both `min` and `max` bounds are specified, the mapping is:

$$\begin{aligned} P_{\text{internal}} &= \arcsin\left(\frac{2(P_{\text{bounded}} - \text{min})}{(\text{max} - \text{min})} - 1\right) \\ P_{\text{bounded}} &= \text{min} + (\sin(P_{\text{internal}}) + 1) \frac{(\text{max} - \text{min})}{2} \end{aligned}$$

With only an upper limit `max` supplied, but `min` left unbounded, the mapping is:

$$\begin{aligned} P_{\text{internal}} &= \sqrt{(\text{max} - P_{\text{bounded}} + 1)^2 - 1} \\ P_{\text{bounded}} &= \text{max} + 1 - \sqrt{P_{\text{internal}}^2 + 1} \end{aligned}$$

With only a lower limit `min` supplied, but `max` left unbounded, the mapping is:

$$\begin{aligned} P_{\text{internal}} &= \sqrt{(P_{\text{bounded}} - \text{min} + 1)^2 - 1} \\ P_{\text{bounded}} &= \text{min} - 1 + \sqrt{P_{\text{internal}}^2 + 1} \end{aligned}$$

With these mappings, the value for the bounded Parameter cannot exceed the specified bounds, though the internally varied value can be freely varied.

It bears repeating that code from `leastsqbound` was adopted to implement the transformation described above. The challenging part (thanks again to Jonathan J. Helmus!) here is to re-transform the covariance matrix so that the uncertainties can be estimated for bounded Parameters. This is included by using the derivative $dP_{\text{internal}}/dP_{\text{bounded}}$ from the equations above to re-scale the Jacobian matrix before constructing the covariance matrix from it. Tests show that this re-scaling of the covariance matrix works quite well, and that uncertainties estimated for bounded are quite reasonable. Of course, if the best fit value is very close to a boundary, the derivative estimated uncertainty and correlations for that parameter may not be reliable.

The [MINUIT](#) documentation recommends caution in using bounds. Setting bounds can certainly increase the number of function evaluations (and so computation time), and in some cases may cause some instabilities, as the range of acceptable parameter values is not fully explored. On the other hand, preliminary tests suggest that using `max` and `min` to set clearly outlandish bounds does not greatly affect performance or results.

USING MATHEMATICAL CONSTRAINTS

Being able to fix variables to a constant value or place upper and lower bounds on their values can greatly simplify modeling real data. These capabilities are key to lmfit's Parameters. In addition, it is sometimes highly desirable to place mathematical constraints on parameter values. For example, one might want to require that two Gaussian peaks have the same width, or have amplitudes that are constrained to add to some value. Of course, one could rewrite the objective or model function to place such requirements, but this is somewhat error-prone, and limits the flexibility so that exploring constraints becomes laborious.

To simplify the setting of constraints, Parameters can be assigned a mathematical expression of other Parameters, builtin constants, and builtin mathematical functions that will be used to determine its value. The expressions used for constraints are evaluated using the `asteval` module, which uses Python syntax, and evaluates the constraint expressions in a safe and isolated namespace.

This approach to mathematical constraints allows one to not have to write a separate model function for two Gaussians where the two `sigma` values are forced to be equal, or where amplitudes are related. Instead, one can write a more general two Gaussian model (perhaps using `GaussianModel`) and impose such constraints on the Parameters for a particular fit.

11.1 Overview

Just as one can place bounds on a Parameter, or keep it fixed during the fit, so too can one place mathematical constraints on parameters. The way this is done with lmfit is to write a Parameter as a mathematical expression of the other parameters and a set of pre-defined operators and functions. The constraint expressions are simple Python statements, allowing one to place constraints like:

```
from lmfit import Parameters

pars = Parameters()
pars.add('frac_curve1', value=0.5, min=0, max=1)
pars.add('frac_curve2', expr='1-frac_curve1')
```

as the value of the `frac_curve1` parameter is updated at each step in the fit, the value of `frac_curve2` will be updated so that the two values are constrained to add to 1.0. Of course, such a constraint could be placed in the fitting function, but the use of such constraints allows the end-user to modify the model of a more general-purpose fitting function.

Nearly any valid mathematical expression can be used, and a variety of built-in functions are available for flexible modeling.

11.2 Supported Operators, Functions, and Constants

The mathematical expressions used to define constrained Parameters need to be valid Python expressions. As you would expect, the operators `+`, `-`, `*`, `/`, and `**`, are supported. In fact, a much more complete set can be used, including Python's bit- and logical operators:

```
+, -, *, /, **, &, |, ^, <<, >>, %, and, or,
==, >, >=, <, <=, !=, ~, not, is, is not, in, not in
```

The values for `e` (2.7182818...) and `pi` (3.1415926...) are available, as are several supported mathematical and trigonometric function:

```
abs, acos, acosh, asin, asinh, atan, atan2, atanh, ceil,
copysign, cos, cosh, degrees, exp, fabs, factorial,
floor, fmod, frexp, fsum, hypot, isinf, isnan, ldexp,
log, log10, log1p, max, min, modf, pow, radians, sin,
sinh, sqrt, tan, tanh, trunc
```

In addition, all Parameter names will be available in the mathematical expressions. Thus, with parameters for a few peak-like functions:

```
pars = Parameters()
pars.add('amp_1', value=0.5, min=0, max=1)
pars.add('cen_1', value=2.2)
pars.add('wid_1', value=0.2)
```

The following expression are all valid:

```
pars.add('amp_2', expr='(2.0 - amp_1**2)')
pars.add('wid_2', expr='sqrt(pi)*wid_1')
pars.add('cen_2', expr='cen_1 * wid_2 / max(wid_1, 0.001)')
```

In fact, almost any valid Python expression is allowed. A notable example is that Python's 1-line *if expression* is supported:

```
pars.add('param_a', value=1)
pars.add('param_b', value=2)
pars.add('test_val', value=100)

pars.add('bounded', expr='param_a if test_val/2. > 100 else param_b')
```

which is equivalent to the more familiar:

```
if pars['test_val'].value/2. > 100:
    bounded = pars['param_a'].value
else:
    bounded = pars['param_b'].value
```

11.3 Using Inequality Constraints

A rather common question about how to set up constraints that use an inequality, say, $x + y \leq 10$. This can be done with algebraic constraints by recasting the problem, as $x + y = \delta$ and $\delta \leq 10$. That is, first, allow x to be held by the freely varying parameter `x`. Next, define a parameter `delta` to be variable with a maximum value of 10, and define parameter `y` as `delta - x`:

```
pars = Parameters()
pars.add('x', value=5, vary=True)
pars.add('delta', value=5, max=10, vary=True)
pars.add('y', expr='delta-x')
```

The essential point is that an inequality still implies that a variable (here, `delta`) is needed to describe the constraint. The secondary point is that upper and lower bounds can be used as part of the inequality to make the definitions more convenient.

11.4 Advanced usage of Expressions in lmfit

The expression used in a constraint is converted to a Python [Abstract Syntax Tree](#), which is an intermediate version of the expression – a syntax-checked, partially compiled expression. Among other things, this means that Python’s own parser is used to parse and convert the expression into something that can easily be evaluated within Python. It also means that the symbols in the expressions can point to any Python object.

In fact, the use of Python’s AST allows a nearly full version of Python to be supported, without using Python’s built-in `eval()` function. The `asteval` module actually supports most Python syntax, including for- and while-loops, conditional expressions, and user-defined functions. There are several unsupported Python constructs, most notably the `class` statement, so that new classes cannot be created, and the `import` statement, which helps make the `asteval` module safe from malicious use.

One important feature of the `asteval` module is that you can add domain-specific functions into the it, for later use in constraint expressions. To do this, you would use the `_asteval` attribute of the `Parameters` class, which contains a complete AST interpreter. The `asteval` interpreter uses a flat namespace, implemented as a single dictionary. That means you can preload any Python symbol into the namespace for the constraints, for example this Lorentzian function:

```
def mylorentzian(x, amp, cen, wid):
    "lorentzian function: wid = half-width at half-max"
    return (amp / (1 + ((x-cen) / wid)**2))
```

You can add this user-defined function to the `asteval` interpreter of the `Parameters` class:

```
from lmfit import Parameters

pars = Parameters()
pars._asteval.symtable['lorentzian'] = mylorentzian
```

and then initialize the `Minimizer` class with this parameter set:

```
from lmfit import Minimizer

def userfcn(x, params):
    pass
```

(continues on next page)

(continued from previous page)

```
fitter = Minimizer(userfcn, pars)
```

Alternatively, one can first initialize the `Minimizer` class and add the function to the `asteval` interpreter of `Minimizer`. `params` afterwards:

```
pars = Parameters()
fitter = Minimizer(userfcn, pars)
fitter.params._asteval.symtable['lorentzian'] = mylorentzian
```

In both cases the user-defined `lorentzian()` function can now be used in constraint expressions.

RELEASE NOTES

This section discusses changes between versions, especially changes significant to the use and behavior of the library. This is not meant to be a comprehensive list of changes. For such a complete record, consult the [lmfit GitHub repository](#).

12.1 Version 1.3.2 Release Notes (July 19, 2024)

Fixes:

- fix typo in restoring a `_buildmodel` dict (PR #957, Issue #956)
- fixes for Numpy2 support (PR #959, Issue #958)
- ensure that correct initial params are used when re-fitting a ModeResult (PR #961, Issue #960)
- make sure that CompositeModels cannot have a prefix (PR #961, Issue #954)

Build, Maintenance:

- update pre-commit hooks, adding codespell exceptions
- update to latest SciPy/NumPy versions, including dependency versions for NumPy 2.
- now require asteval>=1.0 and uncertainties>=3.2.2

12.2 Version 1.3.1 Release Notes (April 19, 2024)

Mostly fixes for bugs introduced in 1.3.0

- allow `Model.eval_uncertainty` to be performed with single points for `x` independent variables (PR #952, Issue #951)
- allow `Model._parse_param` to handle older-style passed-in ‘argnames’ and ‘kwargs’ as for variadic function, add test (PR #950)
- better allow (or re-allow) `Model` function independent variables / keyword arguments to be given non-default values at model creation time
- add `form` as independent variable for builtin Step, Rectangle, Thermal Distribution models.
- use a copy of `sys.modules` when iterating over it. (#949)
- use `Model._repr_string(long=True)` for model `Model.__repr__()`.

12.3 Version 1.3.0 Release Notes (April 4, 2024)

New features:

- add '`min_rel_change`' as optional variable in calculation of confidence intervals with `Model.conf_interval()`. (PR #937).
- `Model.eval_uncertainty` now takes an optional `dscale` parameter (default value of 0.01) to set the step size for calculating derivatives (PR #933).
- add calculation of `predicted_interval` to `Model.eval_uncertainty` (PR #933).

Bug fixes/enhancements:

- restore best-fit parameter values for high accuracy values of constrained values (PR #907)
- improvement to `Model` for the difference between `Parameter`, “independent variable”, and “option”. With this change, keyword arguments to model functions with non-numeric default values such as `do_thing=True`, or `form='linear'` has those arguments become clearly identified as independent variables, and use the provided values as default values. (PR #941)
- better saving/loading saved states of `Model` now use `dill`, have several cleanups, and are now versioned for future-proofing. Also, propagate `funcdets` for `Parameters` when loading a `Model`. (PR #932, PR #934)
- in the `TNC` method, `maxfun` is used instead of `maxiter`.
- fix bug calculating r-squared for fits with weights (PR #921, PR #923)
- fix bug in `modelresult.eval_uncertainty()` after `load_modelresult()` (PR #909)
- use `StringIO` for `pandas.read_json`.
- add test for `MinimizerResult.uvars` after successful fit (PR #913)
- adding an example using `basinhopping`, can take other methods as command-line argument

Maintenance/Deprecations:

- drop support for Python 3.7 that reached EOL on 2023-06-27 (PR #927)
- fix tests for Python 3.12 and Python 3.13-dev
- increase minimum numpy version to 1.23 and scipy to 1.8.
- updates for compatibility with numpy 2.0
- the `dill` package is now required. (#940)
- build switched to use `pyproject.toml` (#928)
- fix broken links in Examples gallery
- fix intersphinx mapping to scipy docs.

12.4 Version 1.2.2 Release Notes (July 14, 2023)

New features:

- add `ModelError.uvars` output to a `ModelError` after a successful fit that contains `ufloats` from the `uncertainties` package which can be used for downstream calculations that propagate the uncertainties (and correlations) of the variable `Parameters`. (PR #888)
- Outputs of residual functions, including `Model._residual`, are more explicitly coerced to 1d-arrays of datatype `Float64`. This decreases the expectation for the user-supplied code to return `ndarrays`, and increases the tolerance for more “array-like” objects or `ndarrays` that are not `Float64` or 1-dimensional. (PR #899)
- `Model.fit` now takes a `coerce_farray` option, defaulting to `True` to control whether to input data and independent variables that are “array-like” are coerced to `ndarrays` of datatype `Float64` or `Complex128`. If set to `False` then independent data that “array-like” (`pandas.Series`, `int32` arrays, etc) will be sent to the model function unaltered. The user may then use other features of these objects, but may also need to explicitly coerce the datatype of the result the change described above about coercing the result causes problems. (Discussion #873; PR #899)

Bug fixes/enhancements:

- fixed bug in `Model.make_params()` for non-composite models that use a prefix (Discussion #892; Issue #893; PR #895)
- fixed bug with aborted fits for several methods having incorrect or invalid fit statistics. (Discussion #894; Issue #896; PR #897)
- `Model.eval_uncertainty` now correctly calculates complex (real/imaginary pairs) uncertainties for Models that generate complex results. (Issue #900; PR #901)
- `Model.eval` now returns and array-like value. This adds to the coercion features above and fixes a bug for composite models that return lists (Issue #875; PR #901)
- the HTML representation for a `ModelError` or `MinimizerResult` are improved, and create fewer entries in the Table of Contents for Jupyter lab. (Issue #884; PR #883; PR #902)

12.5 Version 1.2.1 Release Notes (May 02, 2023)

Bug fixes/enhancements:

- fixed bug in `Model.make_params()` for initial parameter values that were not recognized as floats such as `np.Int64`. (Issue #871; PR #872)
- explicitly set `maxfun` for `1-bfgs-b` method when setting `maxiter`. (Issue #864; Discussion #865; PR #866)

12.6 Version 1.2.0 Release Notes (April 05, 2023)

New features:

- add `create_params` function (PR #844)
- add `chi2_out` and `nsigma` options to `conf_interval2d()`
- add `ModelError.summary()` to return many resulting fit statistics and attributes into a JSON-able dict.
- add `correl_table()` function to `lmfit.printfuncs` and `correl_mode` option to `fit_report()` and `ModelError.fit_report()` to optionally display a RST-formatted table of a correlation matrix.

Bug fixes/enhancements:

- fix bug when setting `param.vary=True` for a constrained parameter (Issue #859; PR #860)
- fix bug in reported uncertainties for constrained parameters by better propagating uncertainties (Issue #855; PR #856)
- Coercing of user input data and independent data for `Model` to `float64` `ndarrays` is somewhat less aggressive and will not increase the precision of `numpy` `ndarrays` (see [Data Types for data and independent data with Model](#) for details). The resulting calculation from a model or objective function is more aggressively coerced to `float64`. (Issue #850; PR #853)
- the default value of `epsfcn` is increased to `1.e-10` to allow for handling of data with precision less than `float64` (Issue #850; PR #853)
- fix `conf_interval2d` to use “increase chi-square by σ^{**2} reduced chi-square” to give the `sigma`-level probabilities (Issue #848; PR #852)
- fix reading of older `ModelResult` (Issue #845; included in PR #844)
- fix `deepcopy` of `Parameters` and user data (mguhyo; PR #837)
- improve `Model.make_params` and `create_params` to take optional dict of `Parameter` attributes (PR #844)
- fix reporting of `nfev` from `least_squares` to better reflect actual number of function calls (Issue #842; PR #844)
- fix bug in `Model.eval` when mixing parameters and keyword arguments (PR #844, #839)
- re-adds `residual` to saved `Model` result (PR #844, #830)
- `ConstantModel` and `ComplexConstantModel` will return an `ndarray` of the same shape as the independent variable `x` (JeppeKlitgaard, Issue #840; PR #841)
- update tests for latest versions of NumPy and SciPy.
- many fixes of doc typos and updates of dependencies, pre-commit hooks, and CI.

12.7 Version 1.1.0 Release Notes (November 27, 2022)

New features:

- add `Pearson4Model` (@lellid; PR #800)
- add `SplineModel` (PR #804)
- add `R^2 rsquared` statistic to fit outputs and reports for `Model` fits (Issue #803; PR #810)
- add calculation of `dely` for model components of composite models (Issue #761; PR #826)

Bug fixes/enhancements:

- make sure variable `spercent` is always defined in `params_html_table` functions (reported by @MySlientWind; Issue #768, PR #770)
- always initialize the variables `success` and `covar` the `MinimizerResult` (reported by Marc W. Pound; PR #771)
- build package following PEP517/PEP518; use `pyproject.toml` and `setup.cfg`; leave `setup.py` for now (PR #777)
- components used to create a `CompositeModel` can now have different independent variables (@Julian-Hochhaus; Discussion #787; PR #788)

- fixed function definition for `StepModel(form='linear')`, was not consistent with the other ones (@matpompili; PR #794)
- fixed height factor for `Gaussian2dModel`, was not correct (@matpompili; PR #795)
- for covariances with negative diagonal elements, we set the covariance to `None` (PR #813)
- fixed linear mode for `RectangleModel` (@arunpersaud; Issue #815; PR #816)
- report correct initial values for parameters with bounds (Issue #820; PR #821)
- allow recalculation of confidence intervals (@jagerber48; PR #798)
- include ‘residual’ in JSON output of `ModelError.dumps` (@mac01021; PR #830)
- supports and is tested against Python 3.11; updated minimum required version of SciPy, NumPy, and asteval (PR #832)

Deprecations:

- remove support for Python 3.6 which reached EOL on 2021-12-23 (PR #790)

12.8 Version 1.0.3 Release Notes (October 14, 2021)

Potentially breaking change:

- argument `x` is now required for the `guess` method of `Models` (Issue #747; PR #748)

To get reasonable estimates for starting values one should always supply both `x` and `y` values; in some cases it would work when only providing `data` (i.e., `y`-values). With the change above, `x` is now required in the `guess` method call, so scripts might need to be updated to explicitly supply `x`.

Bug fixes/enhancements:

- do not overwrite user-specified figure titles in `Model.plot()` functions and allow setting with `title` keyword argument (PR #711)
- preserve `Parameters` subclass in `deepcopy` (@jenshnielsen; PR #719)
- coerce `data` and `indepndent_vars` to NumPy array with `dtype=float64` or `dtype=complex128` where applicable (Issues #723 and #728)
- fix collision between parameter names in built-in models and user-specified parameters (Issue #710 and PR #732)
- correct error message in `PolynomialModel` (@kremeyer; PR #737)
- improved handling of altered JSON data (Issue #739; PR #740, reported by Matthew Giammar)
- map `max_nfev` to `maxiter` when using `differential_evolution` (PR #749, reported by Olivier B.)
- correct use of noise versus experimental uncertainty in the documentation (PR #751, reported by Andrés Zelcer)
- specify return type of `eval` method more precisely and allow for plotting of (Complex)ConstantModel by coercing their `float`, `int`, or `complex` return value to a `numpy.ndarray` (Issue #684 and PR #754)
- fix `dho` (Damped Harmonic Oscillator) lineshape (PR #755; @rayosborn)
- reset `Minimizer._abort` to `False` before starting a new fit (Issue #756 and PR #757; @zelcer)
- fix typo in `guess_from_peak2d` (@ivan-usov; PR #758)

Various:

- update `asteval` dependency to $\geq 0.9.22$ to avoid DeprecationWarnings from NumPy v1.20.0 (PR #707)
- remove incorrectly spelled `DonaichModel` and `donaich` lineshape, deprecated in version 1.0.1 (PR #707)

- remove occurrences of `OrderedDict` throughout the code; `dict` is order-preserving since Python 3.6 (PR #713)
- update the contributing instructions (PR #718; @martin-majlis)
- (again) defer import of `matplotlib` to when it is needed (@zobristnicholas; PR #721)
- fix description of `name` argument in `Parameters.add` (@kristianmeyerr; PR #725)
- update dependencies, make sure a functional development environment is installed on Windows (Issue #712)
- use `setuptools_scm` for version info instead of `versioneer` (PR #729)
- transition to using `f-strings` (PR #730)
- mark `test_manypeaks_speed.py` as flaky to avoid intermittent test failures (repeat up to 5 times; PR #745)
- update `scipy` dependency to $\geq 1.14.0$ (PR #751)
- improvement to output of examples in `sphinx-gallery` and use higher resolution figures (PR #753)
- remove deprecated functions `lmfit.printfuncs.report_errors` and `asteval` argument in `Parameters` class (PR #759)

12.9 Version 1.0.2 Release Notes (February 7, 2021)

Version 1.0.2 officially supports Python 3.9 and has dropped support for Python 3.5. The minimum version of the following dependencies were updated: `asteval` $\geq 0.9.21$, `numpy` ≥ 1.18 , and `scipy` ≥ 1.3 .

New features:

- added two-dimensional Gaussian lineshape and model (PR #642; @mpmdean)
- all built-in models are now registered in `lmfit.models.lmfit_models`; new `Model` class attribute `valid_forms` (PR #663; @rayosborn)
- added a `SineModel` (PR #676; @lneuhau)
- add the `run_mcmc_kwarg` argument to `Minimizer.emcee` to pass to the `emcee.EnsembleSampler.run_mcmc` function (PR #694; @rbnrvw)

Bug fixes:

- `ModelError.eval_uncertainty` should use provided `Parameters` (PR #646)
- center in lognormal model can be negative (Issue #644, PR #645; @YoshieraHuang)
- restore best-fit values after calculation of covariance matrix (Issue #655, PR #657)
- add helper-function `not_zero` to prevent `ZeroDivisionError` in lineshapes and use in exponential lineshape (Issue #631, PR #664; @s-weigand)
- save `last_internal_values` and use to restore internal values if fit is aborted (PR #667)
- dumping a fit using the `lbfgsb` method now works, convert bytes to string if needed (Issue #677, PR #678; @leonfoks)
- fix use of callable Jacobian for scalar methods (PR #681; @msttimberg)
- preserve float/int types when encoding for JSON (PR #696; @jedzill4)
- better support for saving/loading of ExpressionModels and assure that `init_params` and `init_fit` are set when loading a `ModelError` (PR #706)

Various:

- update minimum dependencies (PRs #688, #693)

- improvements in coding style, docstrings, CI, and test coverage (PRs #647, #649, #650, #653, #654; #685, #668, #689)
- fix typo in Oscillator (PR #658; @flothesof)
- add example using SymPy (PR #662)
- allow better custom pool for emcee() (Issue #666, PR #667)
- update NIST Strd reference functions and tests (PR #670)
- make building of documentation cross-platform (PR #673; @s-weigand)
- relax module name check in `test_check_ast_errors` for Python 3.9 (Issue #674, PR #675; @mwhudson)
- fix/update layout of documentation, now uses the sphinx13 theme (PR #687)
- fixed DeprecationWarnings reported by NumPy v1.2.0 (PR #699)
- increase value of `tiny` and check for it in bounded parameters to avoid “parameter not moving from initial value” (Issue #700, PR #701)
- add `max_nfev` to `basinhopping` and `brute` (now supported everywhere in lmfit) and set to more uniform default values (PR #701)
- use Azure Pipelines for CI, drop Travis (PRs #696 and #702)

12.10 Version 1.0.1 Release Notes

Version 1.0.1 is the last release that supports Python 3.5. All newer version will require 3.6+ so that we can use formatting-strings and rely on dictionaries being ordered.

New features:

- added thermal distribution model and lineshape (PR #620; @mpmdean)
- introduced a new argument `max_nfev` to uniformly specify the maximum number of function evaluations (PR #610) **Please note: all other arguments (e.g., ``maxfev``, ``maxiter``, ...) will no longer be passed to the underlying solver. A warning will be emitted stating that one should use ``max_nfev``.**
- the attribute `call_kws` was added to the `MinimizerResult` class and contains the keyword arguments that are supplied to the solver in SciPy.

Bug fixes:

- fixes to the `load` and `__setstate__` methods of the `Parameter` class
- fixed failure of `ModelError.dump()` due to missing attributes (Issue #611, PR #623; @mpmdean)
- `guess_from_peak` function now also works correctly with decreasing x-values or when using pandas (PRs #627 and #629; @mpmdean)
- the `Parameter.set()` method now correctly first updates the boundaries and then the value (Issue #636, PR #637; @arunpersaud)

Various:

- fixed typo for the use of expressions in the documentation (Issue #610; @jkrogager)
- removal of PY2-compatibility and unused code and improved test coverage (PRs #619, #631, and #633)
- removed deprecated `isParameter` function and automatic conversion of an `uncertainties` object (PR #626)
- inaccurate FWHM calculations were removed from built-in models, others labeled as estimates (Issue #616 and PR #630)

- corrected spelling mistake for the Doniach lineshape and model (Issue #634; @rayosborn)
- removed unsupported/untested code for IPython notebooks in lmfit/ui/*

12.11 Version 1.0.0 Release Notes

Version 1.0.0 supports Python 3.5, 3.6, 3.7, and 3.8

New features:

- no new features are introduced in 1.0.0.

Improvements:

- support for Python 2 and use of the `six` package are removed. (PR #612)

Various:

- documentation updates to clarify the use of `emcee`. (PR #614)

12.12 Version 0.9.15 Release Notes

Version 0.9.15 is the last release that supports Python 2.7; it now also fully supports Python 3.8.

New features, improvements, and bug fixes:

- move application of parameter bounds to setter instead of getter (PR #587)
- add support for non-array Jacobian types in `least_squares` (Issue #588, @ezwelty in PR #589)
- add more information (i.e., `acor` and `acceptance_fraction`) about `emcee` fit (@j-zimmermann in PR #593)
- “name” is now a required positional argument for `Parameter` class, update the magic methods (PR #595)
- fix `nvars` count and bound handling in confidence interval calculations (Issue #597, PR #598)
- support Python 3.8; requires `asteval >= 0.9.16` (PR #599)
- only support `emcee` version 3 (i.e., no `PTSampler` anymore) (PR #600)
- fix and refactor `prob_bunc` in confidence interval calculations (PR #604)
- fix adding `Parameters` with custom user-defined symbols (Issue #607, PR #608; thanks to @gbouvignies for the report)

Various:

- bump requirements to LTS version of SciPy/ NumPy and code clean-up (PR #591)
- documentation updates (PR #596, and others)
- improve test coverage and Travis CI updates (PR #595, and others)
- update pre-commit hooks and configuration in `setup.cfg`

To-be deprecated: - function `Parameter.isParameter` and conversion from `uncertainties.core.Variable` to value in `_getval` (PR #595)

12.13 Version 0.9.14 Release Notes

New features:

- the global optimizers `shgo` and `dual_annealing` (new in SciPy v1.2) are now supported (Issue #527; PRs #545 and #556)
- `eval` method added to the `Parameter` class (PR #550 by @zobristnicholas)
- avoid `ZeroDivisionError` in `printfuncs.params_html_table` (PR #552 by @aaristov and PR #559)
- add parallelization to `brute` method (PR #564, requires SciPy v1.3)

Bug fixes:

- consider only varying parameters when reporting potential issues with calculating errorbars (PR #549) and compare value to both `min` and `max` (PR #571)
- guard against division by zero in lineshape functions and `FWHM` and `height` expression calculations (PR #545)
- fix issues with restoring a saved `Model` (Issue #553; PR #554)
- always set `result.method` for `emcee` algorithm (PR #558)
- more careful adding of parameters to handle out-of-order constraint expressions (Issue #560; PR #561)
- make sure all parameters in `Model.guess()` use prefixes (PRs #567 and #569)
- use `inspect.signature` for PY3 to support wrapped functions (Issue #570; PR #576)
- fix `result.nfev`` for `brute` method when using parallelization (Issue #578; PR #579)

Various:

- remove “missing” in the `Model` class (replaced by `nan_policy`) and “drop” as option to `nan_policy` (replaced by `omit`) deprecated since 0.9 (PR #565).
- deprecate ‘`report_errors`’ in `printfuncs.py` (PR #571)
- updates to the documentation to use `jupyter-sphinx` to include examples/output (PRs #573 and #575)
- include a Gallery with examples in the documentation using `sphinx-gallery` (PR #574 and #583)
- improve test-coverage (PRs #571, #572 and #585)
- add/clarify warning messages when NaN values are detected (PR #586)
- several updates to docstrings (Issue #584; PR #583, and others)
- update pre-commit hooks and several docstrings

12.14 Version 0.9.13 Release Notes

New features:

- Clearer warning message in fit reports when uncertainties should but cannot be estimated, including guesses of which Parameters to examine (#521, #543)
- `SplitLorenztianModel` and `split_lorentzian` function (#523)
- HTML representations for `Parameter`, `MinimizerResult`, and `Model` so that they can be printed better with Jupyter (#524, #548)
- support parallelization for differential evolution (#526)

Bug fixes:

- delay import of matplotlib (and so, the selection of its backend) as late as possible (#528, #529)
- fix for saving, loading, and reloading ModelResults (#534)
- fix to leastsq to report the best-fit values, not the values tried last (#535, #536)
- fix synchronization of all parameter values on Model.guess() (#539, #542)
- improve deprecation warnings for outdated nan_policy keywords (#540)
- fix for edge case in gformat() (#547)

Project management:

- using pre-commit framework to improve and enforce coding style (#533)
- added code coverage report to github main page
- updated docs, github templates, added several tests.
- dropped support and testing for Python 3.4.

12.15 Version 0.9.12 Release Notes

Lmfit package is now licensed under BSD-3.

New features:

- SkewedVoigtModel was added as built-in model (Issue #493)
- Parameter uncertainties and correlations are reported for least_squares
- Plotting of complex-valued models is now handled in ModelResult class (PR #503)
- A model's independent variable is allowed to be an object (Issue #492)
- Added usersyms to Parameters() initialization to make it easier to add custom functions and symbols (Issue #507)
- the numdifftools package can be used to calculate parameter uncertainties and correlations for all solvers that do not natively support this (PR #506)
- emcee can now be used as method keyword-argument to Minimizer.minimize and minimize function, which allows for using emcee in the Model class (PR #512; see examples/example_emcee_with_Model.py)

(Bug)fixes:

- asteval errors are now flushed after raising (Issue #486)
- max_time and evaluation time for ExpressionModel increased to 1 hour (Issue #489)
- loading a saved ModelResult now restores all attributes (Issue #491)
- development versions of scipy and emcee are now supported (Issue #497 and PR #496)
- ModelResult.eval() do no longer overwrite the userkws dictionary (Issue #499)
- running the test suite requires pytest only (Issue #504)
- improved FWHM calculation for VoigtModel (PR #514)

12.16 Version 0.9.10 Release Notes

Two new global algorithms were added: basinhopping and AMPGO. Basinhopping wraps the method present in `scipy`, and more information can be found in the documentation ([basinhopping\(\)](#) and `scipy.optimize.basinhopping`). The Adaptive Memory Programming for Global Optimization (AMPGO) algorithm was adapted from Python code written by [Andrea Gavana](#). A more detailed explanation of the algorithm is available in the [AMPGO paper](#) and specifics for lmfit can be found in the `ampgo()` function.

Lmfit uses the external uncertainties (<https://github.com/lebigot/uncertainties>) package (available on PyPI), instead of distributing its own fork.

An `AbortFitException` is now raised when the fit is aborted by the user (i.e., by using `iter_cb`).

Bugfixes:

- all exceptions are allowed when trying to import matplotlib
- simplify and fix corner-case errors when testing closeness of large integers

12.17 Version 0.9.9 Release Notes

Lmfit now uses the asteval (<https://github.com/newville/asteval>) package instead of distributing its own copy. The minimum required asteval version is 0.9.12, which is available on PyPI. If you see import errors related to asteval, please make sure that you actually have the latest version installed.

12.18 Version 0.9.6 Release Notes

Support for SciPy 0.14 has been dropped: SciPy 0.15 is now required. This is especially important for lmfit maintenance, as it means we can now rely on SciPy having code for differential evolution and do not need to keep a local copy.

A brute force method was added, which can be used either with `Minimizer.brute()` or using the `method='brute'` option to `Minimizer.minimize()`. This method requires finite bounds on all varying parameters, or that parameters have a finite `brute_step` attribute set to specify the step size.

Custom cost functions can now be used for the scalar minimizers using the `reduce_fcn` option.

Many improvements to documentation and docstrings in the code were made. As part of that effort, all API documentation in this main Sphinx documentation now derives from the docstrings.

Uncertainties in the resulting best-fit for a model can now be calculated from the uncertainties in the model parameters.

Parameters have two new attributes: `brute_step`, to specify the step size when using the `brute` method, and `user_data`, which is unused but can be used to hold additional information the user may desire. This will be preserved on copy and pickling.

Several bug fixes and cleanups.

Versioneer was updated to 0.18.

Tests can now be run either with nose or pytest.

12.19 Version 0.9.5 Release Notes

Support for Python 2.6 and SciPy 0.13 has been dropped.

12.20 Version 0.9.4 Release Notes

Some support for the new `least_squares` routine from SciPy 0.17 has been added.

Parameters can now be used directly in floating point or array expressions, so that the Parameter value does not need `sigma = params['sigma'].value`. The older, explicit usage still works, but the docs, samples, and tests have been updated to use the simpler usage.

Support for Python 2.6 and SciPy 0.13 is now explicitly deprecated and will be dropped in version 0.9.5.

12.21 Version 0.9.3 Release Notes

Models involving complex numbers have been improved.

The `emcee` module can now be used for uncertainty estimation.

Many bug fixes, and an important fix for performance slowdown on getting parameter values.

ASV benchmarking code added.

12.22 Version 0.9.0 Release Notes

This upgrade makes an important, non-backward-compatible change to the way many fitting scripts and programs will work. Scripts that work with version 0.8.3 will not work with version 0.9.0 and vice versa. The change was not made lightly or without ample discussion, and is really an improvement. Modifying scripts that did work with 0.8.3 to work with 0.9.0 is easy, but needs to be done.

12.22.1 Summary

The upgrade from 0.8.3 to 0.9.0 introduced the `MinimizerResult` class (see [MinimizerResult – the optimization result](#)) which is now used to hold the return value from `minimize()` and `Minimizer.minimize()`. This returned object contains many goodness of fit statistics, and holds the optimized parameters from the fit. Importantly, the parameters passed into `minimize()` and `Minimizer.minimize()` are no longer modified by the fit. Instead, a copy of the passed-in parameters is made which is changed and returns as the `params` attribute of the returned `MinimizerResult`.

12.22.2 Impact

This upgrade means that a script that does:

```
my_pars = Parameters()
my_pars.add('amp', value=300.0, min=0)
my_pars.add('center', value=5.0, min=0, max=10)
my_pars.add('decay', value=1.0, vary=False)

result = minimize(objfunc, my_pars)
```

will still work, but that `my_pars` will **NOT** be changed by the fit. Instead, `my_pars` is copied to an internal set of parameters that is changed in the fit, and this copy is then put in `result.params`. To look at fit results, use `result.params`, not `my_pars`.

This has the effect that `my_pars` will still hold the starting parameter values, while all of the results from the fit are held in the `result` object returned by `minimize()`.

If you want to do an initial fit, then refine that fit to, for example, do a pre-fit, then refine that result different fitting method, such as:

```
result1 = minimize(objfunc, my_pars, method='nelder')
result1.params['decay'].vary = True
result2 = minimize(objfunc, result1.params, method='leastsq')
```

and have access to all of the starting parameters `my_pars`, the result of the first fit `result1`, and the result of the final fit `result2`.

12.22.3 Discussion

The main goal for making this change were to

1. give a better return value to `minimize()` and `Minimizer.minimize()` that can hold all of the information about a fit. By having the return value be an instance of the `MinimizerResult` class, it can hold an arbitrary amount of information that is easily accessed by attribute name, and even be given methods. Using objects is good!
2. To limit or even eliminate the amount of “state information” a `Minimizer` holds. By state information, we mean how much of the previous fit is remembered after a fit is done. Keeping (and especially using) such information about a previous fit means that a `Minimizer` might give different results even for the same problem if run a second time. While it’s desirable to be able to adjust a set of `Parameters` re-run a fit to get an improved result, doing this by changing an internal attribute (`Minimizer.params`) has the undesirable side-effect of not being able to “go back”, and makes it somewhat cumbersome to keep track of changes made while adjusting parameters and re-running fits.

EXAMPLES GALLERY

Below are examples of the different things you can do with lmfit. Click on any image to see the complete source code and output.

We encourage users (i.e., YOU) to submit user-guide-style, documented, and preferably self-contained examples of how you use lmfit for inclusion in this gallery! Please note that many of the examples below currently do *not* follow these guidelines yet.

13.1 Fit with Data in a pandas DataFrame

Simple example demonstrating how to read in the data using pandas and supply the elements of the DataFrame to lmfit.

```
import pandas as pd

from lmfit.models import LorentzianModel
```

read the data into a pandas DataFrame, and use the x and y columns:

```
dframe = pd.read_csv('peak.csv')

model = LorentzianModel()
params = model.guess(dframe['y'], x=dframe['x'])

result = model.fit(dframe['y'], params, x=dframe['x'])
```

and gives the fitting results:

```
print(result.fit_report())
```

```
[[Model]]
  Model(lorentzian)
[[Fit Statistics]]
  # fitting method    = leastsq
  # function evals   = 21
  # data points     = 101
  # variables       = 3
  chi-square        = 13.0737250
  reduced chi-square = 0.13340536
  Akaike info crit  = -200.496119
```

(continues on next page)

(continued from previous page)


```

Bayesian info crit = -192.650757
R-squared          = 0.98351484
[[Variables]]
  amplitude: 39.1530621 +/- 0.62389897 (1.59%) (init = 50.7825)
  center:    9.22379948 +/- 0.01835867 (0.20%) (init = 9.3)
  sigma:     1.15503770 +/- 0.02603721 (2.25%) (init = 1.3)
  fwhm:      2.31007541 +/- 0.05207442 (2.25%) == '2.0000000*sigma'
  height:    10.7899571 +/- 0.17160652 (1.59%) == '0.3183099*amplitude/max(1e-15,_
-sigma)'
[[Correlations]] (unreported correlations are < 0.100)
  C(amplitude, sigma) = +0.7087

```

and plot below:

```
result.plot_fit()
```



Total running time of the script: (0 minutes 0.290 seconds)

13.2 Using an ExpressionModel

ExpressionModels allow a model to be built from a user-supplied expression. See: https://lmfit.github.io/lmfit-py/builtin_models.html#user-defined-models

```
import matplotlib.pyplot as plt
import numpy as np

from lmfit.models import ExpressionModel
```

Generate synthetic data for the user-supplied model:

```
x = np.linspace(-10, 10, 201)
amp, cen, wid = 3.4, 1.8, 0.5

y = amp * np.exp(-(x-cen)**2 / (2*wid**2)) / (np.sqrt(2*np.pi)*wid)
np.random.seed(2021)
y = y + np.random.normal(size=x.size, scale=0.01)
```

Define the ExpressionModel and perform the fit:

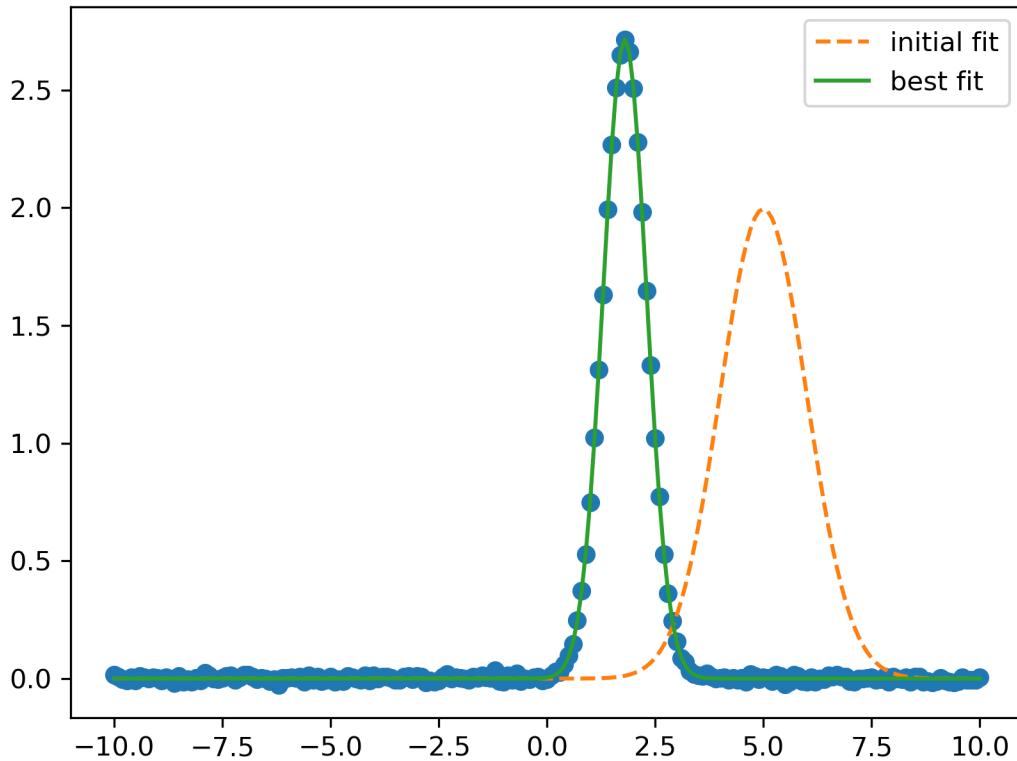
```
gmod = ExpressionModel("amp * exp(-(x-cen)**2 / (2*wid**2)) / (sqrt(2*pi)*wid)")
result = gmod.fit(y, x=x, amp=5, cen=5, wid=1)
```

this results in the following output:

```
print(result.fit_report())
```

```
[[Model]]
<lmfit.ExpressionModel('amp * exp(-(x-cen)**2 / (2*wid**2)) / (sqrt(2*pi)*wid)')>
[[Fit Statistics]]
# fitting method   = leastsq
# function evals  = 52
# data points     = 201
# variables       = 3
chi-square        = 0.01951689
reduced chi-square = 9.8570e-05
Akaike info crit  = -1851.19580
Bayesian info crit = -1841.28588
R-squared          = 0.99967271
[[Variables]]
amp: 3.40625133 +/- 0.00512077 (0.15%) (init = 5)
cen: 1.80121155 +/- 8.6847e-04 (0.05%) (init = 5)
wid: 0.50029616 +/- 8.6848e-04 (0.17%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(amp, wid) = +0.5774
```

```
plt.plot(x, y, 'o')
plt.plot(x, result.init_fit, '--', label='initial fit')
plt.plot(x, result.best_fit, '-', label='best fit')
plt.legend()
```



Total running time of the script: (0 minutes 0.195 seconds)

13.3 Fit Using Inequality Constraint

Sometimes specifying boundaries using `min` and `max` are not sufficient, and more complicated (inequality) constraints are needed. In the example below the center of the Lorentzian peak is constrained to be between 0-5 away from the center of the Gaussian peak.

See also: <https://lmfit.github.io/lmfit-py/constraints.html#using-inequality-constraints>

```
import matplotlib.pyplot as plt
import numpy as np

from lmfit import Minimizer, create_params, report_fit
from lmfit.lineshapes import gaussian, lorentzian

def residual(pars, x, data):
    model = (gaussian(x, pars['amp_g'], pars['cen_g'], pars['wid_g']) +
             lorentzian(x, pars['amp_l'], pars['cen_l'], pars['wid_l']))
    return model - data
```

Generate the simulated data using a Gaussian and Lorentzian lineshape:

```
np.random.seed(0)
x = np.linspace(0, 20.0, 601)

data = (gaussian(x, 21, 6.1, 1.2) + lorentzian(x, 10, 9.6, 1.3) +
        np.random.normal(scale=0.1, size=x.size))
```

Create the fitting parameters and set an inequality constraint for `cen_l`. First, we add a new fitting parameter `peak_split`, which can take values between 0 and 5. Afterwards, we constrain the value for `cen_l` using the expression to be '`peak_split+cen_g`':

```
pfit = create_params(amp_g=10, cen_g=5, wid_g=1, amp_l=10,
                     peak_split=dict(value=2.5, min=0, max=5),
                     cen_l=dict(expr='peak_split+cen_g'),
                     wid_l=dict(expr='wid_g'))

mini = Minimizer(residual, pfit, fcn_args=(x, data))
out = mini.leastsq()
best_fit = data + out.residual
```


Performing a fit, here using the `leastsq` algorithm, gives the following fitting results:

```
report_fit(out.params)
```

```
[[Variables]]
amp_g: 21.2722842 +/- 0.05138772 (0.24%) (init = 10)
cen_g: 6.10496396 +/- 0.00334613 (0.05%) (init = 5)
wid_g: 1.21434954 +/- 0.00327317 (0.27%) (init = 1)
amp_l: 9.46504173 +/- 0.05445415 (0.58%) (init = 10)
peak_split: 3.52163544 +/- 0.01004618 (0.29%) (init = 2.5)
cen_l: 9.62659940 +/- 0.01066172 (0.11%) == 'peak_split+cen_g'
wid_l: 1.21434954 +/- 0.00327317 (0.27%) == 'wid_g'
[[Correlations]] (unreported correlations are < 0.100)
C(amp_g, wid_g) = +0.6199
C(amp_g, peak_split) = +0.3796
C(wid_g, peak_split) = +0.3445
C(amp_g, amp_l) = -0.2951
C(cen_g, amp_l) = -0.2761
C(amp_g, cen_g) = +0.1936
C(wid_g, amp_l) = -0.1651
C(cen_g, wid_g) = +0.1546
```

and figure:

```
plt.plot(x, data, 'o')
plt.plot(x, best_fit, '--', label='best fit')
plt.legend()
plt.show()
```



Total running time of the script: (0 minutes 0.199 seconds)

13.4 Fit Using differential_evolution Algorithm

This example compares the `leastsq` and `differential_evolution` algorithms on a fairly simple problem.

```
import matplotlib.pyplot as plt
import numpy as np

import lmfit

def resid(params, x, ydata):
    decay = params['decay'].value
    offset = params['offset'].value
    omega = params['omega'].value
    amp = params['amp'].value

    y_model = offset + amp * np.sin(x*omega) * np.exp(-x/decay)
    return y_model - ydata
```

Generate synthetic data and set-up Parameters with initial values/boundaries:

```

decay = 5
offset = 1.0
amp = 2.0
omega = 4.0

np.random.seed(2)
x = np.linspace(0, 10, 101)
y = offset + amp*np.sin(omega*x) * np.exp(-x/decay)
yn = y + np.random.normal(size=y.size, scale=0.450)

params = lmfit.Parameters()
params.add('offset', 2.0, min=0, max=10.0)
params.add('omega', 3.3, min=0, max=10.0)
params.add('amp', 2.5, min=0, max=10.0)
params.add('decay', 1.0, min=0, max=10.0)

```

Perform the fits and show fitting results and plot:

```

o1 = lmfit.minimize(resid, params, args=(x, yn), method='leastsq')
print("# Fit using leastsq:")
lmfit.report_fit(o1)

```

```

# Fit using leastsq:
[[Fit Statistics]]
    # fitting method    = leastsq
    # function evals   = 65
    # data points      = 101
    # variables        = 4
    chi-square         = 21.7961792
    reduced chi-square = 0.22470288
    Akaike info crit  = -146.871969
    Bayesian info crit = -136.411487
[[Variables]]
    offset: 0.96333089 +/- 0.04735890 (4.92%) (init = 2)
    omega:  3.98700839 +/- 0.02079709 (0.52%) (init = 3.3)
    amp:    1.80253587 +/- 0.19401928 (10.76%) (init = 2.5)
    decay:  5.76279753 +/- 1.04073348 (18.06%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
    C(amp, decay) = -0.7550

```

```

o2 = lmfit.minimize(resid, params, args=(x, yn), method='differential_evolution')
print("\n\n# Fit using differential_evolution:")
lmfit.report_fit(o2)

```

```

# Fit using differential_evolution:
[[Fit Statistics]]
    # fitting method    = differential_evolution
    # function evals   = 1720
    # data points      = 101
    # variables        = 4
    chi-square         = 21.7961792
    reduced chi-square = 0.22470288

```

(continues on next page)

(continued from previous page)

```

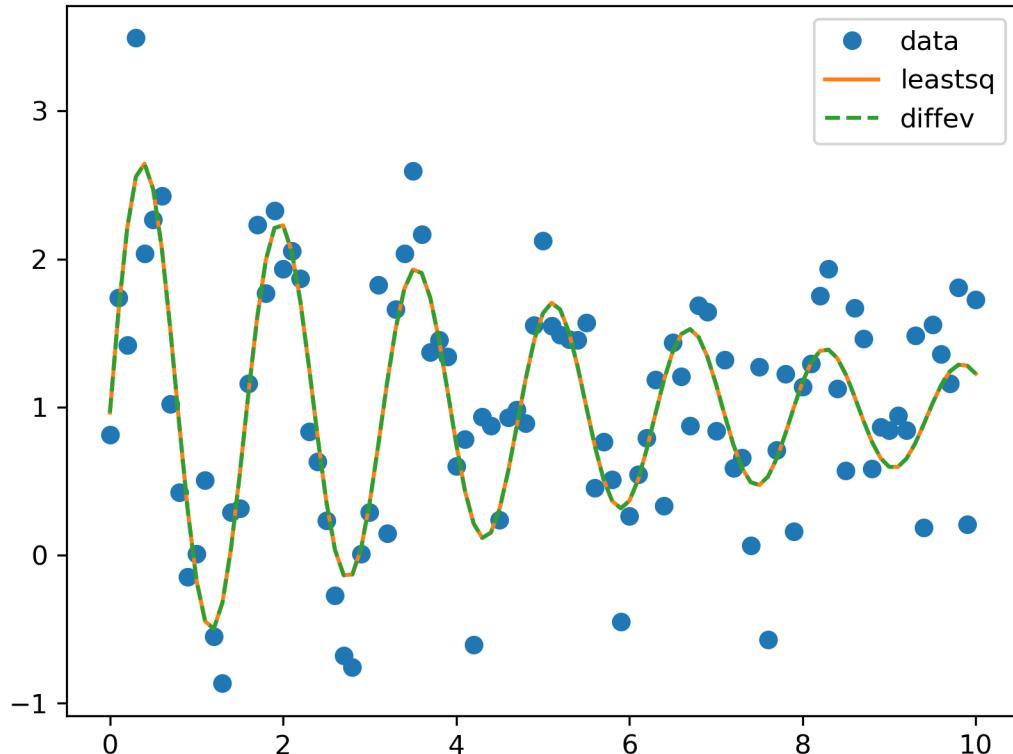
Akaike info crit   = -146.871969
Bayesian info crit = -136.411487
[[Variables]]
  offset: 0.96333168 +/- 0.04735902 (4.92%) (init = 2)
  omega:  3.98700858 +/- 0.02121811 (0.53%) (init = 3.3)
  amp:    1.80253517 +/- 0.19022457 (10.55%) (init = 2.5)
  decay:   5.76281259 +/- 1.00451713 (17.43%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
  C(amp, decay) = -0.7434

```

```

plt.plot(x, yn, 'o', label='data')
plt.plot(x, yn+o1.residual, '--', label='leastsq')
plt.plot(x, yn+o2.residual, '---', label='diffev')
plt.legend()

```



Total running time of the script: (0 minutes 0.377 seconds)

13.5 Fit Using Bounds

A major advantage of using lmfit is that one can specify boundaries on fitting parameters, even if the underlying algorithm in SciPy does not support this. For more information on how this is implemented, please refer to: <https://lmfit.github.io/lmfit-py/bounds.html>

The example below shows how to set boundaries using the `min` and `max` attributes to fitting parameters.

```
import matplotlib.pyplot as plt
from numpy import exp, linspace, pi, random, sign, sin

from lmfit import create_params, minimize
from lmfit.printfuncs import report_fit
```

create the ‘true’ Parameter values and residual function:

```
p_true = create_params(amp=14.0, period=5.4321, shift=0.12345, decay=0.010)

def residual(pars, x, data=None):
    argu = (x * pars['decay'])**2
    shift = pars['shift']
    if abs(shift) > pi/2:
        shift = shift - sign(shift)*pi
    model = pars['amp'] * sin(shift + x/pars['period']) * exp(-argu)
    if data is None:
        return model
    return model - data
```

Generate synthetic data and initialize fitting Parameters:

```
random.seed(0)
x = linspace(0, 250, 1500)
noise = random.normal(scale=2.8, size=x.size)
data = residual(p_true, x) + noise

fit_params = create_params(amp=dict(value=13, max=20, min=0),
                           period=dict(value=2, max=10),
                           shift=dict(value=0, max=pi/2., min=-pi/2.),
                           decay=dict(value=0.02, max=0.1, min=0))
```

Perform the fit and show the results:

```
out = minimize(residual, fit_params, args=(x,), kws={'data': data})
fit = residual(out.params, x)
```

```
report_fit(out, modelpars=p_true, correl_mode='table')
```


```
[[Fit Statistics]]
# fitting method    = leastsq
# function evals   = 79
# data points      = 1500
# variables        = 4
```

(continues on next page)

(continued from previous page)

```
chi-square      = 11301.3646
reduced chi-square = 7.55438813
Akaike info crit  = 3037.18756
Bayesian info crit = 3058.44044
[[Variables]]
amp:    13.8904759 +/- 0.24410753 (1.76%) (init = 13), model_value = 14
period: 5.44026387 +/- 0.01416106 (0.26%) (init = 2), model_value = 5.4321
shift:   0.12464389 +/- 0.02414210 (19.37%) (init = 0), model_value = 0.12345
decay:   0.00996363 +/- 2.0275e-04 (2.03%) (init = 0.02), model_value = 0.01
[[Correlations]]
+-----+-----+-----+-----+-----+
| Variable | amp      | period    | shift     | decay    |
+-----+-----+-----+-----+-----+
| amp      | +1.0000  | -0.0700  | -0.0870  | +0.5757  |
| period    | -0.0700  | +1.0000  | +0.7999  | -0.0404  |
| shift     | -0.0870  | +0.7999  | +1.0000  | -0.0502  |
| decay     | +0.5757  | -0.0404  | -0.0502  | +1.0000  |
+-----+-----+-----+-----+-----+
```

```
plt.plot(x, data, 'o', label='data')
plt.plot(x, fit, label='best fit')
plt.legend()
plt.show()
```



Total running time of the script: (0 minutes 0.208 seconds)

13.6 Fit with Algebraic Constraint

Example on how to use algebraic constraints using the `expr` attribute.

```
import matplotlib.pyplot as plt
from numpy import linspace, random

from lmfit.lineshapes import gaussian, lorentzian
from lmfit.models import GaussianModel, LinearModel, LorentzianModel

random.seed(0)
x = linspace(0.0, 20.0, 601)

data = (gaussian(x, amplitude=21, center=8.1, sigma=1.2) +
        lorentzian(x, amplitude=10, center=9.6, sigma=2.4) +
        0.01 + x*0.05 + random.normal(scale=0.23, size=x.size))

model = GaussianModel(prefix='g_') + LorentzianModel(prefix='l_') + LinearModel(prefix=
    ↪'line_')
```

(continues on next page)

(continued from previous page)

```
params = model.make_params(g_amplitude=10, g_center=9, g_sigma=1,
                           line_slope=0, line_intercept=0)

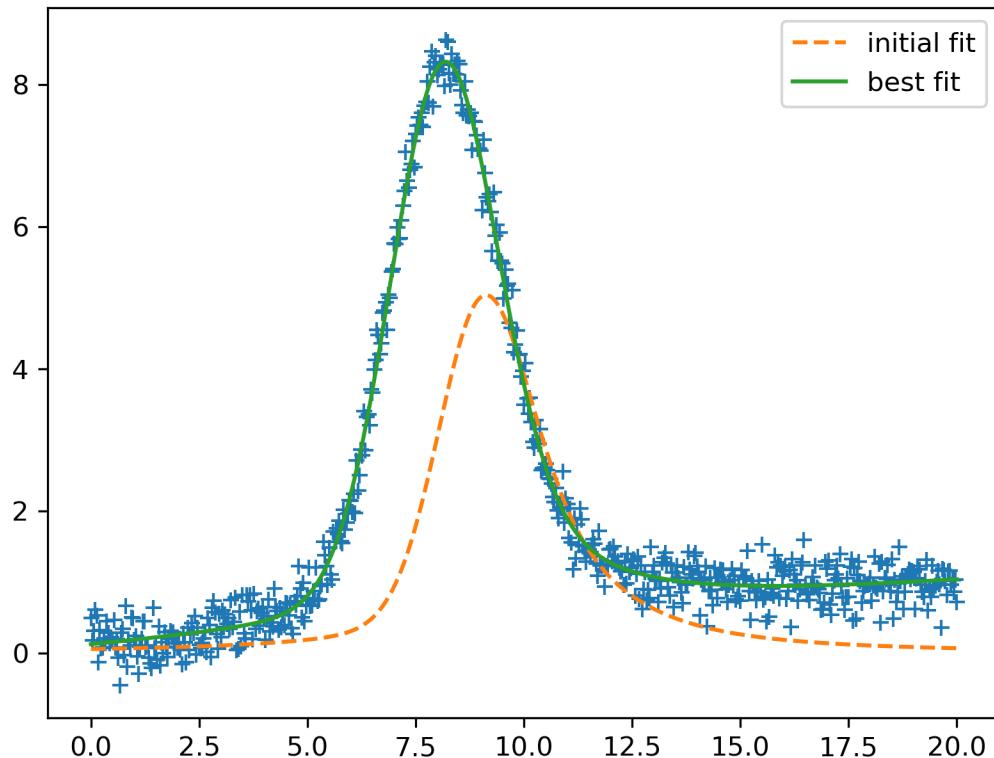
params.add(name='total_amplitude', value=20)
params.set(l_amplitude=dict(expr='total_amplitude - g_amplitude'))
params.set(l_center=dict(expr='1.5+g_center'))
params.set(l_sigma=dict(expr='2*g_sigma'))

data_uncertainty = 0.021 # estimate of data error (for all data points)

init = model.eval(params, x=x)
result = model.fit(data, params, x=x, weights=1.0/data_uncertainty)

print(result.fit_report())

plt.plot(x, data, '+')
plt.plot(x, init, '--', label='initial fit')
plt.plot(x, result.best_fit, '-', label='best fit')
plt.legend()
plt.show()
```



```

[[Model]]
  ((Model(gaussian, prefix='g_') + Model(lorentzian, prefix='l_')) + Model(linear,_
prefix='line_'))
[[Fit Statistics]]
# fitting method = leastsq
# function evals = 65
# data points = 601
# variables = 6
chi-square = 71878.3055
reduced chi-square = 120.803875
Akaike info crit = 2887.26503
Bayesian info crit = 2913.65660
R-squared = 0.98977025
[[Variables]]
g_amplitude: 21.1877635 +/- 0.32192128 (1.52%) (init = 10)
g_center: 8.11125903 +/- 0.01162987 (0.14%) (init = 9)
g_sigma: 1.20925819 +/- 0.01170853 (0.97%) (init = 1)
l_amplitude: 9.41261441 +/- 0.61672968 (6.55%) == 'total_amplitude - g_amplitude'
l_center: 9.61125903 +/- 0.01162987 (0.12%) == '1.5+g_center'
l_sigma: 2.41851637 +/- 0.02341707 (0.97%) == '2*g_sigma'
line_slope: 0.04615727 +/- 0.00170178 (3.69%) (init = 0)
line_intercept: 0.05128584 +/- 0.02448063 (47.73%) (init = 0)
g_fwhm: 2.84758536 +/- 0.02757149 (0.97%) == '2.3548200*g_sigma'
g_height: 6.98998378 +/- 0.05837066 (0.84%) == '0.3989423*g_amplitude/max(1e-
15, g_sigma)'
l_fwhm: 4.83703275 +/- 0.04683414 (0.97%) == '2.0000000*l_sigma'
l_height: 1.23882905 +/- 0.08992735 (7.26%) == '0.3183099*l_amplitude/max(1e-
15, l_sigma)'
total_amplitude: 30.6003779 +/- 0.36481425 (1.19%) (init = 20)
[[Correlations]] (unreported correlations are < 0.100)
C(g_amplitude, g_sigma) = +0.8662
C(g_amplitude, g_center) = +0.7496
C(line_slope, line_intercept) = -0.7144
C(g_center, total_amplitude) = -0.6952
C(g_center, g_sigma) = +0.6227
C(g_amplitude, total_amplitude) = -0.6115
C(line_intercept, total_amplitude) = -0.5883
C(g_sigma, total_amplitude) = -0.4115
C(g_center, line_intercept) = +0.3868
C(g_amplitude, line_intercept) = +0.1834
C(g_amplitude, line_slope) = +0.1825
C(g_sigma, line_slope) = +0.1739

```

Total running time of the script: (0 minutes 0.227 seconds)

13.7 Fit Specifying Different Reduce Function

The `reduce_fcn` specifies how to convert a residual array to a scalar value for the scalar minimizers. The default value is `None` (i.e., “sum of squares of residual”) - alternatives are: `negentropy`, `neglogcauchy`, or a user-specified callable. For more information please refer to: <https://lmfit.github.io/lmfit-py/fitting.html#using-the-minimizer-class>

Here, we use as an example the Student’s t log-likelihood for robust fitting of data with outliers.

```
import matplotlib.pyplot as plt
import numpy as np

import lmfit

def resid(params, x, ydata):
    decay = params['decay'].value
    offset = params['offset'].value
    omega = params['omega'].value
    amp = params['amp'].value

    y_model = offset + amp * np.sin(x*omega) * np.exp(-x/decay)
    return y_model - ydata
```

Generate synthetic data with noise/outliers and initialize fitting Parameters:

```
decay = 5
offset = 1.0
amp = 2.0
omega = 4.0

np.random.seed(2)
x = np.linspace(0, 10, 101)
y = offset + amp * np.sin(omega*x) * np.exp(-x/decay)
yn = y + np.random.normal(size=y.size, scale=0.250)

outliers = np.random.randint(int(len(x)/3.0), len(x), int(len(x)/12))
yn[outliers] += 5*np.random.random(len(outliers))

params = lmfit.create_params(offset=2.0, omega=3.3, amp=2.5,
                             decay=dict(value=1, min=0))
```

Perform fits using the L-BFGS-B method with different `reduce_fcn`:

```
method = 'L-BFGS-B'
o1 = lmfit.minimize(resid, params, args=(x, yn), method=method)
print("# Fit using sum of squares:\n")
lmfit.report_fit(o1)
```

```
# Fit using sum of squares:
```

```
[[Fit Statistics]]
# fitting method = L-BFGS-B
```

(continues on next page)

(continued from previous page)

```

# function evals      = 130
# data points        = 101
# variables          = 4
chi-square           = 32.1674767
reduced chi-square  = 0.33162347
Akaike info crit   = -107.560626
Bayesian info crit = -97.1001440
[[Variables]]
offset: 1.10392445 +/- 0.05751441 (5.21%) (init = 2)
omega: 3.97313428 +/- 0.02073920 (0.52%) (init = 3.3)
amp: 1.69977024 +/- 0.21587469 (12.70%) (init = 2.5)
decay: 7.65901859 +/- 1.87209333 (24.44%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(amp, decay) = -0.7733

```

```

o2 = lmfit.minimize(resid, params, args=(x, yn), method=method,
                     reduce_fcn='neglogcauchy')
print("\n\n# Robust Fit, using log-likelihood with Cauchy PDF:\n")
lmfit.report_fit(o2)

```

```
# Robust Fit, using log-likelihood with Cauchy PDF:
```

```


[[Fit Statistics]]
# fitting method      = L-BFGS-B
# function evals      = 135
# data points         = 101
# variables          = 4
chi-square            = 33.5081285
reduced chi-square   = 0.34544462
Akaike info crit    = -103.436579
Bayesian info crit = -92.9760969
[[Variables]]
offset: 1.02005972 +/- 0.06642640 (6.51%) (init = 2)
omega: 3.98224425 +/- 0.02898701 (0.73%) (init = 3.3)
amp: 1.83231335 +/- 0.27241921 (14.87%) (init = 2.5)
decay: 5.77328069 +/- 1.45141325 (25.14%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(amp, decay) = -0.7584
C(offset, amp) = -0.1067

```

```

plt.plot(x, y, 'o', label='true function')
plt.plot(x, yn, '--*', label='with noise+outliers')
plt.plot(x, yn+o1.residual, '-', label='sum of squares fit')
plt.plot(x, yn+o2.residual, '-', label='robust fit')
plt.legend()
plt.show()

```



Total running time of the script: (0 minutes 0.258 seconds)

13.8 Building a lmfit model with SymPy

Sympy is a Python library for symbolic mathematics. It can be very useful to build a model with SymPy and then apply that model to the data with lmfit. This example shows how to do that. Please note that this example requires both the sympy and matplotlib packages.

```
import matplotlib.pyplot as plt
import numpy as np
import sympy
from sympy.parsing import sympy_parser

import lmfit
```

Instead of creating the SymPy symbols explicitly and building an expression with them, we will use the SymPy parser.

```
gauss_peak1 = sympy_parser.parse_expr('A1*exp(-(x-xc1)**2/(2*sigma1**2))')
gauss_peak2 = sympy_parser.parse_expr('A2*exp(-(x-xc2)**2/(2*sigma2**2))')
exp_back = sympy_parser.parse_expr('B*exp(-x/xw)')

model_list = sympy.Array((gauss_peak1, gauss_peak2, exp_back))
model = sum(model_list)
```

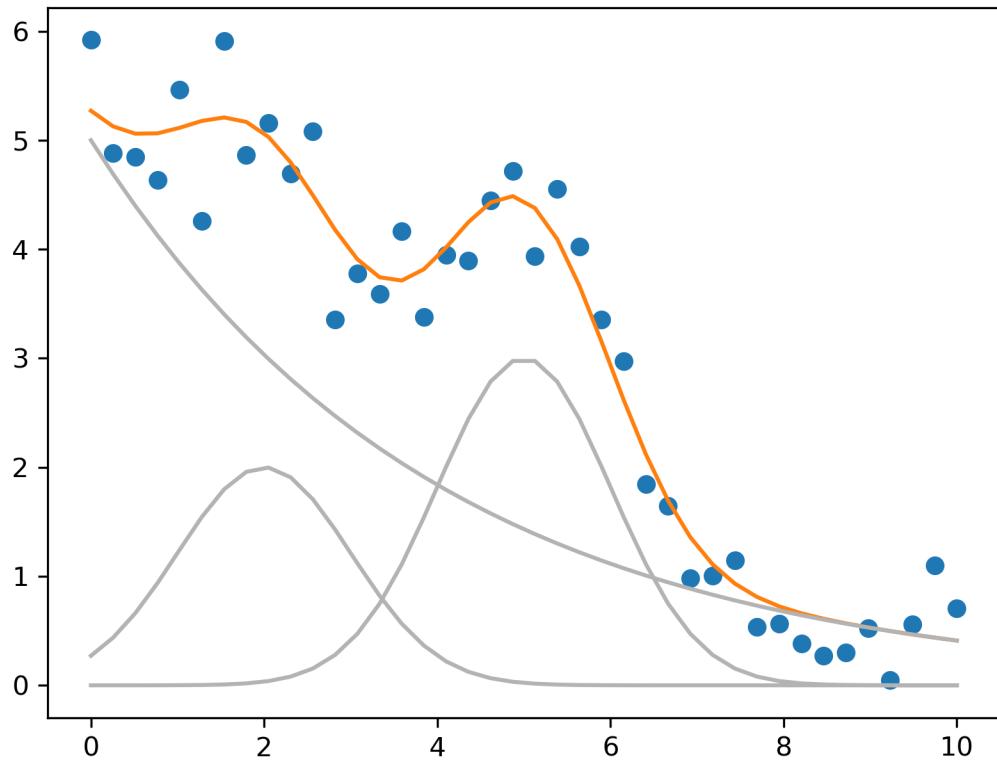
We are using SymPy's lambdify function to make a function from the model expressions. We then use these functions to generate some fake data.

```
model_list_func = sympy.lambdify(list(model_list.free_symbols), model_list)
model_func = sympy.lambdify(list(model.free_symbols), model)
```

Generate synthetic data with noise and plot the data.

```
np.random.seed(1)
x = np.linspace(0, 10, 40)
param_values = dict(x=x, A1=2, sigma1=1, sigma2=1, A2=3, xc1=2, xc2=5, xw=4, B=5)
y = model_func(**param_values)
yi = model_list_func(**param_values)
yn = y + np.random.randn(y.size)*0.4


plt.plot(x, yn, 'o')
plt.plot(x, y)
for c in yi:
    plt.plot(x, c, color='0.7')
```



Next, we will just create a lmfit model from the function and fit the data.

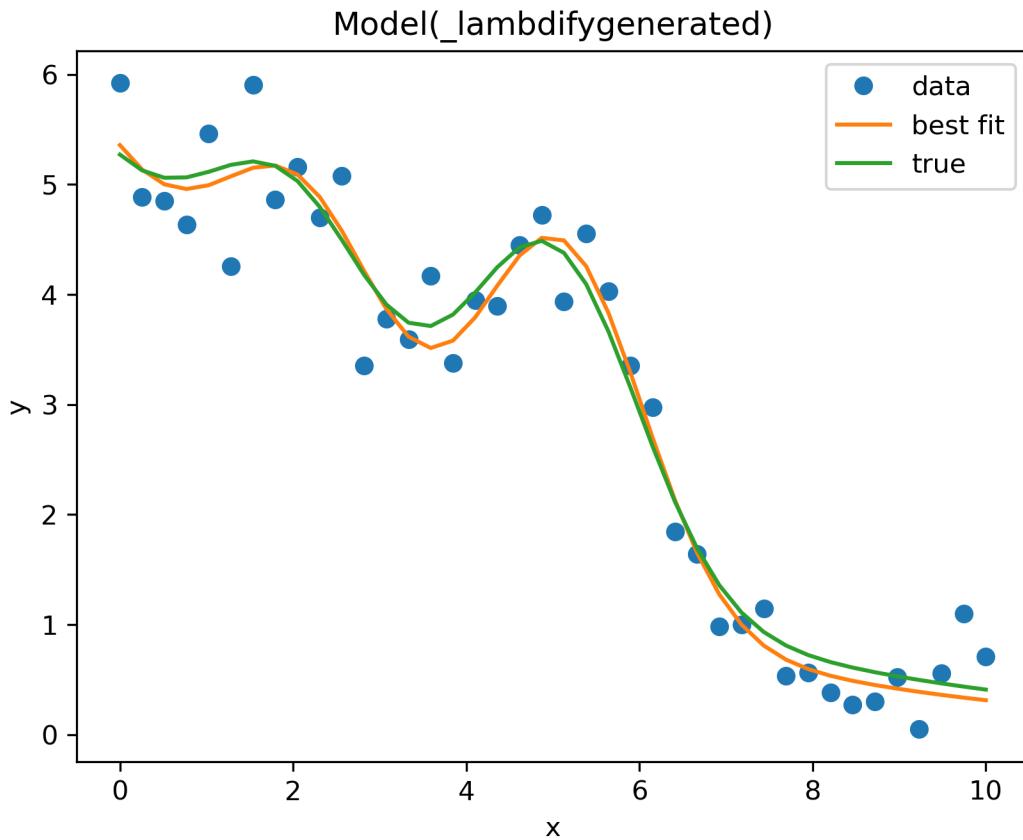
```
lm_mod = lmfit.Model(model_func, independent_vars=('x',))
res = lm_mod.fit(data=yn, **param_values)
```

```
res.plot_fit()
plt.plot(x, y, label='true')
plt.legend()
```



The nice thing of using SymPy is that we can easily modify our fit function. Let's assume we know that the width of both Gaussians are identical. Similarly, we assume that the ratio between both Gaussians is fixed to 3:2 for some reason. Both can be expressed by just substituting the variables.


```
model2 = model.subs('sigma2', 'sigma1').subs('A2', '3/2*A1')
model2_func = sympy.lambdify(list(model2.free_symbols), model2)
lm_mod = lmfit.Model(model2_func, independent_vars=['x'])
param2_values = dict(x=x, A1=2, sigma1=1, xc1=2, xc2=5, xw=4, B=5)
res2 = lm_mod.fit(data=yn, **param2_values)
res2.plot_fit()
plt.plot(x, y, label='true')
plt.legend()
plt.show()
```



Total running time of the script: (0 minutes 1.125 seconds)

13.9 Fit comparing leastsq and basin hopping, or other methods

This example compares the `leastsq` and `basinhopping` algorithms on a decaying sine wave. Note that this can be used to compare other fitting algorithms too.



```

# Fit using leastsq:
[[Model]]
  Model(sine_decay)
[[Fit Statistics]]
  # fitting method   = leastsq
  # function evals  = 31
  # data points    = 201
  # variables      = 4
  chi-square        = 37.4526504
  reduced chi-square = 0.19011498
  Akaike info crit  = -329.725713
  Bayesian info crit = -316.512493
  R-squared         = 0.97868751
[[Variables]]
  amplitude: 12.4411385 +/- 0.18965087 (1.52%) (init = 10)
  frequency: 1.99852115 +/- 0.00324489 (0.16%) (init = 2)
  decay:     4.54866058 +/- 0.09723868 (2.14%) (init = 2)
  offset:    1.25370110 +/- 0.03107114 (2.48%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
  C(amplitude, decay) = -0.7241
  C(amplitude, offset) = -0.1418
#####

```

(continues on next page)

(continued from previous page)

```

# Fit using basinhopping:
[[Model]]
  Model(sine_decay)
[[Fit Statistics]]
  # fitting method = basinhopping
  # function evals = 24924
  # data points   = 201
  # variables     = 4
  chi-square      = 37.4526504
  reduced chi-square = 0.19011498
  Akaike info crit = -329.725713
  Bayesian info crit = -316.512493
  R-squared        = 0.97868751
[[Variables]]
  amplitude: 12.4411365 +/- 0.18862101 (1.52%) (init = 10)
  frequency: 1.99852108 +/- 0.00326218 (0.16%) (init = 2)
  decay:     4.54866141 +/- 0.09621984 (2.12%) (init = 2)
  offset:    1.25370114 +/- 0.03106851 (2.48%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
  C(amplitude, decay) = -0.7205
  C(amplitude, offset) = -0.1412

```

```

import sys

import matplotlib.pyplot as plt
import numpy as np

import lmfit

VALID_METHODS = ['least_squares', 'differential_evolution', 'brute',
                 'basinhopping', 'ampgo', 'nelder', 'lbfgsb', 'powell', 'cg',
                 'newton', 'cobyla', 'bfgs', 'tnc', 'trust-ncg', 'trust-exact',
                 'trust-krylov', 'trust-constr', 'dogleg', 'slsqp', 'emcee',
                 'shgo', 'dual_annealing']

def sine_decay(x, amplitude, frequency, decay, offset):
    return offset + amplitude * np.sin(x*frequency) * np.exp(-x/decay)

x = np.linspace(0, 20, 201)
np.random.seed(2)

ydat = sine_decay(x, 12.5, 2.0, 4.5, 1.25) + np.random.normal(size=len(x), scale=0.40)

model = lmfit.Model(sine_decay)
params = model.make_params(amplitude={'value': 10, 'min': 0, 'max': 1000},

```

(continues on next page)

(continued from previous page)

```

frequency={'value': 2.0, 'min': 0, 'max': 6.0},
decay={'value': 2.0, 'min': 0.001, 'max': 12},
offset=1.0)

# fit with leastsq
result0 = model.fit(ydat, params, x=x, method='leastsq')
print("# Fit using leastsq:")
print(result0.fit_report())

method2 = 'basinhopping'
if len(sys.argv) > 1 and sys.argv[1] in VALID_METHODS:
    method2 = sys.argv[1]

# fit with other method
result = model.fit(ydat, params, x=x, method=method2)
print(f"\n#####\n# Fit using {method2}:")
print(result.fit_report())

# plot comparison
plt.plot(x, ydat, 'o', label='data')
plt.plot(x, result0.best_fit, '+', label='leastsq')
plt.plot(x, result.best_fit, '--', label=method2)
plt.legend()
plt.show()

```

Total running time of the script: (0 minutes 3.289 seconds)

13.10 Fit Multiple Data Sets

Fitting multiple (simulated) Gaussian data sets simultaneously.

All minimizers require the residual array to be one-dimensional. Therefore, in the objective function we need to flatten the array before returning it.

TODO: this could/should be using the Model interface / built-in models!

```

import matplotlib.pyplot as plt
import numpy as np

from lmfit import Parameters, minimize, report_fit

def gauss(x, amp, cen, sigma):
    """Gaussian lineshape."""
    return amp * np.exp(-(x-cen)**2 / (2.*sigma**2))

def gauss_dataset(params, i, x):
    """Calculate Gaussian lineshape from parameters for data set."""
    amp = params[f'amp_{i+1}']

```

(continues on next page)

(continued from previous page)

```

cen = params[f'cen_{i+1}']
sig = params[f'sig_{i+1}']
return gauss(x, amp, cen, sig)

def objective(params, x, data):
    """Calculate total residual for fits of Gaussians to several data sets."""
    ndata, _ = data.shape
    resid = 0.0*data[:]

    # make residual per data set
    for i in range(ndata):
        resid[i, :] = data[i, :] - gauss_dataset(params, i, x)

    # now flatten this to a 1D array, as minimize() needs
    return resid.flatten()

```

Create five simulated Gaussian data sets

```

np.random.seed(2021)
x = np.linspace(-1, 2, 151)
data = []
for _ in np.arange(5):
    amp = 0.60 + 9.50*np.random.rand()
    cen = -0.20 + 1.20*np.random.rand()
    sig = 0.25 + 0.03*np.random.rand()
    dat = gauss(x, amp, cen, sig) + np.random.normal(size=x.size, scale=0.1)
    data.append(dat)
data = np.array(data)

```

Create five sets of fitting parameters, one per data set

```

fit_params = Parameters()
for iy, y in enumerate(data):
    fit_params.add(f'amp_{iy+1}', value=0.5, min=0.0, max=200)
    fit_params.add(f'cen_{iy+1}', value=0.4, min=-2.0, max=2.0)
    fit_params.add(f'sig_{iy+1}', value=0.3, min=0.01, max=3.0)

```

Constrain the values of sigma to be the same for all peaks by assigning sig_2, ..., sig_5 to be equal to sig_1.

```

for iy in (2, 3, 4, 5):
    fit_params[f'sig_{iy}'].expr = 'sig_1'

```

Run the global fit and show the fitting result

```

out = minimize(objective, fit_params, args=(x, data))
report_fit(out.params)

```

```

[[Variables]]
amp_1: 6.32742010 +/- 0.02279089 (0.36%) (init = 0.5)
cen_1: 0.68049261 +/- 0.00126458 (0.19%) (init = 0.4)
sig_1: 0.25755570 +/- 4.9426e-04 (0.19%) (init = 0.3)
amp_2: 6.98604753 +/- 0.02296733 (0.33%) (init = 0.5)

```

(continues on next page)

(continued from previous page)

```

cen_2: 0.50433700 +/- 0.00114536 (0.23%) (init = 0.4)
sig_2: 0.25755570 +/- 4.9426e-04 (0.19%) == 'sig_1'
amp_3: 7.11643510 +/- 0.02300415 (0.32%) (init = 0.5)
cen_3: -0.08260274 +/- 0.00112437 (1.36%) (init = 0.4)
sig_3: 0.25755570 +/- 4.9426e-04 (0.19%) == 'sig_1'
amp_4: 6.10197422 +/- 0.02273421 (0.37%) (init = 0.5)
cen_4: 0.07386098 +/- 0.00131130 (1.78%) (init = 0.4)
sig_4: 0.25755570 +/- 4.9426e-04 (0.19%) == 'sig_1'
amp_5: 9.23910555 +/- 0.02368872 (0.26%) (init = 0.5)
cen_5: 0.34443083 +/- 8.6605e-04 (0.25%) (init = 0.4)
sig_5: 0.25755570 +/- 4.9426e-04 (0.19%) == 'sig_1'
[[Correlations]] (unreported correlations are < 0.100)
C(sig_1, amp_5) = -0.3742
C(sig_1, amp_3) = -0.2968
C(sig_1, amp_2) = -0.2919
C(amp_1, sig_1) = -0.2664
C(sig_1, amp_4) = -0.2575
C(amp_3, amp_5) = +0.1111
C(amp_2, amp_5) = +0.1092


```

Plot the data sets and fits

```

plt.figure()
for i in range(5):
    y_fit = gauss_dataset(out.params, i, x)
    plt.plot(x, data[i, :], 'o', x, y_fit, '-')

```



Total running time of the script: (0 minutes 0.229 seconds)

13.11 Fit using the Model interface

This notebook shows a simple example of using the `lmfit.Model` class. For more information please refer to: <https://lmfit.github.io/lmfit-py/model.html#the-model-class>.

```
import numpy as np
from pandas import Series

from lmfit import Model, Parameter, report_fit
```

The `Model` class is a flexible, concise curve fitter. I will illustrate fitting example data to an exponential decay.

```
def decay(t, N, tau):
    return N*np.exp(-t/tau)
```

The parameters are in no particular order. We'll need some example data. I will use `N=7` and `tau=3`, and add a little noise.

```
t = np.linspace(0, 5, num=1000)
np.random.seed(2021)
data = decay(t, 7, 3) + np.random.randn(t.size)
```

Simplest Usage

```
model = Model(decay, independent_vars=['t'])
result = model.fit(data, t=t, N=10, tau=1)
```

The Model infers the parameter names by inspecting the arguments of the function, `decay`. Then I passed the independent variable, `t`, and initial guesses for each parameter. A residual function is automatically defined, and a least-squared regression is performed.


We can immediately see the best-fit values:

```
print(result.values)
```

```
{'N': 7.146693193035321, 'tau': 2.8980289807074064}
```

and use these best-fit parameters for plotting with the `plot` function:

```
result.plot()
```



We can review the best-fit *Parameters* by accessing `result.params`:

```
result.params.pretty_print()
```

Name	Value	Min	Max	Stderr	Vary	Expr	Brute_Step
N	7.147	-inf	inf	0.0913	True	None	None
tau	2.898	-inf	inf	0.06299	True	None	None

More information about the fit is stored in the `result`, which is an `lmfit.MinimizerResult` object (see: <https://lmfit.github.io/lmfit-py/fitting.html#lmfit.minimizer.MinimizerResult>)

Specifying Bounds and Holding Parameters Constant

Above, the `Model` class implicitly builds `Parameter` objects from keyword arguments of `fit` that match the arguments of `decay`. You can build the `Parameter` objects explicitly; the following is equivalent.

```
result = model.fit(data, t=t,
                    N=Parameter('N', value=10),
                    tau=Parameter('tau', value=1))
report_fit(result.params)
```

```
[[Variables]]
N:    7.14669319 +/- 0.09130428 (1.28%) (init = 10)
tau:  2.89802898 +/- 0.06299118 (2.17%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(N, tau) = -0.7533
```

By building `Parameter` objects explicitly, you can specify bounds (`min`, `max`) and set parameters constant (`vary=False`).

```
result = model.fit(data, t=t,
                    N=Parameter('N', value=7, vary=False),
                    tau=Parameter('tau', value=1, min=0))
report_fit(result.params)
```

```
[[Variables]]
N:    7 (fixed)
tau:  2.97663118 +/- 0.04347476 (1.46%) (init = 1)
```

Defining Parameters in Advance

Passing parameters to `fit` can become unwieldy. As an alternative, you can extract the parameters from `model` like so, set them individually, and pass them to `fit`.

```
params = model.make_params(N=10, tau={'value': 1, 'min': 0})

result = model.fit(data, params, t=t)
report_fit(result.params)
```

```
[[Variables]]
N:    7.14669316 +/- 0.09130423 (1.28%) (init = 10)
tau:  2.89802901 +/- 0.06299127 (2.17%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(N, tau) = -0.7533
```

Keyword arguments override `params`, resetting `value` and all other properties (`min`, `max`, `vary`).

```
result = model.fit(data, params, t=t, tau=1)
report_fit(result.params)
```

```
[[Variables]]
N:    7.14669316 +/- 0.09130423 (1.28%) (init = 10)
tau:  2.89802901 +/- 0.06299127 (2.17%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(N, tau) = -0.7533
```

The input parameters are not modified by `fit`. They can be reused, retaining the same initial value. If you want to use the result of one fit as the initial guess for the next, simply pass `params=result.params`.

#TODO/FIXME: not sure if there ever way a “helpful exception”, but currently it raises a `ValueError`: The input contains `nan` values.

/*A Helpful Exception*

All this implicit magic makes it very easy for the user to neglect to set a `#parameter`. The `fit` function checks for this and raises a helpful exception.

```
# #result = model.fit(data, t=t, tau=1) # N unspecified
```

An `extra` parameter that cannot be matched to the model function will throw a `UserWarning`, but it will not raise, leaving open the possibility of unforeseen extensions calling for some parameters.

Weighted Fits

Use the `sigma` argument to perform a weighted fit. If you prefer to think of the fit in term of `weights`, `sigma=1/weights`.

```
weights = np.arange(len(data))
result = model.fit(data, params, t=t, weights=weights)
report_fit(result.params)
```

```
[[Variables]]
N: 6.98535179 +/- 0.28002384 (4.01%) (init = 10)
tau: 2.97268236 +/- 0.11134755 (3.75%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(N, tau) = -0.9311
```

Handling Missing Data

By default, attempting to fit data that includes a `NaN`, which conventionally indicates a “missing” observation, raises a lengthy exception. You can choose to `omit` (i.e., skip over) missing values instead.

```
data_with_holes = data.copy()
data_with_holes[[5, 500, 700]] = np.nan # Replace arbitrary values with NaN.

model = Model(decay, independent_vars=['t'], nan_policy='omit')
result = model.fit(data_with_holes, params, t=t)
report_fit(result.params)
```

```
[[Variables]]
N: 7.15448795 +/- 0.09181809 (1.28%) (init = 10)
tau: 2.89285089 +/- 0.06306004 (2.18%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(N, tau) = -0.7542
```

If you don’t want to ignore missing values, you can set the model to raise proactively, checking for missing values before attempting the fit.

Uncomment to see the error `#model = Model(decay, independent_vars=['t'], nan_policy='raise') #result = model.fit(data_with_holes, params, t=t)`

The default setting is `nan_policy='raise'`, which does check for `NaNs` and raises an exception when present.

Null-checking relies on `pandas.isnull` if it is available. If `pandas` cannot be imported, it silently falls back on `numpy.isnan`.

Data Alignment

Imagine a collection of time series data with different lengths. It would be convenient to define one sufficiently long array `t` and use it for each time series, regardless of length. `pandas` (<https://pandas.pydata.org/pandas-docs/stable/>) provides tools for aligning indexed data. And, unlike most wrappers to `scipy.leastsq`, `Model` can handle `pandas` objects out of the box, using its data alignment features.

Here I take just a slice of the data and fit it to the full `t`. It is automatically aligned to the correct section of `t` using `Series`' index.

```
model = Model(decay, independent_vars=['t'])
truncated_data = Series(data)[200:800] # data points 200-800
t = Series(t) # all 1000 points
result = model.fit(truncated_data, params, t=t)
report_fit(result.params)
```

```
[[Variables]]
N: 7.10725864 +/- 0.24259071 (3.41%) (init = 10)
tau: 2.92503564 +/- 0.13481789 (4.61%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(N, tau) = -0.9320
```

Data with missing entries and an unequal length still aligns properly.

```
model = Model(decay, independent_vars=['t'], nan_policy='omit')
truncated_data_with_holes = Series(data_with_holes)[200:800]
result = model.fit(truncated_data_with_holes, params, t=t)
report_fit(result.params)
```

```
[[Variables]]
N: 7.11270194 +/- 0.24334895 (3.42%) (init = 10)
tau: 2.92065227 +/- 0.13488230 (4.62%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(N, tau) = -0.9320
```

Total running time of the script: (0 minutes 0.351 seconds)

13.12 Fit Specifying a Function to Compute the Jacobian

Specifying an analytical function to calculate the Jacobian can speed-up the fitting procedure.

```
import matplotlib.pyplot as plt
import numpy as np

from lmfit import Minimizer, Parameters

def func(pars, x, data=None):
    a, b, c = pars['a'], pars['b'], pars['c']
    model = a * np.exp(-b*x) + c
    if data is None:
        return model
    return model - data
```

(continues on next page)

(continued from previous page)

```

def dfunc(pars, x, data=None):
    a, b = pars['a'], pars['b']
    v = np.exp(-b*x)
    return np.array([v, -a*x*v, np.ones(len(x))])

def f(var, x):
    return var[0] * np.exp(-var[1]*x) + var[2]

params = Parameters()
params.add('a', value=10)
params.add('b', value=10)
params.add('c', value=10)

a, b, c = 2.5, 1.3, 0.8
x = np.linspace(0, 4, 50)
y = f([a, b, c], x)
np.random.seed(2021)
data = y + 0.15*np.random.normal(size=x.size)

```

Fit without analytic derivative:

```

min1 = Minimizer(func, params, fcn_args=(x,), fcn_kws={'data': data})
out1 = min1.leastsq()
fit1 = func(out1.params, x)

```

Fit with analytic derivative:

```

min2 = Minimizer(func, params, fcn_args=(x,), fcn_kws={'data': data})
out2 = min2.leastsq(Dfun=dfunc, col_deriv=1)
fit2 = func(out2.params, x)

```

Comparison of fit to exponential decay with/without analytical derivatives to model = $a \cdot \exp(-b \cdot x) + c$:

```

print(f'"true" parameters are: a = {a:.3f}, b = {b:.3f}, c = {c:.3f}\n\n'
      '|=====|\n'
      '| Statistic/Parameter | Without | With |\n'
      '|-----|\n'
      f'| N Function Calls | {out1.nfev:d} | {out2.nfev:d} |\n'
      f'| Chi-square | {out1.chisqr:.4f} | {out2.chisqr:.4f} |\n'
      f'| a | {out1.params['a'].value:.4f} | {out2.params['a'].value:.4f} |\n'
      f'| b | {out1.params['b'].value:.4f} | {out2.params['b'].value:.4f} |\n'
      f'| c | {out1.params['c'].value:.4f} | {out2.params['c'].value:.4f} |\n'
      '|-----|')

```

```
"true" parameters are: a = 2.500, b = 1.300, c = 0.800
```


(continues on next page)

(continued from previous page)

Statistic/Parameter	Without	With
N Function Calls	39	12
Chi-square	1.0920	1.0920
a	2.5635	2.5635
b	1.3585	1.3585
c	0.8241	0.8241

and the best-fit to the synthetic data (with added noise) is the same for both methods:

```
plt.plot(x, data, 'o', label='data')
plt.plot(x, fit1, label='with analytical derivative')
plt.plot(x, fit2, '--', label='without analytical derivative')
plt.legend()
```



Total running time of the script: (0 minutes 0.206 seconds)

13.13 Outlier detection via leave-one-out

Outliers can sometimes be identified by assessing the influence of each datapoint. To assess the influence of one point, we fit the dataset without the point and compare the result with the fit of the full dataset. The code below shows how to do this with lmfit. Note that the presented method is very basic.

```
from collections import defaultdict

import matplotlib.pyplot as plt
import numpy as np

import lmfit
```

Generate test data and model:

```
x = np.linspace(0.3, 10, 100)
np.random.seed(1)
y = 1.0 / (0.1 * x) + 2.0 + 3 * np.random.randn(x.size)

params = lmfit.Parameters()
params.add_many(('a', 0.1), ('b', 1))

def func(x, a, b):
    return 1.0 / (a * x) + b
```

Make five points outliers:


```
idx = np.random.randint(0, x.size, 5)
y[idx] += 10 * np.random.randn(idx.size)
```

Fit the data:

```
model = lmfit.Model(func, independent_vars=['x'])
fit_result = model.fit(y, x=x, a=0.1, b=2)
```

and gives the plot and fitting results below:

```
fit_result.plot_fit()
plt.plot(x[idx], y[idx], 'o', label='outliers')
plt.show()
```



```
print(fit_result.fit_report())
```

```
[[Model]]
  Model(func)
[[Fit Statistics]]
  # fitting method   = leastsq
  # function evals  = 13
  # data points    = 100
  # variables      = 2
  chi-square        = 1338.34458
  reduced chi-square = 13.6565773
  Akaike info crit  = 263.401856
  Bayesian info crit = 268.612196
  R-squared          = 0.70062332
[[Variables]]
  a: 0.08937623 +/- 0.00590174 (6.60%) (init = 0.1)
  b: 1.51298991 +/- 0.46229147 (30.55%) (init = 2)
[[Correlations]] (unreported correlations are < 0.100)
  C(a, b) = +0.6008
```

Fit the dataset while omitting one data point:

```
best_vals = defaultdict(lambda: np.zeros(x.size))
```

(continues on next page)

(continued from previous page)

```

stderrs = defaultdict(lambda: np.zeros(x.size))
chi_sq = np.zeros_like(x)
for i in range(x.size):
    idx2 = np.arange(0, x.size)
    idx2 = np.delete(idx2, i)
    tmp_x = x[idx2]
    tmp = model.fit(y[idx2], x=tmp_x, a=fit_result.params['a'],
                     b=fit_result.params['b'])
    chi_sq[i] = tmp.chisqr
    for p in tmp.params:
        tpar = tmp.params[p]
        best_vals[p][i] = tpar.value
        stderrs[p][i] = (tpar.stderr / fit_result.params[p].stderr)


```

Plot the influence on the red. chisqr of each point:

```

fig, ax = plt.subplots()
ax.plot(x, (fit_result.chisqr - chi_sq) / chi_sq)
ax.scatter(x[idx], fit_result.chisqr / chi_sq[idx] - 1, color='r',
           label='outlier')
ax.set_ylabel(r'Relative red.  $\chi^2$  change')
ax.set_xlabel('x')
ax.legend()

```



Plot the influence on the parameter value and error of each point:


```
fig, axs = plt.subplots(4, figsize=(4, 7), sharex='col')
axs[0].plot(x, best_vals['a'])
axs[0].scatter(x[idx], best_vals['a'][idx], color='r', label='outlier')
axs[0].set_ylabel('best a')

axs[1].plot(x, best_vals['b'])
axs[1].scatter(x[idx], best_vals['b'][idx], color='r', label='outlier')
axs[1].set_ylabel('best b')

axs[2].plot(x, stderrs['a'])
axs[2].scatter(x[idx], stderrs['a'][idx], color='r', label='outlier')
axs[2].set_ylabel('err a change')

axs[3].plot(x, stderrs['b'])
axs[3].scatter(x[idx], stderrs['b'][idx], color='r', label='outlier')
axs[3].set_ylabel('err b change')

axs[3].set_xlabel('x')
```



Total running time of the script: (0 minutes 1.195 seconds)

13.14 Emcee and the Model Interface

```
import corner
import matplotlib.pyplot as plt
import numpy as np

import lmfit
```

Set up a double-exponential function and create a Model:

```
def double_exp(x, a1, t1, a2, t2):
    return a1*np.exp(-x/t1) + a2*np.exp(-(x-0.1) / t2)

model = lmfit.Model(double_exp)
```

Generate some fake data from the model with added noise:

```
truths = (3.0, 2.0, -5.0, 10.0)
x = np.linspace(1, 10, 250)
np.random.seed(0)
y = double_exp(x, *truths)+0.1*np.random.randn(x.size)
```


Create model parameters and give them initial values:

```
p = model.make_params(a1=4, t1=3, a2=4, t2=3)
```

Fit the model using a traditional minimizer, and show the output:

```
result = model.fit(data=y, params=p, x=x, method='Nelder', nan_policy='omit')

lmfit.report_fit(result)
result.plot()
```



```
[[Fit Statistics]]
# fitting method = Nelder-Mead
# function evals = 609
# data points = 250
# variables = 4
chi-square = 2.33333982
reduced chi-square = 0.00948512
Akaike info crit = -1160.54007
Bayesian info crit = -1146.45423
R-squared = 0.94237407
[[Variables]]
a1: 2.98623689 +/- 0.15010519 (5.03%) (init = 4)
t1: 1.30993186 +/- 0.13449656 (10.27%) (init = 3)
a2: -4.33525597 +/- 0.11765824 (2.71%) (init = 4)
```

(continues on next page)

(continued from previous page)

```
t2: 11.8240752 +/- 0.47172610 (3.99%) (init = 3)
[[Correlations]] (unreported correlations are < 0.100)
C(a2, t2) = +0.9876
C(t1, a2) = -0.9278
C(t1, t2) = -0.8852
C(a1, t1) = -0.6093
C(a1, a2) = +0.2973
C(a1, t2) = +0.2319
```

Calculate parameter covariance using emcee:

- start the walkers out at the best-fit values
- set `is_weighted` to `False` to estimate the noise weights
- set some sensible priors on the uncertainty to keep the MCMC in check

```
emcee_kws = dict(steps=5000, burn=500, thin=20, is_weighted=False,
                  progress=False)
emcee_params = result.params.copy()
emcee_params.add('__lnsigma', value=np.log(0.1), min=np.log(0.001), max=np.log(2.0))
```


run the MCMC algorithm and show the results:

```
result_emcee = model.fit(data=y, x=x, params=emcee_params, method='emcee',
                           nan_policy='omit', fit_kws=emcee_kws)
```

```
lmfit.report_fit(result_emcee)
```


```
[[Fit Statistics]]
# fitting method = emcee
# function evals = 500000
# data points = 250
# variables = 5
chi-square = 245.221790
reduced chi-square = 1.00090526
Akaike info crit = 5.17553688
Bayesian info crit = 22.7828415
R-squared = 0.94233862
[[Variables]]
a1: 2.99546858 +/- 0.14834594 (4.95%) (init = 2.986237)
t1: 1.32127391 +/- 0.14079400 (10.66%) (init = 1.309932)
a2: -4.34376940 +/- 0.12389335 (2.85%) (init = -4.335256)
t2: 11.7937607 +/- 0.48879633 (4.14%) (init = 11.82408)
__lnsigma: -2.32712371 +/- 0.04514314 (1.94%) (init = -2.302585)
[[Correlations]] (unreported correlations are < 0.100)
C(a2, t2) = +0.9810
C(t1, a2) = -0.9367
C(t1, t2) = -0.8949
C(a1, t1) = -0.5154
C(a1, a2) = +0.2197
C(a1, t2) = +0.1868
```

```
result_emcee.plot_fit()  
plt.plot(x, model.eval(params=result.params, x=x), '---', label='Nelder')  
plt.legend()
```



Check the acceptance fraction to see whether emcee performed well:

```
plt.plot(result_emcee.acceptance_fraction, 'o')  
plt.xlabel('walker')  
plt.ylabel('acceptance fraction')
```



Try to compute the autocorrelation time:


```
if hasattr(result_emcee, "acor"):
    print("Autocorrelation time for the parameters:")
    print("-----")
    for i, p in enumerate(result.params):
        print(f'{p} = {result_emcee.acor[i]:.3f}')
```

Autocorrelation time for the parameters:

```
-----  
a1 = 61.334  
t1 = 85.867  
a2 = 86.046  
t2 = 84.745
```

Plot the parameter covariances returned by emcee using corner:

```
emcee_corner = corner.corner(result_emcee.flatchain, labels=result_emcee.var_names,
                               truths=list(result_emcee.params.valuesdict().values()))
```



```
print("\nmedian of posterior probability distribution")
print('-----')
lmfit.report_fit(result_emcee.params)
```

```
median of posterior probability distribution
-----
[[Variables]]
 a1: 2.99546858 +/- 0.14834594 (4.95%) (init = 2.986237)
 t1: 1.32127391 +/- 0.14079400 (10.66%) (init = 1.309932)
 a2: -4.34376940 +/- 0.12389335 (2.85%) (init = -4.335256)
 t2: 11.7937607 +/- 0.48879633 (4.14%) (init = 11.82408)
 __lnsigma: -2.32712371 +/- 0.04514314 (1.94%) (init = -2.302585)
[[Correlations]] (unreported correlations are < 0.100)
```

(continues on next page)

(continued from previous page)

```
C(a2, t2) = +0.9810
C(t1, a2) = -0.9367
C(t1, t2) = -0.8949
C(a1, t1) = -0.5154
C(a1, a2) = +0.2197
C(a1, t2) = +0.1868
```

Find the maximum likelihood solution:

```
highest_prob = np.argmax(result_emcee.lnprob)
hp_loc = np.unravel_index(highest_prob, result_emcee.lnprob.shape)
mle_soln = result_emcee.chain[hp_loc]
print("\nMaximum Likelihood Estimation (MLE):")
print('-----')
for ix, param in enumerate(emcee_params):
    print(f" {param}: {mle_soln[ix]:.3f}")

quantiles = np.percentile(result_emcee.flatchain['t1'], [2.8, 15.9, 50, 84.2, 97.7])
print(f"\n1 sigma spread = {0.5 * (quantiles[3] - quantiles[1]):.3f}")
print(f"2 sigma spread = {0.5 * (quantiles[4] - quantiles[0]):.3f}")
```

Maximum Likelihood Estimation (MLE):

```
-----
a1: 2.971
t1: 1.317
a2: -4.336
t2: 11.815
__linsigma: -2.336
```

```
1 sigma spread = 0.141
2 sigma spread = 0.291
```

Total running time of the script: (0 minutes 47.814 seconds)

13.15 Complex Resonator Model

This notebook shows how to fit the parameters of a complex resonator, using *Imfit.Model* and defining a custom *Model* class.

Following Khalil et al. (<https://arxiv.org/abs/1108.3117>), we can model the forward transmission of a microwave resonator with total quality factor Q , coupling quality factor Q_e , and resonant frequency f_0 using:

$$S_{21}(f) = 1 - \frac{QQ_e^{-1}}{1 + 2jQ(f - f_0)/f_0}$$

S_{21} is thus a complex function of a real frequency.

By allowing Q_e to be complex, this model can take into account mismatches in the input and output transmission impedances.

```
import matplotlib.pyplot as plt
import numpy as np

import lmfit
```

Since `scipy.optimize` and `lmfit` require real parameters, we represent Q_e as $Q_e_{\text{real}} + 1j*Q_e_{\text{imag}}$.

```
def linear_resonator(f, f_0, Q, Q_e_real, Q_e_imag):
    Q_e = Q_e_real + 1j*Q_e_imag
    return 1 - (Q * Q_e**-1 / (1 + 2j * Q * (f - f_0) / f_0))
```

The standard practice of defining a `lmfit` model is as follows:

```
class ResonatorModel(lmfit.model.Model):
    __doc__ = "resonator model" + lmfit.models.COMMON_INIT_DOC

    def __init__(self, *args, **kwargs):
        # pass in the defining equation so the user doesn't have to later
        super().__init__(linear_resonator, *args, **kwargs)

        self.set_param_hint('Q', min=0) # enforce Q is positive

    def guess(self, data, f=None, **kwargs):
        verbose = kwargs.pop('verbose', None)
        if f is None:
            return
        argmin_s21 = np.abs(data).argmin()
        fmin = f.min()
        fmax = f.max()
        f_0_guess = f[argmin_s21] # guess that the resonance is the lowest point
        Q_min = 0.1 * (f_0_guess/(fmax-fmin)) # assume the user isn't trying to fit just
        ↪ a small part of a resonance curve
        delta_f = np.diff(f) # assume f is sorted
        min_delta_f = delta_f[delta_f > 0].min()
        Q_max = f_0_guess/min_delta_f # assume data actually samples the resonance
        ↪ reasonably
        Q_guess = np.sqrt(Q_min*Q_max) # geometric mean, why not?
        Q_e_real_guess = Q_guess/(1-np.abs(data[argmin_s21]))
        if verbose:
            print(f"fmin={fmin}, fmax={fmax}, f_0_guess={f_0_guess}")
            print(f"Qmin={Q_min}, Q_max={Q_max}, Q_guess={Q_guess}, Q_e_real_guess={Q_e_
        ↪ real_guess}")
        params = self.make_params(Q=Q_guess, Q_e_real=Q_e_real_guess, Q_e_imag=0, f_0=f_
        ↪ 0_guess)
        params[f'{self.prefix}Q'].set(min=Q_min, max=Q_max)
        params[f'{self.prefix}f_0'].set(min=fmin, max=fmax)
        return lmfit.models.update_param_vals(params, self.prefix, **kwargs)
```

Now let's use the model to generate some fake data:

```
resonator = ResonatorModel()
true_params = resonator.make_params(f_0=100, Q=10000, Q_e_real=9000, Q_e_imag=-9000)
```

(continues on next page)


(continued from previous page)

```

f = np.linspace(99.95, 100.05, 100)
true_s21 = resonator.eval(params=true_params, f=f)
noise_scale = 0.02
np.random.seed(123)
measured_s21 = true_s21 + noise_scale*(np.random.randn(100) + 1j*np.random.randn(100))

plt.plot(f, 20*np.log10(np.abs(measured_s21)))
plt.ylabel('|S21| (dB)')
plt.xlabel('MHz')
plt.title('simulated measurement')

```



Try out the `guess` method we added:

```
guess = resonator.guess(measured_s21, f=f, verbose=True)
```

```

fmin=99.95, fmax=100.05, f_0_guess=100.003535353536
Qmin=100.00353535354105, Q_max=99003.50000055433, Q_guess=3146.537781821432, Q_e_real_
↳ guess=5082.2474265369565

```

And now fit the data using the `guess`-ed values as a starting point:

```
result = resonator.fit(measured_s21, params=guess, f=f, verbose=True)
```

(continues on next page)

(continued from previous page)

```
print(result.fit_report() + '\n')
result.params.pretty_print()
```

```
[[Model]]
  Model(linear_resonator)
[[Fit Statistics]]
  # fitting method    = leastsq
  # function evals   = 41
  # data points      = 200
  # variables        = 4
  chi-square         = 0.08533642
  reduced chi-square = 4.3539e-04
  Akaike info crit  = -1543.89425
  Bayesian info crit = -1530.70099
  R-squared          = (-12528141463701.6+2276419599932.454j)
[[Variables]]
  f_0:      100.000096 +/- 7.0308e-05 (0.00%) (init = 100.0035)
  Q:        10059.4972 +/- 142.294636 (1.41%) (init = 3146.538)
  Q_e_real: 9180.62017 +/- 133.777681 (1.46%) (init = 5082.247)
  Q_e_imag: -9137.03667 +/- 133.769692 (1.46%) (init = 0)
[[Correlations]] (unreported correlations are < 0.100)
  C(Q, Q_e_real) = +0.5175
  C(f_0, Q_e_imag) = +0.5175
  C(f_0, Q_e_real) = +0.5151
  C(Q, Q_e_imag) = -0.5150

```

Name	Value	Min	Max	Stderr	Vary	Expr	Brute_Step
Q	1.006e+04	100	9.9e+04	142.3	True	None	None
Q_e_imag	-9137	-inf	inf	133.8	True	None	None
Q_e_real	9181	-inf	inf	133.8	True	None	None
f_0	100	99.95	100	7.031e-05	True	None	None

Now we'll make some plots of the data and fit. Define a convenience function for plotting complex quantities:

```
def plot_ri(data, *args, **kwargs):
    plt.plot(data.real, data.imag, *args, **kwargs)

fit_s21 = resonator.eval(params=result.params, f=f)
guess_s21 = resonator.eval(params=guess, f=f)


plt.figure()
plot_ri(measured_s21, '.')
plot_ri(fit_s21, '.-', label='best fit')
plot_ri(guess_s21, '--', label='initial fit')
plt.legend()
plt.xlabel('Re(S21)')
plt.ylabel('Im(S21)')


plt.figure()
plt.plot(f, 20*np.log10(np.abs(measured_s21)), '.')
plt.plot(f, 20*np.log10(np.abs(fit_s21)), '.-', label='best fit')
```

(continues on next page)

(continued from previous page)

```
plt.plot(f, 20*np.log10(np.abs(guess_s21)), '--', label='initial fit')
plt.legend()
plt.ylabel('|S21| (dB)')
plt.xlabel('MHz')
```





Total running time of the script: (0 minutes 0.557 seconds)

13.16 Model Selection using lmfit and emcee

FIXME: this is a useful example; however, it doesn't run correctly anymore as the PTSampler was removed in emcee v3...

lmfit.emcee can be used to obtain the posterior probability distribution of parameters, given a set of experimental data. This notebook shows how it can be used for Bayesian model selection.

```
import matplotlib.pyplot as plt
import numpy as np

import lmfit
```

Define a Gaussian lineshape and generate some data:

```
def gauss(x, a_max, loc, sd):
    return a_max * np.exp(-((x - loc) / sd)**2)

x = np.linspace(3, 7, 250)
np.random.seed(0)
y = 4 + 10 * x + gauss(x, 200, 5, 0.5) + gauss(x, 60, 5.8, 0.2)
```

(continues on next page)

(continued from previous page)

```
dy = np.sqrt(y)
y += dy * np.random(y.size)
```

Plot the data:

```
plt.errorbar(x, y)
```

Define the normalised residual for the data:

```
def residual(p, just_generative=False):
    v = p.valuesdict()
    generative = v['a'] + v['b'] * x
    M = 0
    while f'a_max{M}' in v:
        generative += gauss(x, v[f'a_max{M}'], v[f'loc{M}'], v[f'sd{M}'])
        M += 1

    if just_generative:
        return generative
    return (generative - y) / dy
```

Create a Parameter set for the initial guesses:

```
def initial_peak_params(M):
    p = lmfit.Parameters()

    # a and b give a linear background
    a = np.mean(y)
    b = 1

    # a_max, loc and sd are the amplitude, location and SD of each Gaussian
    # component
    a_max = np.max(y)
    loc = np.mean(x)
    sd = (np.max(x) - np.min(x)) * 0.5

    p.add_many(('a', a, True, 0, 10), ('b', b, True, 1, 15))

    for i in range(M):
        p.add_many((f'a_max{i}', 0.5 * a_max, True, 10, a_max),
                   (f'loc{i}', loc, True, np.min(x), np.max(x)),
                   (f'sd{i}', sd, True, 0.1, np.max(x) - np.min(x)))
    return p
```

Solving with *minimize* gives the Maximum Likelihood solution.

```
p1 = initial_peak_params(1)
mi1 = lmfit.minimize(residual, p1, method='differential_evolution')

lmfit.printfuncs.report_fit(mi1.params, min_correl=0.5)
```

From inspection of the data above we can tell that there is going to be more than 1 Gaussian component, but how many are there? A Bayesian approach can be used for this model selection problem. We can do this with *lmfit.emcee*, which uses the *emcee* package to do a Markov Chain Monte Carlo sampling of the posterior probability distribution.

lmfit.emcee requires a function that returns the log-posterior probability. The log-posterior probability is a sum of the log-prior probability and log-likelihood functions.

The log-prior probability encodes information about what you already believe about the system. *lmfit.emcee* assumes that this log-prior probability is zero if all the parameters are within their bounds and *-np.inf* if any of the parameters are outside their bounds. As such it's a uniform prior.

The log-likelihood function is given below. To use non-uniform priors then should include these terms in *lnprob*. This is the log-likelihood probability for the sampling.

```
def lnprob(p):
    resid = residual(p, just_generative=True)
    return -0.5 * np.sum(((resid - y) / dy)**2 + np.log(2 * np.pi * dy**2))
```

To start with we have to create the minimizers and *burn* them in. We create 4 different minimizers representing 0, 1, 2 or 3 Gaussian contributions. To do the model selection we have to integrate the over the log-posterior distribution to see which has the higher probability. This is done using the *thermodynamic_integration_log_evidence* method of the *sampler* attribute contained in the *lmfit.Minimizer* object.

```
# Work out the log-evidence for different numbers of peaks:
total_steps = 310
burn = 300
thin = 10
ntemps = 15
workers = 1 # the multiprocessing does not work with sphinx-gallery
log_evidence = []
res = []

# set up the Minimizers
for i in range(4):
    p0 = initial_peak_params(i)
    # you can't use lnprob as a userfcn with minimize because it needs to be
    # maximised
    mini = lmfit.Minimizer(residual, p0)
    out = mini.minimize(method='differential_evolution')
    res.append(out)

mini = []
# burn in the samplers
for i in range(4):
    # do the sampling
    mini.append(lmfit.Minimizer(lnprob, res[i].params))
    out = mini[i].emcee(steps=total_steps, ntemps=ntemps, workers=workers,
                         reuse_sampler=False, float_behavior='posterior',
                         progress=False)
    # get the evidence
    print(i, total_steps, mini[i].sampler.thermodynamic_integration_log_evidence())
    log_evidence.append(mini[i].sampler.thermodynamic_integration_log_evidence()[0])
```

Once we've burned in the samplers we have to do a collection run. We thin out the MCMC chain to reduce autocorrelation between successive samples.

```
for j in range(6):
    total_steps += 100
    for i in range(4):
```

(continues on next page)

(continued from previous page)

```

# do the sampling
res = mini[i].emcee(burn=burn, steps=100, thin=thin, ntemps=ntemps,
                     workers=workers, reuse_sampler=True, progress=False)
# get the evidence
print(i, total_steps, mini[i].sampler.thermodynamic_integration_log_evidence())
log_evidence.append(mini[i].sampler.thermodynamic_integration_log_evidence()[-1])

plt.plot(log_evidence[-4:])
plt.ylabel('Log-evidence')
plt.xlabel('number of peaks')

```

The Bayes factor is related to the exponential of the difference between the log-evidence values. Thus, 0 peaks is not very likely compared to 1 peak. But 1 peak is not as good as 2 peaks. 3 peaks is not that much better than 2 peaks.

```

r01 = np.exp(log_evidence[-4] - log_evidence[-3])
r12 = np.exp(log_evidence[-3] - log_evidence[-2])
r23 = np.exp(log_evidence[-2] - log_evidence[-1])

print(r01, r12, r23)

```

These numbers tell us that zero peaks is 0 times as likely as one peak. Two peaks is 7e49 times more likely than one peak. Three peaks is 1.1 times more likely than two peaks. With this data one would say that two peaks is sufficient. Caution has to be taken with these values. The log-priors for this sampling are uniform but improper, i.e. they are not normalised properly. Internally the `Inprior` probability is calculated as 0 if all parameters are within their bounds and `-np.inf` if any parameter is outside the bounds. The `Inprob` function defined above is the log-likelihood alone. Remember, that the log-posterior probability is equal to the sum of the log-prior and log-likelihood probabilities. Extra terms can be added to the `Inprob` function to calculate the normalised log-probability. These terms would look something like:

$$\log\left(\prod_i \frac{1}{\max_i - \min_i}\right)$$

where \max_i and \min_i are the upper and lower bounds for the parameter, and the prior is a uniform distribution. Other types of prior are possible. For example, you might expect the prior to be Gaussian.

13.17 Calculate Confidence Intervals

```

import matplotlib.pyplot as plt
from numpy import argsort, exp, linspace, pi, random, sign, sin, unique
from scipy.interpolate import interp1d

from lmfit import Minimizer, conf_interval, conf_interval2d, create_params,
               report_ci, report_fit

```

Define the residual function, specify “true” parameter values, and generate a synthetic data set with some noise:

```

def residual(pars, x, data=None):
    argu = (x*pars['decay'])**2
    shift = pars['shift']
    if abs(shift) > pi/2:
        shift = shift - sign(shift)*pi

```

(continues on next page)

(continued from previous page)

```

model = pars['amp']*sin(shift + x/pars['period']) * exp(-argu)
if data is None:
    return model
return model - data

p_true = create_params(amp=14.0, period=5.33, shift=0.123, decay=0.010)

x = linspace(0.0, 250.0, 2500)
random.seed(2021)
noise = random.normal(scale=0.7215, size=x.size)
data = residual(p_true, x) + noise

```

Create fitting parameters and set initial values:

```
fit_params = create_params(amp=13.0, period=2, shift=0.0, decay=0.020)
```

Set-up the minimizer and perform the fit using leastsq algorithm, and show the report:

```

mini = Minimizer(residual, fit_params, fcn_args=(x,), fcn_kws={'data': data})
out = mini.leastsq()

fit = residual(out.params, x)
report_fit(out)

```

```

[[Fit Statistics]]
# fitting method    = leastsq
# function evals   = 95
# data points      = 2500
# variables        = 4
chi-square         = 1277.24638
reduced chi-square = 0.51171730
Akaike info crit   = -1670.96059
Bayesian info crit = -1647.66441
[[Variables]]
amp:    14.0708269 +/- 0.04936878 (0.35%) (init = 13)
period: 5.32980958 +/- 0.00273143 (0.05%) (init = 2)
shift:  0.12156317 +/- 0.00482312 (3.97%) (init = 0)
decay:  0.01002489 +/- 4.0726e-05 (0.41%) (init = 0.02)
[[Correlations]] (unreported correlations are < 0.100)
C(period, shift) = +0.8002
C(amp, decay)   = +0.5758

```

Calculate the confidence intervals for parameters and display the results:

```

ci, tr = conf_interval(mini, out, trace=True)

report_ci(ci)

names = out.params.keys()
i = 0
gs = plt.GridSpec(4, 4)

```

(continues on next page)


(continued from previous page)

```
sx = []
sy = []
for fixed in names:
    j = 0
    for free in names:
        if j in sx and i in sy:
            ax = plt.subplot(gs[i, j], sharex=sx[j], sharey=sy[i])
        elif i in sy:
            ax = plt.subplot(gs[i, j], sharey=sy[i])
            sx[j] = ax
        elif j in sx:
            ax = plt.subplot(gs[i, j], sharex=sx[j])
            sy[i] = ax
        else:
            ax = plt.subplot(gs[i, j])
            sy[i] = ax
            sx[j] = ax
        if i < 3:
            plt.setp(ax.get_xticklabels(), visible=False)
        else:
            ax.set_xlabel(free)

        if j > 0:
            plt.setp(ax.get_yticklabels(), visible=False)
        else:
            ax.set_ylabel(fixed)

    res = tr[fixed]
    prob = res['prob']
    f = prob < 0.96

    x, y = res[free], res[fixed]
    ax.scatter(x[f], y[f], c=1-prob[f], s=25*(1-prob[f]+0.5))
    ax.autoscale(1, 1)
    j += 1
i += 1
```



	99.73%	95.45%	68.27%	_BEST_	68.27%	95.45%	99.73%
amp :	-0.14795	-0.09863	-0.04934	14.07083	+0.04939	+0.09886	+0.14847
period:	-0.00818	-0.00546	-0.00273	5.32981	+0.00273	+0.00548	+0.00822
shift :	-0.01446	-0.00964	-0.00482	0.12156	+0.00482	+0.00965	+0.01449
decay :	-0.00012	-0.00008	-0.00004	0.01002	+0.00004	+0.00008	+0.00012

It is also possible to calculate the confidence regions for two fixed parameters using the function `conf_interval2d`:

```
names = list(out.params.keys())

plt.figure()
for i in range(4):
    for j in range(4):
        indx = 16-j*4-i
        ax = plt.subplot(4, 4, indx)
        ax.ticklabel_format(style='sci', scilimits=(-2, 2), axis='y')

        # set-up labels and tick marks
        ax.tick_params(labelleft=False, labelbottom=False)
        if indx in (2, 5, 9, 13):
            plt.ylabel(names[j])
            ax.tick_params(labelleft=True)
        if indx == 1:
            ax.tick_params(labelleft=True)
```

(continues on next page)

(continued from previous page)

```
if indx in (13, 14, 15, 16):
    plt.xlabel(names[i])
    ax.tick_params(labelbottom=True)
    [label.set_rotation(45) for label in ax.get_xticklabels()]


if i != j:
    x, y, m = conf_interval2d(mini, out, names[i], names[j], 20, 20)
    plt.contourf(x, y, m, linspace(0, 1, 10))

    x = tr[names[i]][names[i]]
    y = tr[names[i]][names[j]]
    pr = tr[names[i]]['prob']
    s = argsort(x)
    plt.scatter(x[s], y[s], c=pr[s], s=30, lw=1)

else:
    x = tr[names[i]][names[i]]
    y = tr[names[i]]['prob']

    t, s = unique(x, True)
    f = interp1d(t, y[s], 'slinear')
    xn = linspace(x.min(), x.max(), 50)
    plt.plot(xn, f(xn), lw=1)
    plt.ylabel('prob')
    ax.tick_params(labelleft=True)

plt.tight_layout()
plt.show()
```



Total running time of the script: (0 minutes 14.842 seconds)

13.18 Fit Two Dimensional Peaks

This example illustrates how to handle two-dimensional data with lmfit.

```
import matplotlib.pyplot as plt
import numpy as np
from scipy.interpolate import griddata

import lmfit
from lmfit.lineshapes import gaussian2d, lorentzian
```

13.18.1 Two-dimensional Gaussian


We start by considering a simple two-dimensional gaussian function, which depends on coordinates (x, y) . The most general case of experimental data will be irregularly sampled and noisy. Let's simulate some:

```
npoints = 10000
np.random.seed(2021)
x = np.random.rand(npoints)*10 - 4
y = np.random.rand(npoints)*5 - 3
z = gaussian2d(x, y, amplitude=30, centerx=2, centery=-.5, sigmax=.6, sigmay=.8)
z += 2*(np.random.rand(*z.shape)-.5)
error = np.sqrt(z+1)
```

To plot this, we can interpolate the data onto a grid.

```
X, Y = np.meshgrid(np.linspace(x.min(), x.max(), 100),
                    np.linspace(y.min(), y.max(), 100))
Z = griddata((x, y), z, (X, Y), method='linear', fill_value=0)

fig, ax = plt.subplots()
art = ax.pcolor(X, Y, Z, shading='auto')
plt.colorbar(art, ax=ax, label='z')
ax.set_xlabel('x')
ax.set_ylabel('y')
plt.show()
```



In this case, we can use a built-in model to fit

```
model = lmfit.models.Gaussian2dModel()
params = model.guess(z, x, y)
result = model.fit(z, x=x, y=y, params=params, weights=1/error)
lmfit.report_fit(result)
```

```
[[Fit Statistics]]
# fitting method   = leastsq
# function evals  = 87
# data points     = 10000
# variables       = 5
chi-square        = 20618.1774
reduced chi-square = 2.06284916
Akaike info crit  = 7245.87992
Bayesian info crit = 7281.93162
R-squared          = 0.87508800
[[Variables]]
amplitude:  27.4195833 +/- 0.65062974 (2.37%) (init = 16.51399)
centerx:    1.99705425 +/- 0.01405864 (0.70%) (init = 1.940764)
centery:   -0.49516158 +/- 0.01907800 (3.85%) (init = -0.5178641)
sigmax:    0.54740777 +/- 0.01224965 (2.24%) (init = 1.666582)
sigmay:    0.73300589 +/- 0.01617042 (2.21%) (init = 0.8332836)
fwhmx:     1.28904676 +/- 0.02884573 (2.24%) == '2.3548200*sigmax'
```

(continues on next page)

(continued from previous page)

```

fwhmy:      1.72609693 +/- 0.03807842 (2.21%) == '2.3548200*sigmay'
height:     10.8758308 +/- 0.35409534 (3.26%) == '0.1591549*amplitude/(max(1e-15,_
˓→sigmax)*max(1e-15, sigmay))'
[[Correlations]] (unreported correlations are < 0.100)
C(amplitude, sigmay) = +0.2464
C(amplitude, sigmax) = +0.2314

```

To check the fit, we can evaluate the function on the same grid we used before and make plots of the data, the fit and the difference between the two.

```

fig, axs = plt.subplots(2, 2, figsize=(10, 10))

vmax = np.nanpercentile(Z, 99.9)


ax = axs[0, 0]
art = ax.pcolor(X, Y, Z, vmin=0, vmax=vmax, shading='auto')
plt.colorbar(art, ax=ax, label='z')
ax.set_title('Data')

ax = axs[0, 1]
fit = model.func(X, Y, **result.best_values)
art = ax.pcolor(X, Y, fit, vmin=0, vmax=vmax, shading='auto')
plt.colorbar(art, ax=ax, label='z')
ax.set_title('Fit')

ax = axs[1, 0]
fit = model.func(X, Y, **result.best_values)
art = ax.pcolor(X, Y, Z-fit, vmin=0, vmax=10, shading='auto')
plt.colorbar(art, ax=ax, label='z')
ax.set_title('Data - Fit')

for ax in axs.ravel():
    ax.set_xlabel('x')
    ax.set_ylabel('y')
axs[1, 1].remove()
plt.show()

```



13.18.2 Two-dimensional off-axis Lorentzian

We now go on to show a harder example, in which the peak has a Lorentzian profile and an off-axis anisotropic shape. This can be handled by applying a suitable coordinate transform and then using the *lorentzian* function that lmfit provides in the lineshapes module.

```
def lorentzian2d(x, y, amplitude=1., centerx=0., centery=0., sigmax=1., sigmay=1.,
                  rotation=0):
    """Return a two dimensional lorentzian.
```

The maximum of the peak occurs at ``centerx`` and ``centery`` with widths ``sigmax`` and ``sigmay`` in the x and y directions

(continues on next page)

(continued from previous page)

```

respectively. The peak can be rotated by choosing the value of ``rotation``
in radians.
"""
xp = (x - centerx)*np.cos(rotation) - (y - centery)*np.sin(rotation)
yp = (x - centerx)*np.sin(rotation) + (y - centery)*np.cos(rotation)
R = (xp/sigmax)**2 + (yp/sigmay)**2

return 2*amplitude*lorentzian(R)/(np.pi*sigmax*sigmay)

```

Data can be simulated and plotted in the same way as we did before.


```

npoints = 10000
x = np.random.rand(npoints)*10 - 4
y = np.random.rand(npoints)*5 - 3
z = lorentzian2d(x, y, amplitude=30, centerx=2, centery=-.5, sigmax=.6,
                  sigmay=1.2, rotation=30*np.pi/180)
z += 2*(np.random.rand(*z.shape)-.5)
error = np.sqrt(z+1)

X, Y = np.meshgrid(np.linspace(x.min(), x.max(), 100),
                    np.linspace(y.min(), y.max(), 100))
Z = griddata((x, y), z, (X, Y), method='linear', fill_value=0)

fig, ax = plt.subplots()
ax.set_xlabel('x')
ax.set_ylabel('y')
art = ax.pcolor(X, Y, Z, shading='auto')
plt.colorbar(art, ax=ax, label='z')
plt.show()

```



To fit, create a model from the function. Don't forget to tell lmfit that both x and y are independent variables. Keep in mind that lmfit will take the function keywords as default initial guesses in this case and that it will not know that certain parameters only make physical sense over restricted ranges. For example, peak widths should be positive and the rotation can be restricted over a quarter circle.

```
model = lmfit.Model(lorentzian2d, independent_vars=['x', 'y'])
params = model.make_params(amplitude=10, centerx=x[np.argmax(z)],
                           centery=y[np.argmax(z)])
params['rotation'].set(value=.1, min=0, max=np.pi/2)
params['sigmax'].set(value=1, min=0)
params['sigmay'].set(value=2, min=0)

result = model.fit(z, x=x, y=y, params=params, weights=1/error)
lmfit.report_fit(result)
```

```
[[Fit Statistics]]
# fitting method = leastsq
# function evals = 73
# data points = 10000
# variables = 6
chi-square = 11287.3823
reduced chi-square = 1.12941588
Akaike info crit = 1223.00402
Bayesian info crit = 1266.26606
```

(continues on next page)

(continued from previous page)

```
R-squared      = 0.85940121
[[Variables]]
amplitude: 25.6417887 +/- 0.50569636 (1.97%) (init = 10)
centerx:    2.00326033 +/- 0.01163397 (0.58%) (init = 2.051478)
centery:   -0.49692376 +/- 0.01680902 (3.38%) (init = -0.478231)
sigmax:     0.51074241 +/- 0.01027264 (2.01%) (init = 1)
sigmay:     1.11741198 +/- 0.02190772 (1.96%) (init = 2)
rotation:   0.48130689 +/- 0.01542118 (3.20%) (init = 0.1)
[[Correlations]] (unreported correlations are < 0.100)
C(centerx, centery) = +0.5767
C(amplitude, sigmax) = +0.3342
C(amplitude, sigmay) = +0.3046
```

The process of making plots to check it worked is the same as before.

```
fig, axs = plt.subplots(2, 2, figsize=(10, 10))


vmax = np.nanpercentile(Z, 99.9)

ax = axs[0, 0]
art = ax.pcolor(X, Y, Z, vmin=0, vmax=vmax, shading='auto')
plt.colorbar(art, ax=ax, label='z')
ax.set_title('Data')

ax = axs[0, 1]
fit = model.func(X, Y, **result.best_values)
art = ax.pcolor(X, Y, fit, vmin=0, vmax=vmax, shading='auto')
plt.colorbar(art, ax=ax, label='z')
ax.set_title('Fit')

ax = axs[1, 0]
fit = model.func(X, Y, **result.best_values)
art = ax.pcolor(X, Y, Z-fit, vmin=0, vmax=10, shading='auto')
plt.colorbar(art, ax=ax, label='z')
ax.set_title('Data - Fit')

for ax in axs.ravel():
    ax.set_xlabel('x')
    ax.set_ylabel('y')
axs[1, 1].remove()
plt.show()
```



Total running time of the script: (0 minutes 2.648 seconds)

13.19 Global minimization using the brute method (a.k.a. grid search)

This notebook shows a simple example of using `lmfit.minimize.brute` that uses the method with the same name from `scipy.optimize`.

The method computes the function's value at each point of a multidimensional grid of points, to find the global minimum of the function. It behaves identically to `scipy.optimize.brute` in case finite bounds are given on all varying parameters, but will also deal with non-bounded parameters (see below).

```
import copy

from matplotlib.colors import LogNorm
import matplotlib.pyplot as plt
import numpy as np

from lmfit import Minimizer, create_params, fit_report
```

Let's start with the example given in the documentation of SciPy:

“We illustrate the use of brute to seek the global minimum of a function of two variables that is given as the sum of a positive-definite quadratic and two deep “Gaussian-shaped” craters. Specifically, define the objective function f as the sum of three other functions, $f = f1 + f2 + f3$. We suppose each of these has a signature $(z, *params)$, where $z = (x, y)$, and params and the functions are as defined below.”

First, we create a set of Parameters where all variables except x and y are given fixed values. Just as in the documentation we will do a grid search between -4 and 4 and use a stepsize of 0.25. The bounds can be set as usual with the `min` and `max` attributes, and the stepsize is set using `brute_step`.

```
params = create_params(a=dict(value=2, vary=False),
                       b=dict(value=3, vary=False),
                       c=dict(value=7, vary=False),
                       d=dict(value=8, vary=False),
                       e=dict(value=9, vary=False),
                       f=dict(value=10, vary=False),
                       g=dict(value=44, vary=False),
                       h=dict(value=-1, vary=False),
                       i=dict(value=2, vary=False),
                       j=dict(value=26, vary=False),
                       k=dict(value=1, vary=False),
                       l=dict(value=-2, vary=False),
                       scale=dict(value=0.5, vary=False),
                       x=dict(value=0.0, vary=True, min=-4, max=4, brute_step=0.25),
                       y=dict(value=0.0, vary=True, min=-4, max=4, brute_step=0.25))
```

Second, create the three functions and the objective function:

```
def f1(p):
    par = p.valuesdict()
    return (par['a'] * par['x']**2 + par['b'] * par['x'] * par['y'] +
            par['c'] * par['y']**2 + par['d'] * par['x'] + par['e'] * par['y'] +
            par['f'])

def f2(p):
    par = p.valuesdict()
    return (-1.0 * par['g'] * np.exp(-((par['x'] - par['h'])**2 +
                                         (par['y'] - par['i'])**2) / par['scale']))

def f3(p):
    par = p.valuesdict()
    return (-1.0 * par['j'] * np.exp(-((par['x'] - par['k'])**2 +
                                         (par['y'] - par['l'])**2) / par['scale']))
```

(continues on next page)

(continued from previous page)

```
def f(params):
    return f1(params) + f2(params) + f3(params)
```

Performing the actual grid search is done with:

```
fitter = Minimizer(f, params)
result = fitter.minimize(method='brute')
```

, which will increment x and y between -4 in increments of 0.25 until 4 (not inclusive).

```
grid_x, grid_y = (np.unique(par.ravel()) for par in result.brute_grid)
print(grid_x)
```

```
[-4.    -3.75 -3.5  -3.25 -3.  -2.75 -2.5  -2.25 -2.  -1.75 -1.5  -1.25
 -1.    -0.75 -0.5  -0.25  0.   0.25  0.5   0.75  1.   1.25  1.5   1.75
 2.     2.25  2.5   2.75  3.   3.25  3.5   3.75]
```

The objective function is evaluated on this grid, and the raw output from `scipy.optimize.brute` is stored in the `MinimizerResult` as `brute_<parname>` attributes. These attributes are:

`result.brute_x0` – A 1-D array containing the coordinates of a point at which the objective function had its minimum value.

```
print(result.brute_x0)
```

```
[-1.    1.75]
```

`result.brute_fval` – Function value at the point `x0`.

```
print(result.brute_fval)
```

```
-2.8923637137222027
```

`result.brute_grid` – Representation of the evaluation grid. It has the same length as `x0`.

```
print(result.brute_grid)
```

```
[[[-4.    -4.    -4.    ...  -4.    -4.    -4.    ]
 [-3.75 -3.75 -3.75 ...  -3.75 -3.75 -3.75]
 [-3.5   -3.5   -3.5   ...  -3.5   -3.5   -3.5 ]
 ...
 [ 3.25  3.25  3.25 ...  3.25  3.25  3.25]
 [ 3.5   3.5   3.5   ...  3.5   3.5   3.5 ]
 [ 3.75  3.75  3.75 ...  3.75  3.75  3.75]]
 [[-4.    -3.75 -3.5  ...  3.25  3.5   3.75]
 [-4.    -3.75 -3.5  ...  3.25  3.5   3.75]
 [-4.    -3.75 -3.5  ...  3.25  3.5   3.75]
 ...
 [-4.    -3.75 -3.5  ...  3.25  3.5   3.75]]
```

(continues on next page)

(continued from previous page)

```
[-4. -3.75 -3.5 ... 3.25 3.5 3.75]
[-4. -3.75 -3.5 ... 3.25 3.5 3.75]]]
```

result.brute_Jout – Function values at each point of the evaluation grid, i.e., `Jout = func(*grid)`.

```
print(result.brute_Jout)
```

```
[[134. 119.6875 106.25 ... 74.18749997 85.24999999
 97.1875 ]
[129.125 115. 101.75 ... 74.74999948 85.99999987
 98.12499997]
[124.5 110.5625 97.5 ... 75.5624928 86.99999818
 99.31249964]
...
[ 94.12499965 85.24999772 77.24998843 ... 192. 208.5
 225.875 ]
[ 96.49999997 87.81249979 79.99999892 ... 199.8125 216.5
 234.0625 ]
[ 99.125 90.62499998 82.99999992 ... 207.875 224.75
 242.5 ]]
```

Reassuringly, the obtained results are identical to using the method in SciPy directly!


Example 2: fit of a decaying sine wave

In this example, we will explain some of the options of the algorithm.

We start off by generating some synthetic data with noise for a decaying sine wave, define an objective function, and create/initialize a Parameter set.

```
x = np.linspace(0, 15, 301)
np.random.seed(7)
noise = np.random.normal(size=x.size, scale=0.2)
data = (5. * np.sin(2*x - 0.1) * np.exp(-x*x*0.025) + noise)
plt.plot(x, data, 'o')
plt.show()

def fcn2min(params, x, data):
    """Model decaying sine wave, subtract data."""
    amp = params['amp']
    shift = params['shift']
    omega = params['omega']
    decay = params['decay']
    model = amp * np.sin(x*omega + shift) * np.exp(-x*x*decay)
    return model - data
```



In contrast to the implementation in SciPy (as shown in the first example), varying parameters do not need to have finite bounds in lmfit. However, if a parameter does not have finite bounds, then it does need a `brute_step` attribute specified:

```
params = create_params(amp=dict(value=7, min=2.5, brute_step=0.25),
                      decay=dict(value=0.05, brute_step=0.005),
                      shift=dict(value=0.0, min=-np.pi/2., max=np.pi/2),
                      omega=dict(value=3, max=5, brute_step=0.25))
```

Our initial parameter set is now defined as shown below and this will determine how the grid is set-up.

```
params.pretty_print()
```

Name	Value	Min	Max	Stderr	Vary	Expr	Brute_Step
amp	7	2.5	inf	None	True	None	0.25
decay	0.05	-inf	inf	None	True	None	0.005
omega	3	-inf	5	None	True	None	0.25
shift	0	-1.571	1.571	None	True	None	None

First, we initialize a `Minimizer` and perform the grid search:

```
fitter = Minimizer(fcn2min, params, fcn_args=(x, data))
result_brute = fitter.minimize(method='brute', Ns=25, keep=25)
```

(continues on next page)

(continued from previous page)

```
print(fit_report(result_brute))
```

```
[[Fit Statistics]]
# fitting method   = brute
# function evals  = 375000
# data points     = 301
# variables       = 4
chi-square        = 11.9353671
reduced chi-square = 0.04018642
Akaike info crit  = -963.508878
Bayesian info crit = -948.680437
## Warning: uncertainties could not be estimated:
[[Variables]]
amp:   5.00000000 (init = 7)
decay: 0.02500000 (init = 0.05)
shift: -0.13089969 (init = 0)
omega:  2.00000000 (init = 3)
```

We used two new parameters here: `Ns` and `keep`. The parameter `Ns` determines the 'number of grid points along the axes' similarly to its usage in SciPy. Together with `brute_step`, `min` and `max` for a Parameter it will dictate how the grid is set-up:

(1) finite bounds are specified ("SciPy implementation"): uses `brute_step` if present (in the example above) or uses `Ns` to generate the grid. The latter scenario that interpolates `Ns` points from `min` to `max` (inclusive), is here shown for the parameter `shift`:

```
par_name = 'shift'
idx_shift = result_brute.var_names.index(par_name)
grid_shift = np.unique(result_brute.brute_grid[idx_shift].ravel())
print(f"parameter = {par_name}\nnumber of steps = {len(grid_shift)}\ngrid = {grid_shift}\n")
```

```
parameter = shift
number of steps = 25
grid = [-1.57079633 -1.43989663 -1.30899694 -1.17809725 -1.04719755 -0.91629786
-0.78539816 -0.65449847 -0.52359878 -0.39269908 -0.26179939 -0.13089969
0.          0.13089969  0.26179939  0.39269908  0.52359878  0.65449847
0.78539816  0.91629786  1.04719755  1.17809725  1.30899694  1.43989663
1.57079633]
```

If finite bounds are not set for a certain parameter then the user **must** specify `brute_step` - three more scenarios are considered here:

(2) lower bound (`min`) and `brute_step` are specified: `range = (min, min + Ns * brute_step, brute_step)`

```
par_name = 'amp'
idx_shift = result_brute.var_names.index(par_name)
grid_shift = np.unique(result_brute.brute_grid[idx_shift].ravel())
print(f"parameter = {par_name}\nnumber of steps = {len(grid_shift)}\ngrid = {grid_shift}\n")
```

```
parameter = amp
number of steps = 25
```

(continues on next page)

(continued from previous page)

```
grid = [2.5  2.75 3.   3.25 3.5  3.75 4.   4.25 4.5  4.75 5.   5.25 5.5  5.75
       6.   6.25 6.5  6.75 7.   7.25 7.5  7.75 8.   8.25 8.5 ]
```

(3) upper bound (max) and brute_step are specified: range = (max - Ns * brute_step, max, brute_step)

```
par_name = 'omega'
idx_shift = result_brute.var_names.index(par_name)
grid_shift = np.unique(result_brute.brute_grid[idx_shift].ravel())
print(f"parameter = {par_name}\nnumber of steps = {len(grid_shift)}\ngrid = {grid_shift}\n")
```

```
parameter = omega
number of steps = 25
grid = [-1.25 -1.   -0.75 -0.5  -0.25  0.   0.25  0.5   0.75  1.   1.25  1.5
         1.75  2.   2.25  2.5   2.75  3.   3.25  3.5   3.75  4.   4.25  4.5
         4.75]
```

(4) numerical value (value) and brute_step are specified: range = (value - (Ns//2) * brute_step, value + (Ns//2) * brute_step, brute_step)

```
par_name = 'decay'
idx_shift = result_brute.var_names.index(par_name)
grid_shift = np.unique(result_brute.brute_grid[idx_shift].ravel())
print(f"parameter = {par_name}\nnumber of steps = {len(grid_shift)}\ngrid = {grid_shift}\n")
```

```
parameter = decay
number of steps = 24
grid = [-1.00000000e-02 -5.00000000e-03  5.20417043e-18  5.00000000e-03
         1.00000000e-02  1.50000000e-02  2.00000000e-02  2.50000000e-02
         3.00000000e-02  3.50000000e-02  4.00000000e-02  4.50000000e-02
         5.00000000e-02  5.50000000e-02  6.00000000e-02  6.50000000e-02
         7.00000000e-02  7.50000000e-02  8.00000000e-02  8.50000000e-02
         9.00000000e-02  9.50000000e-02  1.00000000e-01  1.05000000e-01]
```

The MinimizerResult contains all the usual best-fit parameters and fitting statistics. For example, the optimal solution from the grid search is given below together with a plot:

```
print(fit_report(result_brute))
```


```
[[Fit Statistics]]
  # fitting method    = brute
  # function evals   = 375000
  # data points      = 301
  # variables        = 4
  chi-square         = 11.9353671
  reduced chi-square = 0.04018642
  Akaike info crit  = -963.508878
  Bayesian info crit = -948.680437
##  Warning: uncertainties could not be estimated:
[[Variables]]
  amp:    5.00000000 (init = 7)
```

(continues on next page)

(continued from previous page)

```
decay: 0.025000000 (init = 0.05)
shift: -0.13089969 (init = 0)
omega: 2.00000000 (init = 3)
```

```
plt.plot(x, data, 'o')
plt.plot(x, data + fcn2min(result_brute.params, x, data), '--')
plt.show()
```



We can see that this fit is already very good, which is what we should expect since our brute force grid is sampled rather finely and encompasses the “correct” values.

In a more realistic, complicated example the brute method will be used to get reasonable values for the parameters and perform another minimization (e.g., using `leastsq`) using those as starting values. That is where the `keep` parameter comes into play: it determines the “number of best candidates from the brute force method that are stored in the `candidates` attribute”. In the example above we store the best-ranking 25 solutions (the default value is 50 and storing all the grid points can be accomplished by choosing `all`). The `candidates` attribute contains the parameters and `chisqr` from the brute force method as a namedtuple, (`'Candidate', ['params', 'score']`), sorted on the (lowest) `chisqr` value. To access the values for a particular candidate one can use `result.candidate[#].params` or `result.candidate[#].score`, where a lower # represents a better candidate. The `show_candidates(#)` uses the `pretty_print()` method to show a specific candidate-# or all candidates when no number is specified.

The optimal fit is, as usual, stored in the `MinimizerResult.params` attribute and is, therefore, identical to `result_brute.show_candidates(1)`.

```
result_brute.show_candidates(1)
```

Candidate #1, chisqr = 11.935							
Name	Value	Min	Max	Stderr	Vary	Expr	Brute_Step
amp	5	2.5	inf	None	True	None	0.25
decay	0.025	-inf	inf	None	True	None	0.005
omega	2	-inf	5	None	True	None	0.25
shift	-0.1309	-1.571	1.571	None	True	None	None

In this case, the next-best scoring candidate has already a chisqr that increased quite a bit:

```
result_brute.show_candidates(2)
```

Candidate #2, chisqr = 13.994							
Name	Value	Min	Max	Stderr	Vary	Expr	Brute_Step
amp	4.75	2.5	inf	None	True	None	0.25
decay	0.025	-inf	inf	None	True	None	0.005
omega	2	-inf	5	None	True	None	0.25
shift	-0.1309	-1.571	1.571	None	True	None	None

and is, therefore, probably not so likely... However, as said above, in most cases you'll want to do another minimization using the solutions from the brute method as starting values. That can be easily accomplished as shown in the code below, where we now perform a `leastsq` minimization starting from the top-25 solutions and accept the solution if the chisqr is lower than the previously 'optimal' solution:

```
best_result = copy.deepcopy(result_brute)

for candidate in result_brute.candidates:
    trial = fitter.minimize(method='leastsq', params=candidate.params)
    if trial.chisqr < best_result.chisqr:
        best_result = trial
```

From the `leastsq` minimization we obtain the following parameters for the most optimal result:

```
print(fit_report(best_result))
```

```
[[Fit Statistics]]
# fitting method    = leastsq
# function evals   = 21
# data points      = 301
# variables        = 4
chi-square         = 10.8653514
reduced chi-square = 0.03658367
Akaike info crit   = -991.780924
Bayesian info crit = -976.952483
[[Variables]]
amp:   5.00323085 +/- 0.03805940 (0.76%) (init = 5)
decay: 0.02563850 +/- 4.4572e-04 (1.74%) (init = 0.03)
shift: -0.09162987 +/- 0.00978382 (10.68%) (init = 0)
omega:  1.99611629 +/- 0.00316225 (0.16%) (init = 2)
[[Correlations]] (unreported correlations are < 0.100)
C(shift, omega) = -0.7855
```


(continues on next page)

(continued from previous page)

```
C(amp, decay) = +0.5838
C(amp, shift) = -0.1208
```

As expected the parameters have not changed significantly as they were already very close to the “real” values, which can also be appreciated from the plots below.

```
plt.plot(x, data, 'o')
plt.plot(x, data + fcn2min(result_brute.params, x, data), ' -',
         label='brute')
plt.plot(x, data + fcn2min(best_result.params, x, data), ' - -',
         label='brute followed by leastsq')
plt.legend()
```



Finally, the results from the brute force grid-search can be visualized using the rather lengthy Python function below (which might get incorporated in lmfit at some point).

```
def plot_results_brute(result, best_vals=True, varlabels=None,
                      output=None):
    """Visualize the result of the brute force grid search.
```

The output file will display the chi-square value per parameter and contour plots for all combination of two parameters.

(continues on next page)

(continued from previous page)

Inspired by the `corner` package (<https://github.com/dfm/corner.py>).

Parameters

`result : :class:`~lmfit.minimizer.MinimizerResult``
Contains the results from the :meth:`brute` method.

`best_vals : bool, optional`
Whether to show the best values from the grid search (default is True).

`varlabels : list, optional`
If None (default), use `result.var_names` as axis labels, otherwise use the names specified in `varlabels`.

`output : str, optional`
Name of the output PDF file (default is 'None')

"""

```
npars = len(result.var_names)
fig, axes = plt.subplots(npars, npars)

if not varlabels:
    varlabels = result.var_names
if best_vals and isinstance(best_vals, bool):
    best_vals = result.params

for i, par1 in enumerate(result.var_names):
    for j, par2 in enumerate(result.var_names):

        # parameter vs chi2 in case of only one parameter
        if npars == 1:
            axes.plot(result.brute_grid, result.brute_Jout, 'o', ms=3)
            axes.set_ylabel(r'$\chi^2$')
            axes.set_xlabel(varlabels[i])
            if best_vals:
                axes.axvline(best_vals[par1].value, ls='dashed', color='r')

        # parameter vs chi2 profile on top
        elif i == j and j < npars-1:
            if i == 0:
                axes[0, 0].axis('off')
            ax = axes[i, j+1]
            red_axis = tuple(a for a in range(npars) if a != i)
            ax.plot(np.unique(result.brute_grid[i]),
                    np.minimum.reduce(result.brute_Jout, axis=red_axis),
                    'o', ms=3)
            ax.set_ylabel(r'$\chi^2$')
            ax.xaxis.set_label_position("right")
            ax.xaxis.set_ticks_position('right')
            ax.set_xticks([])
            if best_vals:
                ax.axvline(best_vals[par1].value, ls='dashed', color='r')
```

(continues on next page)

(continued from previous page)

```

# parameter vs chi2 profile on the left
elif j == 0 and i > 0:
    ax = axes[i, j]
    red_axis = tuple(a for a in range(npars) if a != i)
    ax.plot(np.minimum.reduce(result.brute_Jout, axis=red_axis),
            np.unique(result.brute_grid[i]), 'o', ms=3)
    ax.invert_xaxis()
    ax.set_ylabel(varlabels[i])
    if i != npars-1:
        ax.set_xticks([])
    else:
        ax.set_xlabel(r'$\chi^2$')
    if best_vals:
        ax.axhline(best_vals[par1].value, ls='dashed', color='r')

# contour plots for all combinations of two parameters
elif j > i:
    ax = axes[j, i+1]
    red_axis = tuple(a for a in range(npars) if a not in (i, j))
    X, Y = np.meshgrid(np.unique(result.brute_grid[i]),
                       np.unique(result.brute_grid[j]))
    lvs1 = np.linspace(result.brute_Jout.min(),
                       np.median(result.brute_Jout)/2.0, 7, dtype='int')
    lvs2 = np.linspace(np.median(result.brute_Jout)/2.0,
                       np.median(result.brute_Jout), 3, dtype='int')
    lvs = np.unique(np.concatenate((lvs1, lvs2)))
    ax.contourf(X.T, Y.T, np.minimum.reduce(result.brute_Jout, axis=red_
axis),
                 lvs, norm=LogNorm())
    ax.set_yticks([])
    if best_vals:
        ax.axvline(best_vals[par1].value, ls='dashed', color='r')
        ax.axhline(best_vals[par2].value, ls='dashed', color='r')
        ax.plot(best_vals[par1].value, best_vals[par2].value, 'rs', ms=3)
    if j != npars-1:
        ax.set_xticks([])
    else:
        ax.set_xlabel(varlabels[i])
    if j - i >= 2:
        axes[i, j].axis('off')

if output is not None:
    plt.savefig(output)


```

and finally, to generate the figure:

```

plot_results_brute(result_brute, best_vals=True, varlabels=None)
plt.show()

```



Total running time of the script: (0 minutes 27.423 seconds)

EXAMPLES FROM THE DOCUMENTATION

Below are all the examples that are part of the lmfit documentation.

14.1 Examples from the documentation

Below are all the examples that are part of the lmfit documentation.

14.1.1 Model - savemodel

```
# <examples/doc_model_savemodel.py>
import numpy as np

from lmfit.model import Model, save_model

def mysine(x, amp, freq, shift):
    return amp * np.sin(x*freq + shift)

sinemodel = Model(mysine)
pars = sinemodel.make_params(amp=1, freq=0.25, shift=0)

save_model(sinemodel, 'sinemodel.sav')
# <end examples/doc_model_savemodel.py>
```

Total running time of the script: (0 minutes 0.004 seconds)

14.1.2 Model - savemodelresult

```
[[Model]]
    Model(gaussian)
[[Fit Statistics]]
    # fitting method    = leastsq
    # function evals   = 33
    # data points     = 101
    # variables       = 3
    chi-square        = 3.40883599
```

(continues on next page)

(continued from previous page)

```

reduced chi-square = 0.03478404
Akaike info crit   = -336.263713
Bayesian info crit = -328.418352
R-squared           = 0.98533348
[[Variables]]
  amplitude: 8.88021907 +/- 0.11359530 (1.28%) (init = 5)
  center:    5.65866105 +/- 0.01030493 (0.18%) (init = 5)
  sigma:     0.69765480 +/- 0.01030508 (1.48%) (init = 1)
  fwhm:      1.64285148 +/- 0.02426660 (1.48%) == '2.3548200*sigma'
  height:    5.07800563 +/- 0.06495769 (1.28%) == '0.3989423*amplitude/max(1e-15,_
  sigma)'

[[Correlations]] (unreported correlations are < 0.100)
  C(amplitude, sigma) = +0.5774

```

```

# <examples/doc_model_savemodelresult.py>
import numpy as np

from lmfit.model import save_modelresult
from lmfit.models import GaussianModel

data = np.loadtxt('model1d_gauss.dat')
x = data[:, 0]
y = data[:, 1]

gmodel = GaussianModel()
result = gmodel.fit(y, x=x, amplitude=5, center=5, sigma=1)

save_modelresult(result, 'gauss_modelresult.sav')

print(result.fit_report())
# <end examples/doc_model_savemodelresult.py>

```

Total running time of the script: (0 minutes 0.060 seconds)

14.1.3 Confidence - basic

```

[[Variables]]
  a: 0.09943896 +/- 1.9322e-04 (0.19%) (init = 0.1)
  b: 1.98476942 +/- 0.01222678 (0.62%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
  C(a, b) = +0.6008
    99.73%   95.45%   68.27%   _BEST_   68.27%   95.45%   99.73%
  a: -0.00059  -0.00039  -0.00019   0.09944  +0.00019  +0.00039  +0.00060
  b: -0.03764  -0.02477  -0.01229   1.98477  +0.01229  +0.02477  +0.03764

```

```
# <examples/doc_confidence_basic.py>
import numpy as np

import lmfit

x = np.linspace(0.3, 10, 100)
np.random.seed(0)
y = 1/(0.1*x) + 2 + 0.1*np.random.randn(x.size)

pars = lmfit.create_params(a=0.1, b=1)

def residual(p):
    return 1/(p['a']*x) + p['b'] - y


mini = lmfit.Minimizer(residual, pars)
result = mini.minimize()

print(lmfit.fit_report(result.params))

ci = lmfit.conf_interval(mini, result)
lmfit.printfuncs.report_ci(ci)
# <end examples/doc_confidence_basic.py>
```

Total running time of the script: (0 minutes 0.173 seconds)

14.1.4 Model - loadmodelresult



```
[[Model]]
  Model(gaussian)
[[Fit Statistics]]
  # fitting method = leastsq
  # function evals = 33
  # data points   = 101
  # variables     = 3
  chi-square     = 3.40883599
  reduced chi-square = 0.03478404
  Akaike info crit = -336.263713
  Bayesian info crit = -328.418352
  R-squared       = 0.98533348
[[Variables]]
  amplitude: 8.88021907 +/- 0.11359530 (1.28%) (init = 5)
  center:    5.65866105 +/- 0.01030493 (0.18%) (init = 5)
  sigma:     0.69765480 +/- 0.01030508 (1.48%) (init = 1)
  fwhm:      1.64285148 +/- 0.02426660 (1.48%) == '2.3548200*sigma'
  height:    5.07800563 +/- 0.06495769 (1.28%) == '0.3989423*amplitude/max(1e-15,sigma)'
[[Correlations]] (unreported correlations are < 0.100)
  C(amplitude, sigma) = +0.5774
```

```
# <examples/doc_model_loadmodelresult.py>
import os
import sys

import matplotlib.pyplot as plt
import numpy as np

from lmfit.model import load_modelresult

if not os.path.exists('gauss_modelresult.sav'):
    os.system(f"{sys.executable} doc_model_savemodelresult.py")


data = np.loadtxt('model1d_gauss.dat')
x = data[:, 0]
y = data[:, 1]

result = load_modelresult('gauss_modelresult.sav')
print(result.fit_report())

plt.plot(x, y, 'o')
plt.plot(x, result.best_fit, '-')
plt.show()
# <end examples/doc_model_loadmodelresult.py>
```

Total running time of the script: (0 minutes 0.164 seconds)

14.1.5 Model - loadmodelresult2



```
[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
  ↪ prefix='exp_'))
[[Fit Statistics]]
  # fitting method    = leastsq
  # function evals   = 46
  # data points     = 250
  # variables       = 8
  chi-square        = 1247.52821
  reduced chi-square = 5.15507524
  Akaike info crit  = 417.864631
  Bayesian info crit = 446.036318
  R-squared          = 0.99648654
[[Variables]]
  exp_amplitude: 99.0183278 +/- 0.53748593 (0.54%) (init = 162.2102)
  exp_decay:     90.9508853 +/- 1.10310778 (1.21%) (init = 93.24905)
  g1_amplitude:  4257.77360 +/- 42.3836478 (1.00%) (init = 2000)
  g1_center:     107.030956 +/- 0.15006851 (0.14%) (init = 105)
  g1_sigma:      16.6725772 +/- 0.16048381 (0.96%) (init = 15)
  g1_fwhm:       39.2609181 +/- 0.37791049 (0.96%) == '2.3548200*g1_sigma'
  g1_height:     101.880230 +/- 0.59217173 (0.58%) == '0.3989423*g1_amplitude/max(1e-
```

(continues on next page)

(continued from previous page)

```

→15, g1_sigma)'
g2_amplitude: 2493.41735 +/- 36.1697789 (1.45%) (init = 2000)
g2_center: 153.270102 +/- 0.19466802 (0.13%) (init = 155)
g2_sigma: 13.8069464 +/- 0.18679695 (1.35%) (init = 15)
g2_fwhm: 32.5128735 +/- 0.43987320 (1.35%) == '2.3548200*g2_sigma'
g2_height: 72.0455941 +/- 0.61722243 (0.86%) == '0.3989423*g2_amplitude/max(1e-
→15, g2_sigma)'

[[Correlations]] (unreported correlations are < 0.100)
C(g1_amplitude, g1_sigma) = +0.8243
C(g2_amplitude, g2_sigma) = +0.8154
C(exp_amplitude, exp_decay) = -0.6946
C(g1_sigma, g2_center) = +0.6842
C(g1_center, g2_amplitude) = -0.6689
C(g1_center, g2_sigma) = -0.6520
C(g1_amplitude, g2_center) = +0.6477
C(g1_center, g2_center) = +0.6205
C(g1_center, g1_sigma) = +0.5075
C(exp_decay, g1_amplitude) = -0.5074
C(g1_sigma, g2_amplitude) = -0.4915
C(g2_center, g2_sigma) = -0.4889
C(g1_sigma, g2_sigma) = -0.4826
C(g2_amplitude, g2_center) = -0.4763
C(exp_decay, g2_amplitude) = -0.4270
C(g1_amplitude, g1_center) = +0.4183
C(g1_amplitude, g2_sigma) = -0.4010
C(g1_amplitude, g2_amplitude) = -0.3071
C(exp_amplitude, g2_amplitude) = +0.2821
C(exp_decay, g1_sigma) = -0.2520
C(exp_decay, g2_sigma) = -0.2329
C(exp_amplitude, g2_sigma) = +0.1714
C(exp_decay, g2_center) = -0.1514
C(exp_amplitude, g1_amplitude) = +0.1478
C(exp_decay, g1_center) = +0.1055

```

```

# <examples/doc_model_loadmodelresult2.py>
import os
import sys

import matplotlib.pyplot as plt
import numpy as np

from lmfit.model import load_modelresult

if not os.path.exists('nistgauss_modelresult.sav'):
    os.system(f'{sys.executable} doc_model_savemodelresult2.py')

dat = np.loadtxt('NIST_Gauss2.dat')

```

(continues on next page)

(continued from previous page)

```

x = dat[:, 1]
y = dat[:, 0]


result = load_modelresult('nistgauss_modelresult.sav')
print(result.fit_report())

plt.plot(x, y, 'o')
plt.plot(x, result.best_fit, '-')
plt.show()
# <end examples/doc_model_loadmodelresult2.py>

```

Total running time of the script: (0 minutes 0.182 seconds)

14.1.6 Model - gaussian



```

[[Model]]
  Model(gaussian)
[[Fit Statistics]]
  # fitting method = leastsq
  # function evals = 33
  # data points = 101

```

(continues on next page)

(continued from previous page)

```

# variables      = 3
chi-square      = 3.40883599
reduced chi-square = 0.03478404
Akaike info crit  = -336.263713
Bayesian info crit = -328.418352
R-squared        = 0.98533348
[[Variables]]
amp: 8.88021893 +/- 0.11359522 (1.28%) (init = 5)
cen: 5.65866102 +/- 0.01030495 (0.18%) (init = 5)
wid: 0.69765478 +/- 0.01030505 (1.48%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(amp, wid) = +0.5774

```

```

# <examples/doc_model_gaussian.py>
import matplotlib.pyplot as plt
from numpy import exp, loadtxt, pi, sqrt

from lmfit import Model

data = loadtxt('model1d_gauss.dat')
x = data[:, 0]
y = data[:, 1]

def gaussian(x, amp, cen, wid):
    """1-d gaussian: gaussian(x, amp, cen, wid)"""
    return (amp / (sqrt(2*pi) * wid)) * exp(-(x-cen)**2 / (2*wid**2))

gmodel = Model(gaussian)
result = gmodel.fit(y, x=x, amp=5, cen=5, wid=1)


print(result.fit_report())

plt.plot(x, y, 'o')
plt.plot(x, result.init_fit, '--', label='initial fit')
plt.plot(x, result.best_fit, '-', label='best fit')
plt.legend()
plt.show()
# <end examples/doc_model_gaussian.py>

```

Total running time of the script: (0 minutes 0.181 seconds)

14.1.7 Model - loadmodel



```
[[Model]]
  Model(my sine)
[[Fit Statistics]]
  # fitting method = leastsq
  # function evals = 25
  # data points = 101
  # variables = 3
  chi-square = 7.68903767
  reduced chi-square = 0.07845957
  Akaike info crit = -254.107813
  Bayesian info crit = -246.262451
  R-squared = 0.97252133
[[Variables]]
  amp: 2.32733694 +/- 0.03950824 (1.70%) (init = 3)
  freq: 0.50098739 +/- 5.7726e-04 (0.12%) (init = 0.52)
  shift: 0.53605324 +/- 0.03383110 (6.31%) (init = 0)
[[Correlations]] (unreported correlations are < 0.100)
  C(freq, shift) = -0.8663
```

```

# <examples/doc_model_loadmodel.py>
import os
import sys

import matplotlib.pyplot as plt
import numpy as np

from lmfit.model import load_model

if not os.path.exists('sinemodel.sav'):
    os.system(f"{sys.executable} doc_model_savemodel.py")

def mysine(x, amp, freq, shift):
    return amp * np.sin(x*freq + shift)

data = np.loadtxt('sinedata.dat')
x = data[:, 0]
y = data[:, 1]

model = load_model('sinemodel.sav', funcdefs={'mysine': mysine})
params = model.make_params(amp=dict(value=3, min=0),
                           freq=0.52,
                           shift=dict(value=0, min=-1, max=1))


result = model.fit(y, params, x=x)
print(result.fit_report())

plt.plot(x, y, 'o')
plt.plot(x, result.best_fit, '-')
plt.show()
# <end examples/doc_model_loadmodel.py>

```

Total running time of the script: (0 minutes 0.175 seconds)

14.1.8 Model - with nan policy



```
[[Model]]
  Model(gaussian)
[[Fit Statistics]]
  # fitting method = leastsq
  # function evals = 22
  # data points   = 99
  # variables     = 3
  chi-square     = 3.27990355
  reduced chi-square = 0.03416566
  Akaike info crit = -331.323278
  Bayesian info crit = -323.537918
  R-squared       = 0.98570688
[[Variables]]
  amplitude: 8.82064881 +/- 0.11686114 (1.32%) (init = 5)
  center:    5.65906365 +/- 0.01055590 (0.19%) (init = 6)
  sigma:     0.69165307 +/- 0.01060640 (1.53%) (init = 1)
  fwhm:      1.62871849 +/- 0.02497615 (1.53%) == '2.3548200*sigma'
  height:    5.08770952 +/- 0.06488251 (1.28%) == '0.3989423*amplitude/max(1e-15,_
  sigma)'
[[Correlations]] (unreported correlations are < 0.100)
  C(amplitude, sigma) = +0.6105
```

```
# <examples/doc_model_with_nan_policy.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit.models import GaussianModel

data = np.loadtxt('model1d_gauss.dat')
x = data[:, 0]
y = data[:, 1]

y[44] = np.nan
y[65] = np.nan

# nan_policy = 'raise'
# nan_policy = 'propagate'
nan_policy = 'omit'

gmodel = GaussianModel()
result = gmodel.fit(y, x=x, amplitude=5, center=6, sigma=1,
                     nan_policy=nan_policy)


print(result.fit_report())

# make sure nans are removed for plotting:
x_ = x[np.where(np.isfinite(y))]
y_ = y[np.where(np.isfinite(y))]

plt.plot(x_, y_, 'o')
plt.plot(x_, result.init_fit, '--', label='initial fit')
plt.plot(x_, result.best_fit, '-', label='best fit')
plt.legend()
plt.show()
# <end examples/doc_model_with_nan_policy.py>
```

Total running time of the script: (0 minutes 0.183 seconds)

14.1.9 Builtinmodels - stepmodel



```
[[Model]]
(Model(step, prefix='step_', form='erf') + Model(linear, prefix='line_'))
[[Fit Statistics]]
# fitting method    = leastsq
# function evals   = 55
# data points      = 201
# variables        = 5
chi-square         = 593.709621
reduced chi-square = 3.02913072
Akaike info crit   = 227.700173
Bayesian info crit = 244.216697
R-squared          = 0.99897798
[[Variables]]
line_slope:      1.87162383 +/- 0.09318592 (4.98%) (init = 0)
line_intercept:  12.0964588 +/- 0.27606017 (2.28%) (init = 11.58574)
step_amplitude: 112.858576 +/- 0.65391731 (0.58%) (init = 134.7378)
step_center:    3.13494787 +/- 0.00516602 (0.16%) (init = 2.5)
step_sigma:     0.67393440 +/- 0.01091158 (1.62%) (init = 1.428571)
[[Correlations]] (unreported correlations are < 0.100)
C(line_slope, step_amplitude)   = -0.8791
C(step_amplitude, step_sigma)   = +0.5643
```

(continues on next page)

(continued from previous page)

```

C(line_slope, step_sigma)      = -0.4569
C(line_intercept, step_center) = +0.4269
C(line_slope, line_intercept)  = -0.3093
C(line_slope, step_center)    = -0.2338
C(line_intercept, step_sigma)  = -0.1372
C(line_intercept, step_amplitude) = -0.1173
C(step_amplitude, step_center) = +0.1095

```

```

# <examples/doc_builtinmodels_stepmodel.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit.models import LinearModel, StepModel

x = np.linspace(0, 10, 201)
y = np.ones_like(x)
y[:48] = 0.0
y[48:77] = np.arange(77-48)/(77.0-48)
np.random.seed(0)
y = 110.2 * (y + 9e-3*np.random.randn(x.size)) + 12.0 + 2.22*x

step_mod = StepModel(form='erf', prefix='step_')
line_mod = LinearModel(prefix='line_')

pars = line_mod.make_params(intercept=y.min(), slope=0)
pars += step_mod.guess(y, x=x, center=2.5)

mod = step_mod + line_mod
out = mod.fit(y, pars, x=x)


print(out.fit_report())

plt.plot(x, y)
plt.plot(x, out.init_fit, '--', label='initial fit')
plt.plot(x, out.best_fit, '-', label='best fit')
plt.legend()
plt.show()
# <end examples/doc_builtinmodels_stepmodel.py>

```

Total running time of the script: (0 minutes 0.189 seconds)

14.1.10 Model - uncertainty



```
[[Model]]
  Model(gaussian)
[[Fit Statistics]]
  # fitting method    = leastsq
  # function evals   = 33
  # data points      = 101
  # variables        = 3
  chi-square         = 3.40883599
  reduced chi-square = 0.03478404
  Akaike info crit   = -336.263713
  Bayesian info crit = -328.418352
  R-squared          = 0.98533348
[[Variables]]
  amp: 8.88021893 +/- 0.11359522 (1.28%) (init = 5)
  cen: 5.65866102 +/- 0.01030495 (0.18%) (init = 5)
  wid: 0.69765478 +/- 0.01030505 (1.48%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
  C(amp, wid) = +0.5774
```

```

# <examples/doc_model_uncertainty.py>
import matplotlib.pyplot as plt
from numpy import exp, loadtxt, pi, sqrt

from lmfit import Model

data = loadtxt('model1d_gauss.dat')
x = data[:, 0]
y = data[:, 1]

def gaussian(x, amp, cen, wid):
    """1-d gaussian: gaussian(x, amp, cen, wid)"""
    return (amp / (sqrt(2*pi) * wid)) * exp(-(x-cen)**2 / (2*wid**2))

gmodel = Model(gaussian)
result = gmodel.fit(y, x=x, amp=5, cen=5, wid=1)

print(result.fit_report())


dely = result.eval_uncertainty(sigma=3)

plt.plot(x, y, 'o')
plt.plot(x, result.init_fit, '--', label='initial fit')
plt.plot(x, result.best_fit, '-', label='best fit')
plt.fill_between(x, result.best_fit-dely, result.best_fit+dely,
                 color="#ABABAB", label=r'3-$\sigma$ uncertainty band')
plt.legend()
plt.show()
# <end examples/doc_model_uncertainty.py>

```

Total running time of the script: (0 minutes 0.199 seconds)

14.1.11 Model - two components



```
[[Model]]
(Model(gaussian) + Model(line))
[[Fit Statistics]]
# fitting method = leastsq
# function evals = 55
# data points = 101
# variables = 5
chi-square = 2.57855517
reduced chi-square = 0.02685995
Akaike info crit = -360.457020
Bayesian info crit = -347.381417
R-squared = 0.99194643
[[Variables]]
amp: 8.45930976 +/- 0.12414531 (1.47%) (init = 5)
cen: 5.65547889 +/- 0.00917673 (0.16%) (init = 5)
wid: 0.67545513 +/- 0.00991697 (1.47%) (init = 1)
slope: 0.26484403 +/- 0.00574892 (2.17%) (init = 0)
intercept: -0.96860189 +/- 0.03352202 (3.46%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
C(slope, intercept) = -0.7954
C(amp, wid) = +0.6664
```

(continues on next page)

(continued from previous page)

C(amp, intercept)	= -0.2216
C(amp, slope)	= -0.1692
C(cen, slope)	= -0.1618
C(wid, intercept)	= -0.1477
C(cen, intercept)	= +0.1287
C(wid, slope)	= -0.1127

```

# <examples/doc_model_two_components.py>
import matplotlib.pyplot as plt
from numpy import exp, loadtxt, pi, sqrt

from lmfit import Model

data = loadtxt('model1d_gauss.dat')
x = data[:, 0]
y = data[:, 1] + 0.25*x - 1.0

def gaussian(x, amp, cen, wid):
    """1-d gaussian: gaussian(x, amp, cen, wid)"""
    return (amp / (sqrt(2*pi) * wid)) * exp(-(x-cen)**2 / (2*wid**2))

def line(x, slope, intercept):
    """a line"""
    return slope*x + intercept

mod = Model(gaussian) + Model(line)
pars = mod.make_params(amp=5, cen=5, wid={'value': 1, 'min': 0},
                      slope=0, intercept=1)

result = mod.fit(y, pars, x=x)
print(result.fit_report())

plt.plot(x, y, 'o')
plt.plot(x, result.init_fit, '--', label='initial fit')
plt.plot(x, result.best_fit, ' ', label='best fit')
plt.legend()
plt.show()
# <end examples/doc_model_two_components.py>

```

Total running time of the script: (0 minutes 0.192 seconds)

14.1.12 Fitting - withreport

```
[[Fit Statistics]]
# fitting method = leastsq
# function evals = 83
# data points = 1001
# variables = 4
chi-square = 498.811759
reduced chi-square = 0.50031270
Akaike info crit = -689.222517
Bayesian info crit = -669.587497
[[Variables]]
amp: 13.9121959 +/- 0.14120321 (1.01%) (init = 13)
period: 5.48507038 +/- 0.02666520 (0.49%) (init = 2)
shift: 0.16203673 +/- 0.01405662 (8.67%) (init = 0)
decay: 0.03264539 +/- 3.8015e-04 (1.16%) (init = 0.02)
[[Correlations]] (unreported correlations are < 0.100)
C(period, shift) = +0.7974
C(amp, decay) = +0.5816
C(amp, shift) = -0.2966
C(amp, period) = -0.2432
C(shift, decay) = -0.1819
C(period, decay) = -0.1496
```

```
# <examples/doc_fitting_withreport.py>
from numpy import exp, linspace, pi, random, sign, sin

from lmfit import create_params, fit_report, minimize

p_true = create_params(amp=14.0, period=5.46, shift=0.123, decay=0.032)

def residual(pars, x, data=None):
    """Model a decaying sine wave and subtract data."""
    vals = pars.valuesdict()
    amp = vals['amp']
    per = vals['period']
    shift = vals['shift']
    decay = vals['decay']

    if abs(shift) > pi/2:
        shift = shift - sign(shift)*pi
    model = amp * sin(shift + x/per) * exp(-x*x*decay*decay)
    if data is None:
        return model
    return model - data
```

(continues on next page)

(continued from previous page)

```

random.seed(0)
x = linspace(0.0, 250., 1001)
noise = random.normal(scale=0.7215, size=x.size)
data = residual(p_true, x) + noise

fit_params = create_params(amp=13, period=2, shift=0, decay=0.02)

out = minimize(residual, fit_params, args=(x,), kws={'data': data})

print(fit_report(out))
# <end examples/doc_fitting_withreport.py>

```

Total running time of the script: (0 minutes 0.009 seconds)

14.1.13 Model - savemodelresult2

```

[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
  ↪ prefix='exp_'))
[[Fit Statistics]]
  # fitting method    = leastsq
  # function evals   = 46
  # data points      = 250
  # variables        = 8
  chi-square         = 1247.52821
  reduced chi-square = 5.15507524
  Akaike info crit   = 417.864631
  Bayesian info crit = 446.036318
  R-squared          = 0.99648654
[[Variables]]
  exp_amplitude: 99.0183278 +/- 0.53748593 (0.54%) (init = 162.2102)
  exp_decay:     90.9508853 +/- 1.10310778 (1.21%) (init = 93.24905)
  g1_amplitude:  4257.77360 +/- 42.3836478 (1.00%) (init = 2000)
  g1_center:     107.030956 +/- 0.15006851 (0.14%) (init = 105)
  g1_sigma:       16.6725772 +/- 0.16048381 (0.96%) (init = 15)
  g1_fwhm:        39.2609181 +/- 0.37791049 (0.96%) == '2.3548200*g1_sigma'
  g1_height:     101.880230 +/- 0.59217173 (0.58%) == '0.3989423*g1_amplitude/max(1e-
  ↪ 15, g1_sigma)'
  g2_amplitude:  2493.41735 +/- 36.1697789 (1.45%) (init = 2000)
  g2_center:     153.270102 +/- 0.19466802 (0.13%) (init = 155)
  g2_sigma:       13.8069464 +/- 0.18679695 (1.35%) (init = 15)
  g2_fwhm:        32.5128735 +/- 0.43987320 (1.35%) == '2.3548200*g2_sigma'
  g2_height:     72.0455941 +/- 0.61722243 (0.86%) == '0.3989423*g2_amplitude/max(1e-
  ↪ 15, g2_sigma)'
[[Correlations]] (unreported correlations are < 0.100)
  C(g1_amplitude, g1_sigma)      = +0.8243
  C(g2_amplitude, g2_sigma)      = +0.8154
  C(exp_amplitude, exp_decay)   = -0.6946
  C(g1_sigma, g2_center)         = +0.6842
  C(g1_center, g2_amplitude)     = -0.6689
  C(g1_center, g2_sigma)         = -0.6520

```

(continues on next page)

(continued from previous page)

C(g1_amplitude, g2_center)	= +0.6477
C(g1_center, g2_center)	= +0.6205
C(g1_center, g1_sigma)	= +0.5075
C(exp_decay, g1_amplitude)	= -0.5074
C(g1_sigma, g2_amplitude)	= -0.4915
C(g2_center, g2_sigma)	= -0.4889
C(g1_sigma, g2_sigma)	= -0.4826
C(g2_amplitude, g2_center)	= -0.4763
C(exp_decay, g2_amplitude)	= -0.4270
C(g1_amplitude, g1_center)	= +0.4183
C(g1_amplitude, g2_sigma)	= -0.4010
C(g1_amplitude, g2_amplitude)	= -0.3071
C(exp_amplitude, g2_amplitude)	= +0.2821
C(exp_decay, g1_sigma)	= -0.2520
C(exp_decay, g2_sigma)	= -0.2329
C(exp_amplitude, g2_sigma)	= +0.1714
C(exp_decay, g2_center)	= -0.1514
C(exp_amplitude, g1_amplitude)	= +0.1478
C(exp_decay, g1_center)	= +0.1055

```
# <examples/doc_model_savemodelresult2.py>
import numpy as np

from lmfit.model import save_modelresult
from lmfit.models import ExponentialModel, GaussianModel

dat = np.loadtxt('NIST_Gauss2.dat')
x = dat[:, 1]
y = dat[:, 0]

exp_mod = ExponentialModel(prefix='exp_')
pars = exp_mod.guess(y, x=x)

gauss1 = GaussianModel(prefix='g1_')
pars.update(gauss1.make_params(center=dict(value=105, min=75, max=125),
                               sigma=dict(value=15, min=0),
                               amplitude=dict(value=2000, min=0)))

gauss2 = GaussianModel(prefix='g2_')
pars.update(gauss2.make_params(center=dict(value=155, min=125, max=175),
                               sigma=dict(value=15, min=0),
                               amplitude=dict(value=2000, min=0)))

mod = gauss1 + gauss2 + exp_mod

init = mod.eval(pars, x=x)
```

(continues on next page)

(continued from previous page)

```

result = mod.fit(y, pars, x=x)


save_modelresult(result, 'nistgauss_modelresult.sav')

print(result.fit_report())
# <end examples/doc_model_savemodelresult2.py>

```

Total running time of the script: (0 minutes 0.085 seconds)

14.1.14 Builtinmodels - nistgauss2



```

[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
  ↪ prefix='exp_'))
[[Fit Statistics]]
  # fitting method = leastsq
  # function evals = 37
  # data points   = 250
  # variables     = 8
  chi-square      = 1247.52821
  reduced chi-square = 5.15507524

```

(continues on next page)

(continued from previous page)

```

Akaike info crit    = 417.864631
Bayesian info crit = 446.036318
R-squared           = 0.99648654
[[Variables]]
exp_amplitude: 99.0183265 +/- 0.53748764 (0.54%) (init = 94.53724)
exp_decay:      90.9508884 +/- 1.10310753 (1.21%) (init = 111.1985)
g1_amplitude:   4257.77384 +/- 42.3839276 (1.00%) (init = 3189.648)
g1_center:      107.030957 +/- 0.15006934 (0.14%) (init = 106.5)
g1_sigma:        16.6725783 +/- 0.16048220 (0.96%) (init = 14.5)
g1_fwhm:        39.2609209 +/- 0.37790669 (0.96%) == '2.3548200*g1_sigma'
g1_height:      101.880228 +/- 0.59216965 (0.58%) == '0.3989423*g1_amplitude/max(1e-
˓→15, g1_sigma)'
g2_amplitude:   2493.41698 +/- 36.1699974 (1.45%) (init = 2818.337)
g2_center:      153.270103 +/- 0.19466966 (0.13%) (init = 150)
g2_sigma:        13.8069440 +/- 0.18680331 (1.35%) (init = 15)
g2_fwhm:        32.5128679 +/- 0.43988818 (1.35%) == '2.3548200*g2_sigma'
g2_height:      72.0455954 +/- 0.61722288 (0.86%) == '0.3989423*g2_amplitude/max(1e-
˓→15, g2_sigma)'
[[Correlations]] (unreported correlations are < 0.500)
C(g1_amplitude, g1_sigma)    = +0.8243
C(g2_amplitude, g2_sigma)    = +0.8154
C(exp_amplitude, exp_decay) = -0.6946
C(g1_sigma, g2_center)       = +0.6842
C(g1_center, g2_amplitude)   = -0.6689
C(g1_center, g2_sigma)       = -0.6521
C(g1_amplitude, g2_center)   = +0.6477
C(g1_center, g2_center)      = +0.6205
C(g1_center, g1_sigma)       = +0.5075
C(exp_decay, g1_amplitude)   = -0.5074

```

```

# <examples/doc_nistgauss2.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit.models import ExponentialModel, GaussianModel

dat = np.loadtxt('NIST_Gauss2.dat')
x = dat[:, 1]
y = dat[:, 0]

exp_mod = ExponentialModel(prefix='exp_')
gauss1 = GaussianModel(prefix='g1_')
gauss2 = GaussianModel(prefix='g2_')

def index_of(arrval, value):
    """Return index of array *at or below* value."""

```

(continues on next page)

(continued from previous page)

```
if value < min(arrval):
    return 0
return max(np.where(arrval <= value)[0])

ix1 = index_of(x, 75)
ix2 = index_of(x, 135)
ix3 = index_of(x, 175)

pars1 = exp_mod.guess(y[:ix1], x=x[:ix1])
pars2 = gauss1.guess(y[ix1:ix2], x=x[ix1:ix2])
pars3 = gauss2.guess(y[ix2:ix3], x=x[ix2:ix3])

pars = pars1 + pars2 + pars3
mod = gauss1 + gauss2 + exp_mod


out = mod.fit(y, pars, x=x)

print(out.fit_report(min_correl=0.5))

plt.plot(x, y)
plt.plot(x, out.init_fit, '--', label='initial fit')
plt.plot(x, out.best_fit, '-', label='best fit')
plt.legend()
plt.show()
# <end examples/doc_nistgauss2.py>
```

Total running time of the script: (0 minutes 0.217 seconds)

14.1.15 Model - with iter callback



```

ITER -1 ['peak_amplitude = 3.00000', 'peak_center = 6.00000', 'peak_sigma = 2.00000',
         'bkg_slope = 0.00000', 'bkg_intercept = 0.00000', 'peak_fwhm = 4.70964', 'peak_height_'
         ' = 0.59841']
ITER  0 ['peak_amplitude = 3.00000', 'peak_center = 6.00000', 'peak_sigma = 2.00000',
         'bkg_slope = 0.00000', 'bkg_intercept = 0.00000', 'peak_fwhm = 4.70964', 'peak_height_'
         ' = 0.59841']
ITER  1 ['peak_amplitude = 3.00000', 'peak_center = 6.00000', 'peak_sigma = 2.00000',
         'bkg_slope = 0.00000', 'bkg_intercept = 0.00000', 'peak_fwhm = 4.70964', 'peak_height_'
         ' = 0.59841']
ITER  2 ['peak_amplitude = 3.00004', 'peak_center = 6.00000', 'peak_sigma = 2.00000',
         'bkg_slope = 0.00000', 'bkg_intercept = 0.00000', 'peak_fwhm = 4.70964', 'peak_height_'
         ' = 0.59842']
ITER  3 ['peak_amplitude = 3.00000', 'peak_center = 6.00004', 'peak_sigma = 2.00000',
         'bkg_slope = 0.00000', 'bkg_intercept = 0.00000', 'peak_fwhm = 4.70964', 'peak_height_'
         ' = 0.59841']
ITER  4 ['peak_amplitude = 3.00000', 'peak_center = 6.00000', 'peak_sigma = 2.00003',
         'bkg_slope = 0.00000', 'bkg_intercept = 0.00000', 'peak_fwhm = 4.70970', 'peak_height_'
         ' = 0.59841']
ITER  5 ['peak_amplitude = 3.00000', 'peak_center = 6.00000', 'peak_sigma = 2.00000',
         'bkg_slope = 0.00001', 'bkg_intercept = 0.00000', 'peak_fwhm = 4.70964', 'peak_height_'
         ' = 0.59841']

```

(continues on next page)

(continued from previous page)

```

ITER  6 ['peak_amplitude = 3.00000', 'peak_center = 6.00000', 'peak_sigma = 2.00000',
↪ 'bkg_slope = 0.00000', 'bkg_intercept = 0.00001', 'peak_fwhm = 4.70964', 'peak_height_↪
↪ = 0.59841']

ITER  7 ['peak_amplitude = 28.06096', 'peak_center = 19.99268', 'peak_sigma = 6.62280',
↪ 'bkg_slope = -0.20866', 'bkg_intercept = 3.28814', 'peak_fwhm = 15.59550', 'peak_↪
↪ height = 1.69033']

ITER  8 ['peak_amplitude = 28.06125', 'peak_center = 19.99268', 'peak_sigma = 6.62280',
↪ 'bkg_slope = -0.20866', 'bkg_intercept = 3.28814', 'peak_fwhm = 15.59550', 'peak_↪
↪ height = 1.69034']

ITER  9 ['peak_amplitude = 28.06096', 'peak_center = 19.99267', 'peak_sigma = 6.62280',
↪ 'bkg_slope = -0.20866', 'bkg_intercept = 3.28814', 'peak_fwhm = 15.59550', 'peak_↪
↪ height = 1.69033']

ITER  10 ['peak_amplitude = 28.06096', 'peak_center = 19.99268', 'peak_sigma = 6.62288',
↪ 'bkg_slope = -0.20866', 'bkg_intercept = 3.28814', 'peak_fwhm = 15.59568', 'peak_↪
↪ height = 1.69031']

ITER  11 ['peak_amplitude = 28.06096', 'peak_center = 19.99268', 'peak_sigma = 6.62280',
↪ 'bkg_slope = -0.20866', 'bkg_intercept = 3.28814', 'peak_fwhm = 15.59550', 'peak_↪
↪ height = 1.69033']

ITER  12 ['peak_amplitude = 28.06096', 'peak_center = 19.99268', 'peak_sigma = 6.62280',
↪ 'bkg_slope = -0.20866', 'bkg_intercept = 3.28818', 'peak_fwhm = 15.59550', 'peak_↪
↪ height = 1.69033']

ITER  13 ['peak_amplitude = 69.61513', 'peak_center = 11.64712', 'peak_sigma = 18.26748
↪', 'bkg_slope = 0.15397', 'bkg_intercept = 3.00619', 'peak_fwhm = 43.01664', 'peak_↪
↪ height = 1.52032']

ITER  14 ['peak_amplitude = 26.14334', 'peak_center = 19.99783', 'peak_sigma = 17.92626
↪', 'bkg_slope = -0.18990', 'bkg_intercept = 3.74340', 'peak_fwhm = 42.21311', 'peak_↪
↪ height = 0.58181']

ITER  15 ['peak_amplitude = 26.14361', 'peak_center = 19.99783', 'peak_sigma = 17.92626
↪', 'bkg_slope = -0.18990', 'bkg_intercept = 3.74340', 'peak_fwhm = 42.21311', 'peak_↪
↪ height = 0.58182']

ITER  16 ['peak_amplitude = 26.14334', 'peak_center = 19.99783', 'peak_sigma = 17.92626
↪', 'bkg_slope = -0.18990', 'bkg_intercept = 3.74340', 'peak_fwhm = 42.21311', 'peak_↪
↪ height = 0.58181']

ITER  17 ['peak_amplitude = 26.14334', 'peak_center = 19.99783', 'peak_sigma = 17.92644
↪', 'bkg_slope = -0.18990', 'bkg_intercept = 3.74340', 'peak_fwhm = 42.21355', 'peak_↪
↪ height = 0.58180']

ITER  18 ['peak_amplitude = 26.14334', 'peak_center = 19.99783', 'peak_sigma = 17.92626
↪', 'bkg_slope = -0.18989', 'bkg_intercept = 3.74340', 'peak_fwhm = 42.21311', 'peak_↪
↪ height = 0.58181']

ITER  19 ['peak_amplitude = 26.14334', 'peak_center = 19.99783', 'peak_sigma = 17.92626
↪', 'bkg_slope = -0.18990', 'bkg_intercept = 3.74343', 'peak_fwhm = 42.21311', 'peak_↪
↪ height = 0.58181']

ITER  20 ['peak_amplitude = 28.96026', 'peak_center = 19.98918', 'peak_sigma = 18.29596
↪', 'bkg_slope = -0.27575', 'bkg_intercept = 4.75376', 'peak_fwhm = 43.08370', 'peak_↪
↪ height = 0.63148']

ITER  21 ['peak_amplitude = 28.96056', 'peak_center = 19.98918', 'peak_sigma = 18.29596
↪', 'bkg_slope = -0.27575', 'bkg_intercept = 4.75376', 'peak_fwhm = 43.08370', 'peak_↪
↪ height = 0.63148']

ITER  22 ['peak_amplitude = 28.96026', 'peak_center = 19.98917', 'peak_sigma = 18.29596
↪', 'bkg_slope = -0.27575', 'bkg_intercept = 4.75376', 'peak_fwhm = 43.08370', 'peak_↪
↪ height = 0.63148']

ITER  23 ['peak_amplitude = 28.96026', 'peak_center = 19.98918', 'peak_sigma = 18.29616
↪', 'bkg_slope = -0.27575', 'bkg_intercept = 4.75376', 'peak_fwhm = 43.08370', 'peak_↪
↪ height = 0.63148']

```

(continues on next page)

(continued from previous page)

```

', 'bkg_slope = -0.27575', 'bkg_intercept = 4.75376', 'peak_fwhm = 43.08416', 'peak_
height = 0.63147']
ITER 24 ['peak_amplitude = 28.96026', 'peak_center = 19.98918', 'peak_sigma = 18.29596
', 'bkg_slope = -0.27574', 'bkg_intercept = 4.75376', 'peak_fwhm = 43.08370', 'peak_
height = 0.63148']
ITER 25 ['peak_amplitude = 28.96026', 'peak_center = 19.98918', 'peak_sigma = 18.29596
', 'bkg_slope = -0.27575', 'bkg_intercept = 4.75381', 'peak_fwhm = 43.08370', 'peak_
height = 0.63148']
ITER 26 ['peak_amplitude = 49.50287', 'peak_center = 18.93364', 'peak_sigma = 0.30541',
'bkg_slope = -0.35174', 'bkg_intercept = 4.83315', 'peak_fwhm = 0.71918', 'peak_
height = 64.66412']
ITER 27 ['peak_amplitude = 29.93395', 'peak_center = 19.97768', 'peak_sigma = 17.89130
', 'bkg_slope = -0.29348', 'bkg_intercept = 4.93342', 'peak_fwhm = 42.13079', 'peak_
height = 0.66747']
ITER 28 ['peak_amplitude = 29.93426', 'peak_center = 19.97768', 'peak_sigma = 17.89130
', 'bkg_slope = -0.29348', 'bkg_intercept = 4.93342', 'peak_fwhm = 42.13079', 'peak_
height = 0.66748']
ITER 29 ['peak_amplitude = 29.93395', 'peak_center = 19.97767', 'peak_sigma = 17.89130
', 'bkg_slope = -0.29348', 'bkg_intercept = 4.93342', 'peak_fwhm = 42.13079', 'peak_
height = 0.66747']
ITER 30 ['peak_amplitude = 29.93395', 'peak_center = 19.97768', 'peak_sigma = 17.89149
', 'bkg_slope = -0.29348', 'bkg_intercept = 4.93342', 'peak_fwhm = 42.13124', 'peak_
height = 0.66746']
ITER 31 ['peak_amplitude = 29.93395', 'peak_center = 19.97768', 'peak_sigma = 17.89130
', 'bkg_slope = -0.29347', 'bkg_intercept = 4.93342', 'peak_fwhm = 42.13079', 'peak_
height = 0.66747']
ITER 32 ['peak_amplitude = 29.93395', 'peak_center = 19.97768', 'peak_sigma = 17.89130
', 'bkg_slope = -0.29348', 'bkg_intercept = 4.93347', 'peak_fwhm = 42.13079', 'peak_
height = 0.66747']
ITER 33 ['peak_amplitude = 33.93754', 'peak_center = 19.80974', 'peak_sigma = 14.52919
', 'bkg_slope = -0.31329', 'bkg_intercept = 5.00443', 'peak_fwhm = 34.21363', 'peak_
height = 0.93186']
ITER 34 ['peak_amplitude = 33.93789', 'peak_center = 19.80974', 'peak_sigma = 14.52919
', 'bkg_slope = -0.31329', 'bkg_intercept = 5.00443', 'peak_fwhm = 34.21363', 'peak_
height = 0.93187']
ITER 35 ['peak_amplitude = 33.93754', 'peak_center = 19.80971', 'peak_sigma = 14.52919
', 'bkg_slope = -0.31329', 'bkg_intercept = 5.00443', 'peak_fwhm = 34.21363', 'peak_
height = 0.93186']
ITER 36 ['peak_amplitude = 33.93754', 'peak_center = 19.80974', 'peak_sigma = 14.52935
', 'bkg_slope = -0.31329', 'bkg_intercept = 5.00443', 'peak_fwhm = 34.21399', 'peak_
height = 0.93185']
ITER 37 ['peak_amplitude = 33.93754', 'peak_center = 19.80974', 'peak_sigma = 14.52919
', 'bkg_slope = -0.31328', 'bkg_intercept = 5.00443', 'peak_fwhm = 34.21363', 'peak_
height = 0.93186']
ITER 38 ['peak_amplitude = 33.93754', 'peak_center = 19.80974', 'peak_sigma = 14.52919
', 'bkg_slope = -0.31329', 'bkg_intercept = 5.00448', 'peak_fwhm = 34.21363', 'peak_
height = 0.93186']
ITER 39 ['peak_amplitude = 39.94969', 'peak_center = 11.63266', 'peak_sigma = 17.05427
', 'bkg_slope = -0.30912', 'bkg_intercept = 4.82241', 'peak_fwhm = 40.15974', 'peak_
height = 0.93452']
ITER 40 ['peak_amplitude = 39.95010', 'peak_center = 11.63266', 'peak_sigma = 17.05427
', 'bkg_slope = -0.30912', 'bkg_intercept = 4.82241', 'peak_fwhm = 40.15974', 'peak_
height = 0.93452']

```

(continues on next page)

(continued from previous page)

```

height = 0.93453']
ITER 41 ['peak_amplitude = 39.94969', 'peak_center = 11.63237', 'peak_sigma = 17.05427
', 'bkg_slope = -0.30912', 'bkg_intercept = 4.82241', 'peak_fwhm = 40.15974', 'peak_
height = 0.93452']
ITER 42 ['peak_amplitude = 39.94969', 'peak_center = 11.63266', 'peak_sigma = 17.05445
', 'bkg_slope = -0.30912', 'bkg_intercept = 4.82241', 'peak_fwhm = 40.16016', 'peak_
height = 0.93451']
ITER 43 ['peak_amplitude = 39.94969', 'peak_center = 11.63266', 'peak_sigma = 17.05427
', 'bkg_slope = -0.30911', 'bkg_intercept = 4.82241', 'peak_fwhm = 40.15974', 'peak_
height = 0.93452']
ITER 44 ['peak_amplitude = 39.94969', 'peak_center = 11.63266', 'peak_sigma = 17.05427
', 'bkg_slope = -0.30912', 'bkg_intercept = 4.82246', 'peak_fwhm = 40.15974', 'peak_
height = 0.93452']
ITER 45 ['peak_amplitude = 43.44365', 'peak_center = 7.65269', 'peak_sigma = 0.17915',
'bkg_slope = -0.30320', 'bkg_intercept = 3.85514', 'peak_fwhm = 0.42187', 'peak_height_
= 96.74097']
ITER 46 ['peak_amplitude = 40.21652', 'peak_center = 11.53095', 'peak_sigma = 15.41291
', 'bkg_slope = -0.30665', 'bkg_intercept = 4.70376', 'peak_fwhm = 36.29462', 'peak_
height = 1.04095']
ITER 47 ['peak_amplitude = 40.21693', 'peak_center = 11.53095', 'peak_sigma = 15.41291
', 'bkg_slope = -0.30665', 'bkg_intercept = 4.70376', 'peak_fwhm = 36.29462', 'peak_
height = 1.04096']
ITER 48 ['peak_amplitude = 40.21652', 'peak_center = 11.53065', 'peak_sigma = 15.41291
', 'bkg_slope = -0.30665', 'bkg_intercept = 4.70376', 'peak_fwhm = 36.29462', 'peak_
height = 1.04095']
ITER 49 ['peak_amplitude = 40.21652', 'peak_center = 11.53095', 'peak_sigma = 15.41307
', 'bkg_slope = -0.30665', 'bkg_intercept = 4.70376', 'peak_fwhm = 36.29500', 'peak_
height = 1.04094']
ITER 50 ['peak_amplitude = 40.21652', 'peak_center = 11.53095', 'peak_sigma = 15.41291
', 'bkg_slope = -0.30665', 'bkg_intercept = 4.70376', 'peak_fwhm = 36.29462', 'peak_
height = 1.04095']
ITER 51 ['peak_amplitude = 40.21652', 'peak_center = 11.53095', 'peak_sigma = 15.41291
', 'bkg_slope = -0.30665', 'bkg_intercept = 4.70381', 'peak_fwhm = 36.29462', 'peak_
height = 1.04095']
ITER 52 ['peak_amplitude = 40.73607', 'peak_center = 10.91338', 'peak_sigma = 11.96371
', 'bkg_slope = -0.30679', 'bkg_intercept = 4.50063', 'peak_fwhm = 28.17238', 'peak_
height = 1.35839']
ITER 53 ['peak_amplitude = 40.73648', 'peak_center = 10.91338', 'peak_sigma = 11.96371
', 'bkg_slope = -0.30679', 'bkg_intercept = 4.50063', 'peak_fwhm = 28.17238', 'peak_
height = 1.35840']
ITER 54 ['peak_amplitude = 40.73607', 'peak_center = 10.91308', 'peak_sigma = 11.96371
', 'bkg_slope = -0.30679', 'bkg_intercept = 4.50063', 'peak_fwhm = 28.17238', 'peak_
height = 1.35839']
ITER 55 ['peak_amplitude = 40.73607', 'peak_center = 10.91338', 'peak_sigma = 11.96384
', 'bkg_slope = -0.30679', 'bkg_intercept = 4.50063', 'peak_fwhm = 28.17268', 'peak_
height = 1.35837']
ITER 56 ['peak_amplitude = 40.73607', 'peak_center = 10.91338', 'peak_sigma = 11.96371
', 'bkg_slope = -0.30679', 'bkg_intercept = 4.50063', 'peak_fwhm = 28.17238', 'peak_
height = 1.35839']
ITER 57 ['peak_amplitude = 40.73607', 'peak_center = 10.91338', 'peak_sigma = 11.96371
', 'bkg_slope = -0.30679', 'bkg_intercept = 4.50068', 'peak_fwhm = 28.17238', 'peak_
height = 1.35839']

```

(continues on next page)

(continued from previous page)

```

ITER  58 ['peak_amplitude = 40.86269', 'peak_center = 9.92047', 'peak_sigma = 5.29355',
↪'bkg_slope = -0.31543', 'bkg_intercept = 4.04328', 'peak_fwhm = 12.46536', 'peak_
↪height = 3.07957']
ITER  59 ['peak_amplitude = 40.86311', 'peak_center = 9.92047', 'peak_sigma = 5.29355',
↪'bkg_slope = -0.31543', 'bkg_intercept = 4.04328', 'peak_fwhm = 12.46536', 'peak_
↪height = 3.07960']
ITER  60 ['peak_amplitude = 40.86269', 'peak_center = 9.92015', 'peak_sigma = 5.29355',
↪'bkg_slope = -0.31543', 'bkg_intercept = 4.04328', 'peak_fwhm = 12.46536', 'peak_
↪height = 3.07957']
ITER  61 ['peak_amplitude = 40.86269', 'peak_center = 9.92047', 'peak_sigma = 5.29361',
↪'bkg_slope = -0.31543', 'bkg_intercept = 4.04328', 'peak_fwhm = 12.46550', 'peak_
↪height = 3.07953']
ITER  62 ['peak_amplitude = 40.86269', 'peak_center = 9.92047', 'peak_sigma = 5.29355',
↪'bkg_slope = -0.31542', 'bkg_intercept = 4.04328', 'peak_fwhm = 12.46536', 'peak_
↪height = 3.07957']
ITER  63 ['peak_amplitude = 40.86269', 'peak_center = 9.92047', 'peak_sigma = 5.29355',
↪'bkg_slope = -0.31543', 'bkg_intercept = 4.04332', 'peak_fwhm = 12.46536', 'peak_
↪height = 3.07957']
ITER  64 ['peak_amplitude = 34.49433', 'peak_center = 6.96011', 'peak_sigma = 3.79883',
↪'bkg_slope = -0.21584', 'bkg_intercept = 3.04624', 'peak_fwhm = 8.94557', 'peak_height_
↪= 3.62249']
ITER  65 ['peak_amplitude = 34.49469', 'peak_center = 6.96011', 'peak_sigma = 3.79883',
↪'bkg_slope = -0.21584', 'bkg_intercept = 3.04624', 'peak_fwhm = 8.94557', 'peak_height_
↪= 3.62253']
ITER  66 ['peak_amplitude = 34.49433', 'peak_center = 6.95978', 'peak_sigma = 3.79883',
↪'bkg_slope = -0.21584', 'bkg_intercept = 3.04624', 'peak_fwhm = 8.94557', 'peak_height_
↪= 3.62249']
ITER  67 ['peak_amplitude = 34.49433', 'peak_center = 6.96011', 'peak_sigma = 3.79888',
↪'bkg_slope = -0.21584', 'bkg_intercept = 3.04624', 'peak_fwhm = 8.94568', 'peak_height_
↪= 3.62245']
ITER  68 ['peak_amplitude = 34.49433', 'peak_center = 6.96011', 'peak_sigma = 3.79883',
↪'bkg_slope = -0.21584', 'bkg_intercept = 3.04624', 'peak_fwhm = 8.94557', 'peak_height_
↪= 3.62249']
ITER  69 ['peak_amplitude = 34.49433', 'peak_center = 6.96011', 'peak_sigma = 3.79883',
↪'bkg_slope = -0.21584', 'bkg_intercept = 3.04628', 'peak_fwhm = 8.94557', 'peak_height_
↪= 3.62249']
ITER  70 ['peak_amplitude = 0.08766', 'peak_center = 8.87255', 'peak_sigma = 0.10402',
↪'bkg_slope = -0.26833', 'bkg_intercept = 4.93654', 'peak_fwhm = 0.24494', 'peak_height_
↪= 0.33620']
ITER  71 ['peak_amplitude = 30.67432', 'peak_center = 8.04096', 'peak_sigma = 1.99254',
↪'bkg_slope = -0.19527', 'bkg_intercept = 2.91723', 'peak_fwhm = 4.69208', 'peak_height_
↪= 6.14155']
ITER  72 ['peak_amplitude = 30.67464', 'peak_center = 8.04096', 'peak_sigma = 1.99254',
↪'bkg_slope = -0.19527', 'bkg_intercept = 2.91723', 'peak_fwhm = 4.69208', 'peak_height_
↪= 6.14161']
ITER  73 ['peak_amplitude = 30.67432', 'peak_center = 8.04063', 'peak_sigma = 1.99254',
↪'bkg_slope = -0.19527', 'bkg_intercept = 2.91723', 'peak_fwhm = 4.69208', 'peak_height_
↪= 6.14155']
ITER  74 ['peak_amplitude = 30.67432', 'peak_center = 8.04096', 'peak_sigma = 1.99257',
↪'bkg_slope = -0.19527', 'bkg_intercept = 2.91723', 'peak_fwhm = 4.69214', 'peak_height_
↪= 6.14147']
ITER  75 ['peak_amplitude = 30.67432', 'peak_center = 8.04096', 'peak_sigma = 1.99254',
↪'bkg_slope = -0.19527', 'bkg_intercept = 2.91723', 'peak_fwhm = 4.69214', 'peak_height_
↪= 6.14147']

```

(continues on next page)

(continued from previous page)

```

→ 'bkg_slope = -0.19527', 'bkg_intercept = 2.91723', 'peak_fwhm = 4.69208', 'peak_height_
→ = 6.14155']
ITER 76 ['peak_amplitude = 30.67432', 'peak_center = 8.04096', 'peak_sigma = 1.99254',
→ 'bkg_slope = -0.19527', 'bkg_intercept = 2.91726', 'peak_fwhm = 4.69208', 'peak_height_
→ = 6.14155']
ITER 77 ['peak_amplitude = 21.89977', 'peak_center = 7.48302', 'peak_sigma = 1.18769',
→ 'bkg_slope = -0.19932', 'bkg_intercept = 3.46743', 'peak_fwhm = 2.79680', 'peak_height_
→ = 7.35606']
ITER 78 ['peak_amplitude = 21.90000', 'peak_center = 7.48302', 'peak_sigma = 1.18769',
→ 'bkg_slope = -0.19932', 'bkg_intercept = 3.46743', 'peak_fwhm = 2.79680', 'peak_height_
→ = 7.35614']
ITER 79 ['peak_amplitude = 21.89977', 'peak_center = 7.48269', 'peak_sigma = 1.18769',
→ 'bkg_slope = -0.19932', 'bkg_intercept = 3.46743', 'peak_fwhm = 2.79680', 'peak_height_
→ = 7.35606']
ITER 80 ['peak_amplitude = 21.89977', 'peak_center = 7.48302', 'peak_sigma = 1.18771',
→ 'bkg_slope = -0.19932', 'bkg_intercept = 3.46743', 'peak_fwhm = 2.79684', 'peak_height_
→ = 7.35595']
ITER 81 ['peak_amplitude = 21.89977', 'peak_center = 7.48302', 'peak_sigma = 1.18769',
→ 'bkg_slope = -0.19932', 'bkg_intercept = 3.46743', 'peak_fwhm = 2.79680', 'peak_height_
→ = 7.35606']
ITER 82 ['peak_amplitude = 21.89977', 'peak_center = 7.48302', 'peak_sigma = 1.18769',
→ 'bkg_slope = -0.19932', 'bkg_intercept = 3.46746', 'peak_fwhm = 2.79680', 'peak_height_
→ = 7.35606']
ITER 83 ['peak_amplitude = 24.51986', 'peak_center = 7.64789', 'peak_sigma = 1.23846',
→ 'bkg_slope = -0.20238', 'bkg_intercept = 3.36745', 'peak_fwhm = 2.91635', 'peak_height_
→ = 7.89852']
ITER 84 ['peak_amplitude = 24.52011', 'peak_center = 7.64789', 'peak_sigma = 1.23846',
→ 'bkg_slope = -0.20238', 'bkg_intercept = 3.36745', 'peak_fwhm = 2.91635', 'peak_height_
→ = 7.89861']
ITER 85 ['peak_amplitude = 24.51986', 'peak_center = 7.64756', 'peak_sigma = 1.23846',
→ 'bkg_slope = -0.20238', 'bkg_intercept = 3.36745', 'peak_fwhm = 2.91635', 'peak_height_
→ = 7.89852']
ITER 86 ['peak_amplitude = 24.51986', 'peak_center = 7.64789', 'peak_sigma = 1.23848',
→ 'bkg_slope = -0.20238', 'bkg_intercept = 3.36745', 'peak_fwhm = 2.91639', 'peak_height_
→ = 7.89841']
ITER 87 ['peak_amplitude = 24.51986', 'peak_center = 7.64789', 'peak_sigma = 1.23846',
→ 'bkg_slope = -0.20238', 'bkg_intercept = 3.36745', 'peak_fwhm = 2.91635', 'peak_height_
→ = 7.89852']
ITER 88 ['peak_amplitude = 24.51986', 'peak_center = 7.64789', 'peak_sigma = 1.23846',
→ 'bkg_slope = -0.20238', 'bkg_intercept = 3.36748', 'peak_fwhm = 2.91635', 'peak_height_
→ = 7.89852']
ITER 89 ['peak_amplitude = 24.52133', 'peak_center = 7.63730', 'peak_sigma = 1.22397',
→ 'bkg_slope = -0.20264', 'bkg_intercept = 3.36998', 'peak_fwhm = 2.88223', 'peak_height_
→ = 7.99252']
ITER 90 ['peak_amplitude = 24.52158', 'peak_center = 7.63730', 'peak_sigma = 1.22397',
→ 'bkg_slope = -0.20264', 'bkg_intercept = 3.36998', 'peak_fwhm = 2.88223', 'peak_height_
→ = 7.99260']
ITER 91 ['peak_amplitude = 24.52133', 'peak_center = 7.63697', 'peak_sigma = 1.22397',
→ 'bkg_slope = -0.20264', 'bkg_intercept = 3.36998', 'peak_fwhm = 2.88223', 'peak_height_
→ = 7.99252']
ITER 92 ['peak_amplitude = 24.52133', 'peak_center = 7.63730', 'peak_sigma = 1.22399',
→ 'bkg_slope = -0.20264', 'bkg_intercept = 3.36998', 'peak_fwhm = 2.88227', 'peak_height_
→ = 7.99252']

```

(continues on next page)

(continued from previous page)

```

↳= 7.99240']
ITER 93 ['peak_amplitude = 24.52133', 'peak_center = 7.63730', 'peak_sigma = 1.22397',
↳ 'bkg_slope = -0.20264', 'bkg_intercept = 3.36998', 'peak_fwhm = 2.88223', 'peak_height↳
↳= 7.99252']
ITER 94 ['peak_amplitude = 24.52133', 'peak_center = 7.63730', 'peak_sigma = 1.22397',
↳ 'bkg_slope = -0.20264', 'bkg_intercept = 3.37001', 'peak_fwhm = 2.88223', 'peak_height↳
↳= 7.99252']
ITER 95 ['peak_amplitude = 24.52372', 'peak_center = 7.63753', 'peak_sigma = 1.22415',
↳ 'bkg_slope = -0.20264', 'bkg_intercept = 3.36986', 'peak_fwhm = 2.88265', 'peak_height↳
↳= 7.99214']
ITER 96 ['peak_amplitude = 24.52397', 'peak_center = 7.63753', 'peak_sigma = 1.22415',
↳ 'bkg_slope = -0.20264', 'bkg_intercept = 3.36986', 'peak_fwhm = 2.88265', 'peak_height↳
↳= 7.99222']
ITER 97 ['peak_amplitude = 24.52372', 'peak_center = 7.63720', 'peak_sigma = 1.22415',
↳ 'bkg_slope = -0.20264', 'bkg_intercept = 3.36986', 'peak_fwhm = 2.88265', 'peak_height↳
↳= 7.99214']
ITER 98 ['peak_amplitude = 24.52372', 'peak_center = 7.63753', 'peak_sigma = 1.22416',
↳ 'bkg_slope = -0.20264', 'bkg_intercept = 3.36986', 'peak_fwhm = 2.88269', 'peak_height↳
↳= 7.99202']
ITER 99 ['peak_amplitude = 24.52372', 'peak_center = 7.63753', 'peak_sigma = 1.22415',
↳ 'bkg_slope = -0.20264', 'bkg_intercept = 3.36986', 'peak_fwhm = 2.88265', 'peak_height↳
↳= 7.99214']
ITER 100 ['peak_amplitude = 24.52372', 'peak_center = 7.63753', 'peak_sigma = 1.22415',
↳ 'bkg_slope = -0.20264', 'bkg_intercept = 3.36989', 'peak_fwhm = 2.88265', 'peak_
↳ height = 7.99214']
ITER 101 ['peak_amplitude = 24.52371', 'peak_center = 7.63753', 'peak_sigma = 1.22415',
↳ 'bkg_slope = -0.20264', 'bkg_intercept = 3.36986', 'peak_fwhm = 2.88264', 'peak_
↳ height = 7.99214']
ITER 101 ['peak_amplitude = 24.52371', 'peak_center = 7.63753', 'peak_sigma = 1.22415',
↳ 'bkg_slope = -0.20264', 'bkg_intercept = 3.36986', 'peak_fwhm = 2.88264', 'peak_
↳ height = 7.99214']

Nfev = 101
[[Model]]
    (Model(gaussian, prefix='peak_') + Model(linear, prefix='bkg_'))
[[Fit Statistics]]
    # fitting method    = leastsq
    # function evals   = 101
    # data points      = 401
    # variables        = 5
    chi-square         = 20.0043556
    reduced chi-square = 0.05051605
    Akaike info crit  = -1192.20257
    Bayesian info crit = -1172.23276
    R-squared          = 0.99377421
[[Variables]]
    peak_amplitude: 24.5237052 +/- 0.16281835 (0.66%) (init = 3)
    peak_center:    7.63752785 +/- 0.00746969 (0.10%) (init = 6)
    peak_sigma:     1.22414559 +/- 0.00811005 (0.66%) (init = 2)
    bkg_slope:      -0.20264093 +/- 0.00204346 (1.01%) (init = 0)
    bkg_intercept:  3.36986054 +/- 0.02653942 (0.79%) (init = 0)
    peak_fwhm:      2.88264251 +/- 0.01909771 (0.66%) == '2.3548200*peak_sigma'
    peak_height:    7.99214038 +/- 0.04318559 (0.54%) == '0.3989423*peak_amplitude/'

```

(continues on next page)

(continued from previous page)

```

→max(1e-15, peak_sigma)'
[[Correlations]] (unreported correlations are < 0.100)
C(bkg_slope, bkg_intercept)      = -0.8574
C(peak_amplitude, peak_sigma)    = +0.6681
C(peak_amplitude, bkg_intercept) = -0.5260
C(peak_sigma, bkg_intercept)     = -0.3514
C(peak_amplitude, bkg_slope)     = +0.2858
C(peak_sigma, bkg_slope)         = +0.1909
C(peak_center, bkg_slope)        = -0.1451
C(peak_center, bkg_intercept)    = +0.1244

```

```

# <examples/doc_with_itercb.py>
import matplotlib.pyplot as plt
from numpy import linspace, random

from lmfit.lineshapes import gaussian
from lmfit.models import GaussianModel, LinearModel


def per_iteration(pars, iteration, resid, *args, **kws):
    print(" ITER ", iteration, [f" {p.name} = {p.value:.5f}" for p in pars.values()])

x = linspace(0., 20, 401)
y = gaussian(x, amplitude=24.56, center=7.6543, sigma=1.23)
random.seed(2021)
y = y - .20*x + 3.333 + random.normal(scale=0.23, size=x.size)

mod = GaussianModel(prefix='peak_') + LinearModel(prefix='bkg_')

pars = mod.make_params(peak_amplitude=dict(value=3.0, min=0),
                      peak_center=dict(value=6.0, min=0, max=20),
                      peak_sigma=2.0,
                      bkg_intercept=0,
                      bkg_slope=0)

out = mod.fit(y, pars, x=x, iter_cb=per_iteration)

plt.plot(x, y, '--')


print(f'Nfev = {out.nfev}')
print(out.fit_report())

plt.plot(x, out.best_fit, '-', label='best fit')
plt.legend()
plt.show()
# <end examples/doc_with_itercb.py>

```

Total running time of the script: (0 minutes 0.212 seconds)

14.1.16 Parameters - valuesdict



```
[[Fit Statistics]]
# fitting method    = leastsq
# function evals   = 64
# data points      = 301
# variables        = 4
chi-square         = 12.1867036
reduced chi-square = 0.04103267
Akaike info crit   = -957.236198
Bayesian info crit = -942.407756
[[Variables]]
amp:    5.03088059 +/- 0.04005824 (0.80%) (init = 10)
decay:  0.02495457 +/- 4.5396e-04 (1.82%) (init = 0.1)
omega:  2.00026310 +/- 0.00326183 (0.16%) (init = 3)
shift: -0.10264952 +/- 0.01022294 (9.96%) (init = 0)
[[Correlations]] (unreported correlations are < 0.100)
C(omega, shift) = -0.7852
C(amp, decay)   = +0.5840
C(amp, shift)   = -0.1179
```

```

# <examples/doc_parameters_valuesdict.py>
import numpy as np

from lmfit import Minimizer, create_params, report_fit

# create data to be fitted
x = np.linspace(0, 15, 301)
np.random.seed(2021)
data = (5.0 * np.sin(2.0*x - 0.1) * np.exp(-x*x*0.025) +
        np.random.normal(size=x.size, scale=0.2))

# define objective function: returns the array to be minimized
def fcn2min(params, x, data):
    """Model a decaying sine wave and subtract data."""
    v = params.valuesdict()

    model = v['amp'] * np.sin(x * v['omega'] + v['shift']) * np.exp(-x*x*v['decay'])
    return model - data

# create a set of Parameters
params = create_params(amp=dict(value=10, min=0),
                       decay=0.1,
                       omega=3.0,
                       shift=dict(value=0.0, min=-np.pi/2., max=np.pi/2.))

# do fit, here with the default leastsq algorithm
minner = Minimizer(fcn2min, params, fcn_args=(x, data))
result = minner.minimize()

# calculate final result
final = data + result.residual


# write error report
report_fit(result)


# try to plot results
try:
    import matplotlib.pyplot as plt
    plt.plot(x, data, '+')
    plt.plot(x, final)
    plt.show()
except ImportError:
    pass
# <end of examples/doc_parameters_valuesdict.py>

```


Total running time of the script: (0 minutes 0.175 seconds)

14.1.17 Builtinmodels - peakmodels





•



•

```
[[Model]]
  Model(gaussian)
[[Fit Statistics]]
  # fitting method = leastsq
  # function evals = 25
  # data points   = 401
  # variables     = 3
  chi-square     = 29.9943157
  reduced chi-square = 0.07536260
  Akaike info crit = -1033.77437
```

(continues on next page)

(continued from previous page)

```

Bayesian info crit = -1021.79248
R-squared          = 0.99045513
[[Variables]]
  amplitude: 30.3135789 +/- 0.15712752 (0.52%) (init = 43.62238)
  center:    9.24277046 +/- 0.00737497 (0.08%) (init = 9.25)
  sigma:     1.23218496 +/- 0.00737506 (0.60%) (init = 1.35)
  fwhm:      2.90157379 +/- 0.01736695 (0.60%) == '2.3548200*sigma'
  height:    9.81457271 +/- 0.05087308 (0.52%) == '0.3989423*amplitude/max(1e-15,_
  ~sigma)'
[[Correlations]]
+-----+-----+-----+
| Variable | amplitude | center   | sigma    |
+-----+-----+-----+
| amplitude | +1.0000  | -0.0000  | +0.5774  |
| center    | -0.0000  | +1.0000  | -0.0000  |
| sigma     | +0.5774  | -0.0000  | +1.0000  |
+-----+-----+-----+
[[Model]]
  Model(lorentzian)
[[Fit Statistics]]
  # fitting method = leastsq
  # function evals = 25
  # data points   = 401
  # variables     = 3
  chi-square      = 53.7535387
  reduced chi-square = 0.13505914
  Akaike info crit = -799.830322
  Bayesian info crit = -787.848438
  R-squared        = 0.98289441
[[Variables]]
  amplitude: 38.9726380 +/- 0.31386754 (0.81%) (init = 54.52798)
  center:    9.24439393 +/- 0.00927645 (0.10%) (init = 9.25)
  sigma:     1.15483177 +/- 0.01315708 (1.14%) (init = 1.35)
  fwhm:      2.30966354 +/- 0.02631416 (1.14%) == '2.0000000*sigma'
  height:    10.7421504 +/- 0.08634317 (0.80%) == '0.3183099*amplitude/max(1e-15,_
  ~sigma)'
[[Correlations]]
+-----+-----+-----+
| Variable | amplitude | center   | sigma    |
+-----+-----+-----+
| amplitude | +1.0000  | -0.0002  | +0.7087  |
| center    | -0.0002  | +1.0000  | -0.0002  |
| sigma     | +0.7087  | -0.0002  | +1.0000  |
+-----+-----+-----+
[[Model]]
  Model(voigt)
[[Fit Statistics]]
  # fitting method = leastsq
  # function evals = 25
  # data points   = 401
  # variables     = 3
  chi-square      = 14.5448627

```

(continues on next page)

(continued from previous page)

```

reduced chi-square = 0.03654488
Akaike info crit   = -1324.00615
Bayesian info crit = -1312.02427
R-squared           = 0.99537150
[[Variables]]
  amplitude: 35.7553799 +/- 0.13861559 (0.39%) (init = 65.43358)
  center:    9.24411179 +/- 0.00505496 (0.05%) (init = 9.25)
  sigma:     0.73015485 +/- 0.00368473 (0.50%) (init = 0.8775)
  gamma:     0.73015485 +/- 0.00368473 (0.50%) == 'sigma'
  fwhm:      2.62949983 +/- 0.01326979 (0.50%) == '1.0692*gamma+sqrt(0.
  ↵8664*gamma**2+5.545083*sigma**2)'
  height:    10.2204068 +/- 0.03959933 (0.39%) == '(amplitude/(max(1e-15,
  ↵sigma*sqrt(2*pi)))*real(wofz((1j*gamma)/(max(1e-15, sigma*sqrt(2))))))'
[[Correlations]]
+-----+-----+-----+-----+
| Variable | amplitude | center   | sigma    |
+-----+-----+-----+-----+
| amplitude | +1.0000  | -0.0001  | +0.6513  |
| center    | -0.0001  | +1.0000  | -0.0001  |
| sigma     | +0.6513  | -0.0001  | +1.0000  |
+-----+-----+-----+-----+
[[Model]]
  Model(voigt)
[[Fit Statistics]]
  # fitting method   = leastsq
  # function evals   = 25
  # data points      = 401
  # variables        = 3
  chi-square         = 14.5448627
  reduced chi-square = 0.03654488
  Akaike info crit   = -1324.00615
  Bayesian info crit = -1312.02427
  R-squared           = 0.99537150
[[Variables]]
  amplitude: 35.7553799 +/- 0.13861559 (0.39%) (init = 65.43358)
  center:    9.24411179 +/- 0.00505496 (0.05%) (init = 9.25)
  sigma:     0.73015485 +/- 0.00368473 (0.50%) (init = 0.8775)
  gamma:     0.73015485 +/- 0.00368473 (0.50%) == 'sigma'
  fwhm:      2.62949983 +/- 0.01326979 (0.50%) == '1.0692*gamma+sqrt(0.
  ↵8664*gamma**2+5.545083*sigma**2)'
  height:    10.2204068 +/- 0.03959933 (0.39%) == '(amplitude/(max(1e-15,
  ↵sigma*sqrt(2*pi)))*real(wofz((1j*gamma)/(max(1e-15, sigma*sqrt(2))))))'
[[Correlations]]
+-----+-----+-----+-----+
| Variable | amplitude | center   | sigma    |
+-----+-----+-----+-----+
| amplitude | +1.0000  | -0.0001  | +0.6513  |
| center    | -0.0001  | +1.0000  | -0.0001  |
| sigma     | +0.6513  | -0.0001  | +1.0000  |
+-----+-----+-----+-----+

```

```

# <examples/doc_builtinmodels_peakmodels.py>
import matplotlib.pyplot as plt
from numpy import loadtxt

from lmfit.models import GaussianModel, LorentzianModel, VoigtModel

data = loadtxt('test_peak.dat')
x = data[:, 0]
y = data[:, 1]

# Gaussian model
mod = GaussianModel()
pars = mod.guess(y, x=x)
out = mod.fit(y, pars, x=x)

print(out.fit_report(correl_mode='table'))

plt.plot(x, y)
plt.plot(x, out.best_fit, ' -', label='Gaussian Model')
plt.legend()
plt.show()

# Lorentzian model
mod = LorentzianModel()
pars = mod.guess(y, x=x)
out = mod.fit(y, pars, x=x)

print(out.fit_report(correl_mode='table'))

plt.figure()
plt.plot(x, y, ' -')
plt.plot(x, out.best_fit, ' -', label='Lorentzian Model')
plt.legend()
plt.show()

# Voigt model
mod = VoigtModel()
pars = mod.guess(y, x=x)
out = mod.fit(y, pars, x=x)

print(out.fit_report(correl_mode='table'))

fig, axes = plt.subplots(1, 2, figsize=(12.8, 4.8))

axes[0].plot(x, y, ' -')
axes[0].plot(x, out.best_fit, ' -', label='Voigt Model\\ngamma constrained')
axes[0].legend()

```

(continues on next page)

(continued from previous page)

```

# allow the gamma parameter to vary in the fit
pars['gamma'].vary = True
out_gamma = mod.fit(y, pars, x=x)
print(out.fit_report(correl_mode='table'))


axes[1].plot(x, y, ' ')
axes[1].plot(x, out_gamma.best_fit, ' ', label='Voigt Model\nngamma unconstrained')
axes[1].legend()

plt.show()
# <end examples/doc_builtinmodels_peakmodels.py>

```

Total running time of the script: (0 minutes 0.786 seconds)

14.1.18 Builtinmodels - nistgauss



```

[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
  ↪ prefix='exp_'))
[[Fit Statistics]]
  # fitting method   = leastsq
  # function evals  = 46
  # data points    = 250
  # variables      = 8
  chi-square        = 1247.52821
  reduced chi-square = 5.15507524
  Akaike info crit  = 417.864631
  Bayesian info crit = 446.036318
  R-squared         = 0.99648654
[[Variables]]
  exp_amplitude: 99.0183278 +/- 0.53748593 (0.54%) (init = 162.2102)
  exp_decay:     90.9508853 +/- 1.10310778 (1.21%) (init = 93.24905)
  g1_amplitude:  4257.77360 +/- 42.3836478 (1.00%) (init = 2000)
  g1_center:     107.030956 +/- 0.15006851 (0.14%) (init = 105)

```

(continues on next page)

(continued from previous page)

```

g1_sigma:      16.6725772 +/- 0.16048381 (0.96%) (init = 15)
g1_fwhm:      39.2609181 +/- 0.37791049 (0.96%) == '2.3548200*g1_sigma'
g1_height:    101.880230 +/- 0.59217173 (0.58%) == '0.3989423*g1_amplitude/max(1e-
  15, g1_sigma)'
g2_amplitude: 2493.41735 +/- 36.1697789 (1.45%) (init = 2000)
g2_center:    153.270102 +/- 0.19466802 (0.13%) (init = 155)
g2_sigma:     13.8069464 +/- 0.18679695 (1.35%) (init = 15)
g2_fwhm:      32.5128735 +/- 0.43987320 (1.35%) == '2.3548200*g2_sigma'
g2_height:    72.0455941 +/- 0.61722243 (0.86%) == '0.3989423*g2_amplitude/max(1e-
  15, g2_sigma)'
[[Correlations]]
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
| Variable      | exp_amplitude | exp_decay      | g1_amplitude | g1_center      | g1_
| sigma         | g2_amplitude  | g2_center      | g2_sigma      |               |
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
| exp_amplitude | +1.0000      | -0.6946      | +0.1478      | -0.0467      | +0.
| 0218          | +0.2821      | +0.0331      | +0.1714      | +0.1055      | -0.
| exp_decay     | -0.6946      | +1.0000      | -0.5074      | +0.4183      | +0.
| 2520          | -0.4270      | -0.1514      | -0.2329      | -0.6689      | -0.
| g1_amplitude  | +0.1478      | -0.5074      | +1.0000      | +0.6205      | +0.
| 8243          | -0.3071      | +0.6477      | -0.4010      | -0.4915      | -0.
| g1_center     | -0.0467      | +0.1055      | +0.4183      | -0.4763      | +0.
| 5075          | -0.6689      | +0.6205      | -0.6520      | +0.8154      | +0.
| g1_sigma       | +0.0218      | -0.2520      | +0.8243      | +0.6842      | +0.
| 0000          | -0.4915      | +0.6842      | -0.4826      | -0.4763      | +0.
| g2_amplitude  | +0.2821      | -0.4270      | -0.3071      | +1.0000      | -0.
| 4915          | +1.0000      | -0.4763      | +0.8154      | +0.0331      | +0.
| g2_center     | +0.0331      | -0.1514      | +0.6477      | -0.4763      | +0.
| 6842          | -0.4763      | +1.0000      | -0.4889      | +0.1714      | -0.
| g2_sigma       | +0.1714      | -0.2329      | -0.4010      | +0.8154      | +0.
| 4826          | +0.8154      | -0.4889      | +1.0000      |               | +0.
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+

```

```

# <examples/doc_builtinmodels_nistgauss.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit.models import ExponentialModel, GaussianModel

dat = np.loadtxt('NIST_Gauss2.dat')
x = dat[:, 1]
y = dat[:, 0]

exp_mod = ExponentialModel(prefix='exp_')

```

(continues on next page)

(continued from previous page)

```

pars = exp_mod.guess(y, x=x)

gauss1 = GaussianModel(prefix='g1_')
pars.update(gauss1.make_params(center=dict(value=105, min=75, max=125),
                               sigma=dict(value=15, min=0),
                               amplitude=dict(value=2000, min=0)))

gauss2 = GaussianModel(prefix='g2_')
pars.update(gauss2.make_params(center=dict(value=155, min=125, max=175),
                               sigma=dict(value=15, min=0),
                               amplitude=dict(value=2000, min=0)))

mod = gauss1 + gauss2 + exp_mod

init = mod.eval(pars, x=x)
out = mod.fit(y, pars, x=x)

print(out.fit_report(correl_mode='table'))

fig, axes = plt.subplots(1, 2, figsize=(12.8, 4.8))
axes[0].plot(x, y)
axes[0].plot(x, init, '--', label='initial fit')
axes[0].plot(x, out.best_fit, ' ', label='best fit')
axes[0].legend()


comps = out.eval_components(x=x)
axes[1].plot(x, y)
axes[1].plot(x, comps['g1_'], '--', label='Gaussian component 1')
axes[1].plot(x, comps['g2_'], '--', label='Gaussian component 2')
axes[1].plot(x, comps['exp_'], '--', label='Exponential component')
axes[1].legend()

plt.show()
# <end examples/doc_builtinmodels_nistgauss.py>

```

Total running time of the script: (0 minutes 0.410 seconds)

14.1.19 Parameters - basic



```
[[Fit Statistics]]
# fitting method = leastsq
# function evals = 64
# data points = 301
# variables = 4
chi-square = 12.1867036
reduced chi-square = 0.04103267
Akaike info crit = -957.236198
Bayesian info crit = -942.407756
[[Variables]]
amp: 5.03088059 +/- 0.04005824 (0.80%) (init = 10)
decay: 0.02495457 +/- 4.5396e-04 (1.82%) (init = 0.1)
omega: 2.00026310 +/- 0.00326183 (0.16%) (init = 3)
shift: -0.10264952 +/- 0.01022294 (9.96%) (init = 0)
[[Correlations]] (unreported correlations are < 0.100)
C(omega, shift) = -0.7852
C(amp, decay) = +0.5840
C(amp, shift) = -0.1179
```

```

# <examples/doc_parameters_basic.py>
import numpy as np

from lmfit import Minimizer, Parameters, create_params, report_fit

# create data to be fitted
x = np.linspace(0, 15, 301)
np.random.seed(2021)
data = (5.0 * np.sin(2.0*x - 0.1) * np.exp(-x*x*0.025) +
        np.random.normal(size=x.size, scale=0.2))

# define objective function: returns the array to be minimized
def fcn2min(params, x, data):
    """Model a decaying sine wave and subtract data."""
    amp = params['amp']
    shift = params['shift']
    omega = params['omega']
    decay = params['decay']
    model = amp * np.sin(x*omega + shift) * np.exp(-x*x*decay)
    return model - data

# create a set of Parameters
params = Parameters()
params.add('amp', value=10, min=0)
params.add('decay', value=0.1)
params.add('shift', value=0.0, min=-np.pi/2., max=np.pi/2.)
params.add('omega', value=3.0)

# ... or use
params = create_params(amp=dict(value=10, min=0),
                       decay=0.1,
                       omega=3,
                       shift=dict(value=0, min=-np.pi/2, max=np.pi/2))

# do fit, here with the default leastsq algorithm
minner = Minimizer(fcn2min, params, fcn_args=(x, data))
result = minner.minimize()

# calculate final result
final = data + result.residual

# write error report
report_fit(result)

# try to plot results
try:
    import matplotlib.pyplot as plt
    plt.plot(x, data, '+')
    plt.plot(x, final)
    plt.show()
except ImportError:

```


(continues on next page)

(continued from previous page)

```
pass
# <end of examples/doc_parameters_basic.py>
```

Total running time of the script: (0 minutes 0.171 seconds)

14.1.20 Model - composite



```
[[Model]]
(Model(jump) <function convolve at 0x1852b3060> Model(gaussian))
[[Fit Statistics]]
# fitting method = leastsq
# function evals = 33
# data points = 201
# variables = 3
chi-square = 24.7562335
reduced chi-square = 0.12503148
Akaike info crit = -414.939746
Bayesian info crit = -405.029832
R-squared = 0.99632577
[[Variables]]
mid: 4 (fixed)
amplitude: 0.62508458 +/- 0.00189732 (0.30%) (init = 1)
center: 5.50853669 +/- 0.00973231 (0.18%) (init = 3.5)
sigma: 0.59576097 +/- 0.01348579 (2.26%) (init = 1.5)
[[Correlations]] (unreported correlations are < 0.100)
C(amplitude, center) = +0.3292
C(amplitude, sigma) = +0.2680
```

```
# <examples/doc_model_composite.py>
import matplotlib.pyplot as plt
```

(continues on next page)

(continued from previous page)

```

import numpy as np

from lmfit import CompositeModel, Model
from lmfit.lineshapes import gaussian, step

# create data from broadened step
x = np.linspace(0, 10, 201)
y = step(x, amplitude=12.5, center=4.5, sigma=0.88, form='erf')
np.random.seed(0)
y = y + np.random.normal(scale=0.35, size=x.size)

def jump(x, mid):
    """Heaviside step function."""
    o = np.zeros(x.size)
    imid = max(np.where(x <= mid)[0])
    o[imid:] = 1.0
    return o

def convolve(arr, kernel):
    """Simple convolution of two arrays."""
    npts = min(arr.size, kernel.size)
    pad = np.ones(npts)
    tmp = np.concatenate((pad*arr[0], arr, pad*arr[-1]))
    out = np.convolve(tmp, kernel, mode='valid')
    noff = int((len(out) - npts) / 2)
    return out[noff:noff+npts]

# create Composite Model using the custom convolution operator
mod = CompositeModel(Model(jump), Model(gaussian), convolve)

# create parameters for model. Note that 'mid' and 'center' will be highly
# correlated. Since 'mid' is used as an integer index, it will be very
# hard to fit, so we fix its value
pars = mod.make_params(amplitude=dict(value=1, min=0),
                       center=3.5,
                       sigma=dict(value=1.5, min=0),
                       mid=dict(value=4, vary=False))

# fit this model to data array y
result = mod.fit(y, params=pars, x=x)

print(result.fit_report())

# generate components
comps = result.eval_components(x=x)

# plot results
fig, axes = plt.subplots(1, 2, figsize=(12.8, 4.8))

```

(continues on next page)

(continued from previous page)

```

axes[0].plot(x, y, 'bo')
axes[0].plot(x, result.init_fit, 'k--', label='initial fit')
axes[0].plot(x, result.best_fit, 'r-', label='best fit')
axes[0].legend()


axes[1].plot(x, y, 'bo')
axes[1].plot(x, 10*comps['jump'], 'k--', label='Jump component')
axes[1].plot(x, 10*comps['gaussian'], 'r-', label='Gaussian component')
axes[1].legend()

plt.show()
# <end examples/doc_model_composite.py>

```

Total running time of the script: (0 minutes 0.363 seconds)

14.1.21 Builtinmodels - splinemodel



```

[[Model]]
    (Model(gaussian, prefix='peak_') + Model(spline_model, prefix='bkg_'))
[[Fit Statistics]]
    # fitting method    = leastsq

```

(continues on next page)

(continued from previous page)

```

# function evals      = 92
# data points        = 501
# variables          = 14
chi-square           = 52.6611549
reduced chi-square  = 0.10813379
Akaike info crit   = -1100.61674
Bayesian info crit = -1041.58425
R-squared            = 0.94690612

[[Variables]]
peak_amplitude: 12.2231138 +/- 0.29554074 (2.42%) (init = 8)
peak_center:    16.4280869 +/- 0.01091050 (0.07%) (init = 16)
peak_sigma:     0.72096402 +/- 0.01336666 (1.85%) (init = 1)
peak_fwhm:      1.69774050 +/- 0.03147609 (1.85%) == '2.3548200*peak_sigma'
peak_height:    6.76360675 +/- 0.09854036 (1.46%) == '0.3989423*peak_amplitude/
max(1e-15, peak_sigma)'
bkg_s0:          3.51175736 +/- 0.04941392 (1.41%) (init = 3.787995)
bkg_s1:          3.72930068 +/- 0.09558236 (2.56%) (init = 3.959487)
bkg_s2:          4.26846495 +/- 0.12650286 (2.96%) (init = 4.384009)
bkg_s3:          4.42375491 +/- 0.10170203 (2.30%) (init = 4.431971)
bkg_s4:          4.49590447 +/- 0.10615551 (2.36%) (init = 4.243976)
bkg_s5:          3.96515316 +/- 0.09336554 (2.35%) (init = 4.115153)
bkg_s6:          3.35531898 +/- 0.12669983 (3.78%) (init = 3.965325)
bkg_s7:          2.89909737 +/- 0.16190201 (5.58%) (init = 2.788437)
bkg_s8:          2.82656972 +/- 0.13445491 (4.76%) (init = 2.984317)
bkg_s9:          3.43338674 +/- 0.15987280 (4.66%) (init = 3.383491)
bkg_s10:         3.73024845 +/- 0.12096864 (3.24%) (init = 3.791937)

[[Correlations]] (unreported correlations are < 0.300)
C(bkg_s7, bkg_s8)      = -0.8192
C(peak_amplitude, peak_sigma) = +0.7987
C(bkg_s8, bkg_s9)      = -0.7063
C(bkg_s5, bkg_s6)      = -0.6950
C(peak_amplitude, bkg_s7) = -0.6878
C(bkg_s2, bkg_s3)      = -0.6672
C(bkg_s9, bkg_s10)     = -0.6060
C(bkg_s3, bkg_s4)      = -0.5743
C(bkg_s1, bkg_s2)      = -0.5646
C(bkg_s4, bkg_s5)      = -0.5542
C(bkg_s7, bkg_s9)      = +0.5216
C(peak_sigma, bkg_s7)   = -0.5192
C(peak_amplitude, bkg_s8) = +0.5185
C(bkg_s0, bkg_s1)      = +0.4448
C(peak_sigma, bkg_s8)   = +0.3733
C(peak_center, bkg_s6)  = +0.3599
C(bkg_s4, bkg_s6)      = +0.3597
C(bkg_s0, bkg_s2)      = -0.3595
C(bkg_s2, bkg_s4)      = +0.3504
C(bkg_s8, bkg_s10)     = +0.3455
C(bkg_s6, bkg_s7)      = -0.3332
C(peak_center, bkg_s7)  = -0.3301
C(peak_amplitude, bkg_s9) = -0.3206

```

```

# <examples/doc_builtinmodels_splinemodel.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit.models import GaussianModel, SplineModel

data = np.loadtxt('test_splinepeak.dat')
x = data[:, 0]
y = data[:, 1]

plt.plot(x, y, label='data')

model = GaussianModel(prefix='peak_')
params = model.make_params(amplitude=dict(value=8, min=0),
                           center=dict(value=16, min=5, max=25),
                           sigma=dict(value=1, min=0))

# make a background spline with knots evenly spaced over the background,
# but sort of skipping over where the peak is
knot_xvals3 = np.array([1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25])
knot_xvals2 = np.array([1, 3, 5, 7, 9, 11, 13, 16, 19, 21, 23, 25]) # noqa: E241
knot_xvals1 = np.array([1, 3, 5, 7, 9, 11, 13, 19, 21, 23, 25]) # noqa: E241

bkg = SplineModel(prefix='bkg_', xknots=knot_xvals1)
params.update(bkg.guess(y, x))

model = model + bkg

plt.plot(x, model.eval(params, x=x), label='initial')

out = model.fit(y, params, x=x)
print(out.fit_report(min_correl=0.3))
comps = out.eval_components()

plt.plot(x, out.best_fit, label='best fit')
plt.plot(x, comps['bkg_'], label='background')
plt.plot(x, comps['peak_'], label='peak')

knot_yvals = np.array([o.value for o in out.params.values() if o.name.startswith('bkg')])
plt.plot(knot_xvals1, knot_yvals, 'o', color='black', label='spline knots values')
plt.legend()
plt.show()

# knot positions          | peak amplitude
# 11, 13, 19, 21          | 12.223  0.295
# 11, 13, 16, 19, 21       | 11.746  0.594
# 11, 13, 15, 17, 19, 21   | 12.052  0.872

```

(continues on next page)

(continued from previous page)


```
plt.plot(x, y, 'o', label='data')

for nknots in (10, 15, 20, 25, 30):
    model = SplineModel(prefix='bkg_', xknots=np.linspace(0, 25, nknots))
    params = model.guess(y, x)
    out = model.fit(y, params, x=x)
    plt.plot(x, out.best_fit, label=f'best-fit ({nknots} knots)')
plt.legend()
plt.show()


# <end examples/doc_builtinmodels_splinemodel.py>
```

Total running time of the script: (0 minutes 0.407 seconds)

14.1.22 Confidence - advanced



•



```
[[Variables]]
a1: 2.98622095 +/- 0.14867027 (4.98%) (init = 2.986237)
a2: -4.33526363 +/- 0.11527574 (2.66%) (init = -4.335256)
t1: 1.30994276 +/- 0.13121215 (10.02%) (init = 1.309932)
t2: 11.8240337 +/- 0.46316956 (3.92%) (init = 11.82408)
[[Correlations]] (unreported correlations are < 0.500)
C(a2, t2) = +0.9871
C(a2, t1) = -0.9246
C(t1, t2) = -0.8805
C(a1, t1) = -0.5988
      95.45%   68.27%   _BEST_   68.27%   95.45%
a1: -0.27285  -0.14165  2.98622  +0.16354  +0.36343
a2: -0.30440  -0.13219  -4.33526  +0.10689  +0.19684
t1: -0.23392  -0.12494  1.30994  +0.14660  +0.32369
t2: -1.01937  -0.48813  11.82403 +0.46045  +0.90439
```

```
# <examples/doc_confidence_advanced.py>
import matplotlib.pyplot as plt
import numpy as np
```

(continues on next page)

(continued from previous page)

```

import lmfit

x = np.linspace(1, 10, 250)
np.random.seed(0)
y = 3.0*np.exp(-x/2) - 5.0*np.exp(-(x-0.1)/10.) + 0.1*np.random.randn(x.size)

p = lmfit.create_params(a1=4, a2=4, t1=3, t2=3)

def residual(p):
    return p['a1']*np.exp(-x/p['t1']) + p['a2']*np.exp(-(x-0.1)/p['t2']) - y

# create Minimizer
mini = lmfit.Minimizer(residual, p, nan_policy='propagate')

# first solve with Nelder-Mead algorithm
out1 = mini.minimize(method='Nelder')

# then solve with Levenberg-Marquardt using the
# Nelder-Mead solution as a starting point
out2 = mini.minimize(method='leastsq', params=out1.params)

lmfit.report_fit(out2.params, min_correl=0.5)

ci, trace = lmfit.conf_interval(mini, out2, sigmas=[1, 2], trace=True)
lmfit.printfuncs.report_ci(ci)

# plot data and best fit
plt.figure()
plt.plot(x, y)
plt.plot(x, residual(out2.params) + y, '-')
plt.show()

# plot confidence intervals (a1 vs t2 and a2 vs t2)
fig, axes = plt.subplots(1, 2, figsize=(12.8, 4.8))
cx, cy, grid = lmfit.conf_interval2d(mini, out2, 'a1', 't2', 30, 30)
ctp = axes[0].contourf(cx, cy, grid, np.linspace(0, 1, 11))
fig.colorbar(ctp, ax=axes[0])
axes[0].set_xlabel('a1')
axes[0].set_ylabel('t2')

cx, cy, grid = lmfit.conf_interval2d(mini, out2, 'a2', 't2', 30, 30)
ctp = axes[1].contourf(cx, cy, grid, np.linspace(0, 1, 11))
fig.colorbar(ctp, ax=axes[1])
axes[1].set_xlabel('a2')
axes[1].set_ylabel('t2')
plt.show()

# plot dependence between two parameters
fig, axes = plt.subplots(1, 2, figsize=(12.8, 4.8))

```

(continues on next page)

(continued from previous page)

```

cx1, cy1, prob = trace['a1']['a1'], trace['a1']['t2'], trace['a1']['prob']
cx2, cy2, prob2 = trace['t2']['t2'], trace['t2']['a1'], trace['t2']['prob']

axes[0].scatter(cx1, cy1, c=prob, s=30)
axes[0].set_xlabel('a1')
axes[0].set_ylabel('t2')

axes[1].scatter(cx2, cy2, c=prob2, s=30)
axes[1].set_xlabel('t2')
axes[1].set_ylabel('a1')
plt.show()
# <end examples/doc_confidence_advanced.py>

```

Total running time of the script: (0 minutes 5.638 seconds)

14.1.23 Parameters - uvars

```

[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
  ↪ prefix='exp_'))
[[Fit Statistics]]
  # fitting method    = leastsq
  # function evals   = 46
  # data points      = 250
  # variables        = 8
  chi-square         = 1247.52821
  reduced chi-square = 5.15507524
  Akaike info crit   = 417.864631
  Bayesian info crit = 446.036318
  R-squared          = 0.99648654
[[Variables]]
  g1_amplitude:   4257.77390 +/- 42.3837959 (1.00%) (init = 4000)
  g1_center:      107.030957 +/- 0.15006877 (0.14%) (init = 105)
  g1_sigma:       16.6725785 +/- 0.16048211 (0.96%) (init = 15)
  g2_amplitude:   2493.41716 +/- 36.1697062 (1.45%) (init = 2000)
  g2_center:      153.270104 +/- 0.19466774 (0.13%) (init = 155)
  g2_sigma:       13.8069453 +/- 0.18680153 (1.35%) (init = 15)
  exp_amplitude:  99.0183277 +/- 0.53748650 (0.54%) (init = 100)
  exp_decay:      90.9508838 +/- 1.10310767 (1.21%) (init = 100)
  g1_fwhm:        39.2609214 +/- 0.37790648 (0.96%) == '2.3548200*g1_sigma'
  g1_height:      101.880229 +/- 0.59217051 (0.58%) == '0.3989423*g1_amplitude/max(1e-
  ↪ 15, g1_sigma)'
  g2_fwhm:        32.5128709 +/- 0.43988398 (1.35%) == '2.3548200*g2_sigma'
  g2_height:      72.0455939 +/- 0.61721985 (0.86%) == '0.3989423*g2_amplitude/max(1e-
  ↪ 15, g2_sigma)'
[[Correlations]] (unreported correlations are < 0.500)
  C(g1_amplitude, g1_sigma)    = +0.8243
  C(g2_amplitude, g2_sigma)    = +0.8154
  C(exp_amplitude, exp_decay) = -0.6946
  C(g1_sigma, g2_center)       = +0.6842
  C(g1_center, g2_amplitude)   = -0.6689

```

(continues on next page)

(continued from previous page)

```

C(g1_center, g2_sigma)      = -0.6520
C(g1_amplitude, g2_center)  = +0.6477
C(g1_center, g2_center)     = +0.6205
C(g1_center, g1_sigma)      = +0.5075
C(g1_amplitude, exp_decay) = -0.5074
### Peak Area:
Peak1: 4257.77390 +/- 42.38380
Peak2: 2493.41716 +/- 36.16971
Sum : 6751.19106 +/- 46.50802
Correlation of peak1 and peak2 areas:      -0.30712
Stderr Quadrature sum for peak1 and peak2 area2: 55.71924
Stderr of peak1 + peak2 area, with correlation: 46.50802
### Peak Centroid:
Peak Centroid: 124.10846 +/- 0.14074
[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
  ↪ prefix='exp_'))
[[Fit Statistics]]
  # fitting method = leastsq
  # function evals = 46
  # data points   = 250
  # variables     = 8
  chi-square      = 1247.52821
  reduced chi-square = 5.15507524
  Akaike info crit = 417.864631
  Bayesian info crit = 446.036318
  R-squared        = 0.99648654
[[Variables]]
  g1_amplitude: 4257.77390 +/- 42.3837959 (1.00%) (init = 4000)
  g1_center:    107.030957 +/- 0.15006877 (0.14%) (init = 105)
  g1_sigma:     16.6725785 +/- 0.16048211 (0.96%) (init = 15)
  g2_amplitude: 2493.41716 +/- 36.1697062 (1.45%) (init = 2000)
  g2_center:    153.270104 +/- 0.19466774 (0.13%) (init = 155)
  g2_sigma:     13.8069453 +/- 0.18680153 (1.35%) (init = 15)
  exp_amplitude: 99.0183277 +/- 0.53748650 (0.54%) (init = 100)
  exp_decay:    90.9508838 +/- 1.10310767 (1.21%) (init = 100)
  g1_fwhm:      39.2609214 +/- 0.37790648 (0.96%) == '2.3548200*g1_sigma'
  g1_height:    101.880229 +/- 0.59217051 (0.58%) == '0.3989423*g1_amplitude/max(1e-
  ↪ 15, g1_sigma)'
  g2_fwhm:      32.5128709 +/- 0.43988398 (1.35%) == '2.3548200*g2_sigma'
  g2_height:    72.0455939 +/- 0.61721985 (0.86%) == '0.3989423*g2_amplitude/max(1e-
  ↪ 15, g2_sigma)'
  post_area:    6751.19106 +/- 46.5080249 (0.69%) == 'g1_amplitude + g2_amplitude'
  post_centroid: 124.108459 +/- 0.14074482 (0.11%) == '(g1_amplitude*g1_center+ g2_
  ↪ amplitude*g2_center)/(g1_amplitude + g2_amplitude)'
[[Correlations]] (unreported correlations are < 0.500)
  C(g1_amplitude, g1_sigma) = +0.8243
  C(g2_amplitude, g2_sigma) = +0.8154
  C(exp_amplitude, exp_decay) = -0.6946
  C(g1_sigma, g2_center) = +0.6842
  C(g1_center, g2_amplitude) = -0.6689
  C(g1_center, g2_sigma) = -0.6520

```

(continues on next page)

(continued from previous page)

```

C(g1_amplitude, g2_center)  = +0.6477
C(g1_center, g2_center)     = +0.6205
C(g1_center, g1_sigma)      = +0.5075
C(g1_amplitude, exp_decay) = -0.5074
[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
  ↪ prefix='exp_'))
[[Fit Statistics]]
  # fitting method   = leastsq
  # function evals  = 46
  # data points    = 250
  # variables      = 8
  chi-square        = 1247.52821
  reduced chi-square = 5.15507524
  Akaike info crit  = 417.864631
  Bayesian info crit = 446.036318
  R-squared         = 0.99648654
[[Variables]]
  g1_amplitude: 4257.77390 +/- 42.3837959 (1.00%) (init = 4000)
  g1_center:    107.030957 +/- 0.15006877 (0.14%) (init = 105)
  g1_sigma:     16.6725785 +/- 0.16048211 (0.96%) (init = 15)
  g2_amplitude: 2493.41716 +/- 36.1697062 (1.45%) (init = 2000)
  g2_center:    153.270104 +/- 0.19466774 (0.13%) (init = 155)
  g2_sigma:     13.8069453 +/- 0.18680153 (1.35%) (init = 15)
  exp_amplitude: 99.0183277 +/- 0.53748650 (0.54%) (init = 100)
  exp_decay:    90.9508838 +/- 1.10310767 (1.21%) (init = 100)
  g1_fwhm:      39.2609214 +/- 0.37790648 (0.96%) == '2.3548200*g1_sigma'
  g1_height:    101.880229 +/- 0.59217051 (0.58%) == '0.3989423*g1_amplitude/max(1e-
  ↪ 15, g1_sigma)'
  g2_fwhm:      32.5128709 +/- 0.43988398 (1.35%) == '2.3548200*g2_sigma'
  g2_height:    72.0455939 +/- 0.61721985 (0.86%) == '0.3989423*g2_amplitude/max(1e-
  ↪ 15, g2_sigma)'
  area:         6751.19106 +/- 46.5080249 (0.69%) == 'g1_amplitude + g2_amplitude'
  centroid:    124.108459 +/- 0.14074482 (0.11%) == '(g1_amplitude*g1_center+ g2_
  ↪ amplitude*g2_center)/(g1_amplitude + g2_amplitude)'
  post_area:    6751.19106 +/- 46.5080249 (0.69%) == 'g1_amplitude + g2_amplitude'
  post_centroid: 124.108459 +/- 0.14074482 (0.11%) == '(g1_amplitude*g1_center+ g2_
  ↪ amplitude*g2_center)/(g1_amplitude + g2_amplitude)'
[[Correlations]] (unreported correlations are < 0.500)
  C(g1_amplitude, g1_sigma)  = +0.8243
  C(g2_amplitude, g2_sigma)  = +0.8154
  C(exp_amplitude, exp_decay) = -0.6946
  C(g1_sigma, g2_center)     = +0.6842
  C(g1_center, g2_amplitude) = -0.6689
  C(g1_center, g2_sigma)     = -0.6520
  C(g1_amplitude, g2_center) = +0.6477
  C(g1_center, g2_center)    = +0.6205
  C(g1_center, g1_sigma)     = +0.5075
  C(g1_amplitude, exp_decay) = -0.5074

```

```

# <examples/doc_parameters_uvars.py>
import numpy as np

from lmfit.models import ExponentialModel, GaussianModel

dat = np.loadtxt('NIST_Gauss2.dat')
x = dat[:, 1]
y = dat[:, 0]

bkg = ExponentialModel(prefix='exp_')
gauss1 = GaussianModel(prefix='g1_')
gauss2 = GaussianModel(prefix='g2_')

mod = gauss1 + gauss2 + bkg

pars = mod.make_params(g1_center=105, g1_amplitude=4000, g1_sigma={'value': 15, 'min': 0}
                       ,
                       g2_center=155, g2_amplitude=2000, g2_sigma={'value': 15, 'min': 0}
                       ,
                       exp_amplitude=100, exp_decay=100)

out = mod.fit(y, pars, x=x)

print(out.fit_report(min_correl=0.5))

# Area and Centroid of the combined peaks
# option 1: from output uncertainties values

uvar_g1amp = out.uvars['g1_amplitude']
uvar_g2amp = out.uvars['g2_amplitude']
uvar_g1cen = out.uvars['g1_center']
uvar_g2cen = out.uvars['g2_center']

print("### Peak Area: ")
print(f"Peak1: {out.params['g1_amplitude'].value:.5f}+/-{out.params['g1_amplitude'].stderr:.5f}")
print(f"Peak2: {out.params['g2_amplitude'].value:.5f}+/-{out.params['g2_amplitude'].stderr:.5f}")
print(f"Sum : {(uvar_g1amp + uvar_g2amp):.5f}")

area_quadrature = np.sqrt(out.params['g1_amplitude'].stderr**2 + out.params['g2_amplitude'].stderr**2)

print(f"Correlation of peak1 and peak2 areas: {out.params['g1_amplitude'].correl['g2_amplitude']:.5f}")
print(f"Stderr Quadrature sum for peak1 and peak2 area2: {area_quadrature:.5f}")
print(f"Stderr of peak1 + peak2 area, with correlation: {(uvar_g1amp+uvar_g2amp).stdDev:.5f}")

print("### Peak Centroid: ")

centroid = (uvar_g1amp * uvar_g1cen + uvar_g2amp * uvar_g2cen) / (uvar_g1amp + uvar_g2amp)

```

(continues on next page)

(continued from previous page)

```
↳g2amp)

print(f"Peak Centroid: {centroid:.5f}")

# option 2: define corresponding variables after fit

def post_fit(result):
    "example post fit function"
    result.params.add('post_area', expr='g1_amplitude + g2_amplitude')
    result.params.add('post_centroid', expr='(g1_amplitude*g1_center+ g2_amplitude*g2_
↳center)/(g1_amplitude + g2_amplitude)')

mod.post_fit = post_fit


out = mod.fit(y, pars, x=x)
print(out.fit_report(min_correl=0.5))

# option 3: define corresponding variables before fit
pars.add('area', expr='g1_amplitude + g2_amplitude')
pars.add('centroid', expr='(g1_amplitude*g1_center+ g2_amplitude*g2_center)/(g1_
↳amplitude + g2_amplitude)')

out = mod.fit(y, pars, x=x)
print(out.fit_report(min_correl=0.5))
# <end examples/doc_parameters_uvvars.py>
```

Total running time of the script: (0 minutes 0.102 seconds)

14.1.24 Model - uncertainty2



```
[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) + Model(exponential,
  ↪ prefix='bkg_'))
[[Fit Statistics]]
  # fitting method = leastsq
  # function evals = 55
  # data points   = 250
  # variables     = 8
  chi-square      = 1247.52821
  reduced chi-square = 5.15507524
  Akaike info crit = 417.864631
  Bayesian info crit = 446.036318
  R-squared        = 0.99648654
[[Variables]]
  g1_amplitude: 4257.77399 +/- 42.3838008 (1.00%) (init = 3000)
  g1_center:    107.030957 +/- 0.15006868 (0.14%) (init = 100)
  g1_sigma:     16.6725789 +/- 0.16048222 (0.96%) (init = 10)
  g2_amplitude: 2493.41715 +/- 36.1696228 (1.45%) (init = 3000)
  g2_center:    153.270104 +/- 0.19466723 (0.13%) (init = 150)
  g2_sigma:     13.8069453 +/- 0.18680099 (1.35%) (init = 10)
  bkg_amplitude: 99.0183280 +/- 0.53748639 (0.54%) (init = 100)
```

(continues on next page)

(continued from previous page)

```

bkg_decay:      90.9508824 +/- 1.10310769 (1.21%) (init = 80)
g1_fwhm:        39.2609222 +/- 0.37790675 (0.96%) == '2.3548200*g1_sigma'
g1_height:      101.880228 +/- 0.59217122 (0.58%) == '0.3989423*g1_amplitude/max(1e-
˓→15, g1_sigma)'
g2_fwhm:        32.5128710 +/- 0.43988270 (1.35%) == '2.3548200*g2_sigma'
g2_height:      72.0455936 +/- 0.61721901 (0.86%) == '0.3989423*g2_amplitude/max(1e-
˓→15, g2_sigma)'

[[Correlations]] (unreported correlations are < 0.500)
C(g1_amplitude, g1_sigma) = +0.8243
C(g2_amplitude, g2_sigma) = +0.8154
C(bkg_amplitude, bkg_decay) = -0.6946
C(g1_sigma, g2_center) = +0.6842
C(g1_center, g2_amplitude) = -0.6689
C(g1_center, g2_sigma) = -0.6520
C(g1_amplitude, g2_center) = +0.6477
C(g1_center, g2_center) = +0.6205
C(g1_center, g1_sigma) = +0.5075
C(g1_amplitude, bkg_decay) = -0.5074

```

```

# <examples/doc_model_uncertainty2.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit.models import ExponentialModel, GaussianModel

dat = np.loadtxt('NIST_Gauss2.dat')
x = dat[:, 1]
y = dat[:, 0]

model = (GaussianModel(prefix='g1_') +
         GaussianModel(prefix='g2_') +
         ExponentialModel(prefix='bkg_'))

params = model.make_params(bkg_amplitude=100, bkg_decay=80,
                           g1_amplitude=3000,
                           g1_center=100,
                           g1_sigma=10,
                           g2_amplitude=3000,
                           g2_center=150,
                           g2_sigma=10)

result = model.fit(y, params, x=x)
print(result.fit_report(min_correl=0.5))

comps = result.eval_components(x=x)
dely = result.eval_uncertainty(sigma=3)

```

(continues on next page)

(continued from previous page)

```

fig, axes = plt.subplots(2, 2, figsize=(12.8, 9.6))

axes[0][0].plot(x, y, 'o', color='#99002299', markersize=3, label='data')
axes[0][0].plot(x, result.best_fit, '--', label='best fit')
axes[0][0].plot(x, result.init_fit, '---', label='initial fit')
axes[0][0].set_title('data, initial fit, and best-fit')
axes[0][0].legend()

axes[0][1].plot(x, y, 'o', color='#99002299', markersize=3, label='data')
axes[0][1].plot(x, result.best_fit, '--', label='best fit')
axes[0][1].fill_between(x, result.best_fit-dely, result.best_fit+dely,
                       color="#8A8A8A", label=r'3-$\sigma$ band')
axes[0][1].set_title('data, best-fit, and uncertainty band')
axes[0][1].legend()

axes[1][0].plot(x, result.best_fit, '--', label=r'best fit, 3-$\sigma$ band')
axes[1][0].fill_between(x,
                       result.best_fit-dely,
                       result.best_fit+dely,
                       color="#8A8A8A")

axes[1][0].plot(x, comps['bkg_'], label=r'background, 3-$\sigma$ band')
axes[1][0].fill_between(x,
                       comps['bkg_']-result.dely_comps['bkg_'],
                       comps['bkg_']+result.dely_comps['bkg_'],
                       color="#8A8A8A")

axes[1][0].plot(x, comps['g1_'], label=r'Gaussian #1, 3-$\sigma$ band')
axes[1][0].fill_between(x,
                       comps['g1_']-result.dely_comps['g1_'],
                       comps['g1_']+result.dely_comps['g1_'],
                       color="#8A8A8A")

axes[1][0].plot(x, comps['g2_'], label=r'Gaussian #2, 3-$\sigma$ band')
axes[1][0].fill_between(x,
                       comps['g2_']-result.dely_comps['g2_'],
                       comps['g2_']+result.dely_comps['g2_'],
                       color="#8A8A8A")

axes[1][0].set_title('model components with uncertainty bands')
axes[1][0].legend()


axes[1][1].plot(x, result.best_fit, '--', label='best fit')
axes[1][1].plot(x, 10*result.dely, label=r'3-$\sigma$ total (x10)')
axes[1][1].plot(x, 10*result.dely_comps['bkg_'], label=r'3-$\sigma$ background (x10)')
axes[1][1].plot(x, 10*result.dely_comps['g1_'], label=r'3-$\sigma$ Gaussian #1 (x10)')
axes[1][1].plot(x, 10*result.dely_comps['g2_'], label=r'3-$\sigma$ Gaussian #2 (x10)')
axes[1][1].set_title('uncertainties for model components')
axes[1][1].legend()


plt.show()
# <end examples/doc_model_uncertainty2.py>


```

Total running time of the script: (0 minutes 0.898 seconds)

14.1.25 Fitting - emcee







```
•
[[Variables]]
a1: 2.98623689 +/- 0.15010519 (5.03%) (init = 4)
a2: -4.33525597 +/- 0.11765824 (2.71%) (init = 4)
t1: 1.30993186 +/- 0.13449656 (10.27%) (init = 3)
t2: 11.8240752 +/- 0.47172610 (3.99%) (init = 3)
[[Correlations]] (unreported correlations are < 0.500)
C(a2, t2) = +0.9876
C(a2, t1) = -0.9278
C(t1, t2) = -0.8852
C(a1, t1) = -0.6093
The chain is shorter than 50 times the integrated autocorrelation time for 5 parameter(s). Use this estimate with caution and run a longer chain!
N/50 = 20;
tau: [42.15955322 47.347426 48.71211873 46.7985718 40.89881208]
```

median of posterior probability distribution

```
-----[[Variables]]
a1: 2.98945718 +/- 0.14033921 (4.69%) (init = 2.986237)
a2: -4.34687243 +/- 0.12131092 (2.79%) (init = -4.335256)
t1: 1.32883916 +/- 0.13766047 (10.36%) (init = 1.309932)
t2: 11.7836194 +/- 0.47719763 (4.05%) (init = 11.82408)
__lnsigma: -2.32559226 +/- 0.04542650 (1.95%) (init = -2.302585)
[[Correlations]] (unreported correlations are < 0.100)
```

(continues on next page)

(continued from previous page)

```
C(a2, t2) = +0.9811
C(a2, t1) = -0.9377
C(t1, t2) = -0.8943
C(a1, t1) = -0.5076
C(a1, a2) = +0.2140
C(a1, t2) = +0.1777
```

Maximum Likelihood Estimation from emcee

Parameter	MLE Value	Median Value	Uncertainty
a1	2.93839	2.98946	0.14034
a2	-4.35274	-4.34687	0.12131
t1	1.34310	1.32884	0.13766
t2	11.78782	11.78362	0.47720

Error Estimates from emcee

Parameter	-2sigma	-1sigma	median	+1sigma	+2sigma
a1	-0.2656	-0.1362	2.9895	0.1445	0.3141
a2	-0.3209	-0.1309	-4.3469	0.1118	0.1985
t1	-0.2377	-0.1305	1.3288	0.1448	0.3278
t2	-1.0677	-0.4807	11.7836	0.4739	0.8990

```
# <examples/doc_fitting_emcee.py>
import numpy as np

import lmfit

try:
    import matplotlib.pyplot as plt
    HASPYLAB = True
except ImportError:
    HASPYLAB = False

try:
    import corner
    HASCORNER = True
except ImportError:
    HASCORNER = False

x = np.linspace(1, 10, 250)
np.random.seed(0)
y = (3.0*np.exp(-x/2) - 5.0*np.exp(-(x-0.1) / 10.) +
     0.1*np.random.randn(x.size))

p = lmfit.Parameters()
p.add_many(('a1', 4), ('a2', 4), ('t1', 3), ('t2', 3., True))
```

(continues on next page)

(continued from previous page)

```

def residual(p):
    v = p.valuesdict()
    return v['a1']*np.exp(-x/v['t1']) + v['a2']*np.exp(-(x-0.1) / v['t2']) - y

mi = lmfit.minimize(residual, p, method='nelder', nan_policy='omit')
lmfit.printfuncs.report_fit(mi.params, min_correl=0.5)
if HASPYLAB:
    plt.figure()
    plt.plot(x, y, 'o')
    plt.plot(x, residual(mi.params) + y, label='best fit')
    plt.legend()
    plt.show()

# Place bounds on the ln(sigma) parameter that emcee will automatically add
# to estimate the true uncertainty in the data since is_weighted=False
mi.params.add('__lnsigma', value=np.log(0.1), min=np.log(0.001), max=np.log(2))

res = lmfit.minimize(residual, method='emcee', nan_policy='omit', burn=300,
                      steps=1000, thin=20, params=mi.params, is_weighted=False,
                      progress=False)

if HASPYLAB and HASCORNER:
    emcee_corner = corner.corner(res.flatchain, labels=res.var_names,
                                  truths=list(res.params.valuesdict().values()))
    plt.show()

if HASPYLAB:
    plt.plot(res.acceptance_fraction, 'o')
    plt.xlabel('walker')
    plt.ylabel('acceptance fraction')
    plt.show()

if hasattr(res, "acor"):
    print("Autocorrelation time for the parameters:")
    print("-----")
    for i, par in enumerate(p):
        print(par, res.acor[i])

print("\nmedian of posterior probability distribution")
print('-----')
lmfit.report_fit(res.params)

# find the maximum likelihood solution
highest_prob = np.argmax(res.lnprob)
hp_loc = np.unravel_index(highest_prob, res.lnprob.shape)
mle_soln = res.chain[hp_loc]
for i, par in enumerate(p):
    p[par].value = mle_soln[i]

```

(continues on next page)

(continued from previous page)

```

print('\nMaximum Likelihood Estimation from emcee      ')
print('-----')
print('Parameter  MLE Value  Median Value  Uncertainty')
fmt = '  {:5s}  {:11.5f}  {:11.5f}  {:11.5f}'.format
for name, param in p.items():
    print(fmt(name, param.value, res.params[name].value,
              res.params[name].stderr))

if HASPYLAB:
    plt.figure()
    plt.plot(x, y, 'o')
    plt.plot(x, residual(mi.params) + y, label='Nelder-Mead')
    plt.plot(x, residual(res.params) + y, '--', label='emcee')
    plt.legend()
    plt.show()

print('\nError Estimates from emcee      ')
print('-----')
print('Parameter  -2sigma  -1sigma  median  +1sigma  +2sigma')

for name in p.keys():
    quantiles = np.percentile(res.flatchain[name],
                               [2.275, 15.865, 50, 84.135, 97.275])
    median = quantiles[2]
    err_m2 = quantiles[0] - median
    err_m1 = quantiles[1] - median
    err_p1 = quantiles[3] - median
    err_p2 = quantiles[4] - median
    fmt = '  {:5s}  {:8.4f}  {:8.4f}  {:8.4f}  {:8.4f}  {:8.4f}'.format
    print(fmt(name, err_m2, err_m1, median, err_p1, err_p2))

```

Total running time of the script: (0 minutes 9.813 seconds)

14.1.26 Confidence - chi2 maps

```

# <examples/doc_confidence_chi2_maps.py>

import matplotlib.pyplot as plt
import numpy as np

from lmfit import conf_interval, conf_interval2d, report_ci
from lmfit.lineshapes import gaussian
from lmfit.models import GaussianModel, LinearModel

sigma_levels = [1, 2, 3]

rng = np.random.default_rng(seed=102)

```

set up data – deliberately adding imperfections and a small amount of non-Gaussian noise

```

npts = 501
x = np.linspace(1, 100, num=npts)

noise = rng.normal(scale=0.3, size=npts) + 0.2*rng.f(3, 9, size=npts)

y = (gaussian(x, amplitude=83, center=47., sigma=5.)
     + 0.02*x + 4 + 0.25*np.cos((x-20)/8.0) + noise)

mod = GaussianModel() + LinearModel()
params = mod.make_params(amplitude=100, center=50, sigma=5,
                        slope=0, intercept=2)

out = mod.fit(y, params, x=x)
print(out.fit_report(correl_mode='table'))

```

```

[[Model]]
    (Model(gaussian) + Model(linear))
[[Fit Statistics]]
# fitting method    = leastsq
# function evals   = 31
# data points      = 501
# variables        = 5
chi-square         = 103.861381
reduced chi-square = 0.20939794
Akaike info crit   = -778.348033
Bayesian info crit = -757.265003
R-squared          = 0.93782756
[[Variables]]
amplitude: 78.8171374 +/- 1.21910939 (1.55%) (init = 100)
center:    47.0751649 +/- 0.07576660 (0.16%) (init = 50)
sigma:     4.93298753 +/- 0.07984021 (1.62%) (init = 5)
slope:     0.01839006 +/- 7.1957e-04 (3.91%) (init = 0)
intercept: 4.39234411 +/- 0.04420227 (1.01%) (init = 0)
fwhm:      11.6162977 +/- 0.18800933 (1.62%) == '2.3548200*sigma'
height:    6.37412722 +/- 0.08603873 (1.35%) == '0.3989423*amplitude/max(1e-15,_
sigma)'
[[Correlations]]
+-----+-----+-----+-----+-----+-----+
| Variable | amplitude | center | sigma | slope | intercept |
+-----+-----+-----+-----+-----+-----+
| amplitude | +1.0000 | -0.0074 | +0.6371 | +0.0721 | -0.3373 |
| center | -0.0074 | +1.0000 | -0.0048 | -0.1026 | +0.0864 |
| sigma | +0.6371 | -0.0048 | +1.0000 | +0.0459 | -0.2149 |
| slope | +0.0721 | -0.1026 | +0.0459 | +1.0000 | -0.8421 |
| intercept | -0.3373 | +0.0864 | -0.2149 | -0.8421 | +1.0000 |
+-----+-----+-----+-----+-----+

```

run conf_intervale, print report

```

ci = conf_interval(out, out, sigmas=sigma_levels)

print("## Confidence Report:")
report_ci(ci)

```

```
## Confidence Report:
      99.73%   95.45%   68.27%   _BEST_   68.27%   95.45%   99.73%
amplitude: -3.62610 -2.41983 -1.21237 78.81714 +1.22111 +2.45479 +3.70515
center    : -0.22849 -0.15214 -0.07584 47.07516 +0.07587 +0.15225 +0.22873
sigma     : -0.23335 -0.15640 -0.07870 4.93299 +0.08000 +0.16158 +0.24509
slope     : -0.00217 -0.00144 -0.00072 0.01839 +0.00072 +0.00144 +0.00217
intercept: -0.13326 -0.08860 -0.04423 4.39234 +0.04421 +0.08854 +0.13312
```

plot initial fit

```
colors = ('#2030b0', '#b02030', '#207070')
fig, axes = plt.subplots(2, 3, figsize=(15, 9.5))

axes[0, 0].plot(x, y, 'o', markersize=3, label='data', color=colors[0])
axes[0, 0].plot(x, out.best_fit, label='fit', color=colors[1])
axes[0, 0].set_xlabel('x')
axes[0, 0].set_ylabel('y')
axes[0, 0].legend()

aix, aiy = 0, 0
nsamples = 50
explicitly_calculate_sigma = True

for pairs in (('sigma', 'amplitude'), ('intercept', 'amplitude'),
              ('slope', 'intercept'), ('slope', 'center'), ('sigma', 'center')):

    xpar, ypar = pairs
    if explicitly_calculate_sigma:
        print(f"Generating chi-square map for {pairs}")
        c_x, c_y, chi2_mat = conf_interval2d(out, out, xpar, ypar,
                                              nsamples, nsamples, nsigma=3.5,
                                              chi2_out=True)
        # explicitly calculate sigma matrix: sigma increases chi_square
        # from chi_square_best
        # to chi_square + sigma**2 * reduced_chi_square
        # so: sigma = sqrt((chi2-chi2_best)/ reduced_chi_square)
        chi2_min = chi2_mat.min()
        sigma_mat = np.sqrt((chi2_mat-chi2_min)/out.redchi)
    else:
        print(f"Generating sigma map for {pairs}")
        # or, you could just calculate the matrix of probabilities as:
        c_x, c_y, sigma_mat = conf_interval2d(out, out, xpar, ypar,
                                              nsamples, nsamples, nsigma=3.5)

    aix += 1
    if aix == 2:
        aix = 0
        aiy += 1
    ax = axes[aix, aiy]

    cnt = ax.contour(c_x, c_y, sigma_mat, levels=sigma_levels, colors=colors,
```

(continues on next page)

(continued from previous page)

```

        linestyles='-' )
ax.clabel(cnt, inline=True, fmt=r"$\sigma=%.\mathbf{0f}$", fontsize=13)

# draw boxes for estimated uncertainties:
# dotted : scaled stderr from initial fit
# dashed : values found from conf_interval()
xv = out.params[xpar].value
xs = out.params[xpar].stderr
yv = out.params[ypar].value
ys = out.params[ypar].stderr

cix = ci[xpar]
ciy = ci[ypar]


nc = len(sigma_levels)
for i in sigma_levels:
    # dotted line: scaled stderr
    ax.plot((xv-i*xs, xv+i*xs, xv+i*xs, xv-i*xs),
             (yv-i*ys, yv-i*ys, yv+i*ys, yv+i*ys),
             linestyle='dotted', color=colors[i-1])

    # dashed line: refined uncertainties from conf_interval
    xsm, xsp = cix[nc+i][1], cix[nc-i][1]
    ysm, ysp = ciy[nc+i][1], ciy[nc-i][1]
    ax.plot((xsm, xsp, xsp, xsm), (ysm, ysm, ysp, ysp),
             linestyle='dashed', color=colors[i-1])

ax.set_xlabel(xpar)
ax.set_ylabel(ypar)
ax.grid(True, color='#d0d0d0')

plt.show()
# <end examples/doc_confidence_chi2_maps.py>


```



```
Generating chi-square map for ('sigma', 'amplitude')
Generating chi-square map for ('intercept', 'amplitude')
Generating chi-square map for ('slope', 'intercept')
Generating chi-square map for ('slope', 'center')
Generating chi-square map for ('sigma', 'center')
```

Total running time of the script: (0 minutes 57.484 seconds)

14.1.27 Model - uncertainty pred



```
[[Model]]
  Model(logistic_func)
[[Fit Statistics]]
  # fitting method    = leastsq
  # function evals   = 60
  # data points      = 272
  # variables        = 4
  chi-square         = 8469.42359
  reduced chi-square = 31.6023268
  Akaike info crit   = 943.249061
  Bayesian info crit = 957.672270
  R-squared          = 0.83090615
[[Variables]]
  amp:    82.4660404 +/- 0.99770817 (1.21%) (init = 90)
  off:    51.3218514 +/- 1.83125060 (3.57%) (init = 50)
  tau:    3.05525924 +/- 0.11068420 (3.62%) (init = 2)
  gamma:  2.25386376 +/- 0.43544396 (19.32%) (init = 2)
[[Correlations]] (unreported correlations are < 0.100)
  C(off, gamma) = +0.8634
  C(amp, gamma) = -0.7706
  C(off, tau)   = +0.7277
```

(continues on next page)

(continued from previous page)

```
C(amp, off)  = -0.5381
C(tau, gamma) = +0.4699
```

```
# <examples/doc_model_uncertainty_pred.py>

# Here we reproduce the results presented in Section 5.1 of the
# nonlinear regression lecture by Julien Chiquet
# https://jchiquet.github.io/MAP566/docs/regression/map566-lecture-nonlinear-regression.
#html

import matplotlib.pyplot as plt
import numpy as np

from lmfit import Model

# data from the R base package datasets for the waiting time between eruptions
# and the duration of the eruption for the Old Faithful geyser in Yellowstone
# National Park, Wyoming, USA.

tlen = np.array([3.6, 1.8, 3.333, 2.283, 4.533, 2.883, 4.7, 3.6, 1.95, 4.35,
                 1.833, 3.917, 4.2, 1.75, 4.7, 2.167, 1.75, 4.8, 1.6, 4.25,
                 1.8, 1.75, 3.45, 3.067, 4.533, 3.6, 1.967, 4.083, 3.85, 4.433,
                 4.3, 4.467, 3.367, 4.033, 3.833, 2.017, 1.867, 4.833, 1.833,
                 4.783, 4.35, 1.883, 4.567, 1.75, 4.533, 3.317, 3.833, 2.1,
                 4.633, 2, 4.8, 4.716, 1.833, 4.833, 1.733, 4.883, 3.717,
                 1.667, 4.567, 4.317, 2.233, 4.5, 1.75, 4.8, 1.817, 4.4, 4.167,
                 4.7, 2.067, 4.7, 4.033, 1.967, 4.5, 4, 1.983, 5.067, 2.017,
                 4.567, 3.883, 3.6, 4.133, 4.333, 4.1, 2.633, 4.067, 4.933,
                 3.95, 4.517, 2.167, 4, 2.2, 4.333, 1.867, 4.817, 1.833, 4.3,
                 4.667, 3.75, 1.867, 4.9, 2.483, 4.367, 2.1, 4.5, 4.05, 1.867,
                 4.7, 1.783, 4.85, 3.683, 4.733, 2.3, 4.9, 4.417, 1.7, 4.633,
                 2.317, 4.6, 1.817, 4.417, 2.617, 4.067, 4.25, 1.967, 4.6,
                 3.767, 1.917, 4.5, 2.267, 4.65, 1.867, 4.167, 2.8, 4.333,
                 1.833, 4.383, 1.883, 4.933, 2.033, 3.733, 4.233, 2.233, 4.533,
                 4.817, 4.333, 1.983, 4.633, 2.017, 5.1, 1.8, 5.033, 4, 2.4,
                 4.6, 3.567, 4, 4.5, 4.083, 1.8, 3.967, 2.2, 4.15, 2, 3.833,
                 3.5, 4.583, 2.367, 5, 1.933, 4.617, 1.917, 2.083, 4.583,
                 3.333, 4.167, 4.333, 4.5, 2.417, 4, 4.167, 1.883, 4.583, 4.25,
                 3.767, 2.033, 4.433, 4.083, 1.833, 4.417, 2.183, 4.8, 1.833,
                 4.8, 4.1, 3.966, 4.233, 3.5, 4.366, 2.25, 4.667, 2.1, 4.35,
                 4.133, 1.867, 4.6, 1.783, 4.367, 3.85, 1.933, 4.5, 2.383, 4.7,
                 1.867, 3.833, 3.417, 4.233, 2.4, 4.8, 2, 4.15, 1.867, 4.267,
                 1.75, 4.483, 4, 4.117, 4.083, 4.267, 3.917, 4.55, 4.083,
                 2.417, 4.183, 2.217, 4.45, 1.883, 1.85, 4.283, 3.95, 2.333,
                 4.15, 2.35, 4.933, 2.9, 4.583, 3.833, 2.083, 4.367, 2.133,
                 4.35, 2.2, 4.45, 3.567, 4.5, 4.15, 3.817, 3.917, 4.45, 2,
                 4.283, 4.767, 4.533, 1.85, 4.25, 1.983, 2.25, 4.75, 4.117,
```

(continues on next page)

(continued from previous page)

```

2.15, 4.417, 1.817, 4.467])

delay = np.array([79, 54, 74, 62, 85, 55, 88, 85, 51, 85, 54, 84, 78, 47, 83,
                  52, 62, 84, 52, 79, 51, 47, 78, 69, 74, 83, 55, 76, 78, 79,
                  73, 77, 66, 80, 74, 52, 48, 80, 59, 90, 80, 58, 84, 58, 73,
                  83, 64, 53, 82, 59, 75, 90, 54, 80, 54, 83, 71, 64, 77, 81,
                  59, 84, 48, 82, 60, 92, 78, 78, 65, 73, 82, 56, 79, 71, 62,
                  76, 60, 78, 76, 83, 75, 82, 70, 65, 73, 88, 76, 80, 48, 86,
                  60, 90, 50, 78, 63, 72, 84, 75, 51, 82, 62, 88, 49, 83, 81,
                  47, 84, 52, 86, 81, 75, 59, 89, 79, 59, 81, 50, 85, 59, 87,
                  53, 69, 77, 56, 88, 81, 45, 82, 55, 90, 45, 83, 56, 89, 46,
                  82, 51, 86, 53, 79, 81, 60, 82, 77, 76, 59, 80, 49, 96, 53,
                  77, 77, 65, 81, 71, 70, 81, 93, 53, 89, 45, 86, 58, 78, 66,
                  76, 63, 88, 52, 93, 49, 57, 77, 68, 81, 81, 73, 50, 85, 74,
                  55, 77, 83, 83, 51, 78, 84, 46, 83, 55, 81, 57, 76, 84, 77,
                  81, 87, 77, 51, 78, 60, 82, 91, 53, 78, 46, 77, 84, 49, 83,
                  71, 80, 49, 75, 64, 76, 53, 94, 55, 76, 50, 82, 54, 75, 78,
                  79, 78, 78, 70, 79, 70, 54, 86, 50, 90, 54, 54, 77, 79, 64,
                  75, 47, 86, 63, 85, 82, 57, 82, 67, 74, 54, 83, 73, 73, 88,
                  80, 71, 83, 56, 79, 78, 84, 58, 83, 43, 60, 75, 81, 46, 90,
                  46, 74])

```

```

# model with a modified logistic function + offset
def logistic_func(t, amp, off, tau, gamma):
    return (amp-off)/(1+np.exp(-gamma*(t-tau))) + off

# set up model and initial parameter values
mod = Model(logistic_func)
pars = mod.make_params(amp=90, gamma=2, tau=2, off=50)

# run fit, and print report
result = mod.fit(delay, pars, t=tlen)
print(result.fit_report())

# make finely spaced grid of duration values, extending past data range
tfine = np.linspace(1, 6, 101)
yfine = result.eval(t=tfine)

# now calculate uncertainty interval and predicted interval for sigma=2, 95% level
efine = result.eval_uncertainty(t=tfine, sigma=2)
pfine = result.dely_predicted

plt.plot(tlen, delay, '.', label='observations')
plt.plot(tlen, result.best_fit, 'o', label='best fit')
plt.plot(tfine, yfine, '-', label='fine grid model')
plt.fill_between(tfine, yfine-efine, yfine+efine,
                 color="#c0c0c0", label=r'$2\sigma$ confidence interval')

plt.fill_between(tfine, yfine-pfine, yfine+pfine,
                 color="#d0d0a0", label=r'$2\sigma$ predicted interval')

```

(continues on next page)

(continued from previous page)

```
plt.xlabel('duration (minutes)')
plt.ylabel('delay to next eruption (minutes)')
plt.legend()
plt.show()

# <end examples/doc_model_uncertainty_pred.py>
```

Total running time of the script: (0 minutes 0.236 seconds)

PYTHON MODULE INDEX

|

`lmfit.confidence`, 155
`lmfit.minimizer`, 33
`lmfit.model`, 65
`lmfit.models`, 113
`lmfit.parameter`, 23

INDEX

A

aborted (*in module lmfit.model*), 91
aborted (*MinimizerResult attribute*), 40
add() (*Parameters method*), 25
add_many() (*Parameters method*), 26
aic (*in module lmfit.model*), 92
aic (*MinimizerResult attribute*), 41
ampgo() (*Minimizer method*), 53

B

basinhopping() (*Minimizer method*), 52
best_fit (*in module lmfit.model*), 91
best_values (*in module lmfit.model*), 90
bic (*in module lmfit.model*), 92
bic (*MinimizerResult attribute*), 41
BreitWignerModel (*class in lmfit.models*), 120
brute() (*Minimizer method*), 51

C

call_kw() (*in module lmfit.model*), 92
call_kw() (*MinimizerResult attribute*), 40
chisqr (*in module lmfit.model*), 92
chisqr (*MinimizerResult attribute*), 41
ci_out (*in module lmfit.model*), 92
ci_report() (*in module lmfit*), 167
ci_report() (*ModelResult method*), 85
components (*in module lmfit.model*), 91
Composite models, 106
CompositeModel (*class in lmfit.model*), 108
conf_interval() (*in module lmfit*), 165
conf_interval() (*ModelResult method*), 85
conf_interval2d() (*in module lmfit*), 166
ConstantModel (*class in lmfit.models*), 127
correl (*Parameter attribute*), 24
covar (*in module lmfit.model*), 91
covar (*MinimizerResult attribute*), 39
create_params() (*in module lmfit.parameter*), 29
create_uvars() (*Parameters method*), 27

D

DampedHarmonicOscillatorModel (*class in lmfit.models*), 122

DampedOscillatorModel (*class in lmfit.models*), 121
data (*in module lmfit.model*), 91
dely (*in module lmfit.model*), 91
dely_comps (*in module lmfit.model*), 91
dely_predicted (*in module lmfit.model*), 91
DoniachModel (*class in lmfit.models*), 126
dual_annealing() (*Minimizer method*), 54
dump() (*Parameters method*), 27
dumps() (*Parameters method*), 27

E

emcee() (*Minimizer method*), 55
errorbars (*in module lmfit.model*), 91
errorbars (*MinimizerResult attribute*), 40
eval() (*Model method*), 70
eval() (*ModelResult method*), 84
eval() (*Parameters method*), 28
eval_components() (*ModelResult method*), 84
eval_uncertainty() (*ModelResult method*), 86
ExponentialGaussianModel (*class in lmfit.models*), 123
ExponentialModel (*class in lmfit.models*), 132
ExpressionModel (*class in lmfit.models*), 135

F

fit() (*Model method*), 70
fit() (*ModelResult method*), 84
fit_report() (*in module lmfit.printfuncs*), 44
fit_report() (*ModelResult method*), 84
flatchain (*in module lmfit.model*), 91
flatchain (*MinimizerResult attribute*), 40
func (*in module lmfit.model*), 74

G

Gaussian2dModel (*class in lmfit.models*), 134
GaussianModel (*class in lmfit.models*), 114
guess() (*Model method*), 71

I

ier (*in module lmfit.model*), 91
ier (*MinimizerResult attribute*), 40
independent_vars (*in module lmfit.model*), 74

`init_fit (in module lmfit.model)`, 91
`init_params (in module lmfit.model)`, 90
`init_vals (in module lmfit.model)`, 90
`init_vals (MinimizerResult attribute)`, 39
`init_values (in module lmfit.model)`, 90
`init_values (MinimizerResult attribute)`, 40
`iter_cb() (ModelResult method)`, 90

J

`jacfcn() (ModelResult method)`, 90

L

`least_squares() (Minimizer method)`, 49
`leastsq() (Minimizer method)`, 49
`LinearModel (class in lmfit.models)`, 127
`lmdif_message (in module lmfit.model)`, 91
`lmdif_message (MinimizerResult attribute)`, 40
`lmfit.confidence`
 `module`, 155
`lmfit.minimizer`
 `module`, 31
`lmfit.model`
 `module`, 65
`lmfit.models`
 `module`, 113
`lmfit.parameter`
 `module`, 21
`load() (Parameters method)`, 28
`load_model() (in module lmfit.model)`, 81
`load_modelresult() (in module lmfit.model)`, 104
`loads() (Parameters method)`, 28
`LognormalModel (class in lmfit.models)`, 121
`LorentzianModel (class in lmfit.models)`, 114

M

`make_params() (Model method)`, 72
`message (in module lmfit.model)`, 91
`message (MinimizerResult attribute)`, 40
`method (in module lmfit.model)`, 91
`minimize() (in module lmfit.minimizer)`, 33
`minimize() (Minimizer method)`, 48
`Minimizer (class in lmfit.minimizer)`, 46
`MinimizerResult (class in lmfit.minimizer)`, 39
`Model (class in lmfit.model)`, 69
`model (in module lmfit.model)`, 92
`ModelResult (class in lmfit.model)`, 83
`module`
 `lmfit.confidence`, 155
 `lmfit.minimizer`, 31
 `lmfit.model`, 65
 `lmfit.models`, 113
 `lmfit.parameter`, 21
`MoffatModel (class in lmfit.models)`, 117

N

`name (in module lmfit.model)`, 74
`nan_policy (in module lmfit.model)`, 74
`ndata (in module lmfit.model)`, 92
`ndata (MinimizerResult attribute)`, 41
`nfev (in module lmfit.model)`, 92
`nfev (MinimizerResult attribute)`, 41
`nfree (in module lmfit.model)`, 92
`nfree (MinimizerResult attribute)`, 41
`nvarys (in module lmfit.model)`, 92
`nvarys (MinimizerResult attribute)`, 41

O

`opts (in module lmfit.model)`, 74

P

`param_hints (in module lmfit.model)`, 74
`param_names (in module lmfit.model)`, 74
`Parameter (class in lmfit.parameter)`, 23
`Parameters (class in lmfit.parameter)`, 25
`params (in module lmfit.model)`, 90
`params (MinimizerResult attribute)`, 39
`Pearson4Model (class in lmfit.models)`, 118
`Pearson7Model (class in lmfit.models)`, 119
`plot() (ModelResult method)`, 86
`plot_fit() (ModelResult method)`, 88
`plot_residuals() (ModelResult method)`, 89
`PolynomialModel (class in lmfit.models)`, 128
`post_fit() (Model method)`, 73
`PowerLawModel (class in lmfit.models)`, 133
`prefix (in module lmfit.model)`, 74
`prepare_fit() (Minimizer method)`, 51
`pretty_print() (Parameters method)`, 26
`print_param_hints() (Model method)`, 73
`PseudoVoigtModel (class in lmfit.models)`, 117

Q

`QuadraticModel (class in lmfit.models)`, 128

R

`RectangleModel (class in lmfit.models)`, 131
`redchi (in module lmfit.model)`, 92
`redchi (MinimizerResult attribute)`, 41
`Removing a Constraint Expression`, 24
`residual (in module lmfit.model)`, 91
`residual (MinimizerResult attribute)`, 39
`rsquared (in module lmfit.model)`, 92

S

`save_model() (in module lmfit.model)`, 81
`save_modelresult() (in module lmfit.model)`, 104
`scalar_minimize() (Minimizer method)`, 50
`scale_covar (in module lmfit.model)`, 92

`set()` (*Parameter method*), 24
`set_param_hint()` (*Model method*), 73
`shgo()` (*Minimizer method*), 54
`show_candidates()` (*MinimizerResult method*), 41
`SineModel` (*class in lmfit.models*), 130
`SkewedGaussianModel` (*class in lmfit.models*), 123
`SkewedVoigtModel` (*class in lmfit.models*), 124
`SplineModel` (*class in lmfit.models*), 129
`SplitLorentzianModel` (*class in lmfit.models*), 115
`status` (*MinimizerResult attribute*), 40
`stderr` (*Parameter attribute*), 24
`StepModel` (*class in lmfit.models*), 130
`StudentsTModel` (*class in lmfit.models*), 120
`success` (*in module lmfit.model*), 92
`success` (*MinimizerResult attribute*), 40
`summary()` (*ModelResult method*), 85

T

`ThermalDistributionModel` (*class in lmfit.models*),
125

U

`userargs` (*in module lmfit.model*), 92
`userkws` (*in module lmfit.model*), 92
`uvars` (*in module lmfit.model*), 90
`uvars` (*MinimizerResult attribute*), 39

V

`valuesdict()` (*Parameters method*), 27
`var_names` (*in module lmfit.model*), 90
`var_names` (*MinimizerResult attribute*), 39
`VoigtModel` (*class in lmfit.models*), 116

W

`weights` (*in module lmfit.model*), 91